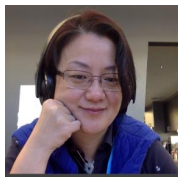


# AWS re:Invent

## Getting started with Amazon Aurora



Debanjan Saha  
GM, Amazon Aurora



Linda Xu  
Principal Architect, Ticketmaster

11/29/2016

# Outline

---

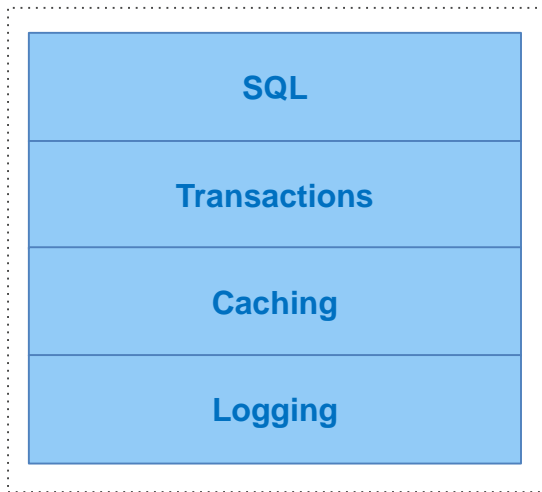
- ▶ What is Amazon Aurora
  - Background and history
- ▶ What you need to know about Aurora
  - Differentiators; use cases; cost of ownership
- ▶ Hear directly from one of our customers
  - Ticketmaster will share their experience

## A bit of history ...

Re-imagining relational databases for the cloud era

# Relational databases were not designed for the cloud

---



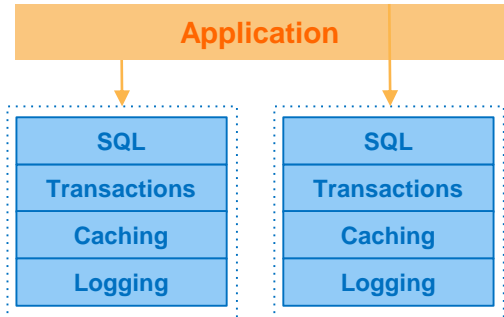
Multiple layers of  
functionality all in a  
monolithic stack

# Not much has changed in last 20 years

---

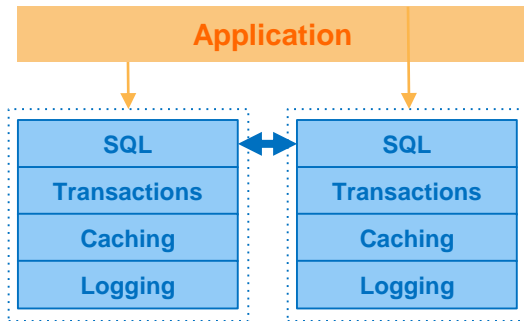
## Sharding

*Coupled at the application layer*



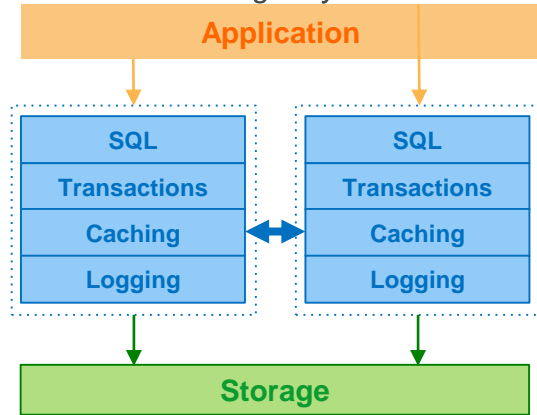
## Shared Nothing

*Coupled at the SQL layer*



## Shared Disk

*Coupled at the caching and storage layer*



Even when you scale it out, you're still replicating the same stack

# Re-imagining relational database

---

**1**

**Scale-out and distributed design**

**2**

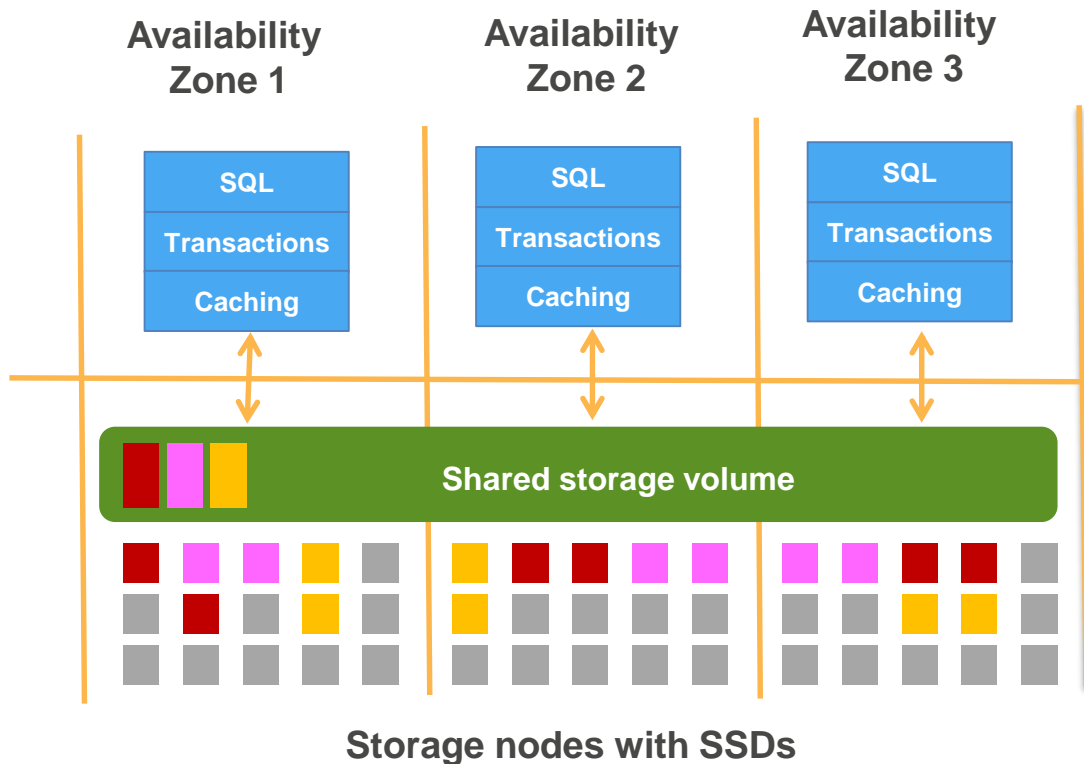
**Service-oriented architecture leveraging AWS services**

**3**

**Automate administrative tasks – fully managed service**

# Scale-out, distributed, multi-tenant architecture

- Purpose-built log-structured distributed storage system designed for databases
- Storage volume is striped across hundreds of storage nodes distributed over 3 different availability zones
- Six copies of data, two copies in each availability zone to protect against AZ+1 failures
- Plan to apply same principles to other layers of the stack



# Leveraging cloud ecosystem

---

Lambda



Invoke Lambda events from stored procedures/triggers.

S3



Load data from S3, store snapshots and backups in S3.

IAM



Use IAM roles to manage database access control.

CloudWatch

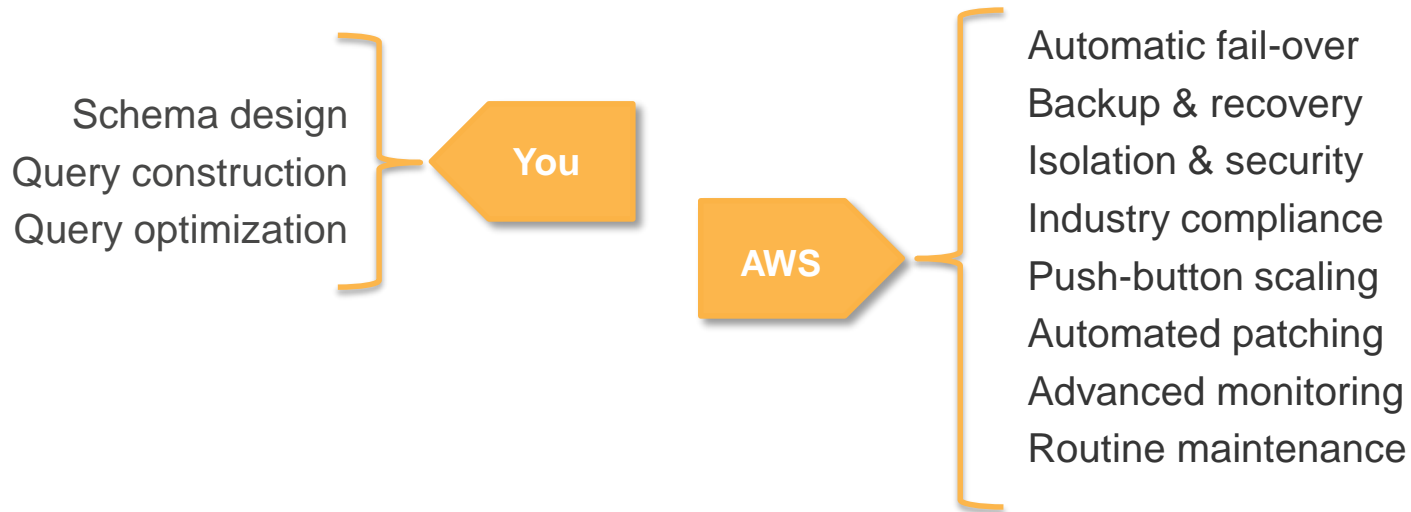


Upload systems metrics and audit logs to CloudWatch.



# Automate administrative tasks

---



Takes care of your time-consuming database management tasks,  
freeing you to focus on your applications and business

# Meet Amazon Aurora .....

## Database reimagined for the cloud

---



- ✓ **Speed** and **availability** of high-end commercial databases
- ✓ **Simplicity** and **cost-effectiveness** of open source databases
- ✓ Drop-in **compatibility** with MySQL
- ✓ Simple **pay as you go** pricing

Delivered as a **managed** service

# Aurora customer adoption

---

Aurora is used by:

**2/3** of top 100 AWS customers

**8** of top 10 gaming customers



**Fastest growing service in AWS history**

# Who are moving to Aurora and why?

---

## Customers using MySQL engines

- Higher performance – up to 5x
- Better availability and durability
- Reduces cost – up to 60%
- Easy migration; no application change

## Customers using commercial engines

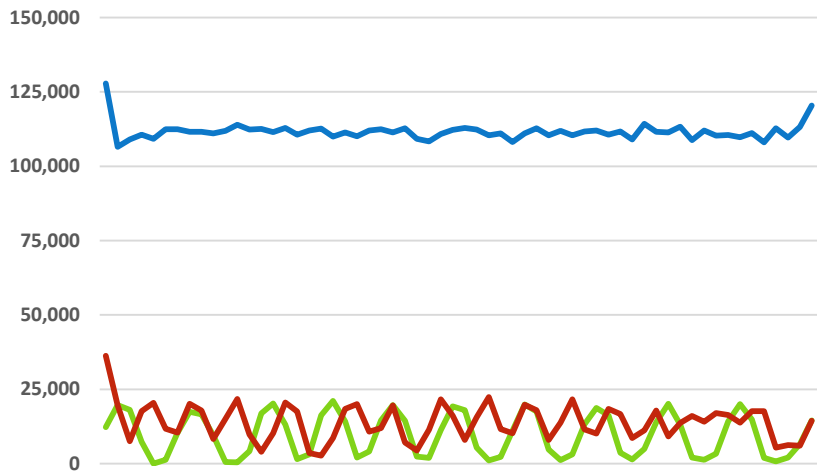
- One tenth of the cost; no licenses
- Integration with cloud ecosystem
- Comparable performance and availability
- Migration tooling and services

**Amazon Aurora is fast ...**

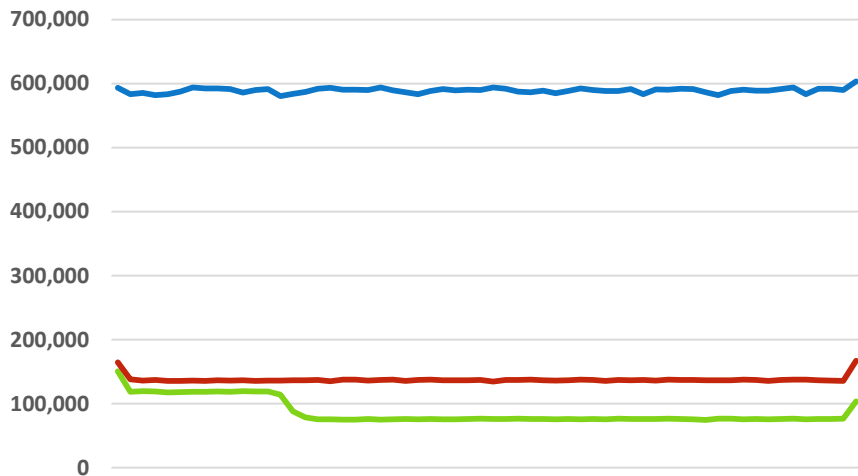
**5x faster than MySQL**

# 5X faster than RDS MySQL 5.6 & 5.7

## WRITE PERFORMANCE



## READ PERFORMANCE



MySQL SysBench results

R3.8XL: 32 cores / 244 GB RAM

**Aurora** 

**MySQL 5.6** 

**MySQL 5.7** 

**Five times higher throughput than stock MySQL  
based on industry standard benchmarks.**

# Aurora Scaling

## With user connection

Connections	Amazon Aurora	RDS MySQL w/ 30K IOPS
50	40,000	10,000
500	71,000	21,000
5,000	110,000	13,000

UP TO  
**8x**  
FASTER

## With number of tables

Tables	Amazon Aurora	MySQL I2.8XL local SSD	RDS MySQL w/ 30K IOPS (single AZ)
10	60,000	18,000	25,000
100	66,000	19,000	23,000
1,000	64,000	7,000	8,000
10,000	54,000	4,000	5,000

UP TO  
**11x**  
FASTER

## With database size - SYSBENCH

DB Size	Amazon Aurora	RDS MySQL w/ 30K IOPS
1GB	107,000	8,400
10GB	107,000	2,400
100GB	101,000	1,500
1TB	26,000	1,200

UP TO  
**21**  
FASTER

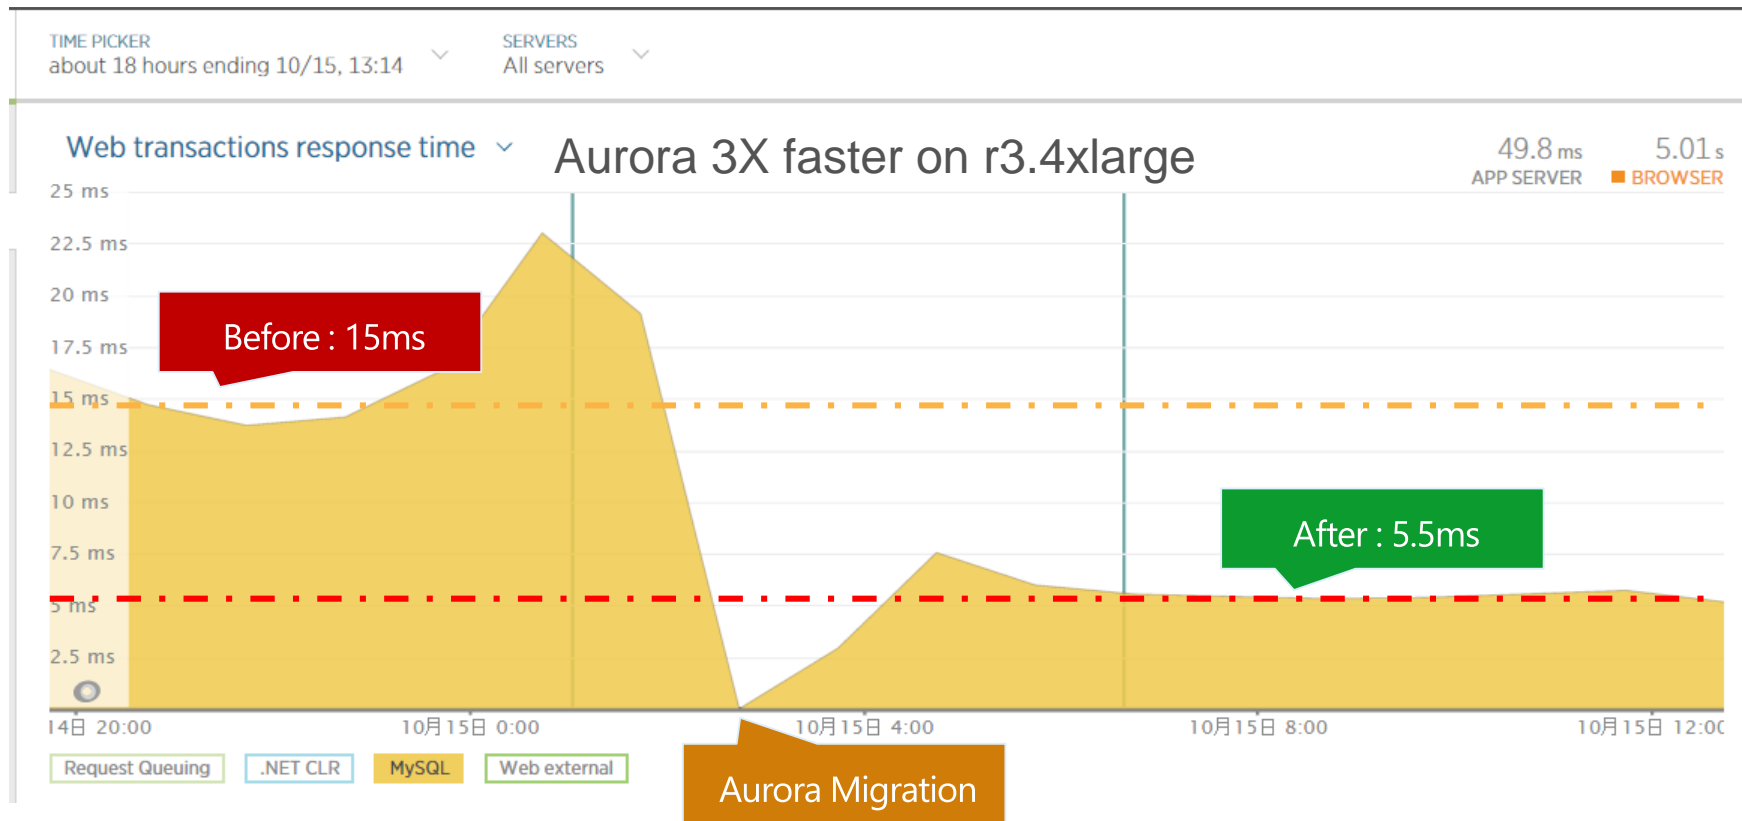
## With database size - TPCC

DB Size	Amazon Aurora	RDS MySQL w/ 30K IOPS
80GB	12,582	585
800GB	9,406	69

UP TO  
**136x**  
FASTER

# Real-life data – gaming workload

Aurora vs. RDS MySQL – r3.4XL, MAZ







# New performance enhancements

---

## Read performance

- ▶ Smart selector
- ▶ Logical read ahead
- ▶ Read views

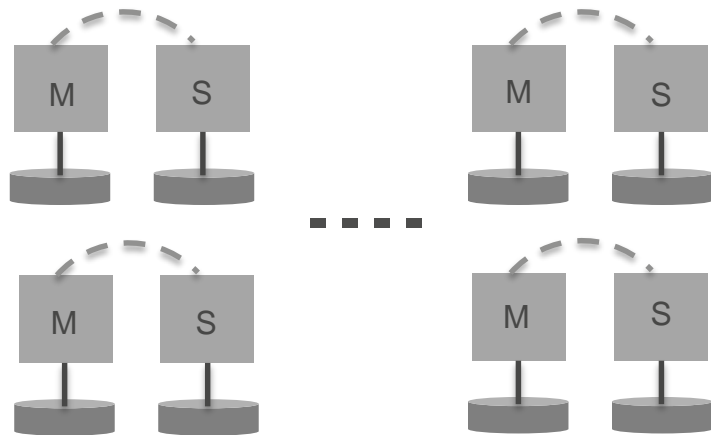
## Write performance

- ▶ NUMA aware scheduler
- ▶ Latch-free lock manager

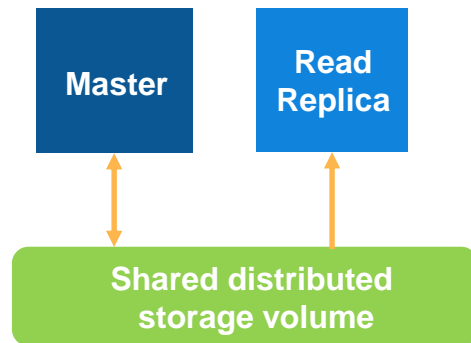
## Meta-data access

- ▶ Instant schema update
- ▶ B-Tree concurrency
- ▶ Catalog concurrency
- ▶ Faster index build

# Use case: MySQL shard consolidation



MySQL shards



Aurora cluster

Customer, a global SAAS provider, was using hundreds of MySQL shards in order to avoid MySQL performance and connection scalability bottlenecks

- Consolidated multiple 29 MySQL shards to single r3.4xlarge Aurora cluster
- Even after consolidation cluster utilization is still 30% with plenty of headroom to grow.

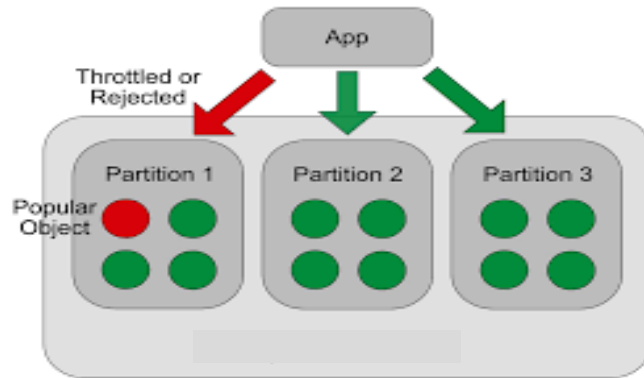
# Use case: Massively concurrent event store

## For messaging, gaming, IoT

---

New Aurora-backed data store reduces operational costs by **40%**

- The cost of reading data (**70%** of user traffic) almost eliminated due to memory-bound nature of the workload.
- Only pay for IO used, not provisioned. Also, Aurora does automatic hot spot management. So, no need to over provision IOPS based on IO requirements of hottest partition.



Customer, a global mobile messaging platform, was using NoSQL key-value database for user messages:

- **~22 million** accesses per hour (70% read, 30% write) - billing grows linearly with the traffic.
- Scalability bottleneck where certain portions (partitions) of data became “hot” and overloaded with requests.

## What about availability

**“Performance only matters if your database is up”**

# 6-way replicated storage

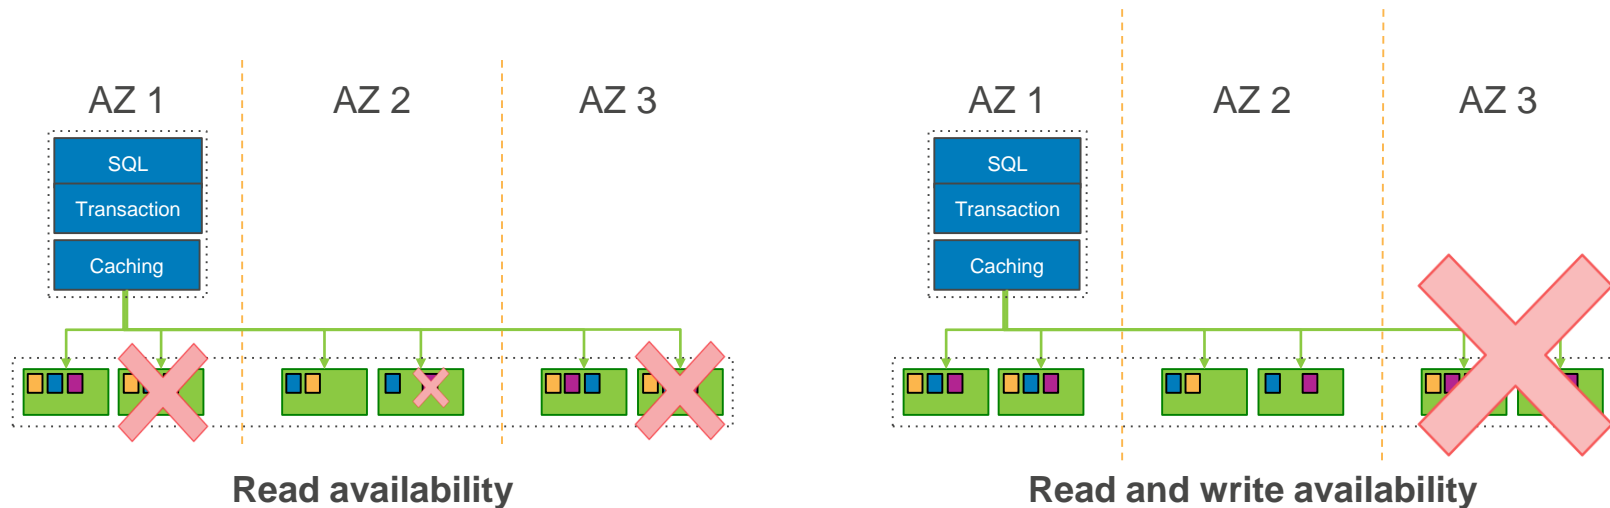
Survives catastrophic failures

Six copies across three availability zones

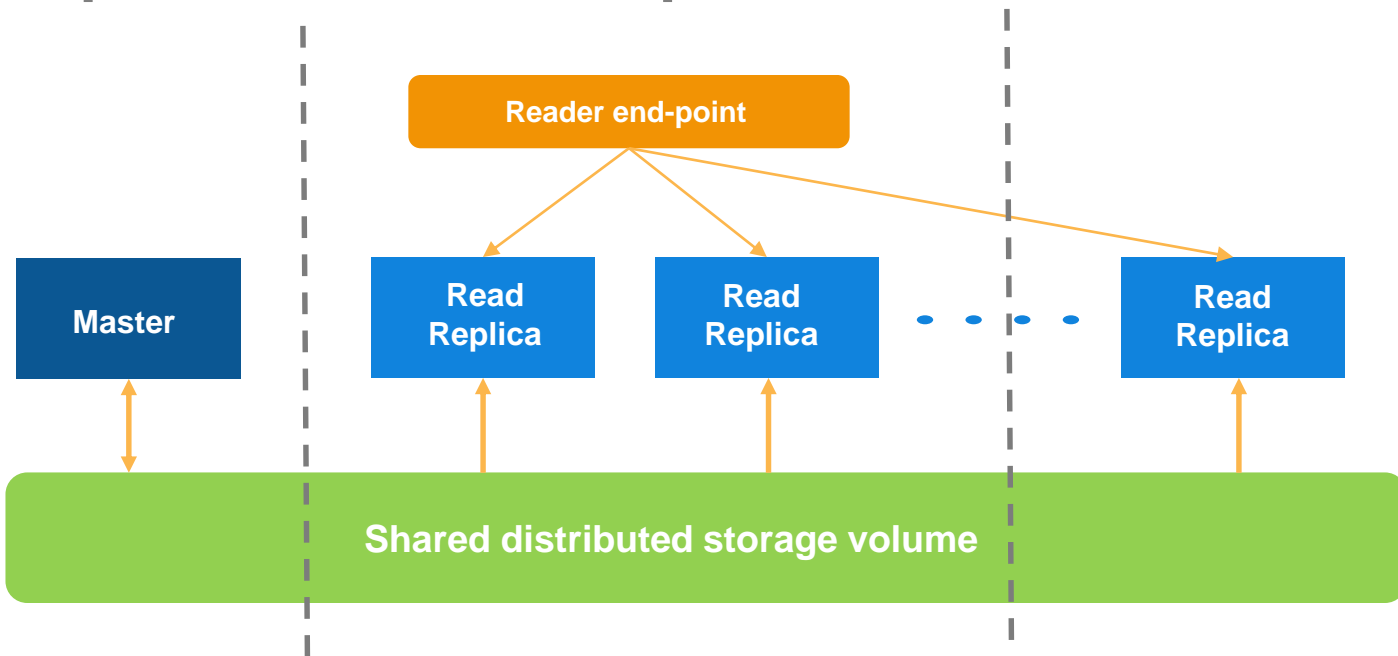
4 out 6 write quorum; 3 out of 6 read quorum

Peer-to-peer replication for repairs

Volume striped across hundreds of storage nodes



# Up to 15 promotable read replicas

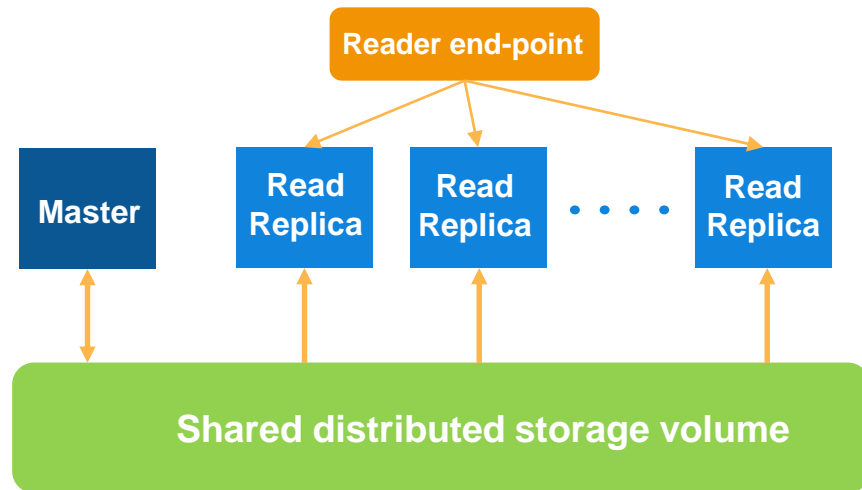


- ▶ Up to 15 promotable read replicas across multiple availability zones
- ▶ Re-do log based replication leads to low replica lag – typically < 10ms
- ▶ Reader end-point with load balancing; customer specifiable failover order

# Use case: Near real-time analytics and reporting

A customer in the travel industry migrated to Aurora for their core reporting application, which is accessed by ~1,000 internal users.

- **Fast provisioning:** replicas can be created, deleted and scaled within minutes based on load.
- **Load balancing:** read-only queries are load balanced across replica fleet through a DNS endpoint – no application configuration needed when replicas are added or removed.
- **Low replication lag:** allows mining for fresh data with no delays, immediately after the data is loaded.
- **Faster, concurrent access:** significant performance gains for core analytics queries - some of the queries executing in 1/100<sup>th</sup> the original time.

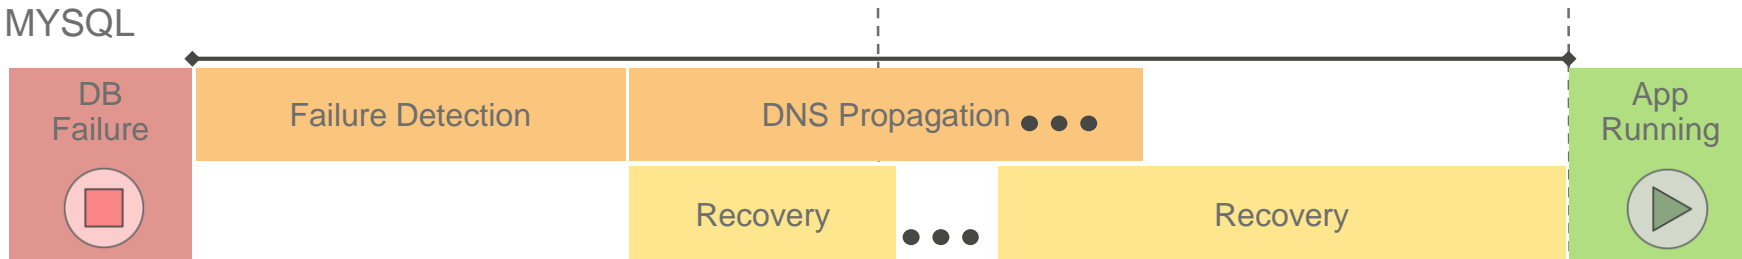


- ▶ Up to 15 promotable read replicas
- ▶ Low replica lag – typically < 10ms
- ▶ Reader end-point with load balancing

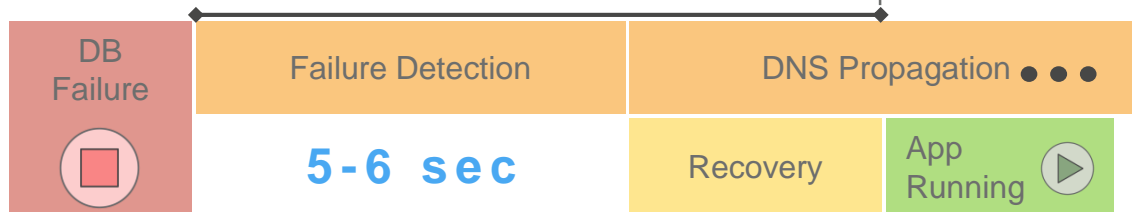


# Automated failover in 15 secs

MYSQL



AURORA WITH MARIADB DRIVER



**5-6 sec**

**5-10 sec**

# Cross-region read replicas

Faster disaster recovery and enhanced data locality

Promote read-replica to a master for faster recovery in the event of disaster

Bring data close to your customer's applications in different regions

Promote to a master for easy migration



# New availability features

---

## Read replicas

- ▶ Read replica end-point
- ▶ Specifiable fail-over order
- ▶ Faster fail-overs < 15 secs

## X-region DR

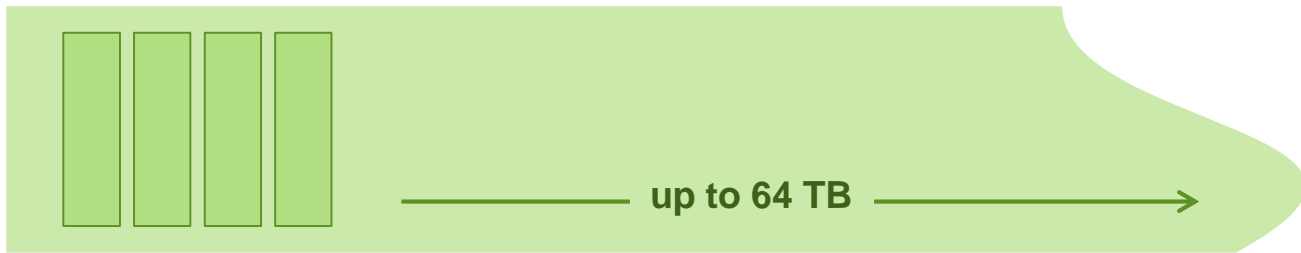
- ▶ Cross-region replication
- ▶ Cross-region snapshot copy **\*Coming soon\***
- ▶ Cross-account snapshot sharing **\*Coming soon\***

## **Amazon Aurora is easy to use**

**Automated storage management, security and compliance,  
advanced monitoring, database migration.**

# Simplify storage management

---

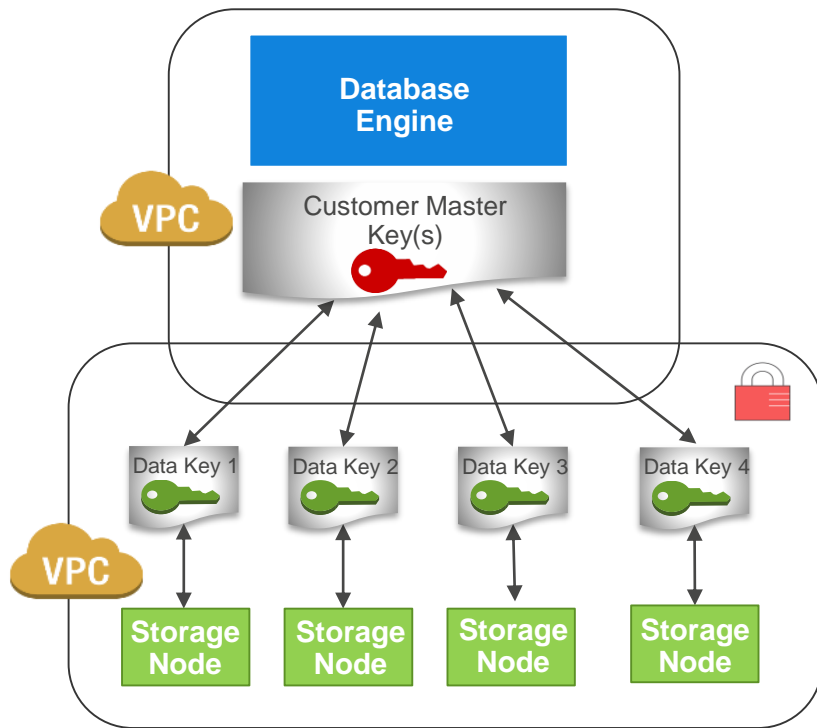


Up to 64TB of storage – auto-incremented in 10GB units

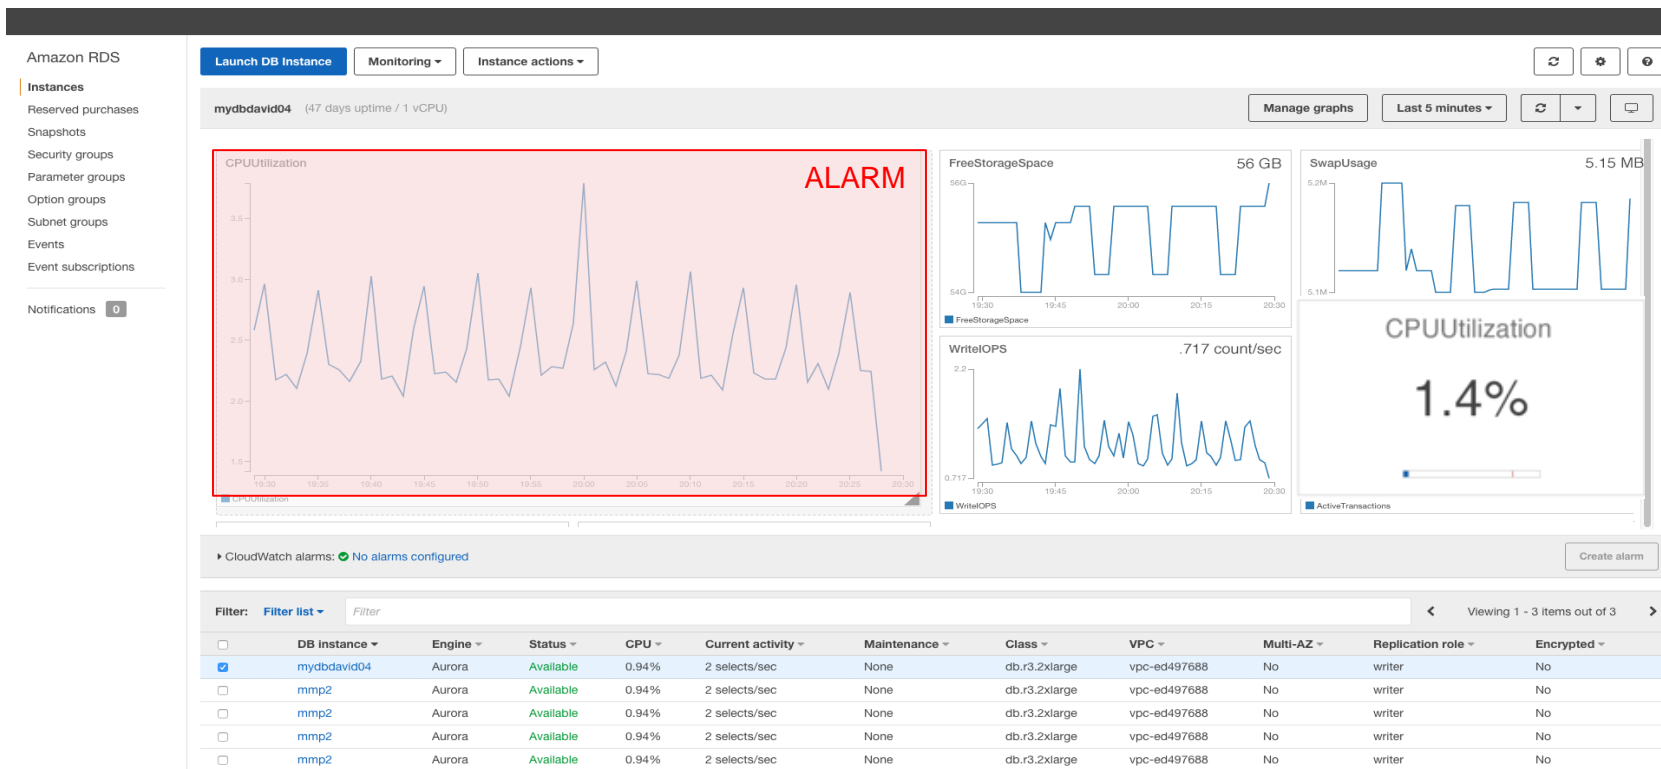
- Continuous, incremental backups to Amazon S3
- Instantly create user snapshots—no performance impact
- Automatic storage scaling up to 64 TB—no performance impact
- Automatic restriping, mirror repair, hot spot management, encryption

# Security and compliance

- ✓ Encryption to secure data at rest using customer managed keys
  - AES-256; hardware accelerated
  - All blocks on disk and in Amazon S3 are encrypted
  - Key management via AWS KMS
- ✓ Encrypted cross-region replication, snapshot copy - SSL to secure data in transit
- ✓ Advanced auditing and logging without any performance impact
- ✓ Industry standard security and data protection – **SOC, ISO, PCI/DSS, HIPAA/BAA**









# Advanced monitoring



50+ system/OS metrics | sorted process list view | 1-60 sec granularity










alarms on specific metrics | egress to CloudWatch Logs | integration with 3<sup>rd</sup>-party tools

# Amazon Aurora migration options

Source database	From where	Recommended option
	RDS	Console based automated snapshot ingestion and catch up via binlog replication.
  	EC2, on premise	Binary snapshot ingestion through S3 and catch up via binlog replication.
 	EC2, on premise, RDS	Schema conversion using SCT and data migration via DMS.

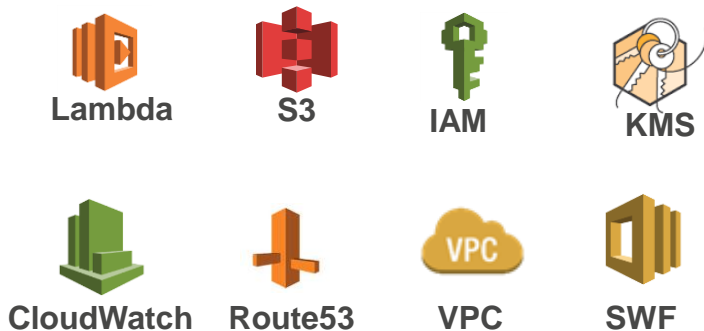


# Leverage MySQL and AWS ecosystems

Business Intelligence	Data Integration	Query and Monitoring
  	  	  

Source: Amazon

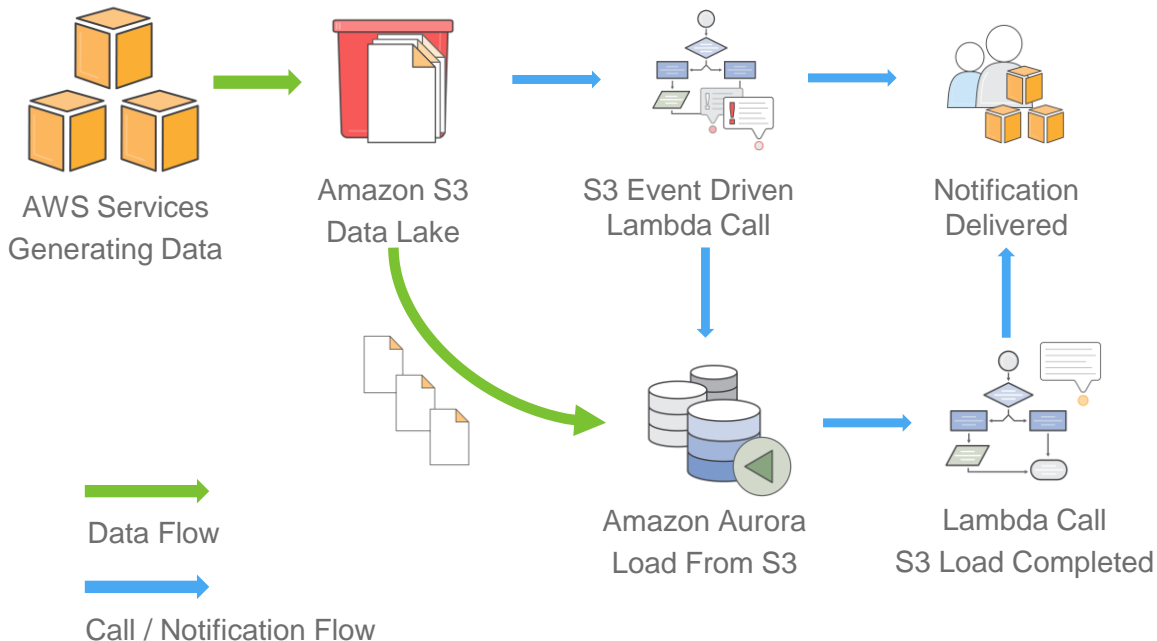
## AWS Ecosystem



**“We ran our compatibility test suites against Amazon Aurora and everything just worked.”** - Dan Jewett, Vice President of Product Management at Tableau

# Use case: Event driven data pipeline

- Simplify custom database logic by moving it from functions, stored procedures etc. to cloud-based code repository.
  - Enable database developers to create rich software features accessible from SQL layer.
- Run code in response to ad-hoc requests, triggers or scheduled database events in any language supported by AWS Lambda (Java, Node.js, Python).
- Accelerate the migration from any programmable database

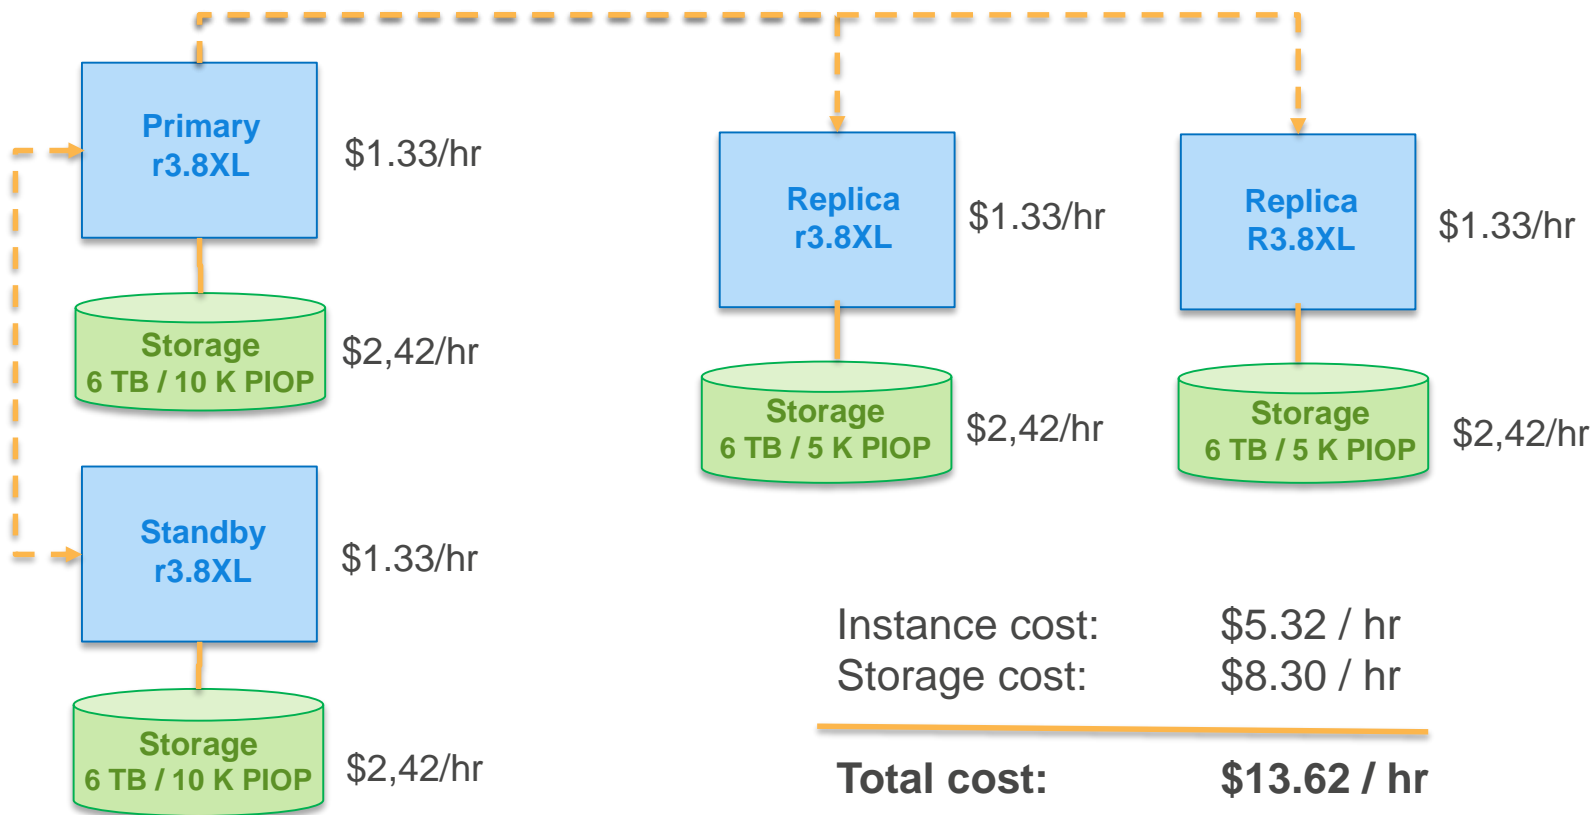


# Amazon Aurora saves you money

**1/10<sup>th</sup> of the cost of commercial databases**  
**Cheaper than even MySQL**

# Cost of ownership: Aurora vs. MySQL

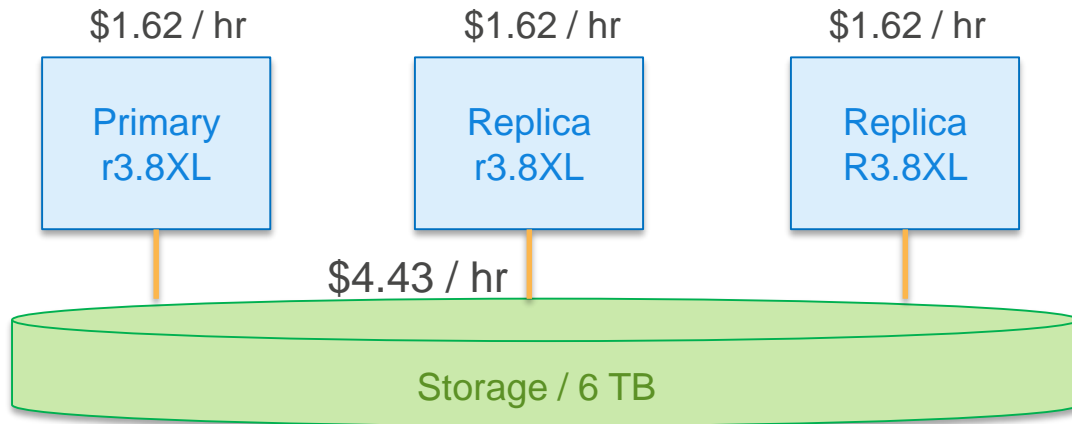
## MySQL configuration hourly cost



# Cost of ownership: Aurora vs. MySQL

## Aurora configuration hourly cost

- No idle standby instance
- Single shared storage volume
- No PIOPs – pay for use I/O
- Reduction in overall IOP



Instance cost:	\$4.86 / hr
Storage cost:	\$4.43 / hr

---

<b>Total cost:</b>	<b>\$9.29 / hr</b>
--------------------	--------------------

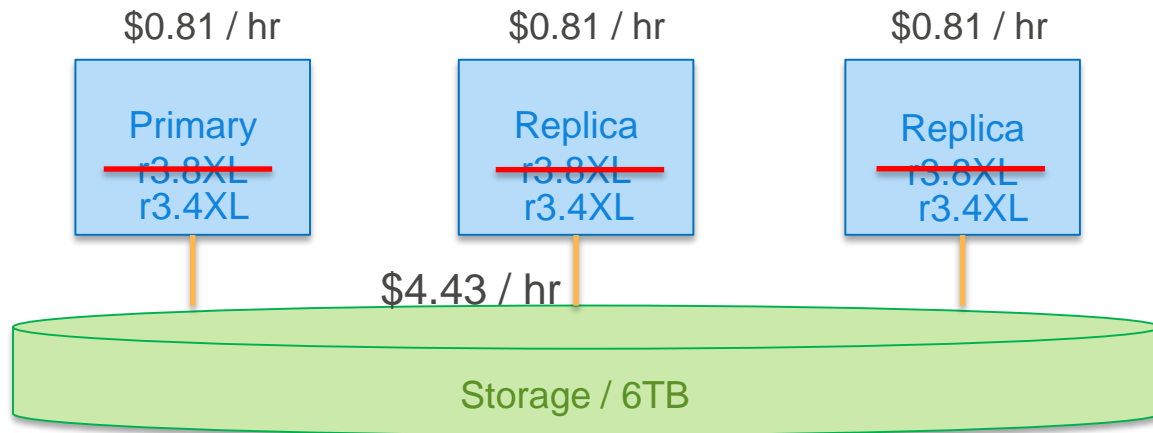
**31.8%  
Savings**

\*At a macro level Aurora saves over 50% in storage cost compared to RDS MySQL.

# Cost of ownership: Aurora vs. MySQL

## Further opportunity for saving

- Use smaller instance size
- Pay-as-you-go storage



Instance cost:	\$2.43 / hr
Storage cost:	\$4.43 / hr

---

<b>Total cost:</b>	<b>\$6.86 / hr</b>
--------------------	--------------------

**49.6%**  
**Savings**

Storage IOPs assumptions:

1. Average IOPs is 50% of Max IOPs
2. 50% savings from shipping logs vs. full pages

# Higher performance, lower Cost

---



Safe.com lowered their bill by **40%** by switching from sharded MySQL to a single Aurora instance.



Double Down Interactive (gaming) lowered their bill by **67%** while also achieving better latencies (most queries ran faster) and lower CPU utilization.

- Fewer instances needed
- Smaller instances can be used
- No need to pre-provision storage
- No additional storage for read replicas

# R3.large is too expensive for dev/test?

We just introduced t2.medium

---

	vCPU	Mem	Hourly Price
db.t2.medium	2	4	\$0.082
db.r3.large	2	15.25	\$0.29
db.r3.xlarge	4	30.5	\$0.58
db.r3.2xlarge	8	61	\$1.16
db.r3.4xlarge	16	122	\$2.32
db.r3.8xlarge	32	244	\$4.64

**\*NEW\***

## t2 RI discounts

Up to 34% with a 1-year RI

Up to 57% with a 3-year RI

\*Prices are for Virginia



# Other Aurora sessions at re:Invent

---

## **DAT303** - Deep Dive on Amazon Aurora

Thu 11:30-12:30, Venetian, Level 4, Delfino 4004

## **DAT301** - Amazon Aurora Best Practices: Getting the Best Out of Your Databases

Wed 5:30-6:30, Venetian, Level 4, Lando 4205

## **DAT322** - Workshop: Stretching Scalability: Doing more with Amazon Aurora

Option 1: Wed 2:00-4:30, Mirage, Trinidad B,

Option 2: Thu 2:30-5:00, Mirage, Antigua B

# Have more questions

---



**Rory Richardson, BD**  
**roryrich@amazon.com**



**Linda Wu, PM**  
**wulind@amazon.com**

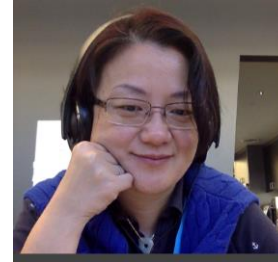


**David Lang, PM**  
**davelang@amazon.com**



**Debanjan Saha, GM**  
**deban@amazon.com**

**Reach out to your super friendly Aurora team**



**Linda Xu**

Principal Architect, Data @ Ticketmaster

20 years of database experience

<https://www.linkedin.com/in/linda-yun-xu-8a53034>

# ABOUT US

## HISTORY

- 1976 - Founded at Arizona State University
  - 1996 - Ticketmaster.com launched
  - 2010 - Live Nation and Ticketmaster join forces to power live experiences
  - 2013 - Transformation journey begins...
- 
- Publicly Traded Company (LYV)
  - \$7.6B Revenue
  - \$25B in GTV (Gross Transaction Value)
  - Top 5 eCommerce site

***ticketmaster***<sup>®</sup>



A full-page background image of a performer in a blue jacket and jeans jumping high on a stage. The stage is covered in a thick shower of multi-colored confetti. In the background, a large crowd of people is visible with their hands raised, and stadium lights illuminate the scene. Various sponsor banners like Allianz, Sony, and ICOF are visible on the stadium walls.

# We power unforgettable moments of joy!

SOMEWHERE IN THE WORLD  
**EVERY 20 MINUTES**  
IS A LIVE NATION EVENT

# What we do ...



**EVERY 20 MINUTES**

Somewhere In The World, There Is A Live Nation Event

**70M**

Unique Monthly Visitors On The Live Nation Family Of Websites



**LIVE NATION BY THE NUMBERS**



**530M**

Fans Touched By Our Platform

**25,500**

Concert Events Worldwide



**465M**

Tickets Sold Globally By Ticketmaster

**37**

Countries Host Live Nation Events



**350+**

Global Icons Under Artist Nation Management

**145M**

Consumer Database Profile



**21M**

Mobile App Downloads

**70+**

Global Festivals



**900**

Brand Partners

**100**

Artist Nation Managers





# Our IT infrastructure Today

---

## Our private cloud is made up of:

- ▶ Hundreds of Ticketmaster products
- ▶ Tens of thousands of servers
- ▶ Multiple Data Centers, globally
- ▶ More than 1,000 databases
  - Oracle, MySQL, MongoDB, Teradata, Microsoft SQL Server and PostgreSQL

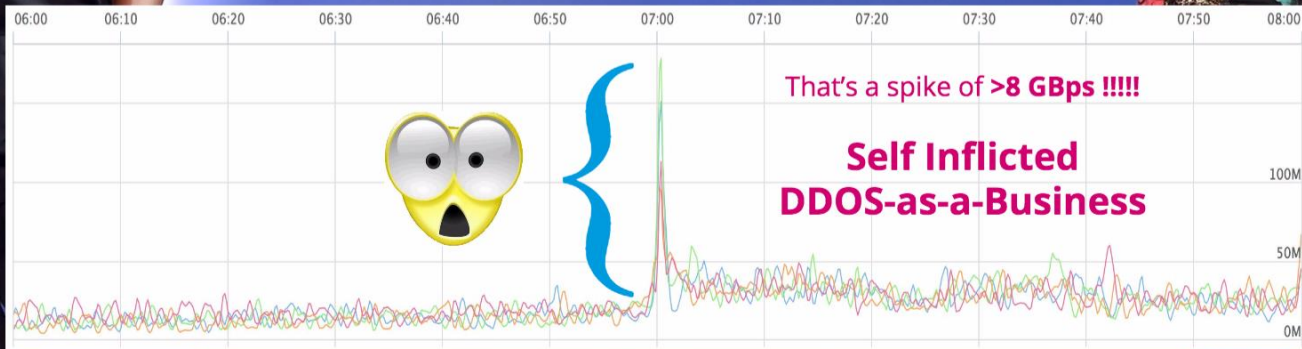


# BIG SCALE, BIG CHALLENGES

## Onsales = Black Friday every day!

- Huge spikes / demand for tickets
- Global company = across time zones
- Limited inventory (Beyonce Tickets!)
- Multiple sales channels

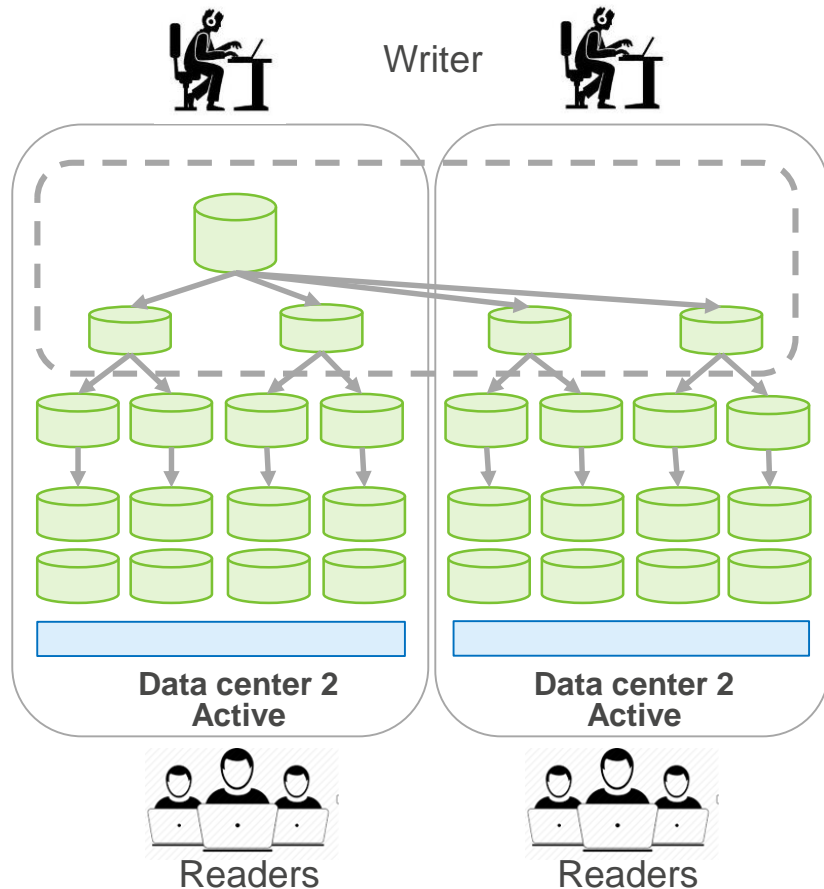
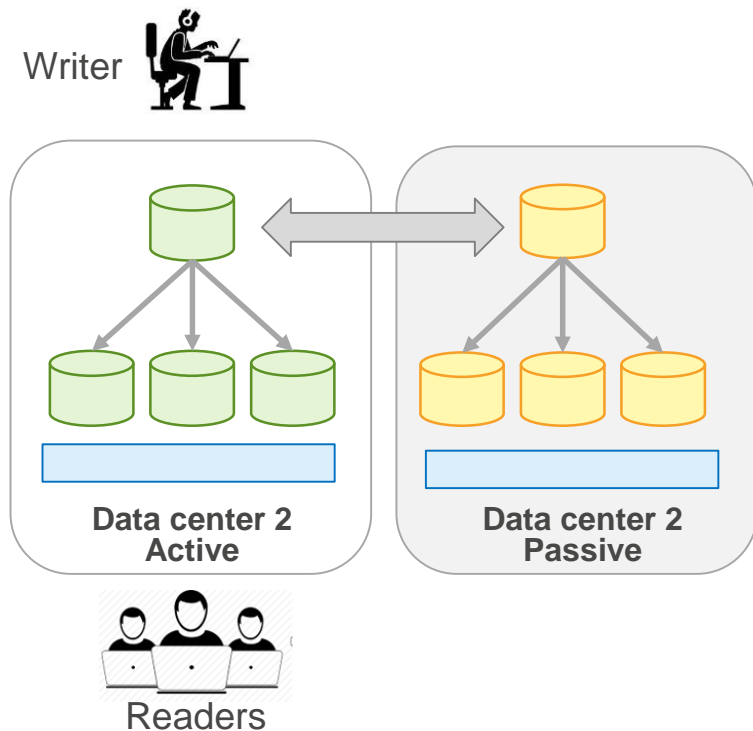
0 to 150M transactions in minutes!





# Challenges we face with MySQL

Ticketmaster has heavy MySQL use cases  
Examples of our MySQL setup



# Challenges we face with MySQL

---

## Scalability

- Scaled for max capacity at all times

## Replication Lag

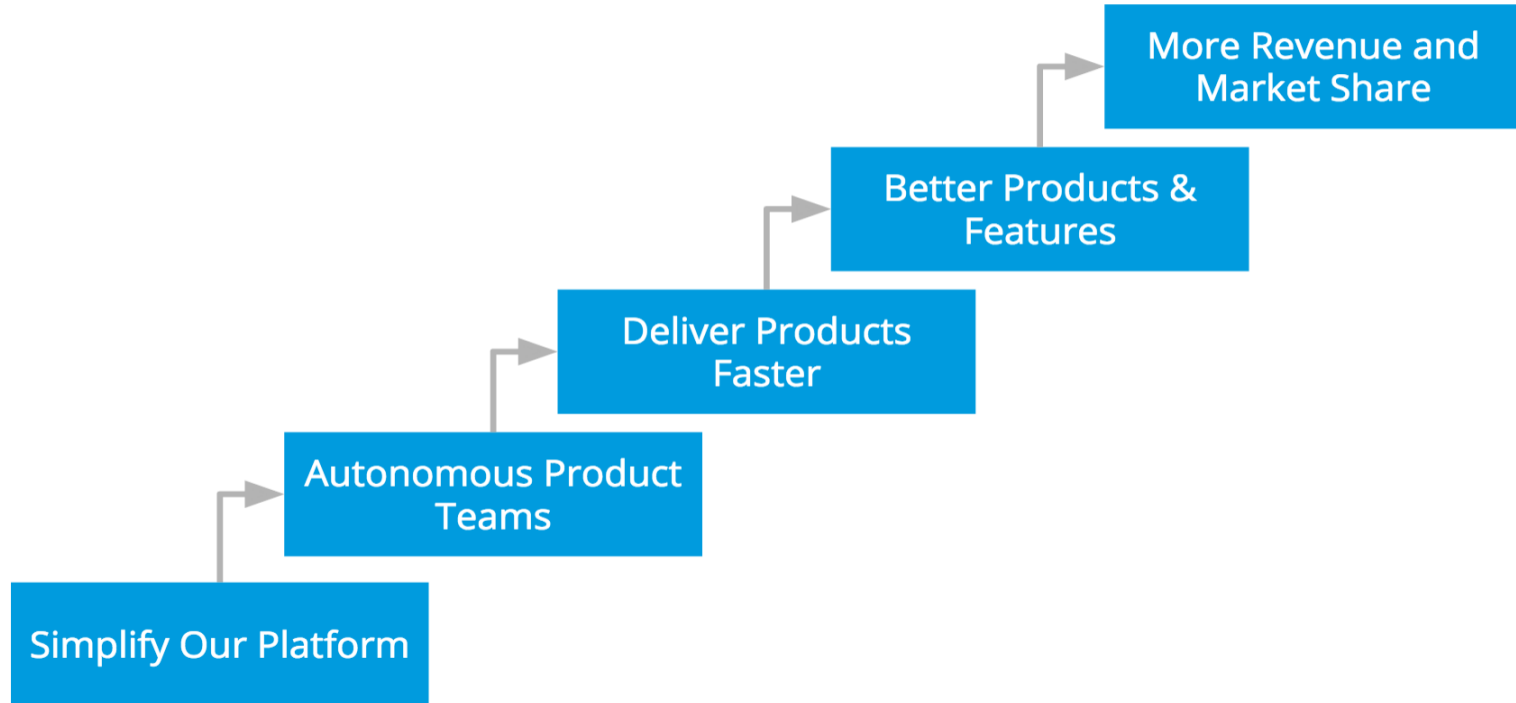
- Read consistency cross slaves
- Single thread replication
- Application can cause high replication lags

## Operations

- Manual failover
- Maintenance on multi-master setup
- Version updates

Where are we headed..

## SIMPLIFY OUR PLATFORM



# Where are we headed..

---

.... We are moving to AWS and to Aurora!

# Why did we choose Aurora?

## Compatibility

- Fully compatible with MySQL 5.6
- Migration from MySQL using MySQL replication
- Migration from Oracle leverages DMS and SCT

## Ease of Operation

- Replication Setup - Auto Failover
- Reader and Writer Endpoint
- Tuned Parameters
- Auto minor version upgrade

## Scalability

- Vertically: Add more CPU and memory per instance
- Horizontally: Add more database instances
- Storage: Automatically add space and IOPs

# Our success story

---

***Account Manager***, our first product migrated to AWS.

- Migrated from MySQL 5.6 to Aurora
- 0 issue be reported since launch
- Replication lag is only 20ms
- 12 MySQL servers to 5 Aurora nodes
- Deployment time took from 1~2 hours to 20 minutes in AWS
- Build a new test environment took from 1 week to 30 minutes in AWS

**Our MySQL application worked on Aurora without any code changes**

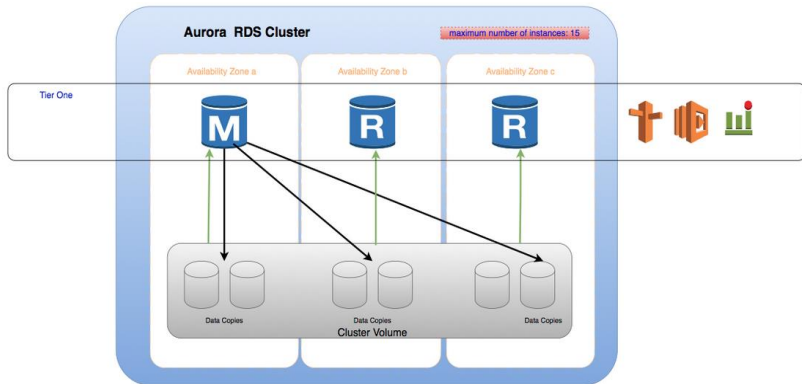
# Our Tools: Terraformer



Terraform

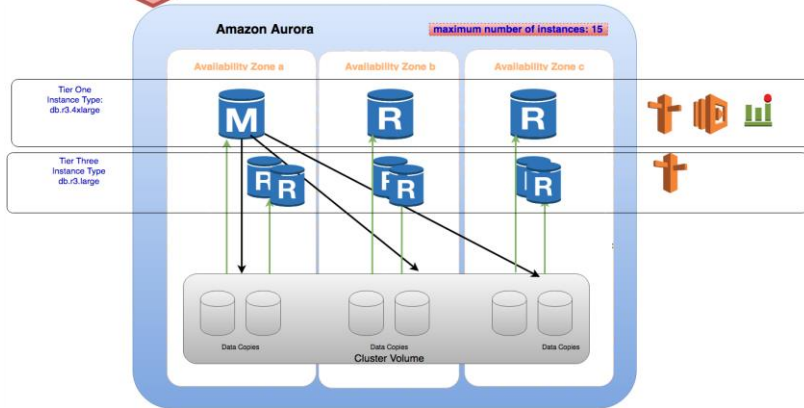
*ticketmaster*<sup>®</sup>

## Regular Business Traffic



## Peak Traffic

OnSale Traffic



- Infrastructure in code including database!
- We use terraformer to manage Aurora
  - Scale in/out horizontally within **10 minutes**, without downtime
  - Scale up/down vertically with **30s** of failover downtime

The background features a large, abstract graphic with blue and orange wavy, ribbon-like shapes. These shapes are overlaid with a pattern of concentric dotted circles in light gray and orange, creating a sense of motion and depth.

**AWS  
re:Invent**

**Thank you!**