CON314

Deploy a Deep Learning Framework on Amazon ECS

Chad Schmutzer, Solutions Architect
Hubert Cheung, Solutions Architect
David Kuo, Solutions Architect
Andy Mui, Solutions Architect

November 2016

amazon web services

What to expect from this workshop

- Workshop goals
- Overview of ECS + ECR
- Overview of AWS CloudFormation
- Overview of EC2 Spot Instances
- Hands on workshop
- Wrap-up

Workshop Goals

Demonstrate ECS value

- Increase infrastructure utilization
- environment isolation
- placing mixed applications in same environment
- easy deployment
- We just chose to use MXNet as an example application to containerize

What's MXNet?

- MXNet is an open-source deep learning framework that allows you to define, train, and deploy deep neural networks on a wide array of devices, from cloud infrastructure to mobile devices. It is highly scalable, allowing for fast model training, and supports a flexible programming model and multiple languages.
- http://mxnet.io/

ECS + ECR

ECS Benefits

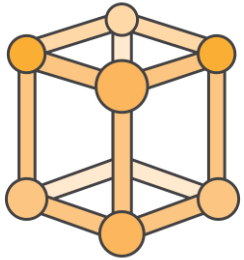

Cluster management made easy

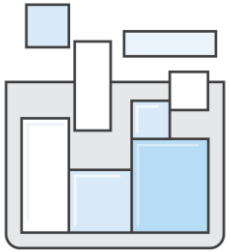

Flexible scheduling

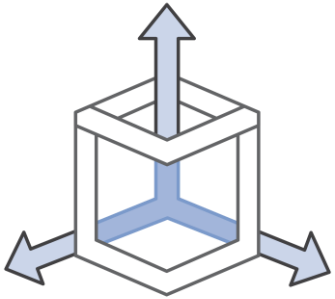

Integrated and extensible

Security

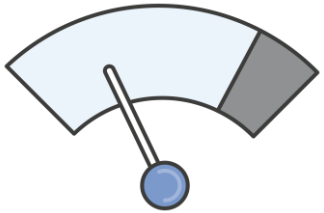

Performance at scale

ECS Architecture

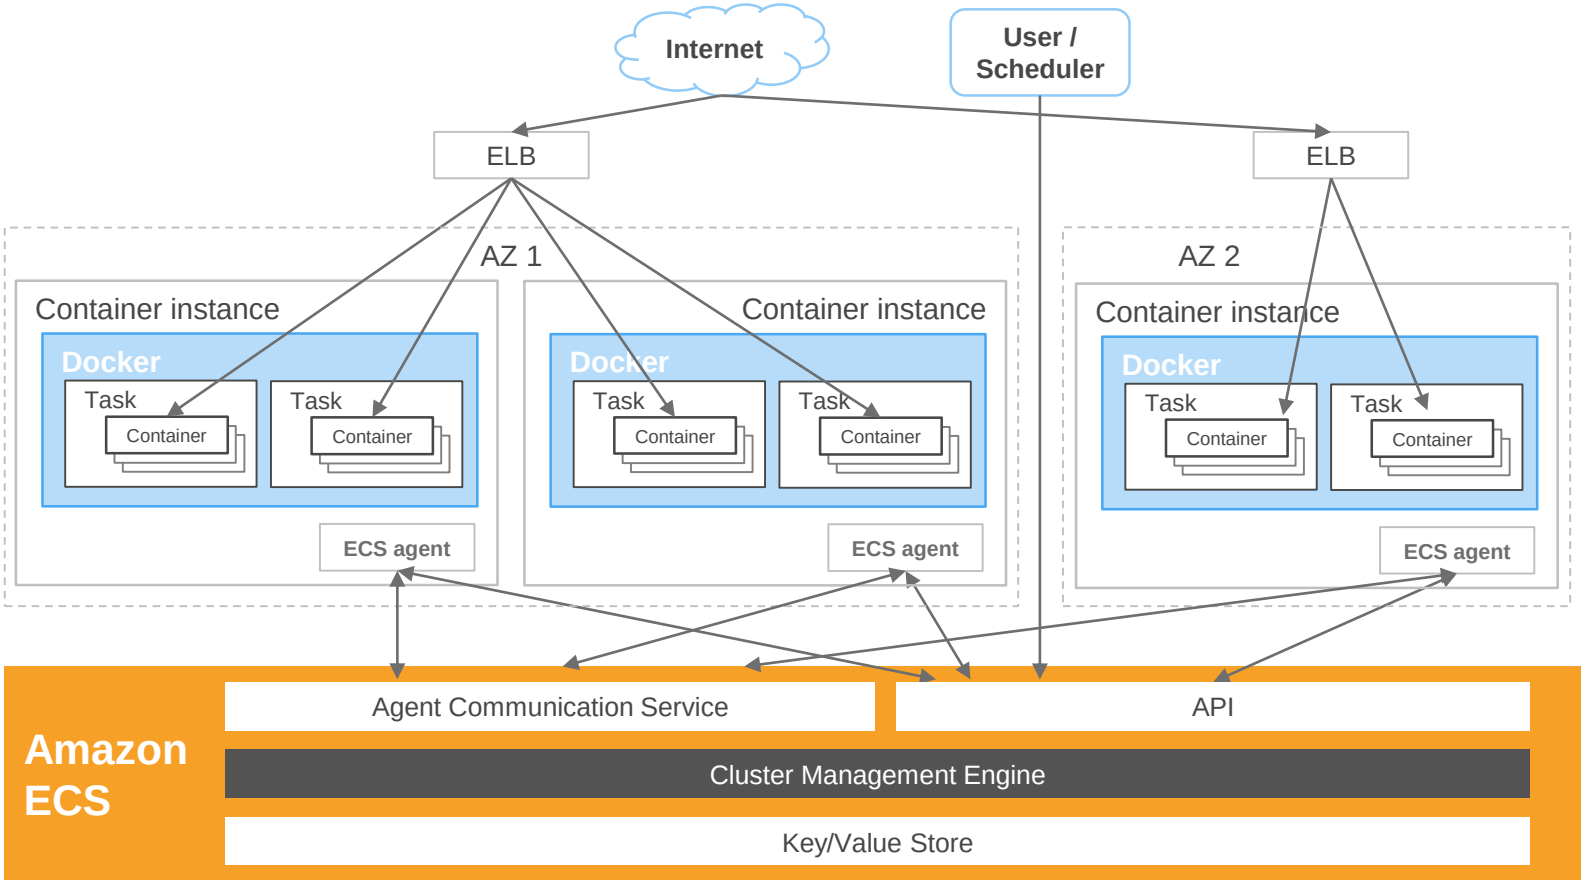

ECS Common Use Cases

Applications and services

- Configuration and deployment
- Microservices

Batch processing

What’s ECR?

- Amazon EC2 Container Registry (Amazon ECR) is a fully-managed Docker container registry that makes it easy for developers to store, manage, and deploy Docker container images. Amazon ECR is integrated with Amazon EC2 Container Service (Amazon ECS), simplifying your development to production workflow.
- Learn more: https://aws.amazon.com/ecr/

How does ECS use ECR?

Amazon ECR

```
{
  "containerDefinitions": [
    {
      "memory": 128,
      "portMappings": [
        {
          "hostPort": 80,
          "containerPort": 80,
          "protocol": "tcp"
        }
      ],
      "essential": true,
      "name": "nginx-container",
      "image": "nginx",
      "logConfiguration": {
        "logDriver": "awslogs",
        "options": {
          "awslogs-group": "ecs-log-streaming",
          "awslogs-region": "us-west-2"
        }
      },
      "cpu": 0
    }
  ],
  "family": "example_task_1"
}
```

ECS Task Definition

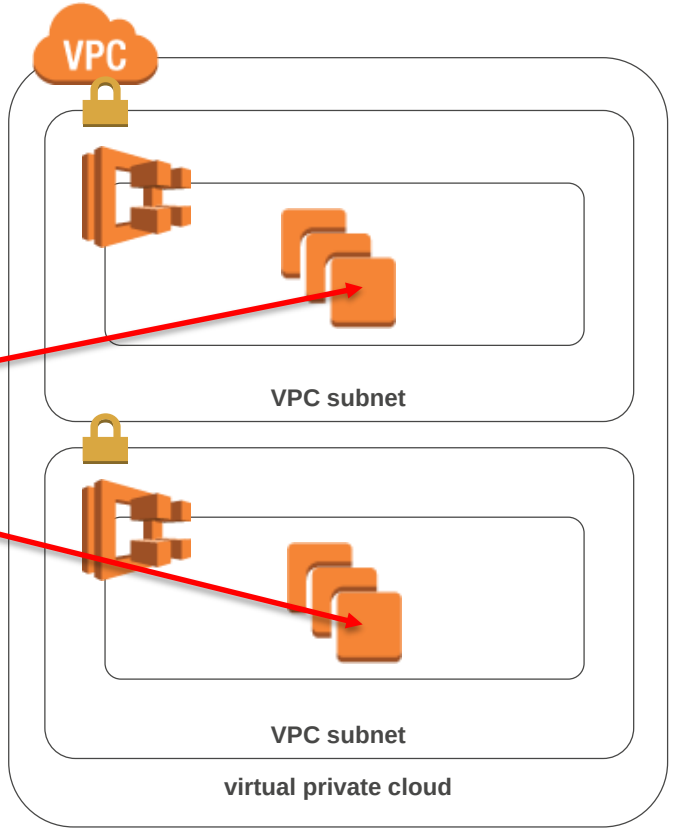

AWS CloudFormation

CloudFormation – Components & Technology

Template

JSON formatted file

Parameter definition

Resource creation

Configuration actions

CloudFormation

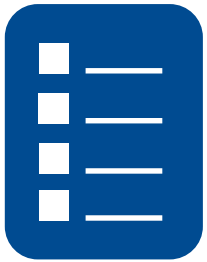

Framework

Stack creation

Stack updates

Error detection and rollback

Stack

Configured AWS resources

Comprehensive service support

Service event aware

Customizable

CloudFormation Benefits

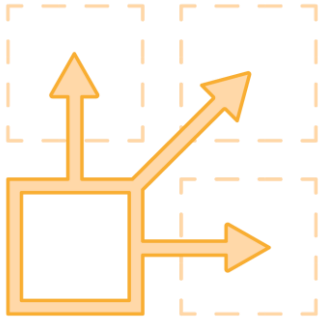

Templated resource provisioning

Infrastructure as code

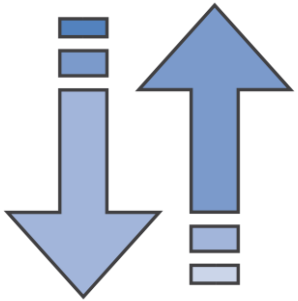

Declarative and flexible

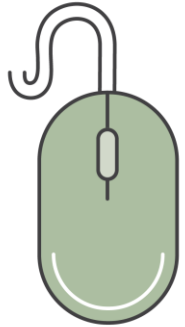

Easy to use

CloudFormation Use Cases

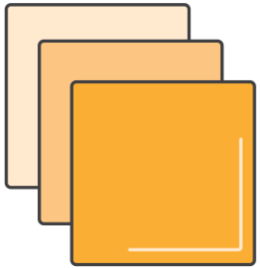

Stack replication

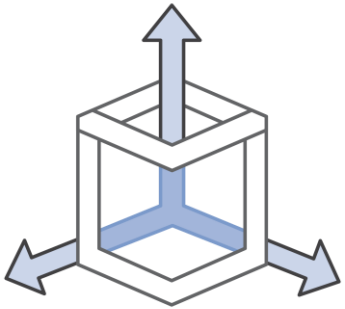

Infrastructure
scale out

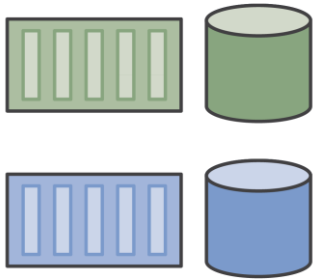

Blue-green
deployments

Infrastructure
as code

Why do customers use CloudFormation?

Developers/DevOps teams value CloudFormation for its ability to treat infrastructure as code, allowing them to apply software engineering principles, such as SOA, revision control, code reviews, integration testing to infrastructure.

IT Admins and MSPs value CloudFormation as a platform to enable standardization, managed consumption, and role-specialization.

ISVs value CloudFormation for its ability to support scaling out of multi-tenant SaaS products by quickly replicating or updating stacks. ISVs also value CloudFormation as a way to package and deploy their software in their customer accounts on AWS.

EC2 Spot Instances

Amazon EC2 Consumption Models

On-Demand

Pay for compute capacity by the hour with no long-term commitments

For spiky workloads, or to define needs

Reserved

Make a low, one-time payment and receive a significant discount on the hourly charge

For committed utilization

Spot

Bid for unused capacity, charged at a Spot price which fluctuates based on supply and demand

For time-insensitive or transient workloads

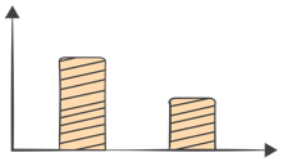

With Spot, the rules are simple

Markets where the price of compute changes based on supply and demand

You'll never pay more than your bid. When the market exceeds your bid you get 2 minutes to wrap up your work

Show me the markets!

C3	1a	1b	1c	On Demand
8XL	$0.50	$0.27	$0.29	$1.76
4XL	$0.21	$0.30	$0.16	$0.88
2XL	$0.08	$0.07	$0.08	$0.44
XL	$0.04	$0.05	$0.04	$.22
L	$0.01	$0.01	$0.04	$0.11

Each instance family

Each instance size

Each Availability Zone

In every region

Is a separate ***Spot Market***

Spot Fleet

Spot Fleet helps you

Launch Thousands of Spot Instances

with one **RequestSpotFleet** call.

Get Best Price

Find the lowest priced horsepower that works for you.

or

Get Diversified Resources

Diversify your fleet. Grow your availability.

And

Apply Custom Weighting

Create your own capacity unit based on your application needs

Diversification with EC2 Spot Fleet

Multiple EC2 Spot Instances selected

Multiple Availability Zones selected

Pick the instances with similar performance characteristics, e.g., c3.large, m3.large, m4.large, r3.large, c4.large.

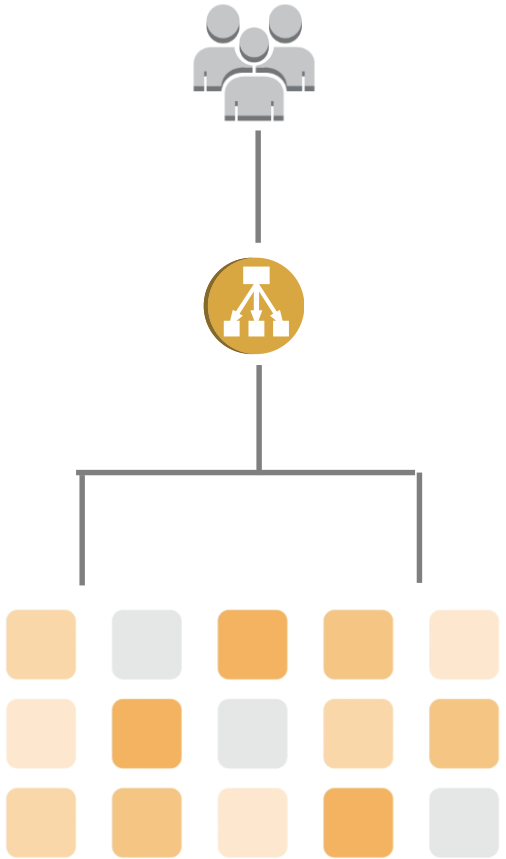

Workshop: Image Classification

Overall Architecture

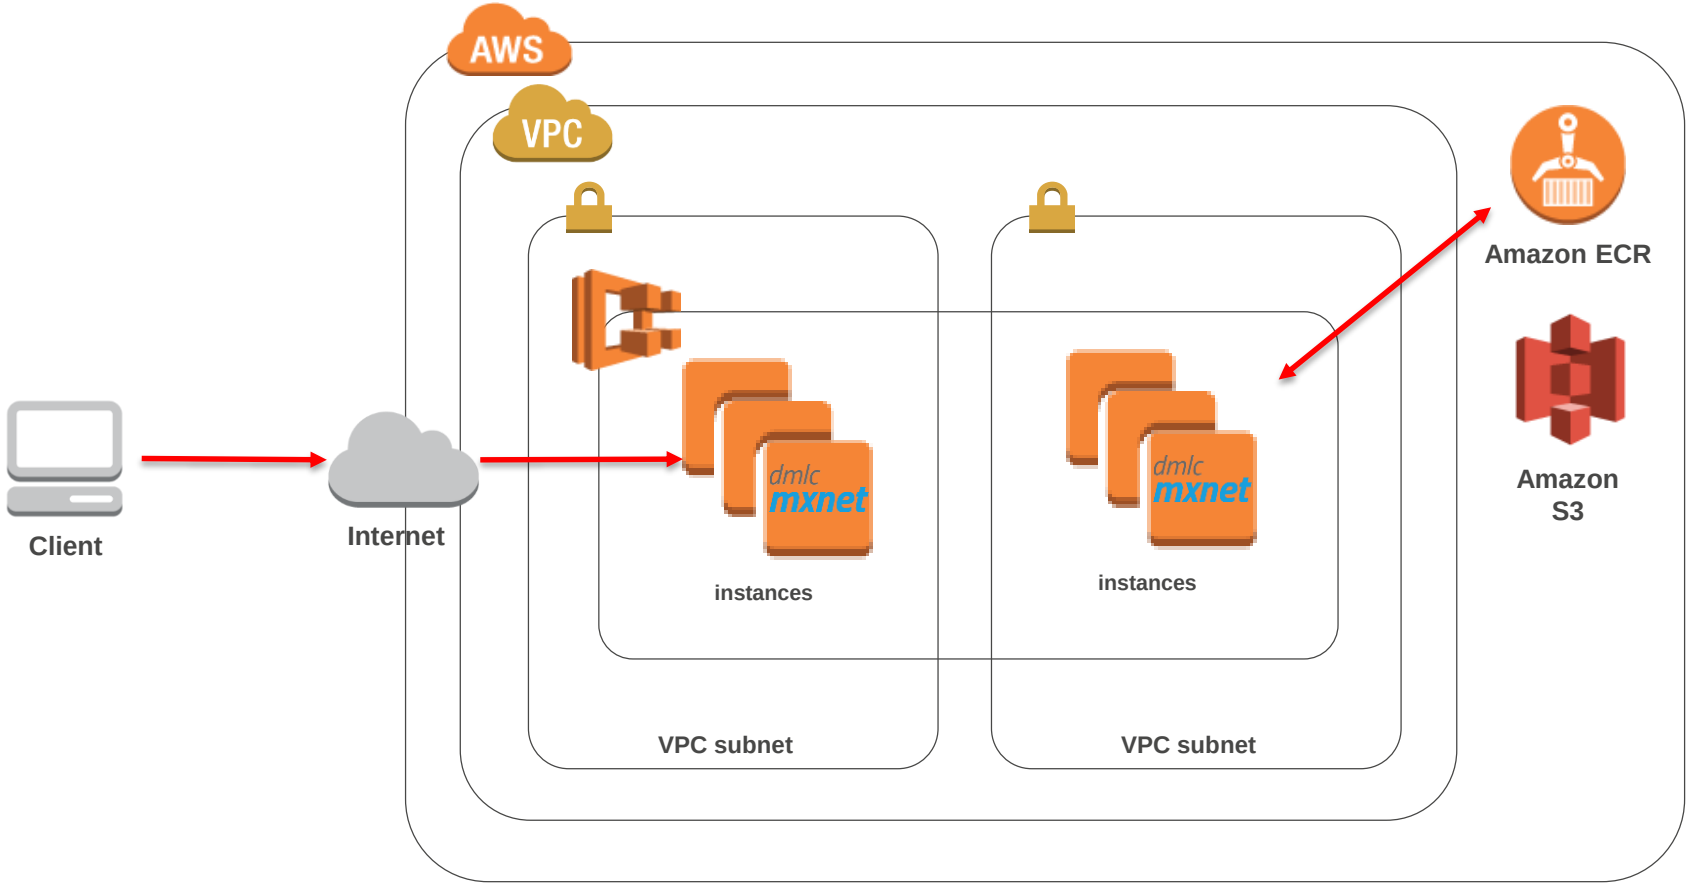

Lab 1: Getting Started

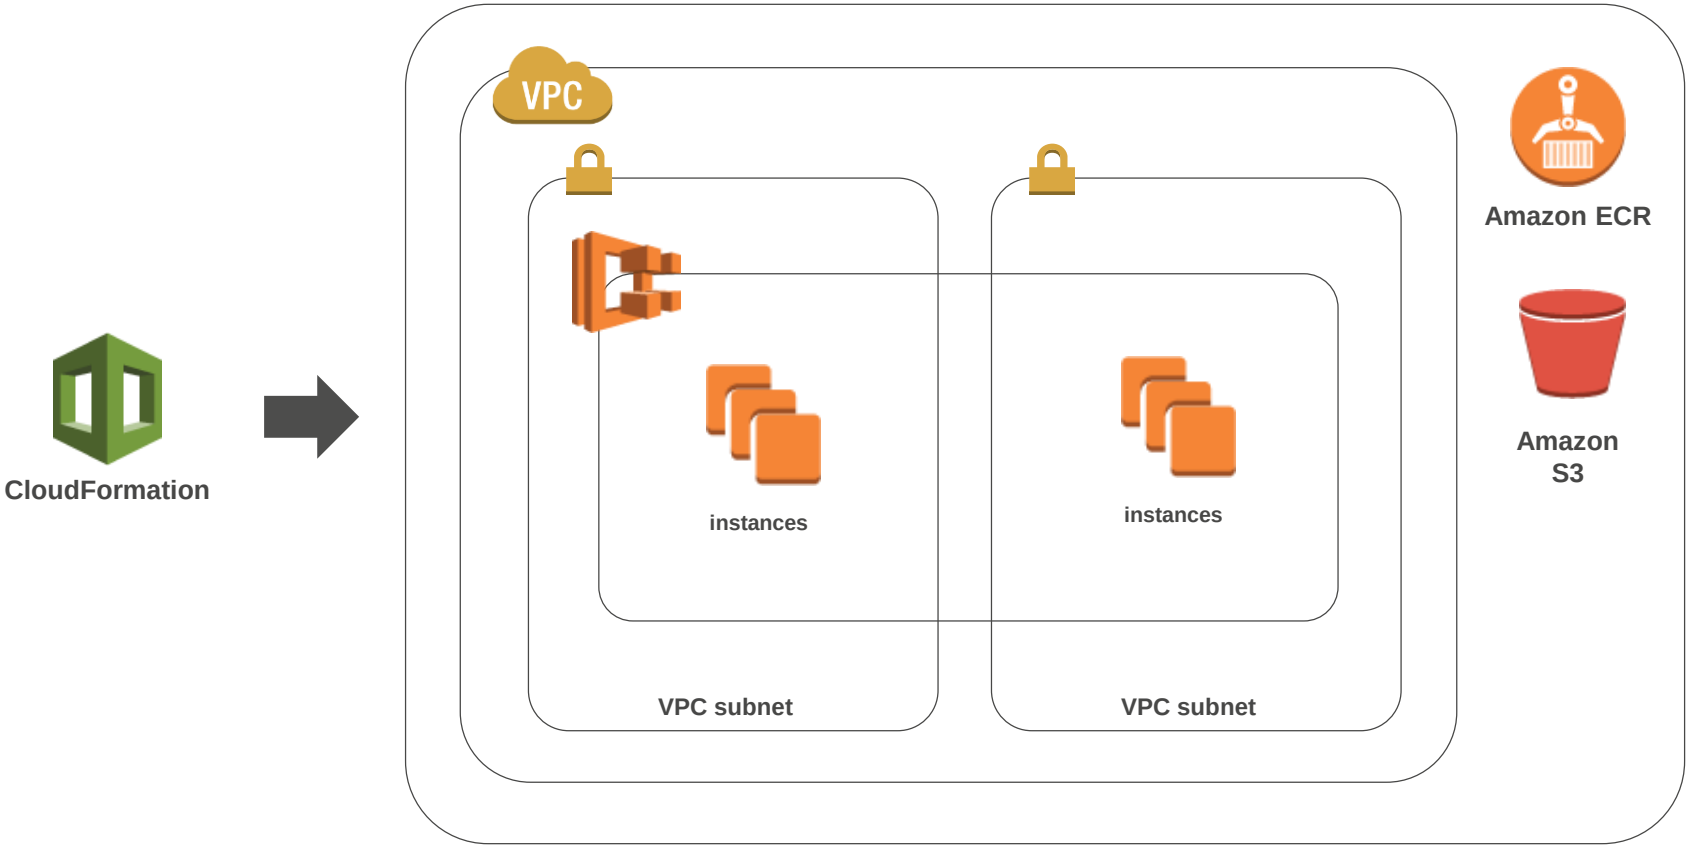

Lab 2: Build MXNet on a Docker Container

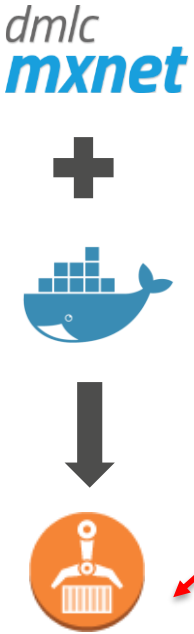

Amazon ECR

```
{
  "containerDefinitions": [
    {
      "memory": 128,
      "portMappings": [
        {
          "hostPort": 80,
          "containerPort": 80,
          "protocol": "tcp"
        }
      ],
      "essential": true,
      "name": "nginx-container",
      "image": "nginx",
      "logConfiguration": {
        "logDriver": "awslogs",
        "options": {
          "awslogs-group": "ecs-log-streaming",
          "awslogs-region": "us-west-2"
        }
      },
      "cpu": 0
    }
  ],
  "family": "example_task_1"
}
```

ECS Task Definition

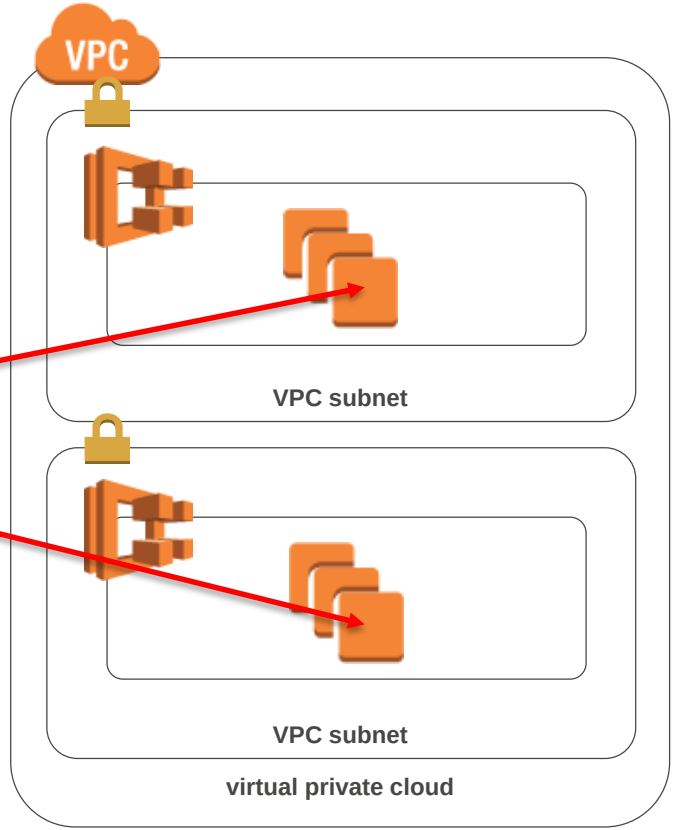

Lab 3: Launch MXNet with ECS

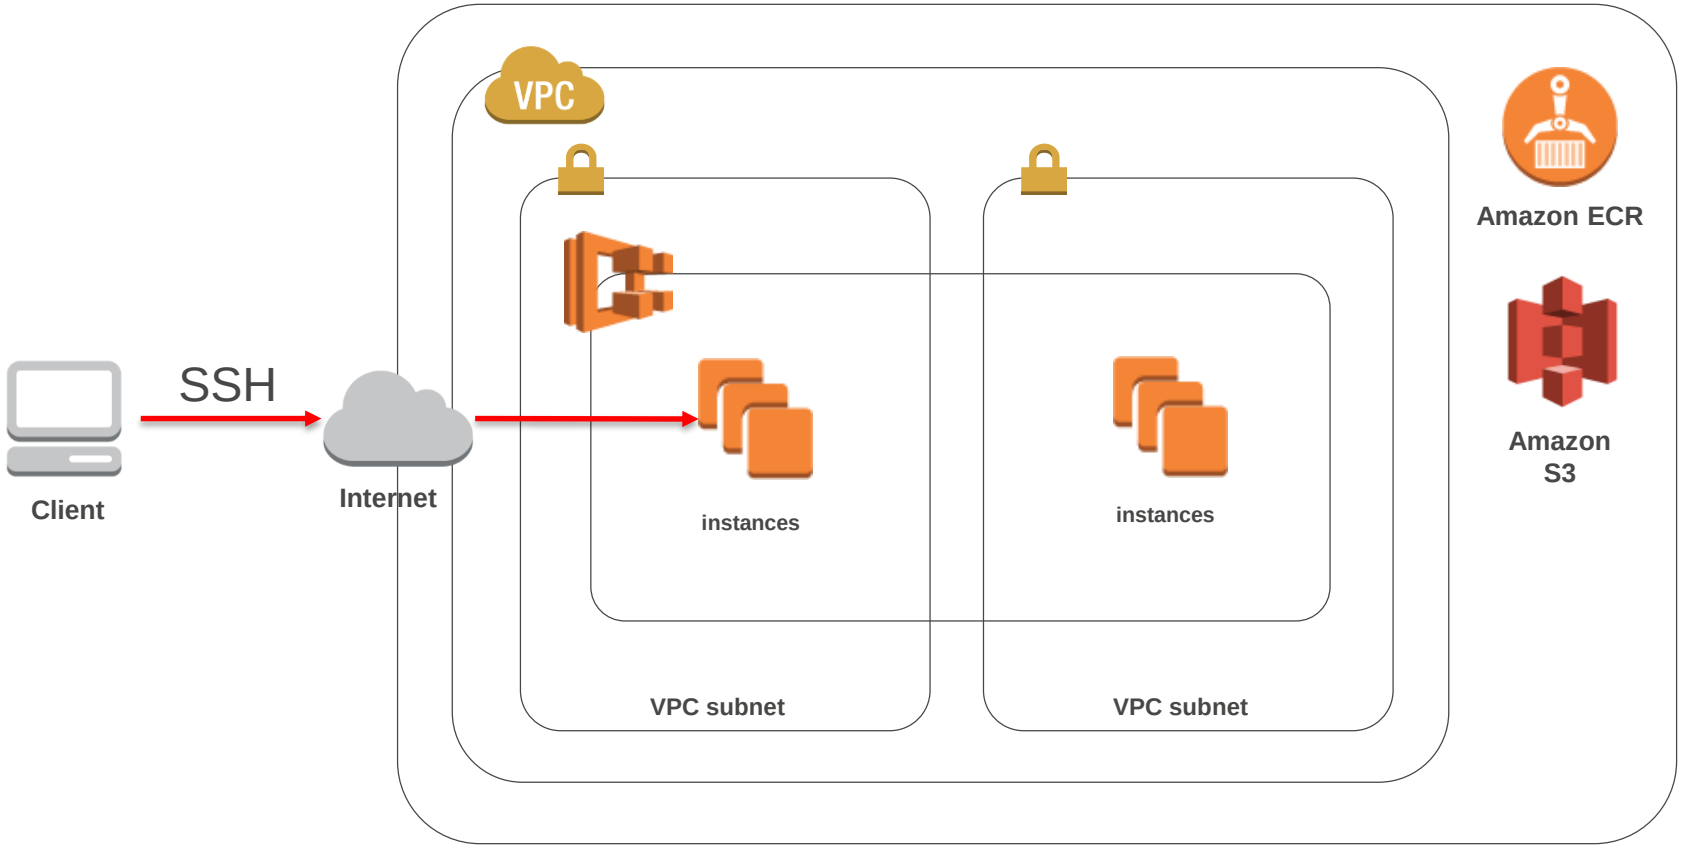

Lab 4: Image Classification Demo

```
In [4]: url = 'http://writm.com/wp-content/uploads/2016/08/Cat-hd-wallpapers.jpg'
        predict(get_image(url), mod, synsets)
```

```
probability=0.692329, class=n02122948 kitten, kitty
probability=0.043847, class=n01323155 kit
probability=0.030002, class=n01318894 pet
probability=0.029693, class=n02122878 tabby, queen
probability=0.026972, class=n01322221 baby
```

Wrap-up

- ECS makes cluster management easy
- ECS has flexible scheduling
- ECS has enables strong security posture
- What other applications can you containerize to deploy and manage at scale?

Thank you!

Remember to complete your evaluations!

Related Sessions

- MAC306 - Using MXNet for Recommendation Modeling at Scale
- CON301 - Operations Management with Amazon ECS
- CON302 - Development Workflow with Docker and Amazon ECS
- CON401 - Amazon ECR Deep Dive on Image Optimization

Appendix

Estimated Workshop Costs