



**AWS
re:Invent**

FIN301

Fraud Detection and Machine Learning on AWS


Pawan Agnihotri, Principal Solutions Architect, Global Financial Services

November 28, 2016



Payments fraud is an ongoing concerns for Financial Services (FS) organizations.

In 2015...



\$8.5b of fraud losses
in the US*



\$21.8b of fraud losses
globally*



\$31.7b of **projected**
fraud losses globally in 2020*

* From The Nilson Report (https://www.nilsonreport.com/constant_contact_promo.php?id_promo=8)

Solving for Multiple Layers Simultaneously

Layer 1

Addresses endpoint authentication

Layer 2

Analyzes session behavior

Layer 3

Monitors account behavior for a channel


Layer 4

Monitoring account behavior across multiple channels


Layer 5

Monitoring multiple account behaviors across multiple channels

Rule-Based Fraud Detection




Rules Fall Short for Fraud Detection




Static set of rules



Human errors & bias

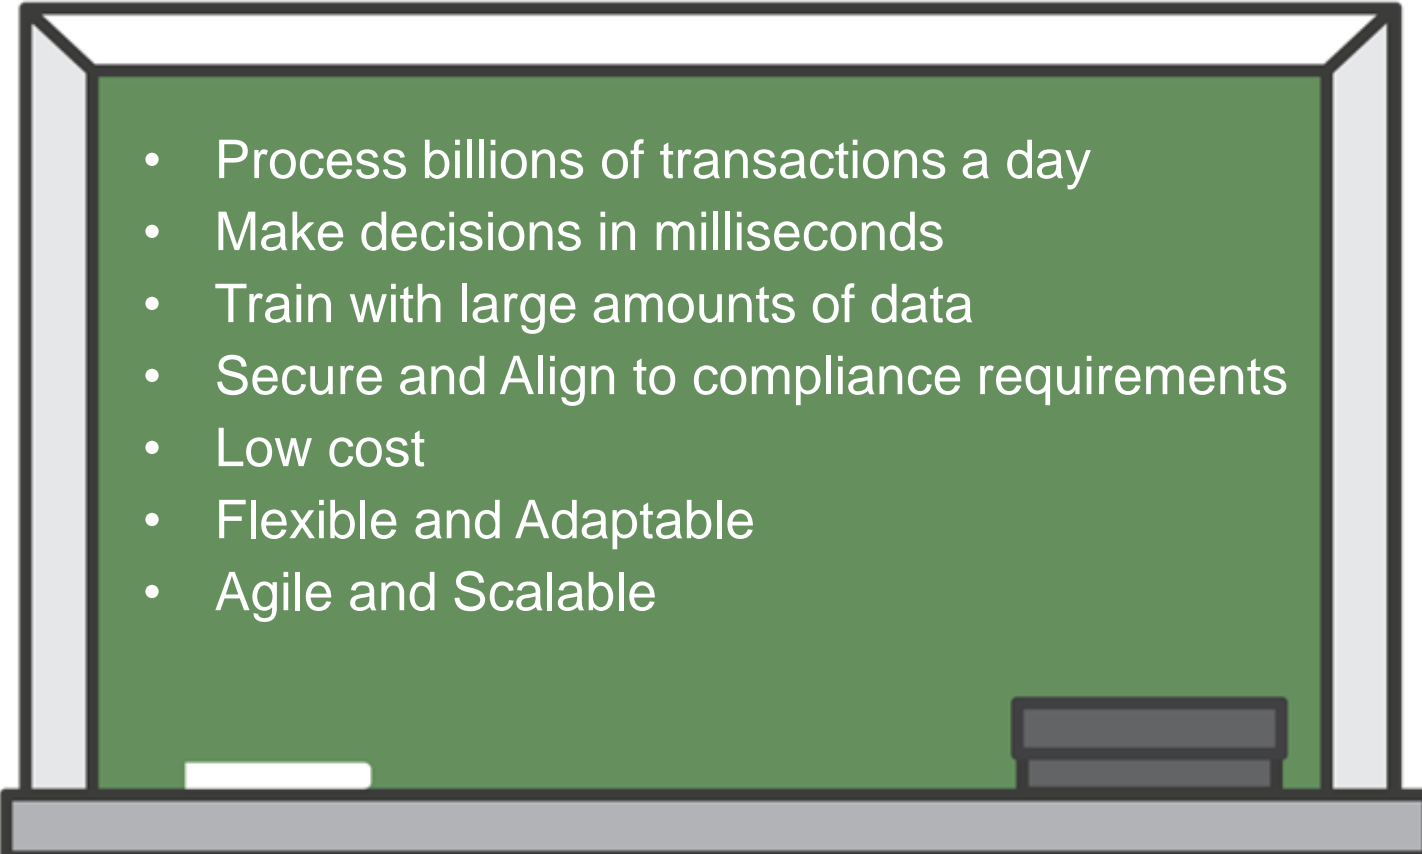


Difficult to manage




Cannot scale

Solution Requirements

- 
- Process billions of transactions a day
 - Make decisions in milliseconds
 - Train with large amounts of data
 - Secure and Align to compliance requirements
 - Low cost
 - Flexible and Adaptable
 - Agile and Scalable


Overview of Machine Learning

Supervised Learning



Supervised Learning

Input



Outcome

Supervised Learning

Input




Input



Input



Input



Outcome



Outcome




Outcome




Outcome

Supervised Learning

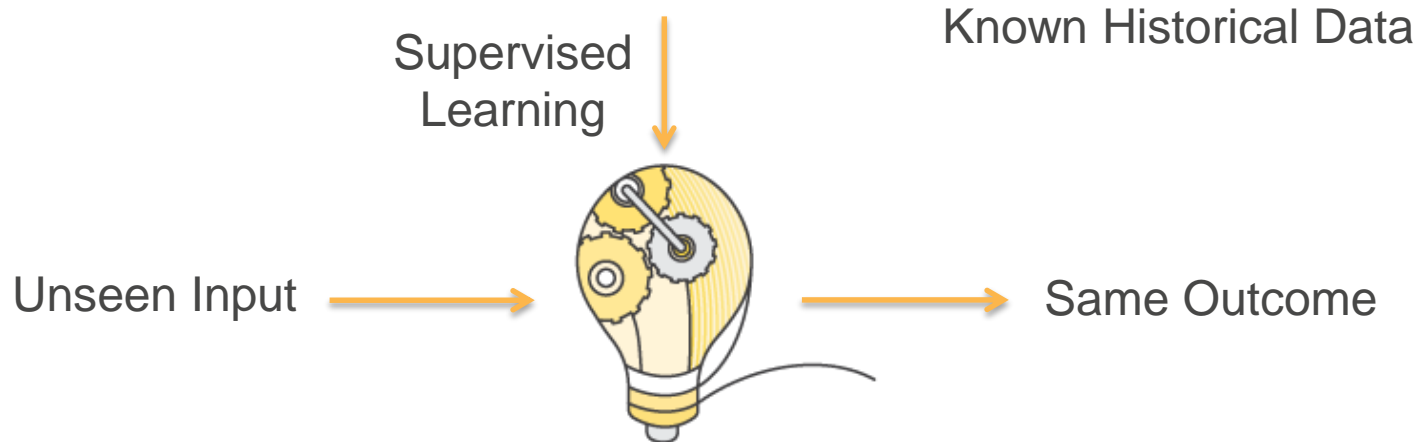



Supervised
Learning

Known Historical Data



Supervised Learning



Tools of the Trade

Amazon Simple Storage Service (S3)



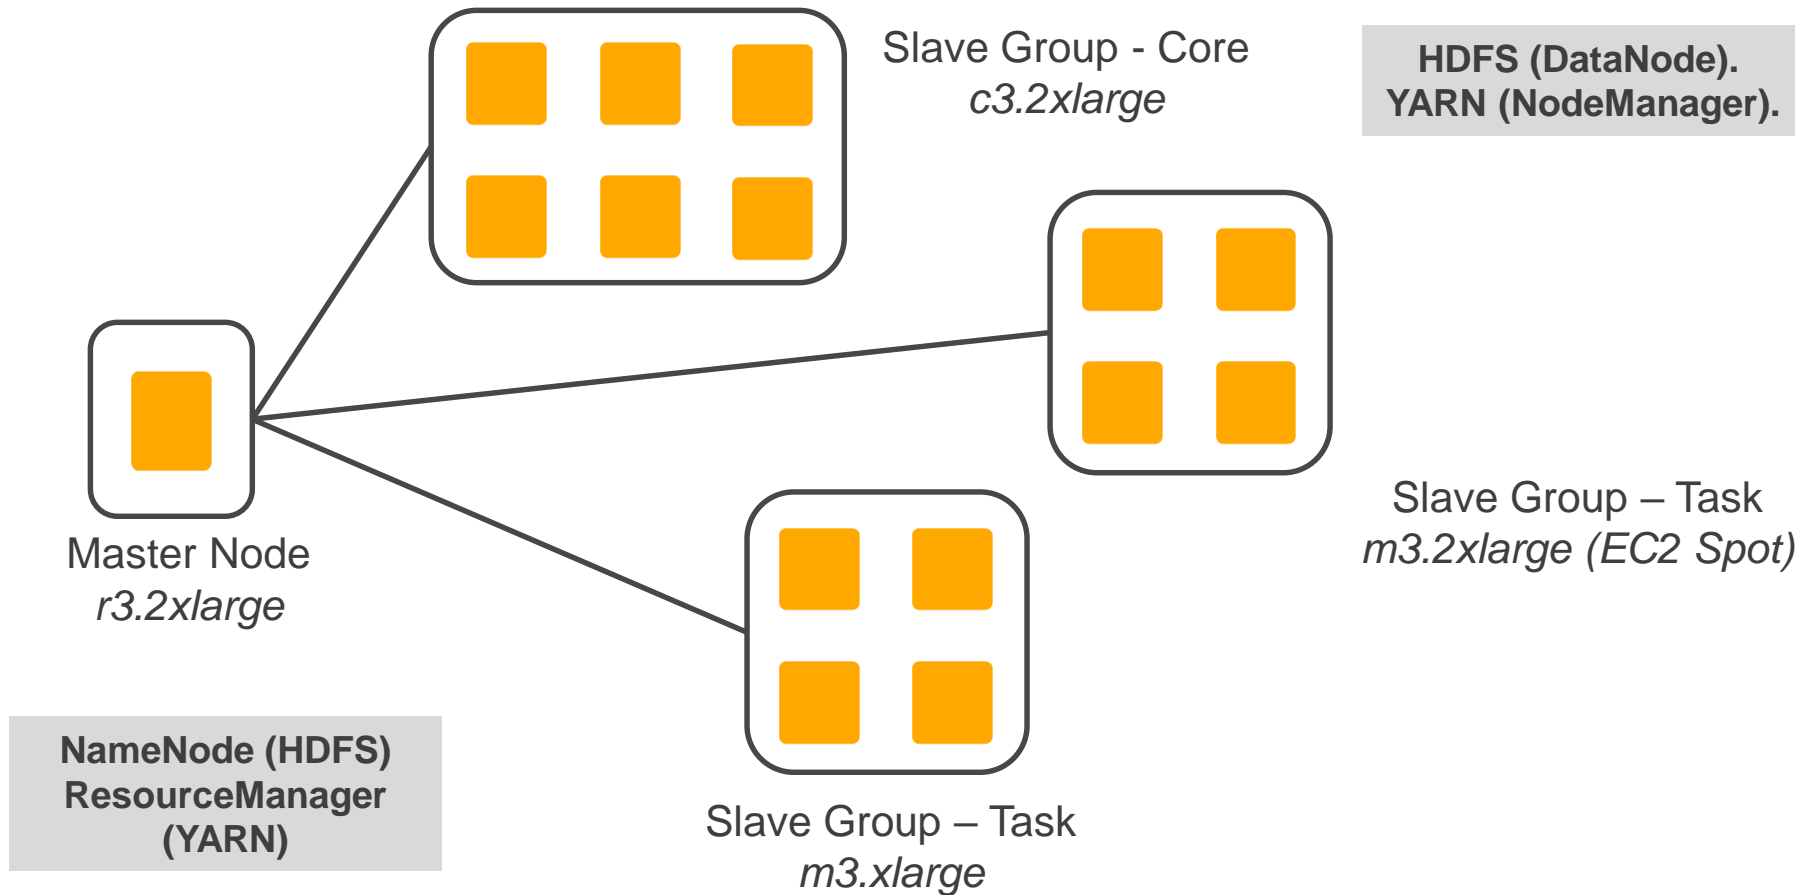
- Highly scalable object storage
- Files are stored as objects and organized into high-level folders called buckets
- Store and retrieve data from anywhere on the web
- Native support data encryption at rest
- Data in transit to and from the service is encrypted using SSL.
- Pay for exactly what you use
- Highly durable (99.999999999% design)
- Limitlessly scalable

Amazon Elastic Map Reduce (EMR)



- *Managed* platform
- MapReduce, Apache Spark, Presto
- Launch a cluster in minutes
- Open source distribution & MapR distribution
- Elasticity of the cloud
- Built in security features
- Support for encryption of data at rest and in transit
- Pay by the hour and save with Spot
- Flexibility to customize

An Example EMR Cluster



Flexibility to add Hadoop applications to Amazon EMR

Processing



Flink



Databases

presto




APACHE
HBASE



Analytics




Amazon Machine Learning






- Easy-to-use service built for developers
- Robust, powerful, and technology-based
- Ability to create models using your data
- Deployable to production in seconds






Amazon Machine Learning Service











Amazon Machine Learning Service



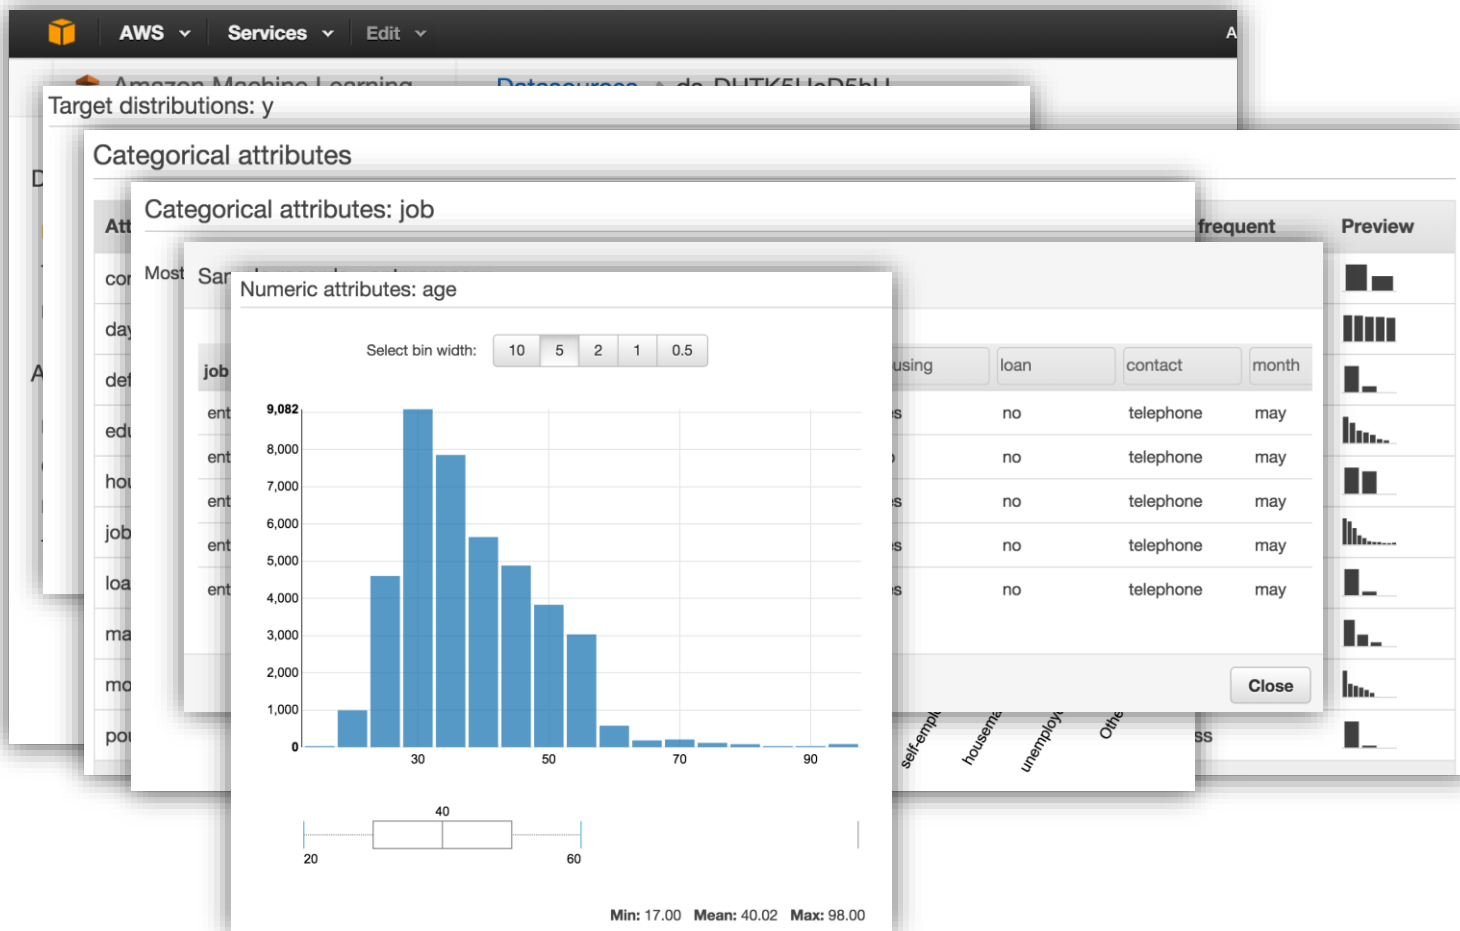
Amazon Machine Learning Service



Amazon Machine Learning Service




Explore and Understand Your Data



Evaluate and Explore Model Performance

of Records



Predicted most likely to be "0"

True negative (86.9%)

This chart shows the distribution of the residuals of your ML model.

Select bin width:

10 5 2 1 0.5

1,335

1,200

1,000

800

600

400

200

0



On your most recent evaluation, **ev-NqLivRuTZiZ**, the ML model's quality score is **better** than the baseline. ⓘ


Average F1 score: 0.64

Baseline F1 score: 0.02

Difference: 0.62


Explore model performance

Total (329) (921) (919) (876) (767) (382) (336) (4530) 0.64



Putting It Together

Credit Card Transaction Dataset



**Customer
Profile**



**Store
Profile**



**Transaction
Details**

Model Creation and Training – Reference Architecture



AWS

VPC

PCI Data Security Standard – High Level Overview



AWS IAM



Amazon CloudWatch



AWS Config



AWS CloudTrail



AWS KMS



Amazon S3

Build and Maintain a Secure Network and Systems

1. Install and maintain a firewall configuration to protect cardholder data
2. Do not use vendor-supplied defaults for system passwords and other security parameters

Protect Cardholder Data

3. Protect stored cardholder data
4. Encrypt transmission of cardholder data across open, public networks

Maintain a Vulnerability Management Program

5. Protect all systems against malware and regularly update anti-malware software or programs
6. Develop and maintain secure systems and applications

Implement Strong Access Control Measures

7. Restrict access to cardholder data by business personnel to the minimum needed to know
8. Identify and authenticate access to system components
9. Restrict physical access to cardholder data

Regularly Monitor and Test Networks

10. Track and monitor all access to network resources and cardholder data
11. Regularly test security systems and processes

Maintain an Information Security Policy

12. Maintain a policy that addresses information security for all personnel




EMR



SSL/TLS



Amazon RDS




Spark

ML lib



Amazon Machine Learning



HBase




H2O.ai

AWS Direct Connect

Corporate Data Center


Online Fraud Detection – Reference Architecture




The Outcomes of the AWS Solution



Cost: Solution price down from \$100K to \$10K



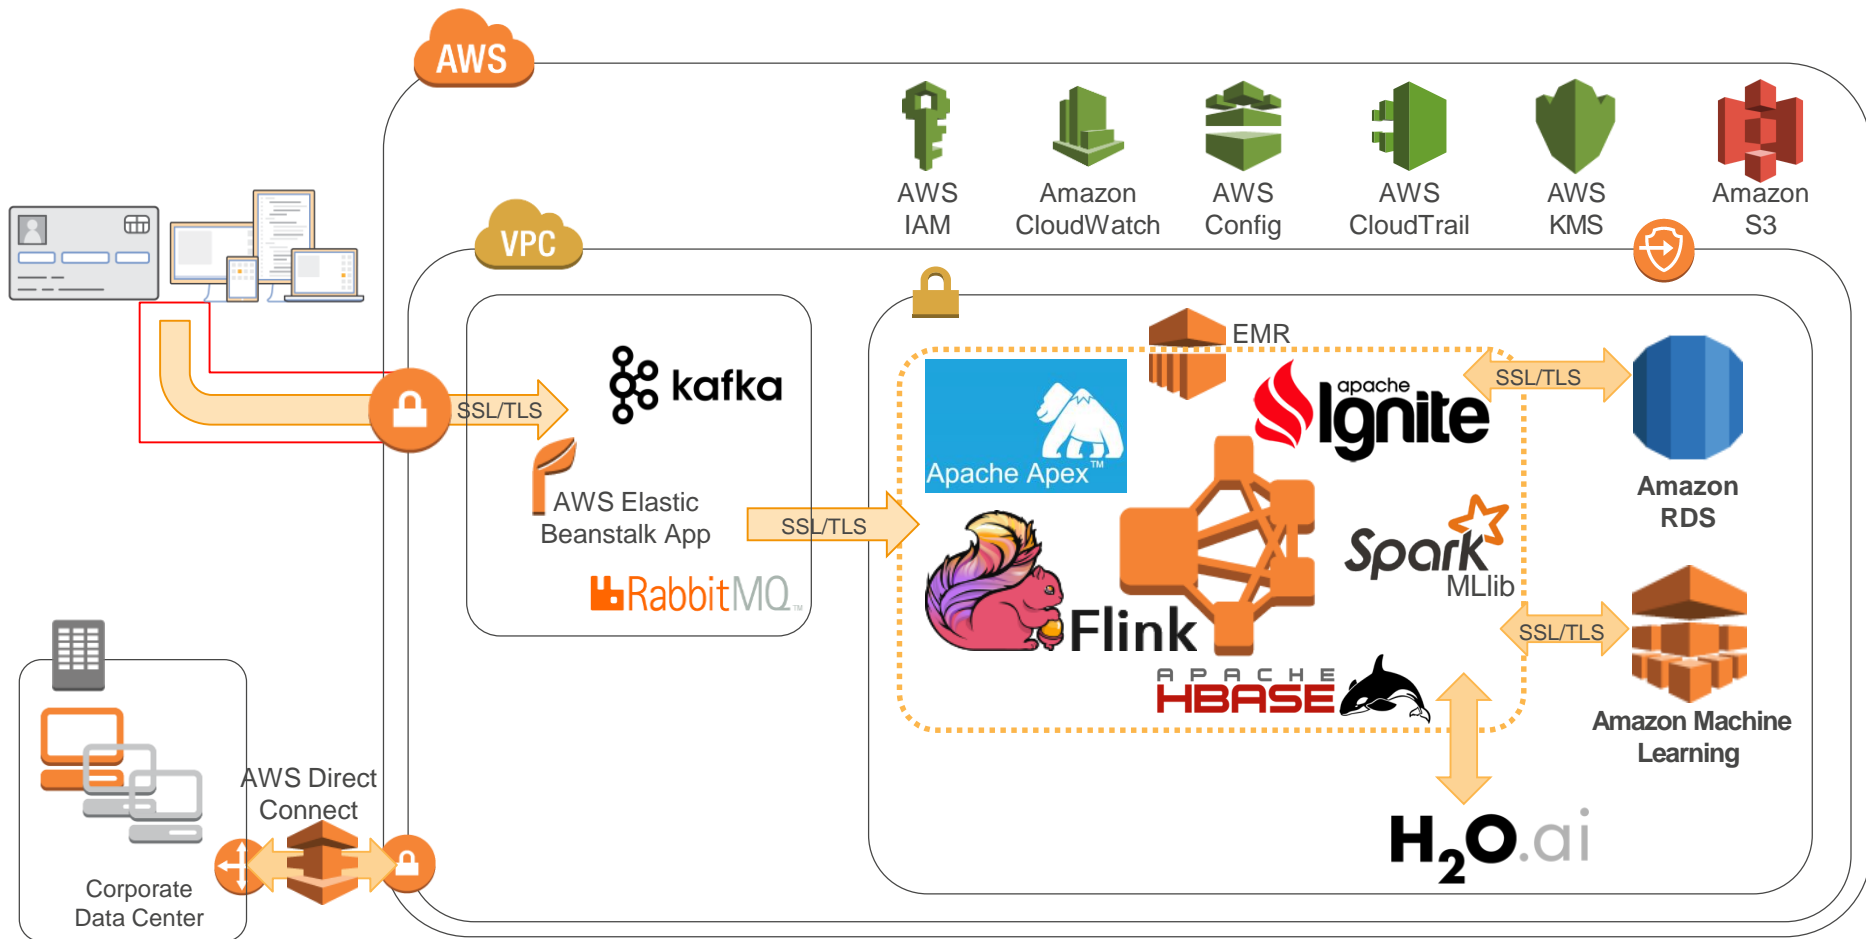
Speed: Development down from months to days



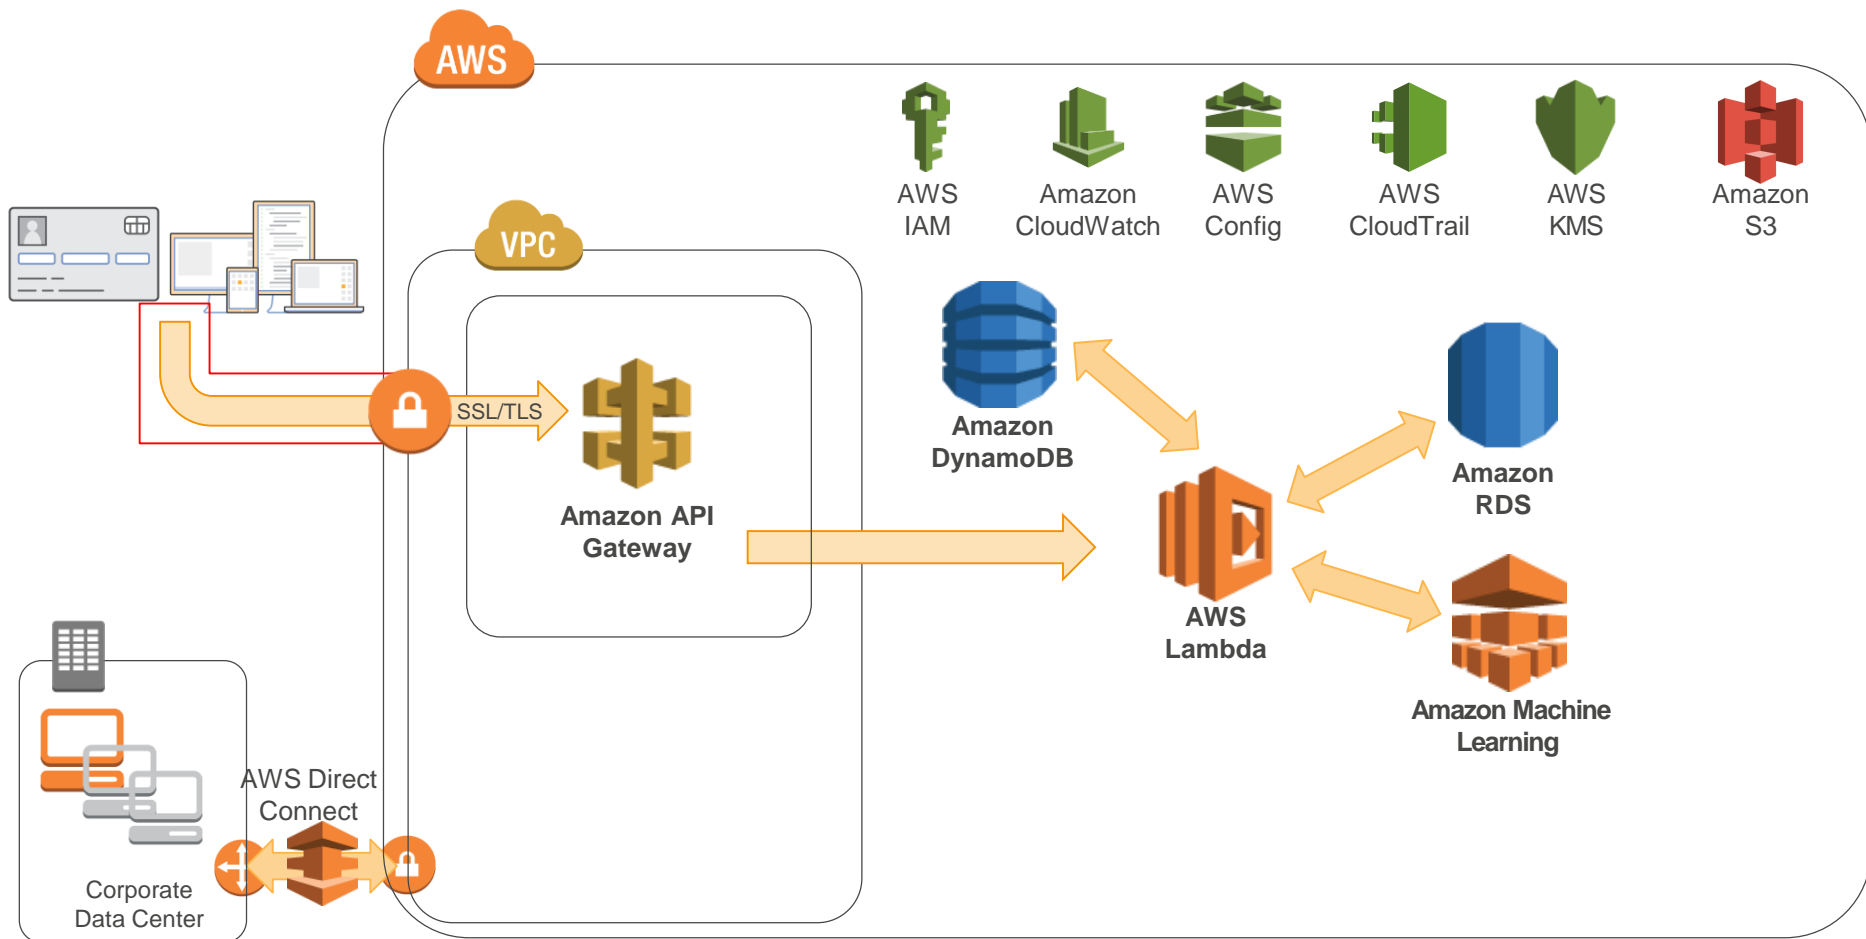
Resources: Focus shift from management to development

Next Evolution of the Platform

Online Fraud Detection – Reference Architecture



Online Fraud Detection – Future State



Other Sessions on Machine Learning

CMP314 - Bringing Deep Learning to the Cloud with Amazon EC2

MAC206 - Machine Learning State of the Union

MAC303 - Developing Classification and Recommendation Engines with Amazon EMR and Apache Spark



**AWS
re:Invent**

Thank you!



**Remember to complete
your evaluations!**