



**AWS
re:Invent**

MAC306

Using MXNet for Recommendation Modeling at Scale

Leo Dirac, Principal Engineer, AWS Deep Learning

November 30, 2016

What to Expect from the Session

Background on recommender systems and machine learning.

Learn how to implement them on MXNet using p2 instances and the AWS Deep Learning AMI.

Explore several types of recommender systems, including advanced deep learning ideas.

Learn tricks for handling sparse data in MXNet.

Background: Recommender Systems & Machine Learning

Netflix Prize: 2006-2009
\$1,000,000

The Netflix logo, consisting of the word "NETFLIX" in a bold, white, sans-serif font with a 3D effect, set against a solid red rectangular background.

NETFLIX

Recommending Movies

```
/* Predict what Star Rating will user u give  
movie m */
```


```
float predictRating(User u, Movie m) {  
    // How???  
}
```

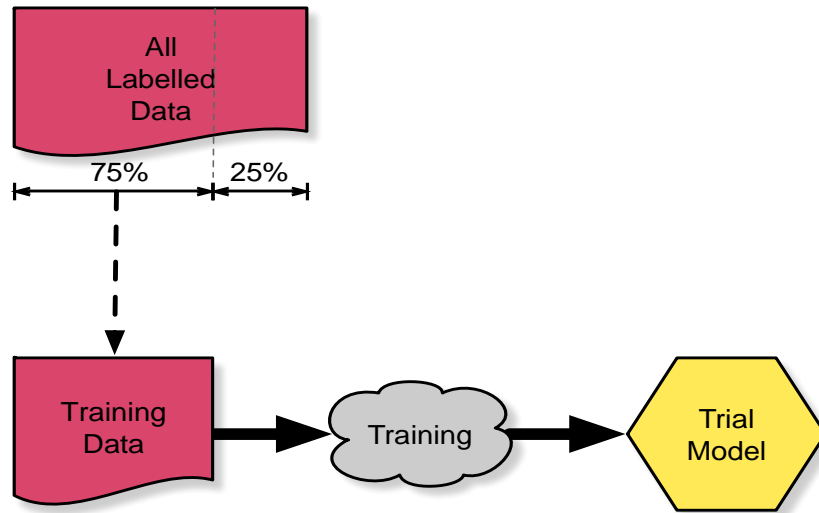
Q: How???

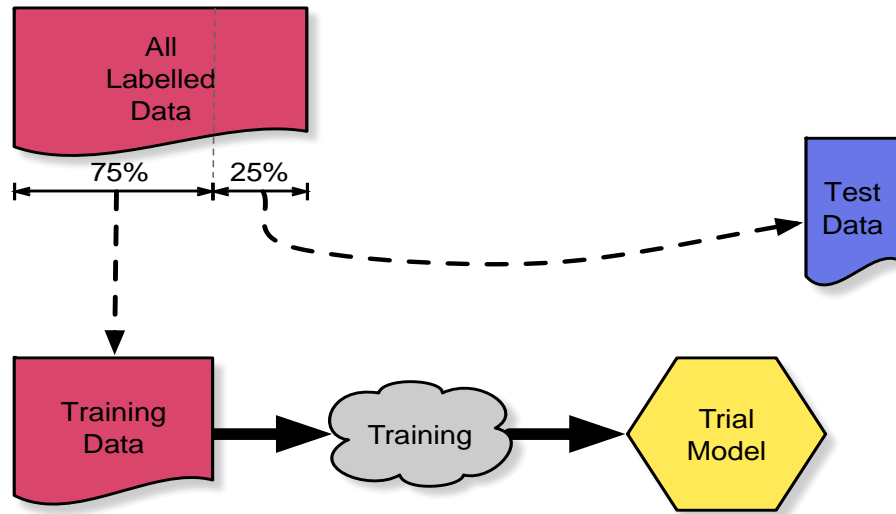
A: Machine Learning: Learn code from data

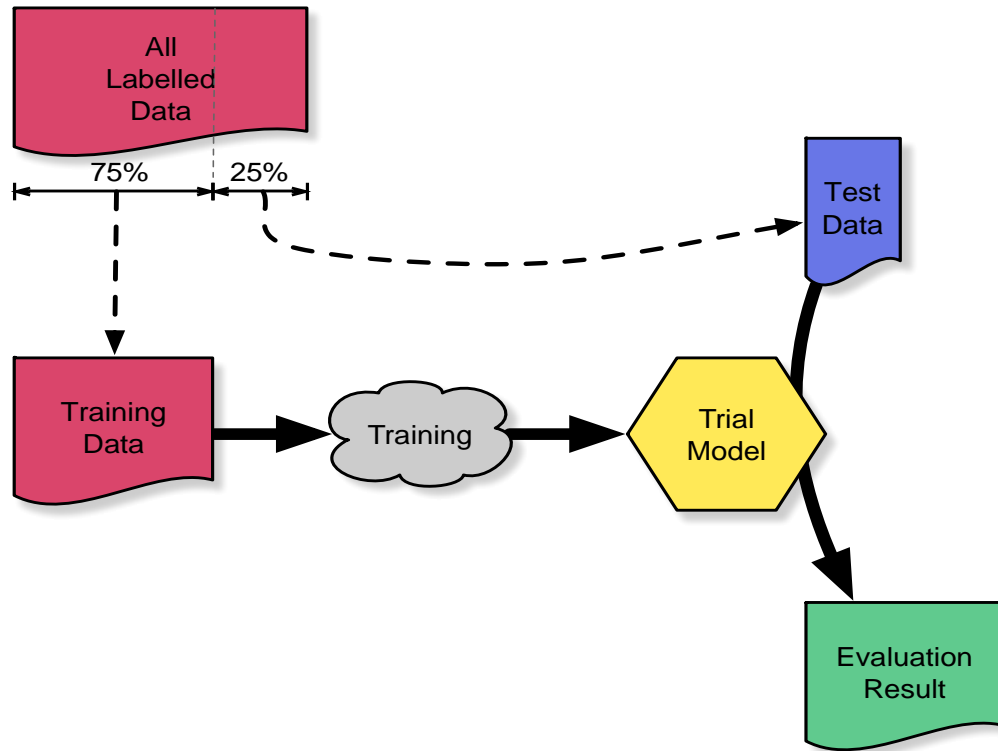
```
float predictRating(User u, Movie m) {  
    return mlModel.run(u,m);  
}
```

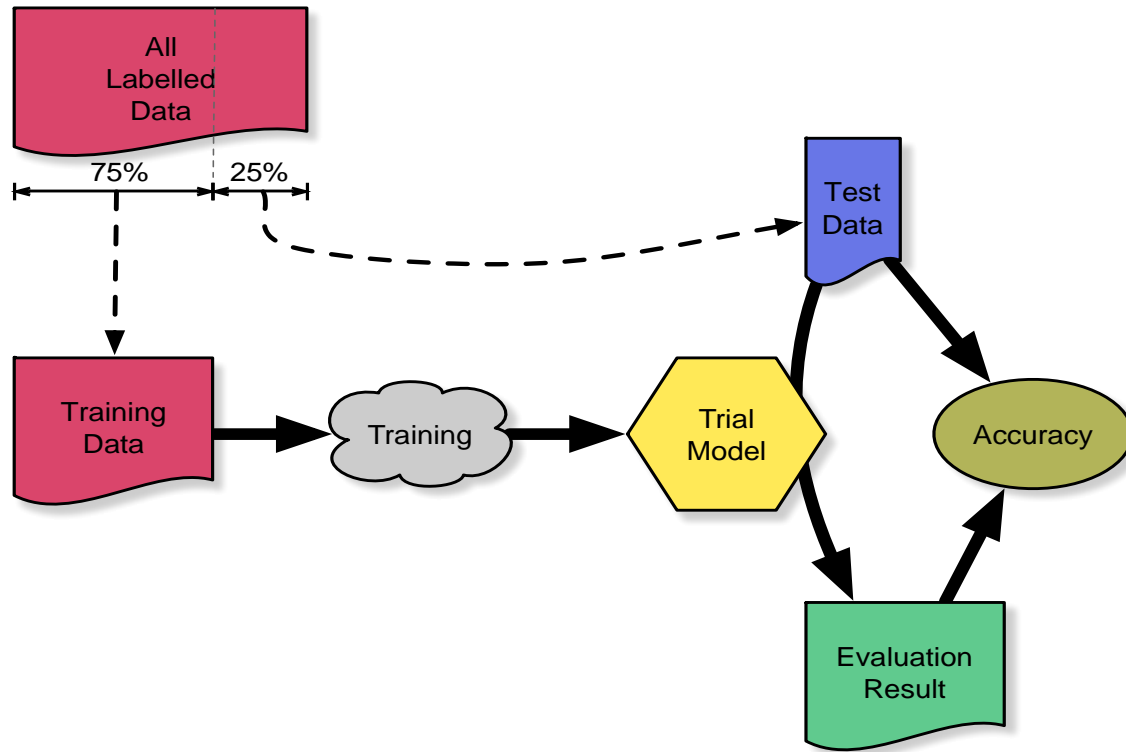












Sparse Data

User-Item Ratings Matrix

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z
1		u001	u002	u003	u004	u005	u006	u007	u008	u009	u010	u011	u012	u013	u014	u015	u016	u017	u018	u019	u020	u021	u022	u023	u024	u025
2	Clockstoppers																									
3	Sabrina: The Teenage Wit																				5					
4	The Stone Boy	5							1																	
5	Heart Condition										5															
6	The Lord of the Rings																			2						
7	Common Wealth															5										
8	Larva				3																					
9	Duel in the Sun													4												
10	Man on the Train																									
11	Treasure Island																									
12	Fighting for Love																									
13	Merce Cunningham: A Life																									
14	Cream: Farewell Concert																									
15	13 Going on 30																									
16	Classic Albums: Stevie Nicks																									1
17	Dark Waters											3														
18	The Lifestyle: Swinging in																	4								
19	Opinion																									
20	Dazed and Confused																				4					
21	Coast to Coast																			1						
22	Since You Went Away																									
23	The Town Is Quiet																									
24	Tornado!																									
25	Sightings: Heartland Ghos																									
26	The Thing				4																					
27	Automotive Series: Porsc															5										
28	Family Guy Presents: Stev							1					5													
29	Love After Death																									
30	F Troop: TV Favorites																									
31	Escape Velocity																									
32	Alphaville																									
33	Classic Albums: Metallica			3																						
34	Monster Man																									
35	WWE: King of the Ring 20																4					3			1	
36	Shaft in Africa																									

Size of user-item ratings matrix

Sample dataset: MovieLens 20M

$(27,000 \text{ movies}) * (138,000 \text{ users})$
 $= 3,700,000,000$ possible ratings

But only 20,000,000 ratings available.

99.5% of ratings are unknown.

Storing the matrix

Dense

3.7B entries

Each entry:

- Rating: 1 byte

3.7 GB

Sparse

20M non-zero entries

Each entry:


- Rating: 1 byte
- Movie_id: 32-bit integer
- User_id: 32-bit integer

180 MB

Sparse is 20x smaller

Matrix Factorization

MF as Math



Embeddings

Emb("The Karate Kid") = $\begin{bmatrix} -3.168 \\ -0.136 \\ 3.770 \\ 4.767 \\ 3.558 \\ -4.168 \\ 0.464 \\ 2.034 \\ 3.411 \\ \dots \\ 0.866 \end{bmatrix}$

Amazon

Embeddings

Emb("The Karate Kid") = $\begin{bmatrix} -3.168 \\ -0.136 \\ 3.770 \\ 4.767 \\ 3.558 \\ -4.168 \\ 0.464 \\ 2.034 \\ 3.411 \\ \dots \\ 0.866 \end{bmatrix}$

Emb("Ferris Bueller") = $\begin{bmatrix} -3.101 \\ -0.057 \\ 3.800 \\ 4.862 \\ 3.632 \\ -4.157 \\ 0.549 \\ 2.064 \\ 3.428 \\ \dots \\ 0.884 \end{bmatrix}$

$$D(\text{Emb}(\text{"K.Kid"}) - \text{Emb}(\text{"Ferris"})) = 0.138$$

$$D(\text{Emb}(\text{"K.Kid"}) - \text{Emb}(\text{"My Little Pony"})) = 1.572$$

MXNet

Flexible and Efficient Library for Deep Learning

Get Started



Star 6,454



Fork 2,391

Flexible

Supports both imperative and symbolic programming

Portable

Runs on CPUs or GPUs, on clusters, servers, desktops, or mobile phones

Multiple Languages

Supports multiple languages, including C++, Python, R, Scala, Julia, Matlab and Javascript – All with the same amazing performance.

Auto-Differentiation

Distributed on Cloud

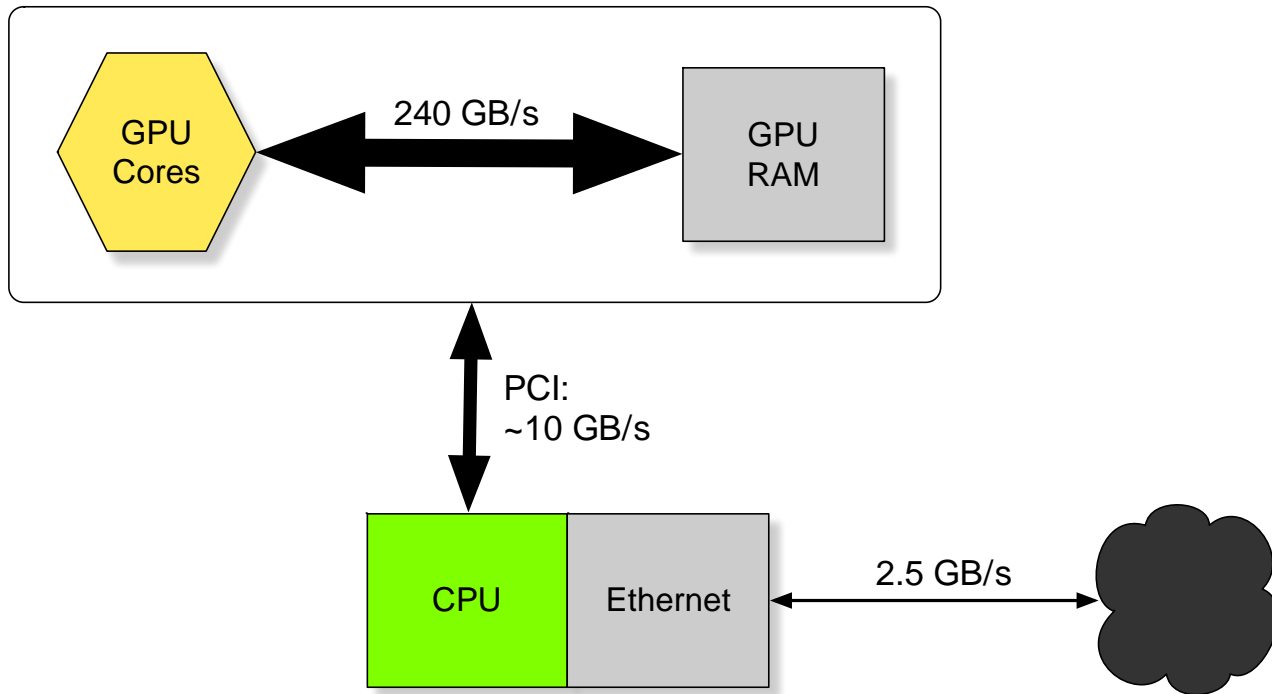
Performance

p2.xlarge

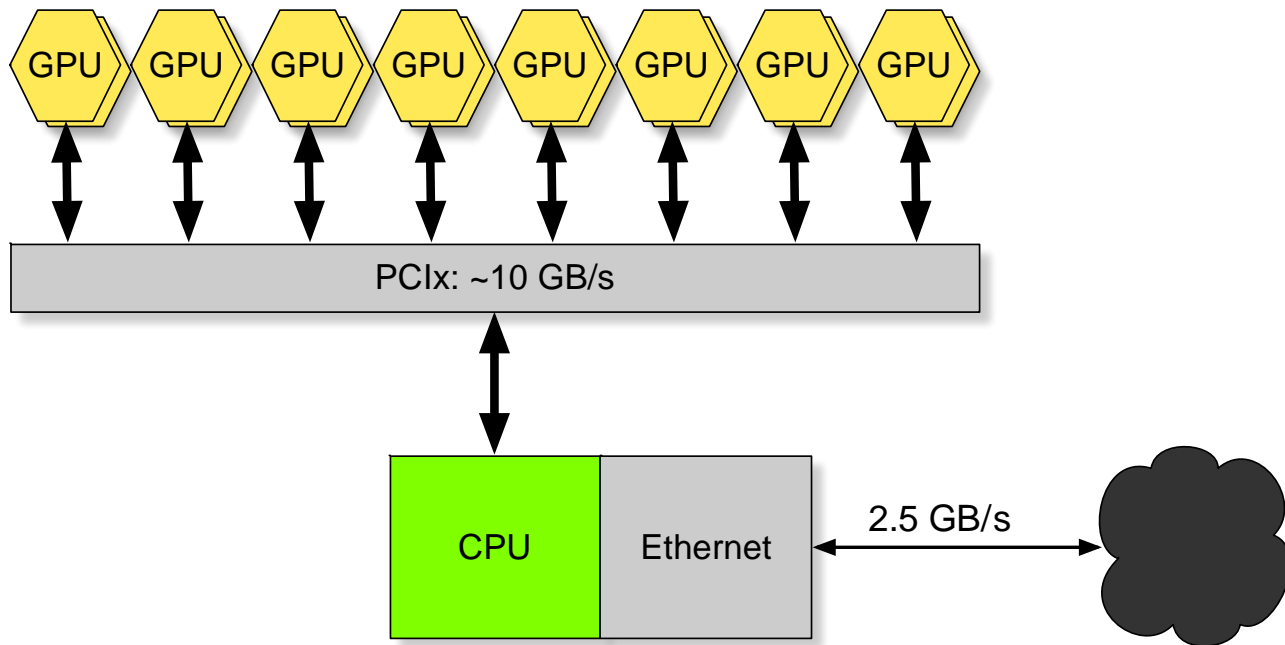
4,300,000,000,000

32-bit floating point
operations/second

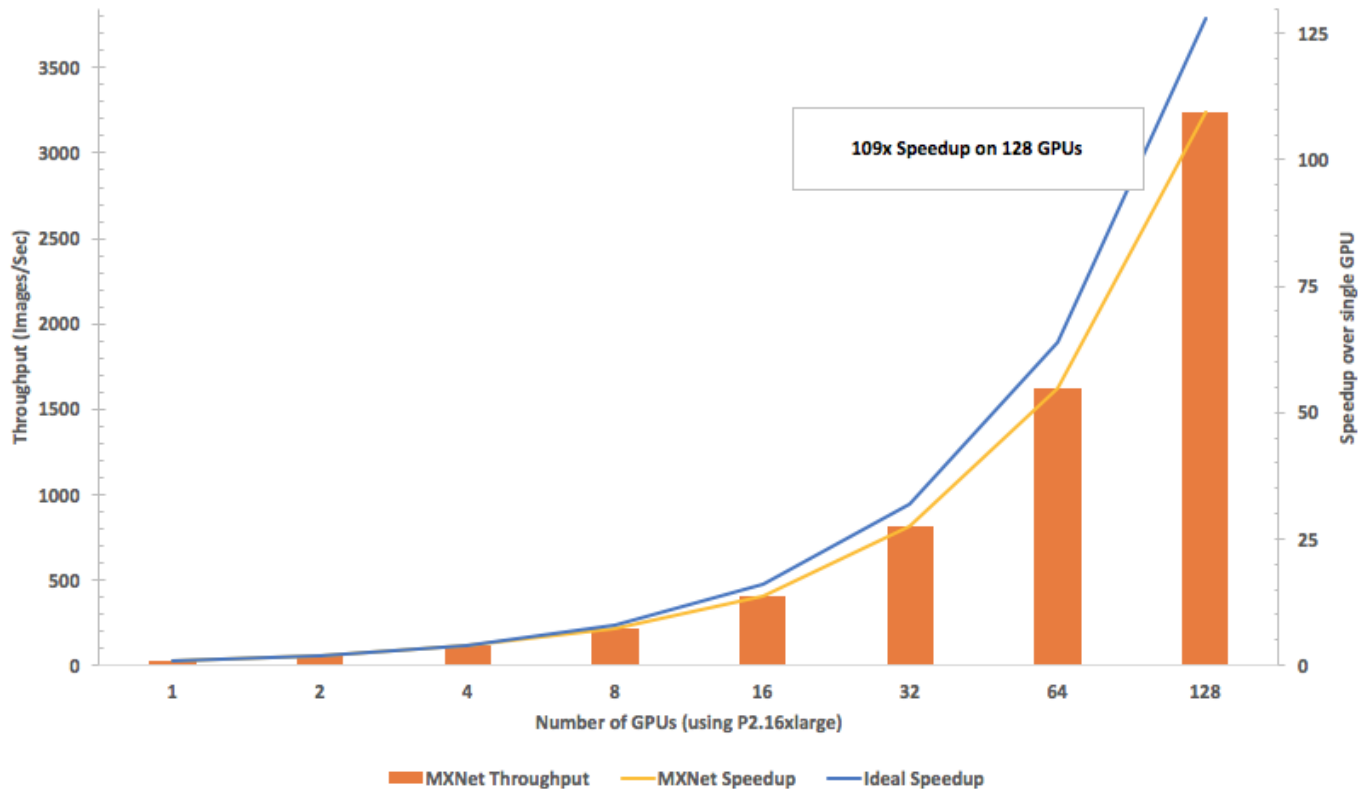
GPUs: Feeding the beast



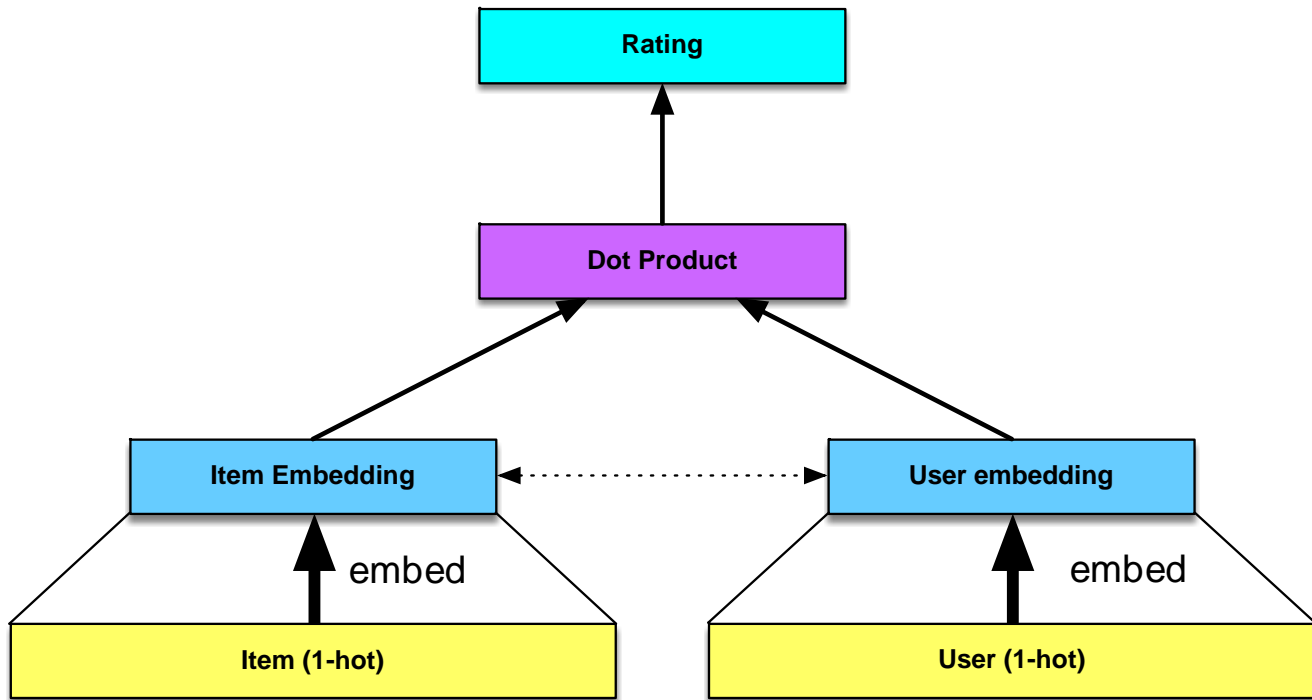
p2.16xlarge



MXNet scaling



MF as a neural network (NN)



Deep Learning AML with p2

Pre-installed:

- MXNet & other popular deep learning frameworks
- GPU Drivers, CUDA, cuDNN
- Jupyter notebook & python libraries

MF Demo in MXNet

[demo1-MF.ipynb](#)

Binary Predictions

Why binary?



Binary user-item matrix

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z	AA	AB	AC	AD	AE	AF	AG	AH	AI
1		u001	u002	u003	u004	u005	u006	u007	u008	u009	u010	u011	u012	u013	u014	u015	u016	u017	u018	u019	u020	u021	u022	u023	u024	u025	u026	u027	u028	u029	u030	u031	u032	u033	u034
2	Clockstoppers																										1								
3	Sabrina: The Teenage Wit																																		
4	The Stone Boy																																1		
5	Heart Condition		1																																
6	The Lord of the Rings																																		
7	Common Wealth																									1									
8	Larva																																		
9	Duel in the Sun	1																																	
10	Man on the Train																																		
11	Treasure Island																																		
12	Fighting for Love																																		
13	Merce Cunningham: A Life																										1								
14	Cream: Farewell Concert																																		
15	13 Going on 30						1																												
16	Classic Albums: Stevie Nicks																									1							1		
17	Dark Waters																																		
18	The Lifestyle: Swinging in													1																					
19	Opinion																									1									
20	Dazed and Confused																																1	1	
21	Coast to Coast					1															1														
22	Since You Went Away								1					1																					1
23	The Town is Quiet			1																													1		
24	Tornado!																																1		
25	Sightings: Heartland Ghos										1																								
26	The Thing																	1																	
27	Automotive Series: Porsc																																		
28	Family Guy Presents: Stev																								1										
29	Love After Death																																		
30	F Troop: TV Favorites	1																																	
31	Escape Velocity															1																			
32	Alphaville											1											1												
33	Classic Albums: Metallica																																		
34	Monster Man																																		1
35	WWE: King of the Ring 20																														1				
36	Shaft in Africa																								1										
37	When We Were Kings																1																		
38	Tai-Pan																																		
39	When a Stranger Calls																		1								1								
40	Baby Einstein: Numbers N																																		
41	Sandbaggers: Collection 2																1																		
42	Allman Brothers: Live at t																															1			
43	WWE: Vengeance: Hell in																																		
44	Midsomer Murders: Faith																																		
45	New Kids on the Block: Gr												1																						
46	Lone Star											1																							

Original data

	A	B	C
1	<u>User</u>	<u>Movie</u>	<u>Label</u>
2	Leo	Xanadu	1
3	Leo	Caddy Shack	1
4	Petra	Dora the Explorer	1
5	Petra	My Little Pony	1

Predicting binary

```
float predictScore(User u, Movie m) {  
    return 1.0;  
}
```

Original data

	A	B	C
1	<u>User</u>	<u>Movie</u>	<u>Label</u>
2	Leo	Xanadu	1
3	Leo	Caddy Shack	1
4	Petra	Dora the Explorer	1
5	Petra	My Little Pony	1

Negative sampling

	A	B	C
1	<u>User</u>	<u>Movie</u>	<u>Label</u>
2	Leo	Xanadu	1
3	Leo	Caddy Shack	1
4	Petra	Dora the Explorer	1
5	Petra	My Little Pony	1
6	Leo	Dora the Explorer	0
7	Leo	My Little Pony	0
8	Petra	Xanadu	0
9	Petra	Caddy Shack	0
10			

Negative sampling

```
from mxreco import NegativeSamplingDataIter
train_data = NegativeSamplingDataIter(
    train_data,
    sample_ratio=5)
```

More details: [BlackOut: Speeding up RNNLM w/ Very Large Vocabularies](#)

Shihao Ji, S. V. N. Vishwanathan, Nadathur Satish, Michael J. Anderson, Pradeep Dubey

Negative Sampling Demo

[demo2-binary.ipynb](#)

Content Features

What do we know?

Behavioral interactions between users & items

Names of items

Pictures of items

What users searched for

How to represent these in NN?

Unique Identifier: Embedding

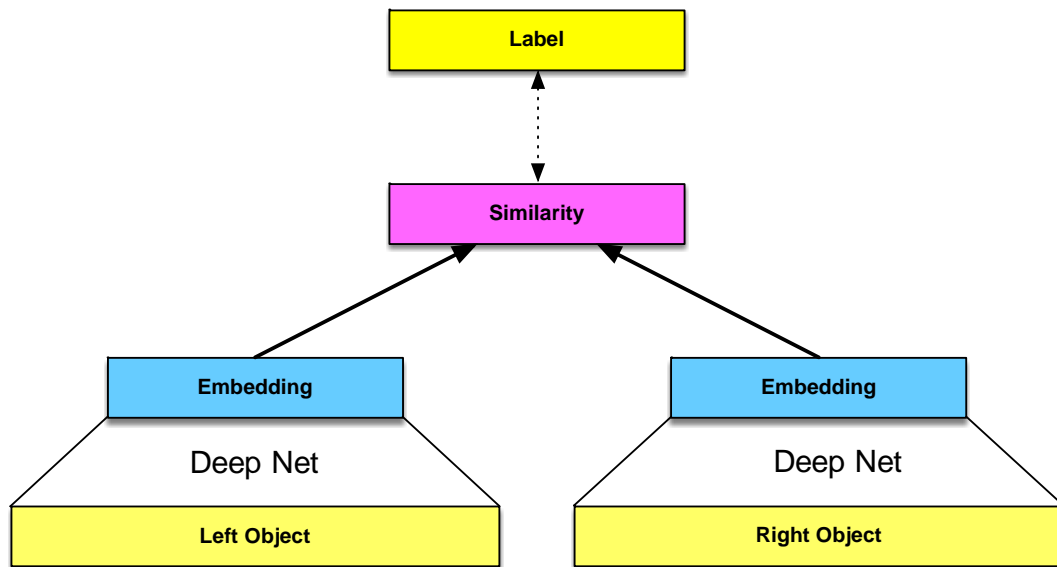
Images: ConvNet (a.k.a. CNN)

Text: LSTM

Text: Bag of Words

DSSM




Deep Structured Semantic Model



CosineLoss layer

```
import mxreco  
pred = mxreco.CosineLoss(a=user, b=item,  
                          label=label)
```

$$\mathcal{L}(\mathbf{A}, \mathbf{B}) = 1 - \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|}$$



Content Features DSSM Demo

[demo3-dssm.ipynb](#)

Inspirational References

[Learning Deep Structured Semantic Models for Web Search using Clickthrough Data](#)

- Po-Sen Huang, Xiaodong He, Jianfeng Gao, Li Deng, Alex Acero, Larry Heck, October, 2013

[Deep Neural Networks for YouTube Recommendations](#)

- Paul Covington and Jay Adams and Emre Sargin, 2016

[Order-Embeddings of Images and Language](#)

- Ivan Vendrov, Ryan Kiros, Sanja Fidler, Raquel Urtasun, March 2016


User-Level Models

Predicting with embeddings

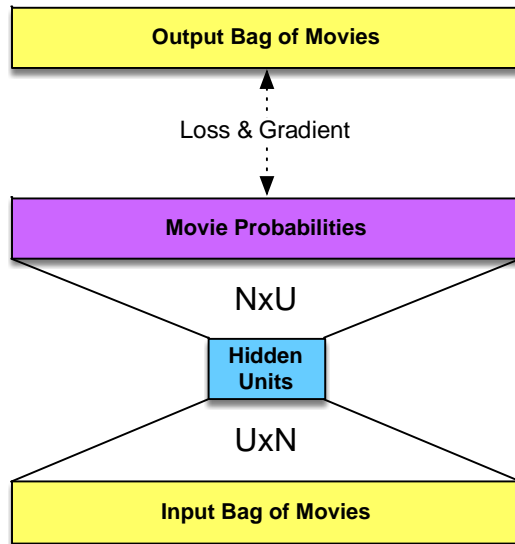
```
def movies_for_user(u):  
    scores = {}  
    for m in movies:  
        score[m.id] = predictScore(u,m)  
    top_movies = sorted(scores.items()...)  
    return top_moves
```

All content at once

```
def movies_for_user(u):  
    scores = userModel.predict(u)  
    top_movies = sorted(scores.items())...  
    return top_movies
```



Multi-label neural network



Sparse output

Sparse input

Storing indexes

Conceptually:

- Predict: 1882, 2808, 24, 160, 1831, 2668
- Inputs: 2986, 329, 2012, 442, 512, 1544, 2615, 1037, 1876, 1917, 2532, 196, 1375, 1779, 2054, 2530, 2628, 1909, 2407, 316, 1356, 1603, 2046, 2428

Storing sparse data

Simpler if fixed width

Pad to end with “-1”

- Predict: 1882, 2808, 24, 160, 1831, 2668,-1,-1,-1,-1
- Inputs: 2986, 329, 2012, 442, 512, 1544, 2615, 1037, 1876,
1917, 2532, 196, 1375, 1779, 2054, 2530, 2628, 1909, 2407,
316, 1356, 1603, 2046, 2428,-1,-1,-1,-1,-1,-1,-1,-1

Trying It Yourself

Trying it yourself

Launch Deep Learning AMI

<https://aws.amazon.com/marketplace/pp/B01M0AXXQB>

Try examples in

<https://github.com/dmlc/mxnet/example/recommender>

The background features a large, abstract graphic with blue and orange wavy, ribbon-like shapes. These shapes are overlaid with a pattern of concentric dotted circles in light gray and orange, creating a sense of motion and depth.

**AWS
re:Invent**

Thank you!



**Remember to complete
your evaluations!**