

AWS re:Invent

How Toyota Racing Development Makes Racing Decisions in Real Time with AWS

The Power of DynamoDB and DynamoDB Streams

Jason Chambers, Toyota Racing Development
Philip Loh, Toyota Racing Development
Martin Sirull, AWS

November 30, 2016

What to Expect from the Session

- How to efficiently manage and hash time-series data in Amazon DynamoDB
- How to stream data using DynamoDB streams
- How to play back DynamoDB data in real time
- Turbo mode: How to mix DynamoDB and Amazon Firehose to get the best of each technology

Who are we?

TRD



Pre-race data

- Use data and models to help predict optimal adjustments for our cars prior to race start
- Provide and present data to drivers for immediate course feedback
- Need data within minutes



Live data

- Provide real-time information and predictions to crews
- Data is used for split-second decisions about pitting, tire changes, and other strategic race strategy decisions
- Need data within seconds



Data architecture

How we laid it out

Our tech stack

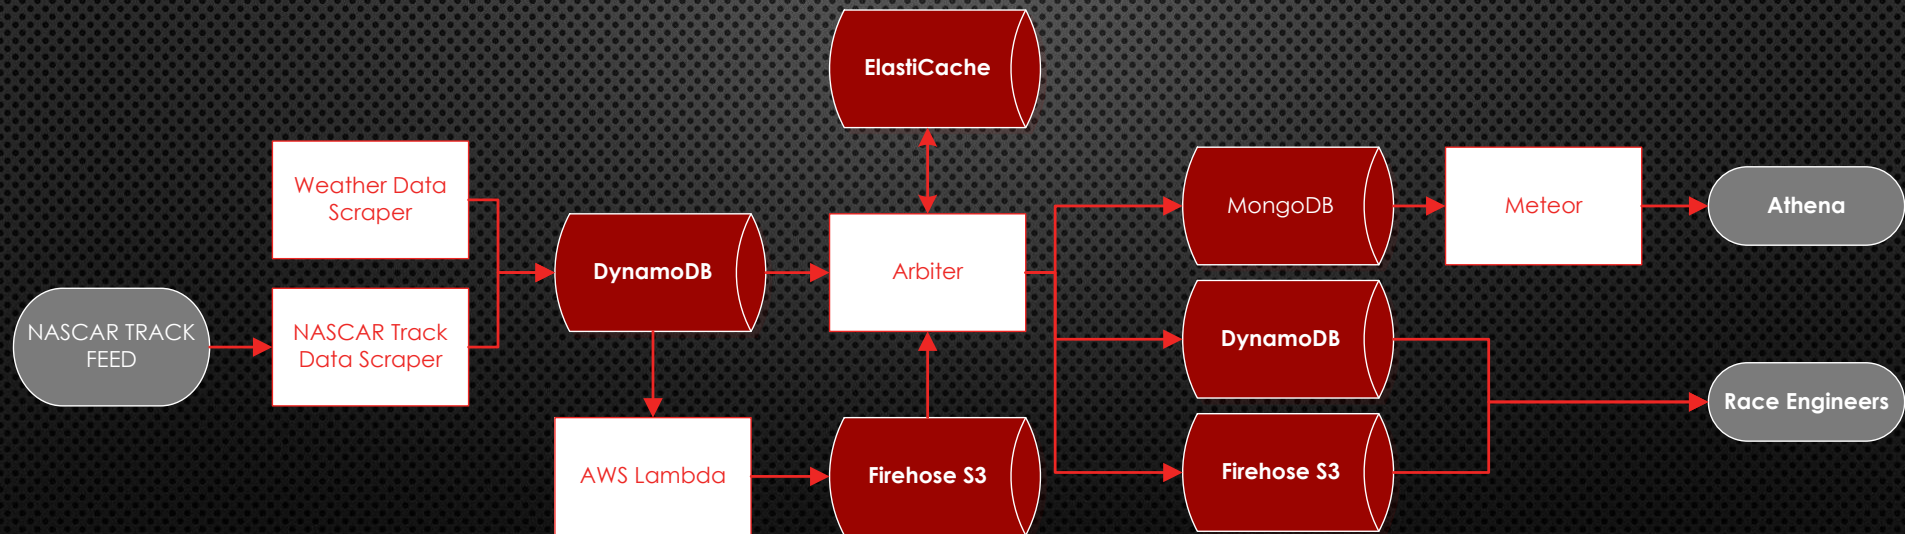
Backend

- Built in JavaScript on Node.js
- Dynamic data – DynamoDB streams & Mongo
- Persistent data - DynamoDB and Amazon S3 services
- Cache – Amazon ElastiCache

Frontend

- Built on JavaScript in Polymer
- Dynamic data served to clients via meteor

Real-time data flow



The power of DynamoDB

Our real-time workhorse

Why we chose DynamoDB

- Decoupled analysis from data collection
- Multiple data sources can write to DynamoDB
- Data processors can listen for this data via streams in real time
- Data is persisted permanently for future queries
- Managed service means no maintenance
- Easy access control with AWS IAM

Hash & range

Choosing the proper hash & range has enormous performance implications

Hash

- Determines partition where data is stored
- Records with equal hashes are stored on same partition
- To query data, you must know at least the record's hash

Range

- Data sharing the same hash is sorted within the partition by its range
- The combination of a record's hash and range constitute its primary key and must be unique

Hash for time-series data

Partition key value	Uniformity
User ID, where the application has many users.	Good
Status code, where there are only a few possible status codes.	Bad
Item creation date, rounded to the nearest time period (e.g. day, hour, minute)	Bad
Device ID, where each device accesses data at relatively similar intervals	Good
Device ID, where even if there are a lot of devices being tracked, one is by far more popular than all the others.	Bad

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/GuidelinesForTables.html>

Time-series data: hot key mitigation

- Separate tables
- Manual partitioning
- Composite keys (TRD solution)

Time-series data: separate tables

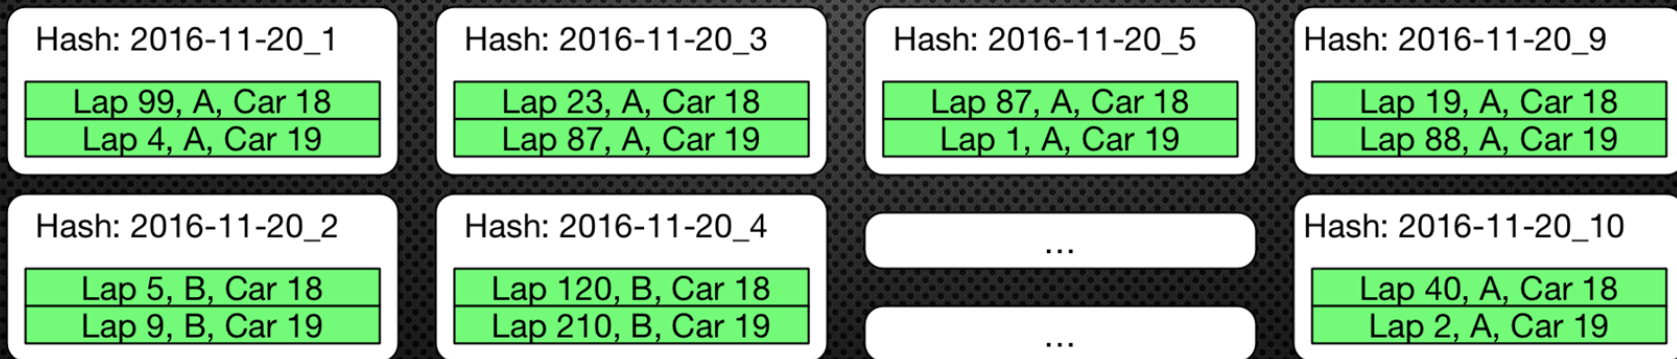
Hash: 2016-11-13

Lap 1, A, Car 18	Lap 3, A, Car 18
Lap 1, B, Car 18	Lap 3, B, Car 18
Lap 1, A, Car 18	...
Lap 1, B, Car 19	...
Lap 2, A, Car 18	Lap 300, A, Car 18
Lap 2, B, Car 18	Lap 300, B, Car 18
Lap 2, A, Car 19	Lap 300, A, Car 19
Lap 2, B, Car 19	Lap 300, B, Car 19

Hash: 2016-11-20

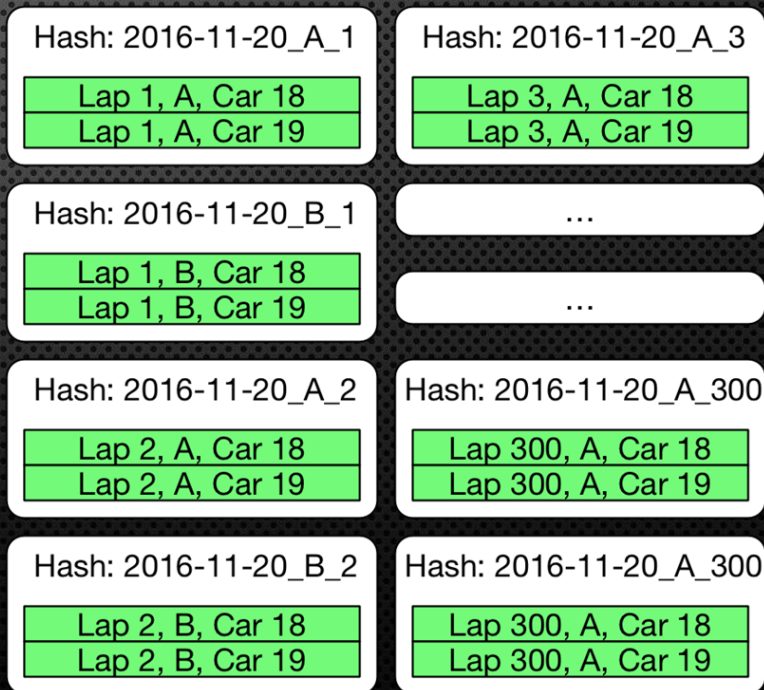
Lap 1, A, Car 18	Lap 3, A, Car 18
Lap 1, B, Car 18	Lap 3, B, Car 18
Lap 1, A, Car 18	...
Lap 1, B, Car 19	...
Lap 2, A, Car 18	Lap 300, A, Car 18
Lap 2, B, Car 18	Lap 300, B, Car 18
Lap 2, A, Car 19	Lap 300, A, Car 19
Lap 2, B, Car 19	Lap 300, B, Car 19

Time-series data: manual partitioning



Time-series data: composite keys

- Our solution to Hot Hashes
- Date, Type, and Lap as hash
- Data potentially in up to 10,000 partitions
- Can query in less than a second
- Could further hash by car number if needed
- Fetching data for entire race can be in parallel



DynamoDB streams

Real-time data feed

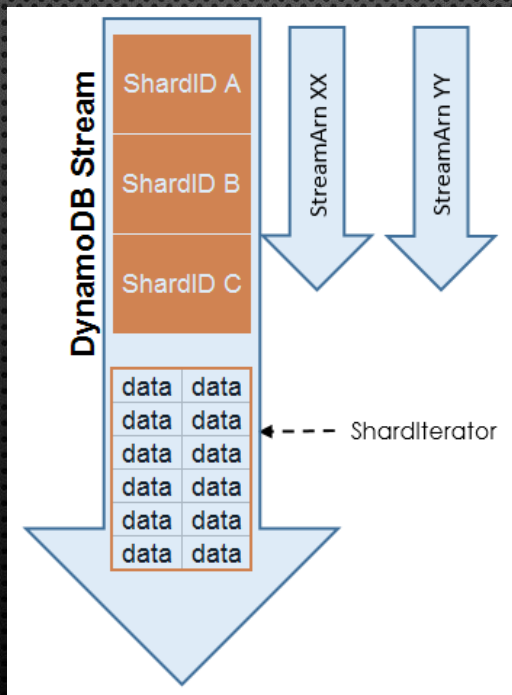
Data streaming

- The Arbiter is set up to monitor DynamoDB tables as a stream
- AWS API calls can be used to “tail” the stream for updates
- TRD has written a JavaScript wrapper to automate the handling of these calls and is shared across multiple services
- For native Java users, AWS has made available the Amazon Kinesis Client Library, which does the same thing

Data streaming

AWS API Calls:

1. ListStream
2. Describe Stream
3. GetShardIterator
4. GetRecords

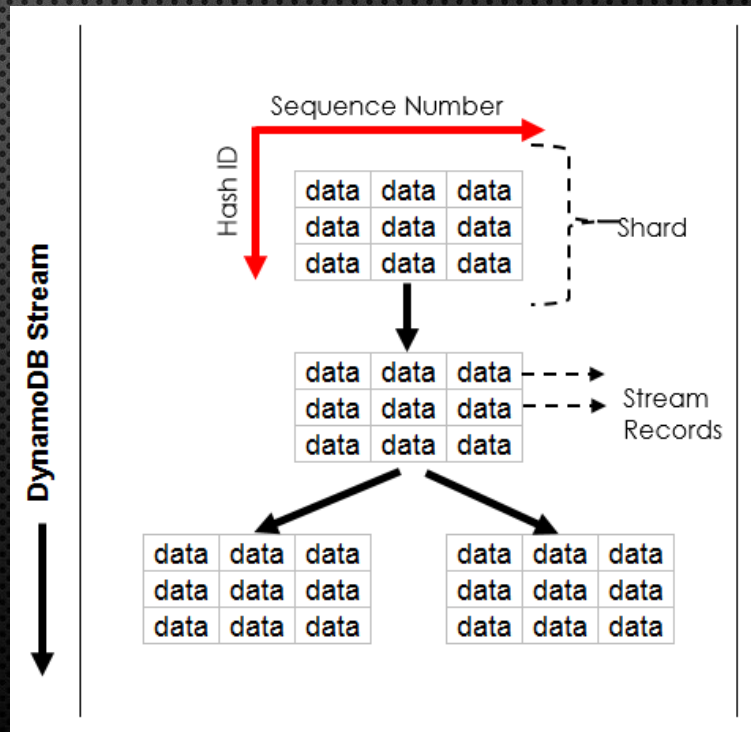


GetShardIterator
Options:

- AFTER/AT_SEQUENCE NUMBER
 - Recover from a system crash
- TRIM_HORIZON
 - Resume from the beginning of the race
- LATEST
 - Start of the race

Sharding

- Tables have partitions and streams have shards
- Shards are bounded by:
 - Min/max hash id
 - Min/max sequence number
- New shards will be created when:
 - Boundaries of the current shard have been exceeded
 - Increased throughput requirements



Sharding

Problems:

- When streaming, we had to actively monitor shard status to determine when it expired and to scraper for new active iterators to stream data continuously

Gotchas:

- Sharding is not instantaneous, hence new shards aren't available immediately upon the active shard expiring
- Need to constantly hunt for new active shards

Testing DynamoDB streams

- Can be challenging for functional testing because we cannot control when a stream re-shards
- API is the same as Kinesis. Can use Kinesis to simulate streams and force re-shards during functional tests
- For node.js we can use Kinesalite for functional testing
- Our most devastating bugs were due to not handling stream re-shards correctly

DynamoDB streams vs. Kinesis streams

	DynamoDB Stream	Kinesis Stream
Persistence	Stream: 24 hours Table: As long as required	24 Hours – 7 days
Sharding	Automatic	Manual
Querying	Stream: As series Table: Ad Hoc	As series
Stream Latency*	~ 1 second	~200-400 ms
Billing	Per Read	Per Shard & Write

*Unofficial. Derived from basic benchmarking, and could change

DynamoDB + Firehose

The perfect match

Turbo mode

Scenario

- Change of prediction model mid-race
- Need to re-process entire data set as fast as possible, then resume from stream where we left off

Problem

- Read limits on DynamoDB streams

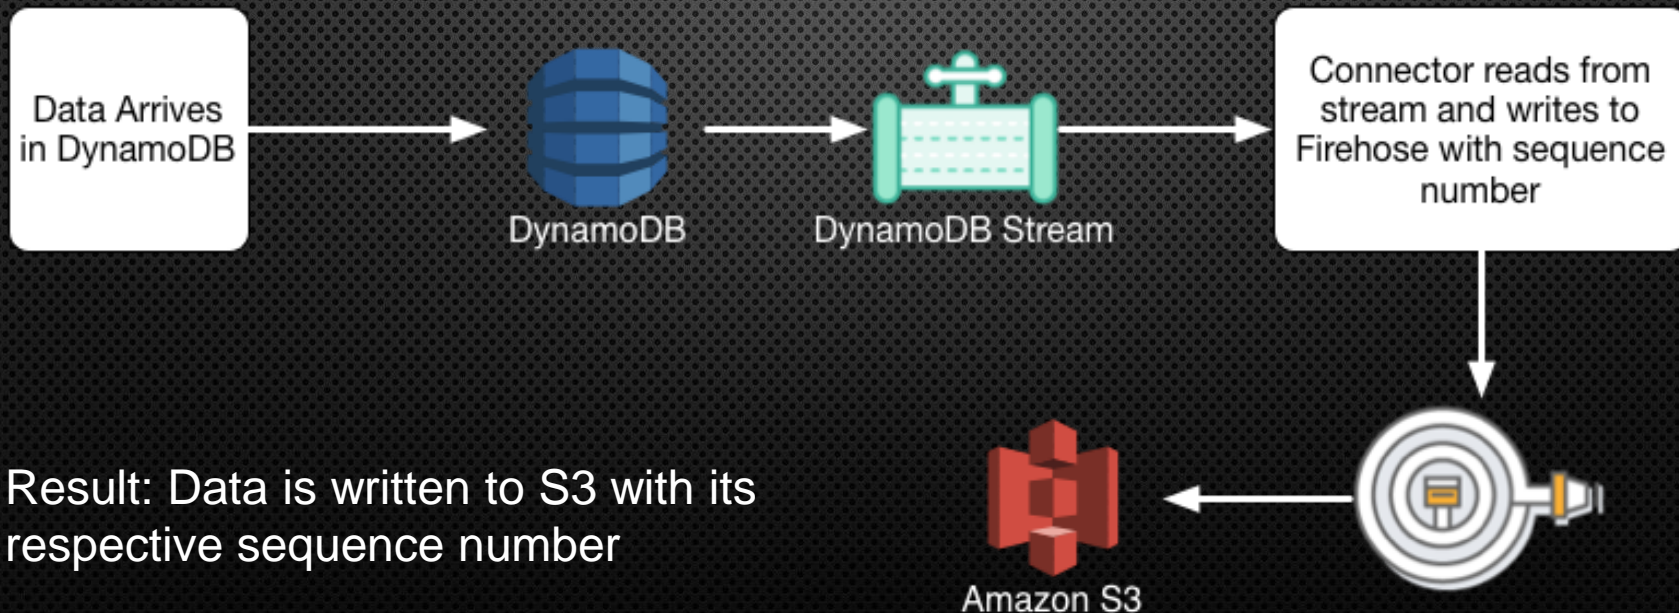
Solution

- Read from S3, then resume from stream

Prerequisite

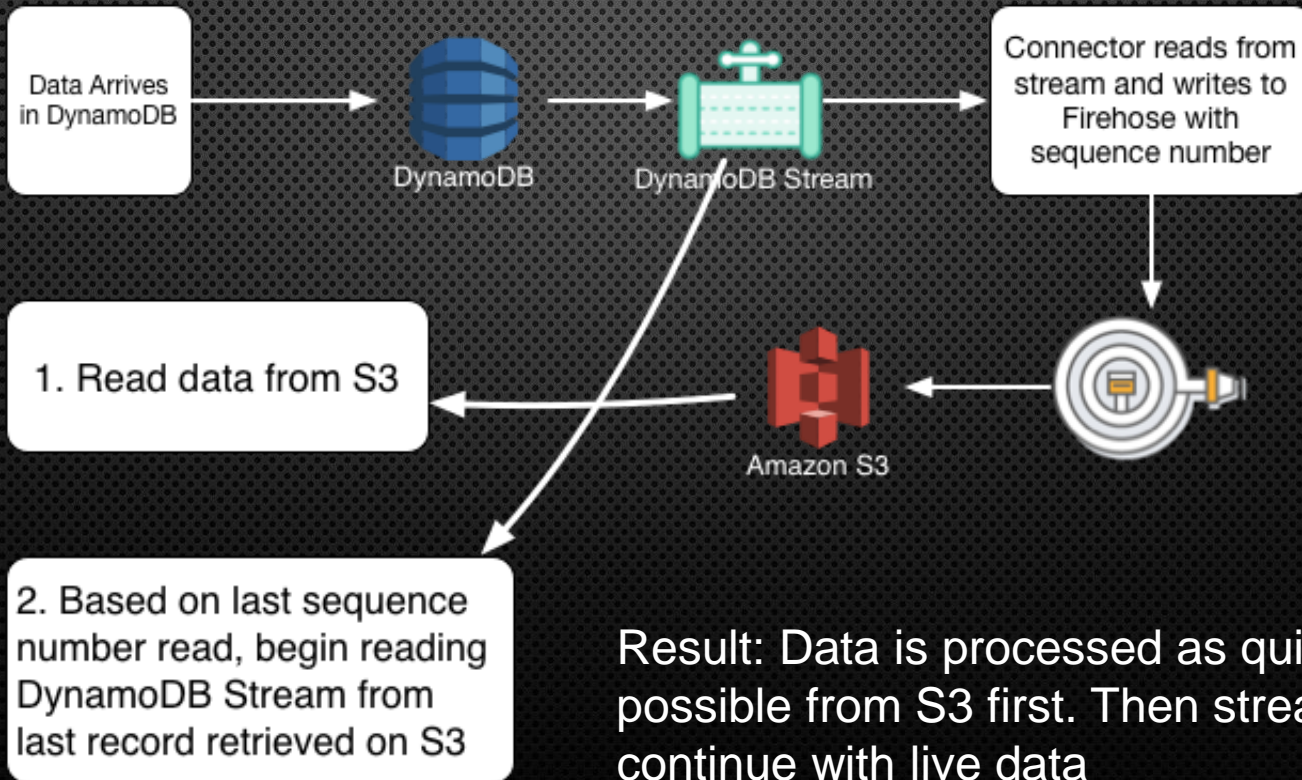
- Data on S3 must contain sequence number for each record for use in resumption of stream

Capture sequence number



Result: Data is written to S3 with its respective sequence number

Turbo read




In action

DynamoDB to Firehose

In action

Athena demo

The screenshot displays the 'Race Positions' interface. At the top, a status bar shows 'Lap: 67', 'Leader: 20', and 'Leader Laptime: 9.830'. The main area is a top-down view of an oval track with a green center. Numerous cars, represented by numbered circles, are positioned around the track. Car 20 is the leader, highlighted in red. Other cars are in various colors (blue, orange, yellow, green, purple). A 'Leader' label is placed in the center of the track. A 'Pitted/Off-Track' box contains cars 1 and 43. A graph at the bottom shows lap times for several cars, with a vertical axis ranging from 32.0 to 32.5. The graph has a 'Lap Time' header and a 'Cumulative Delta' header. The interface includes a search icon and a settings icon in the top right corner.



AWS
re:Invent

Thank you!



**Remember to complete
your evaluations!**