ABD315

Serverless ETL with AWS Glue

Mehul A. Shah

Software Manager, AWS Glue

November 27, 2017

AWS re:Invent

aws

Today's Agenda

Intro to AWS Glue

Construct an ETL flow in 4 steps

Under the hood: customize AWS Glue scripts

Merck – customer testimonial

What is AWS Glue?

Fully-managed, serverless extract-transform-load (ETL) service

for developers, built by developers

1000s of customers and jobs

Select AWS Glue customers

There are many tools already

Amazon Redshift Partner Page for Data Integration

Still, ETL developers hand-code

Canvas-based tools are hard to extend

Code is flexible, powerful, and easy to share

Familiar tools and development pipelines

IDEs, version control, testing, continuous integration

This talk is for developers!

Hand-coding is laborious

schemas change

data formats change

add or change sources

data volume grows

makes hand-coding
error-prone & brittle

AWS Glue does the undifferentiated heavy lifting
so developers can easily customize

AWS Glue Components

Data Catalog

Discover

- Automatic crawling
- Apache Hive Metastore compatible
- Integrated with AWS analytic services

Job Authoring

Develop

- Auto-generates ETL code
- Python and Apache Spark
- Edit, Debug, and Explore

Job Execution

Deploy

- Serverless execution
- Flexible scheduling
- Monitoring and alerting

Common use-cases

Load data warehouses

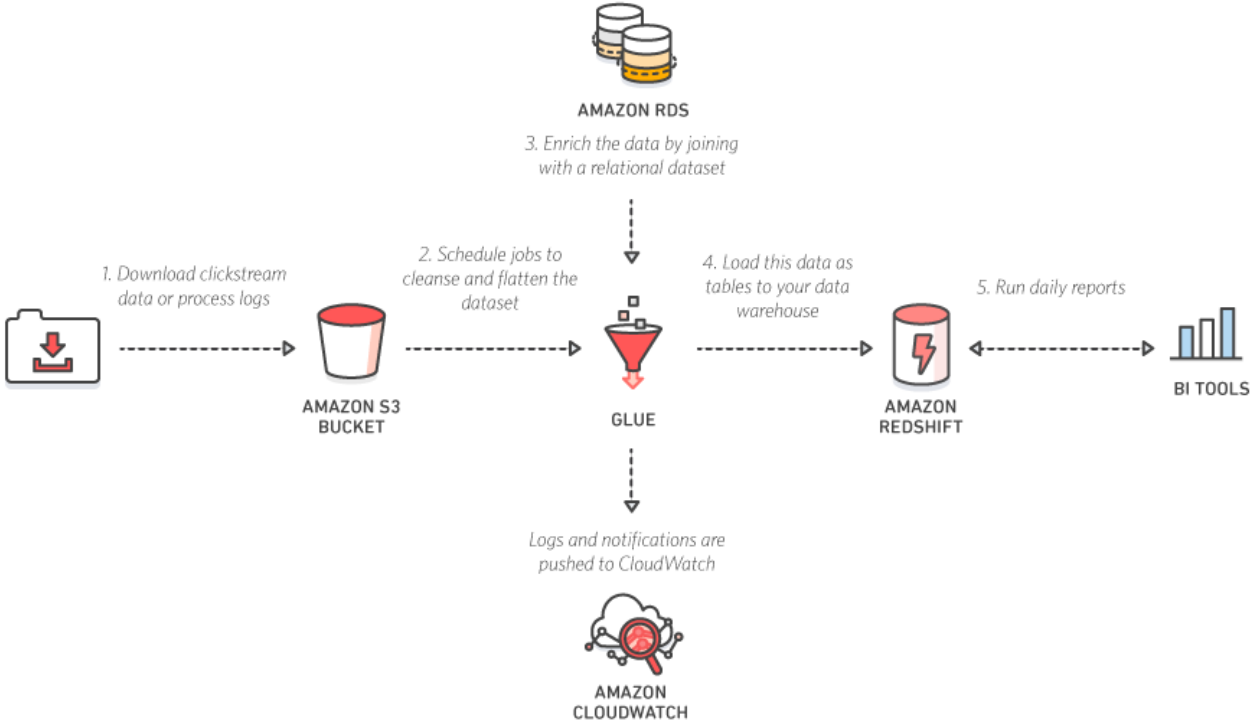

Build a data lake on Amazon S3

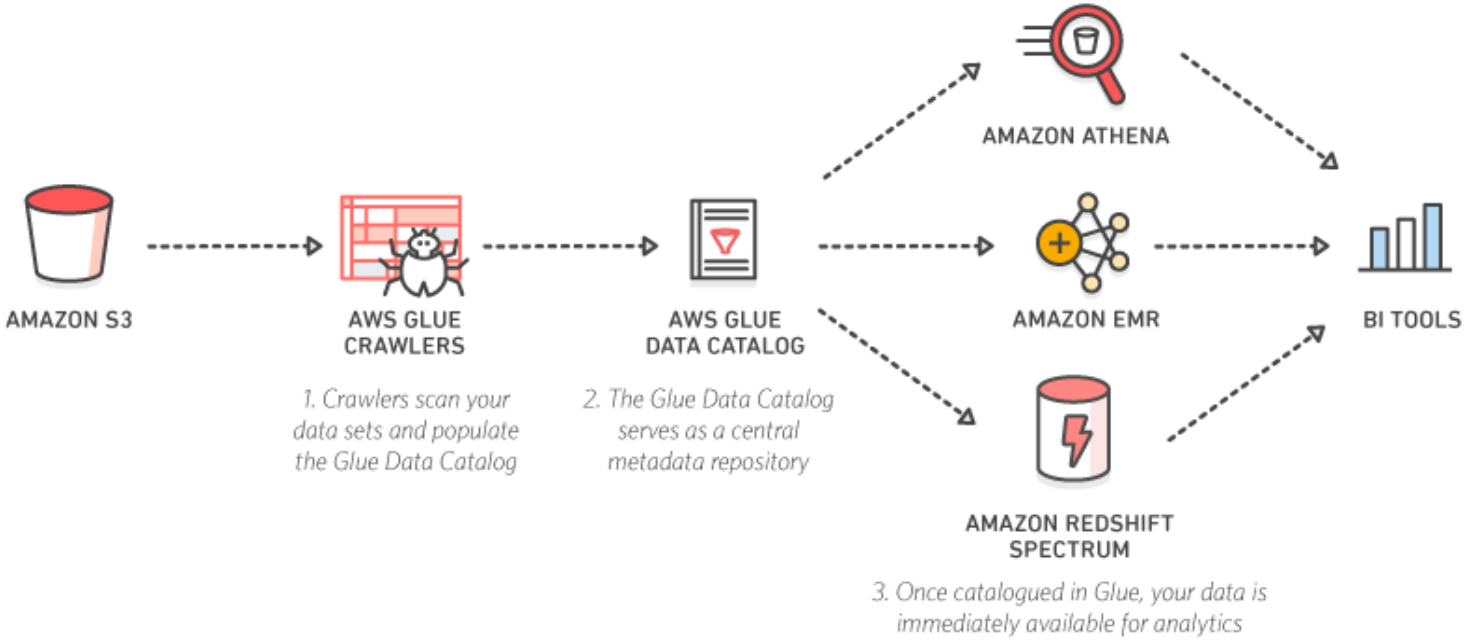

Construct an ETL flow in 4 steps

The 4 Steps

1. Crawl and catalogue your data
2. Specify mappings to generate scripts
3. Interactively edit and explore with dev-endpoints
4. Schedule a job for running in production

ETL example

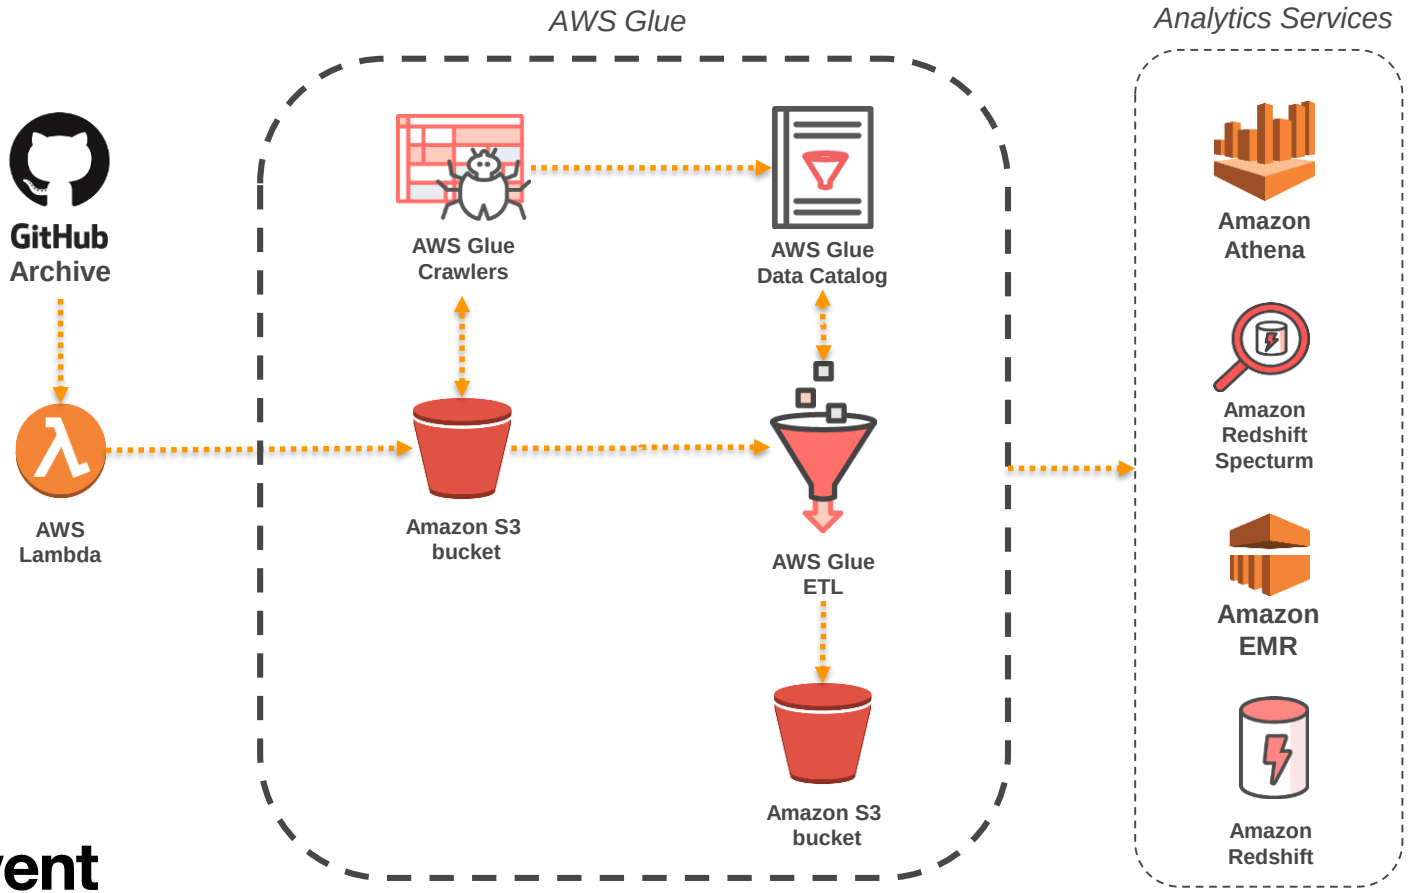

ETL example (con't)

Organize data in Apache Hive-style partitions

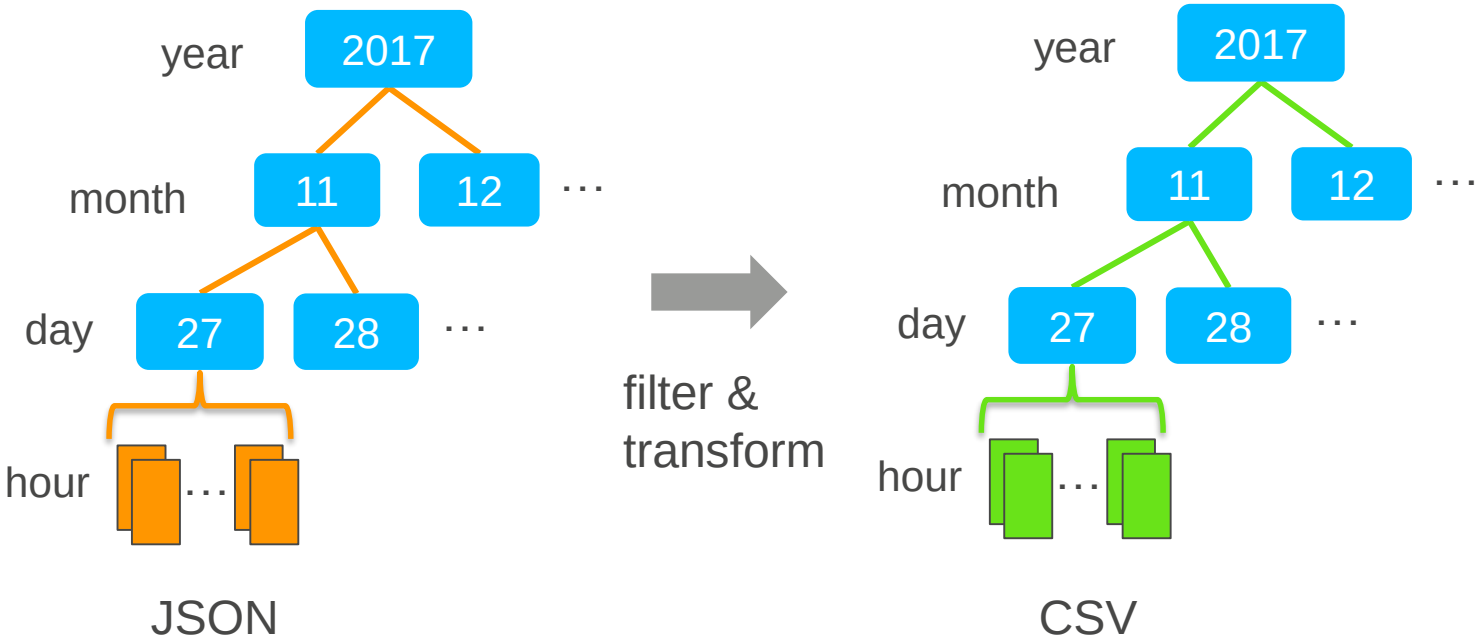

Public GitHub timeline

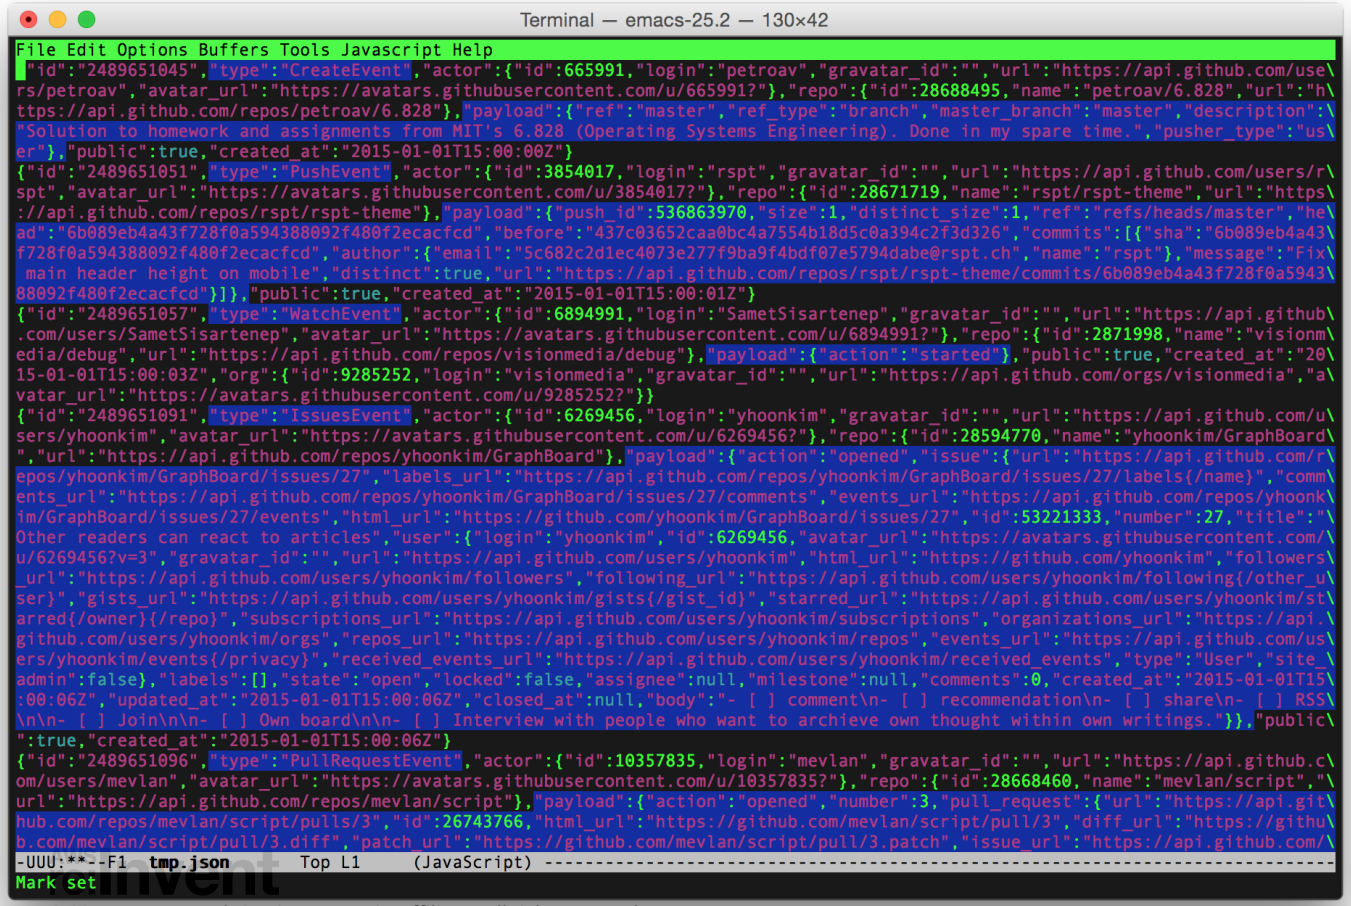

githubarchive.org

35+ event types

unique payload per event type

Step 1: Run crawler

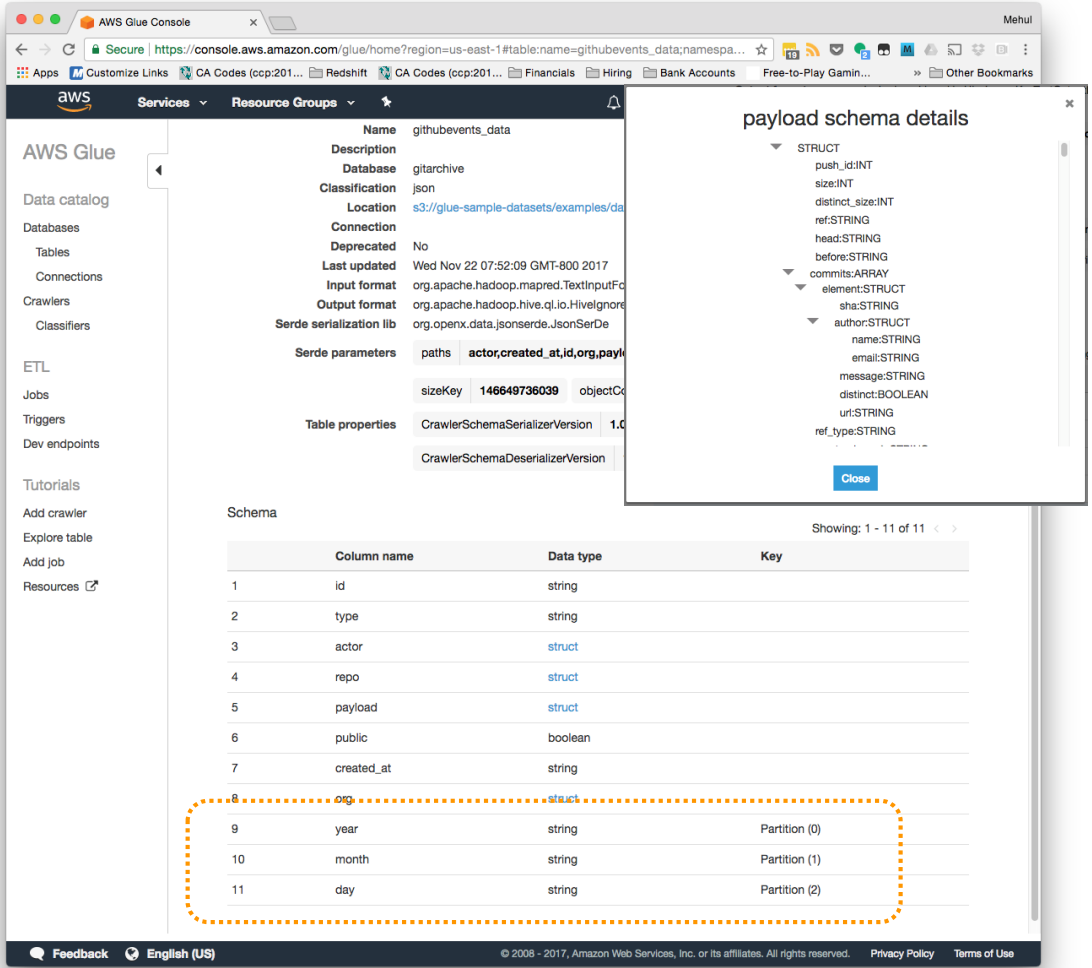

200+ fields

Groups files into
Apache Hive-style partitions

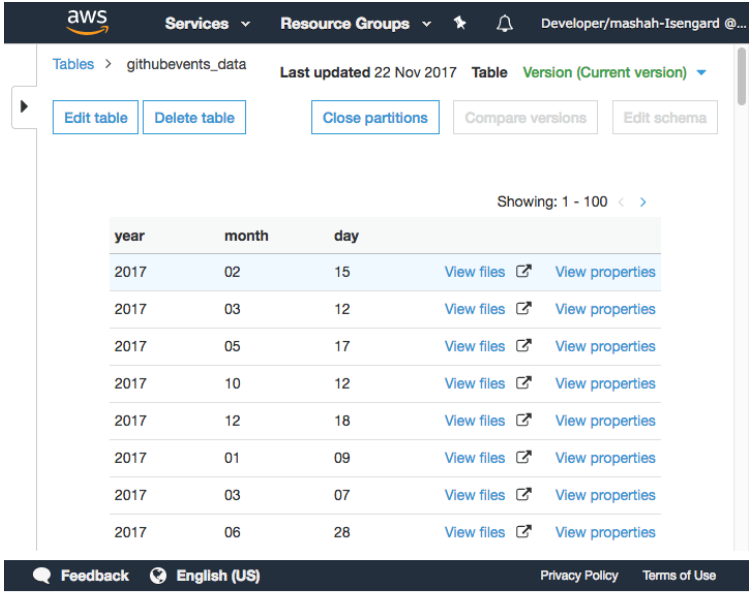

Step 2: Specify mappings

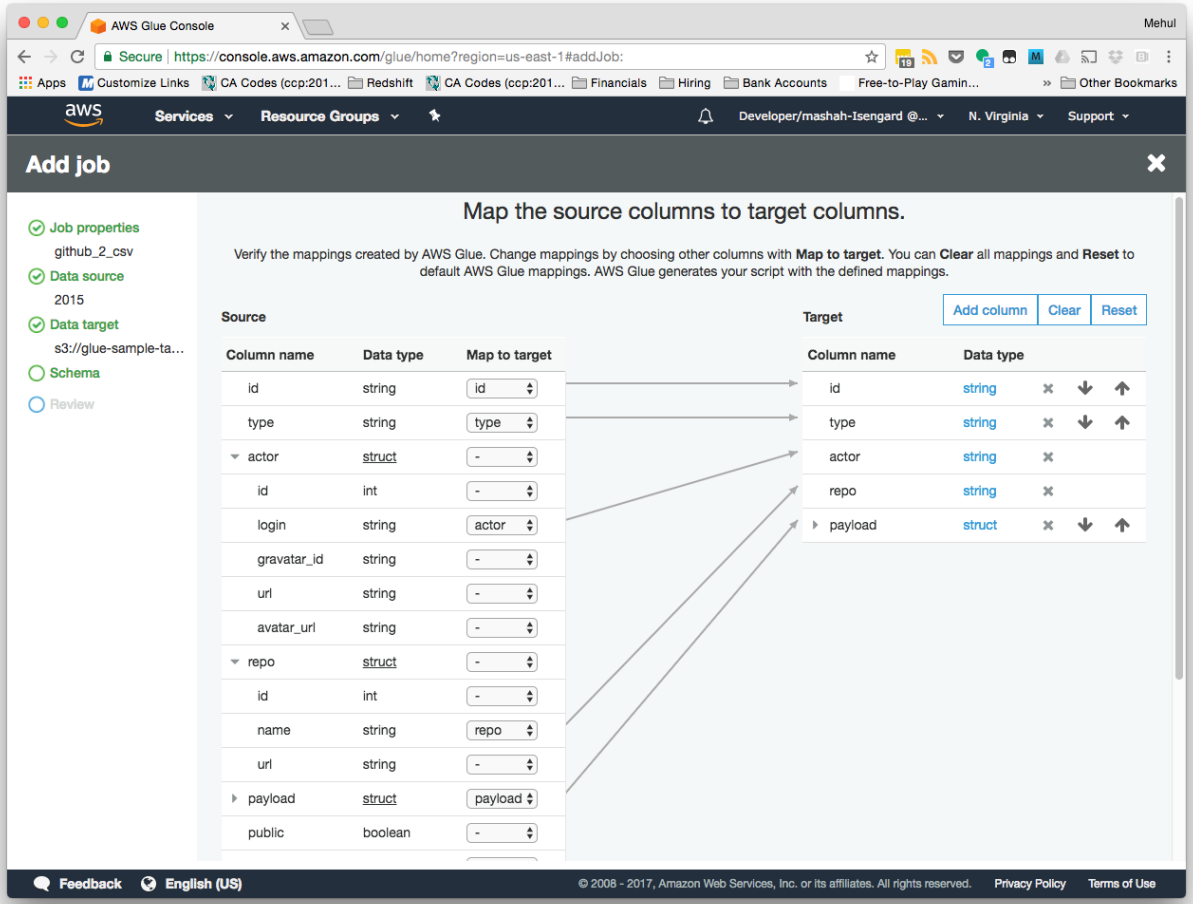

Anatomy of a generated script

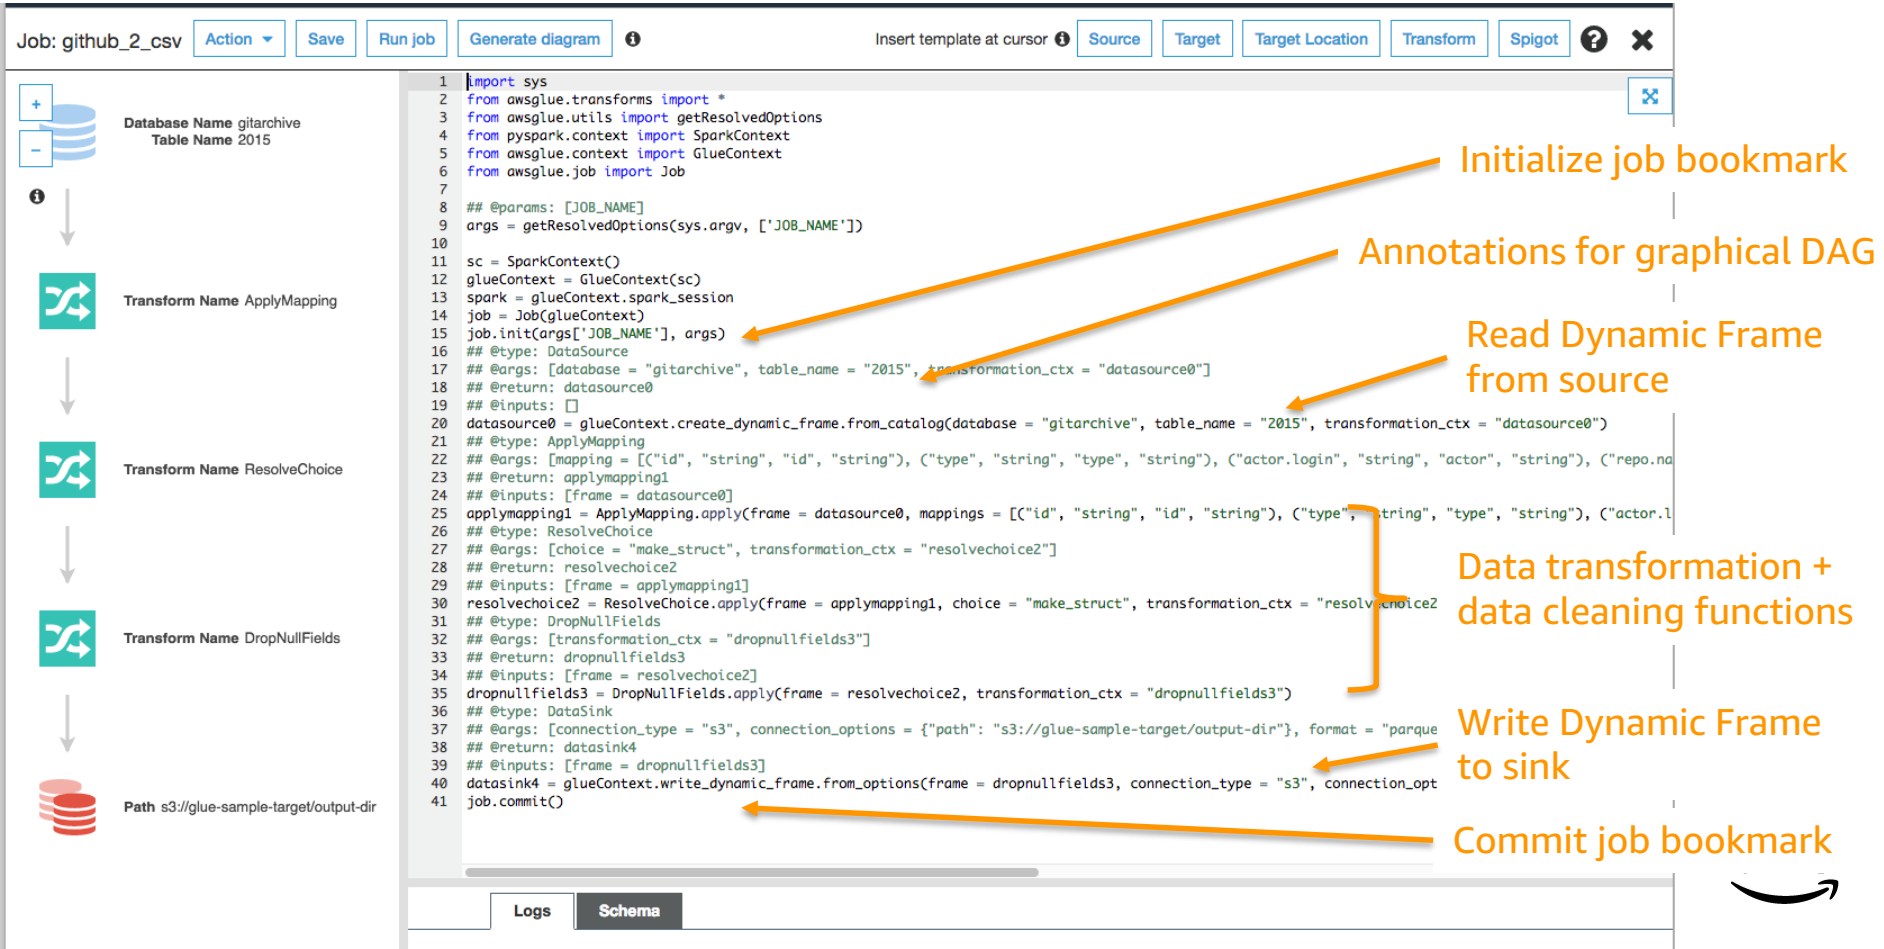

Step 3: Edit + Test with Dev-Endpoints

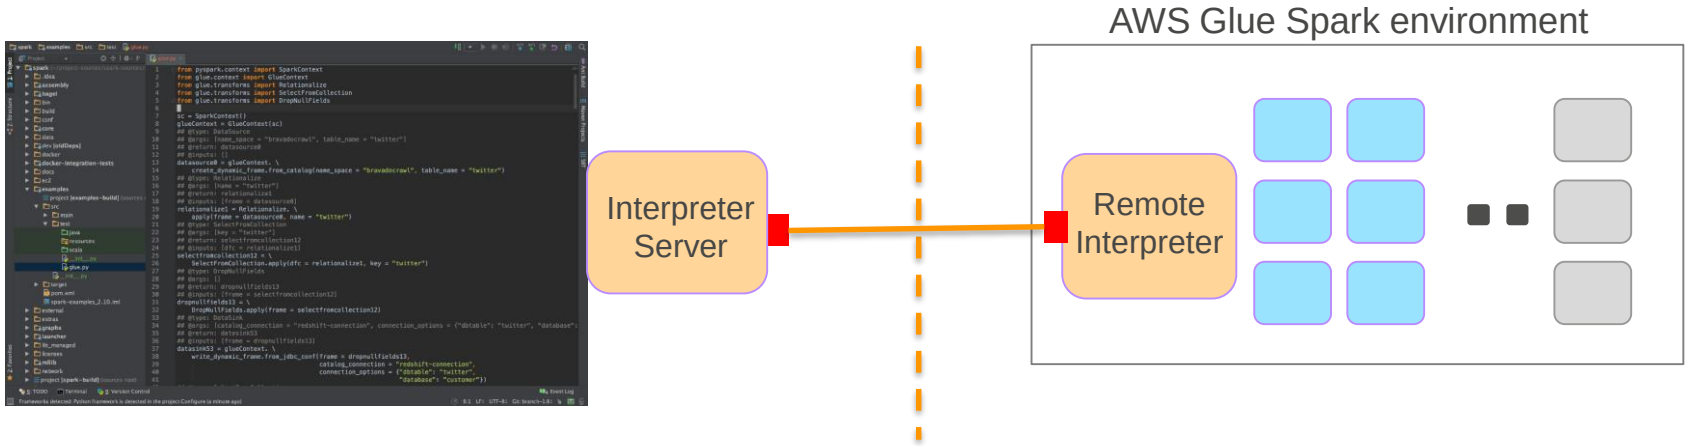

Connect your IDE to an AWS Glue development endpoint.

Environment to **interactively develop,** debug, and test ETL code.

Step 3: Explore and experiment with data

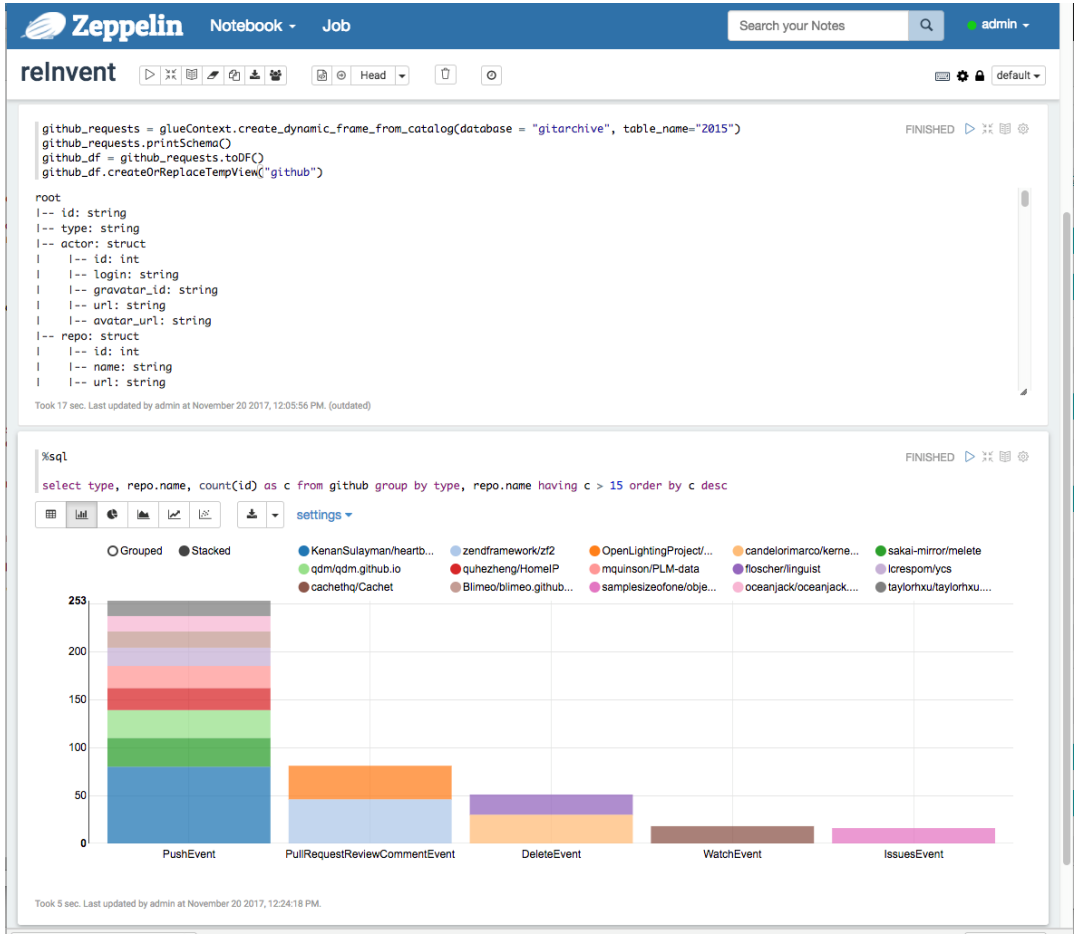

Connect your notebook (e.g. **Zeppelin**) to an AWS Glue development endpoint.

Interactively experiment and **explore** datasets and data sources

Deploy to production

- Push scripts to S3
- Create or register with ETL job

aws

Step 4: Schedule a job

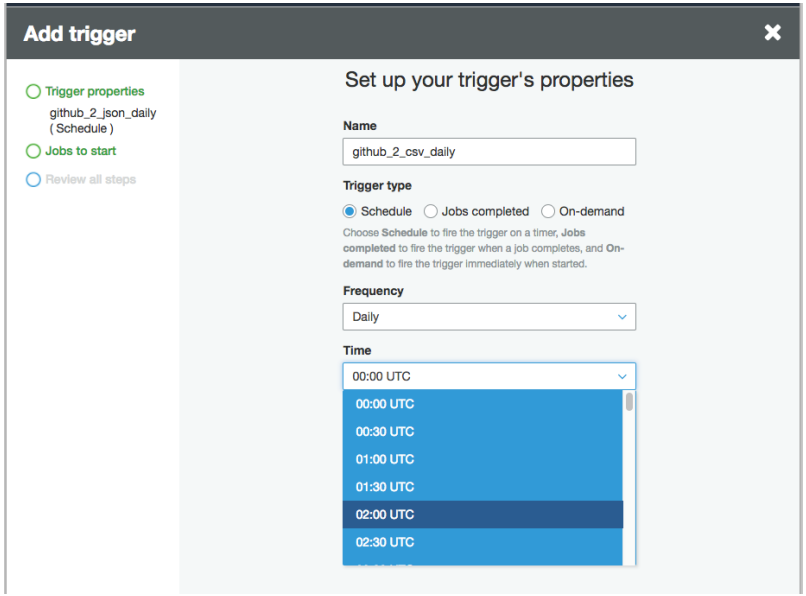

several event types

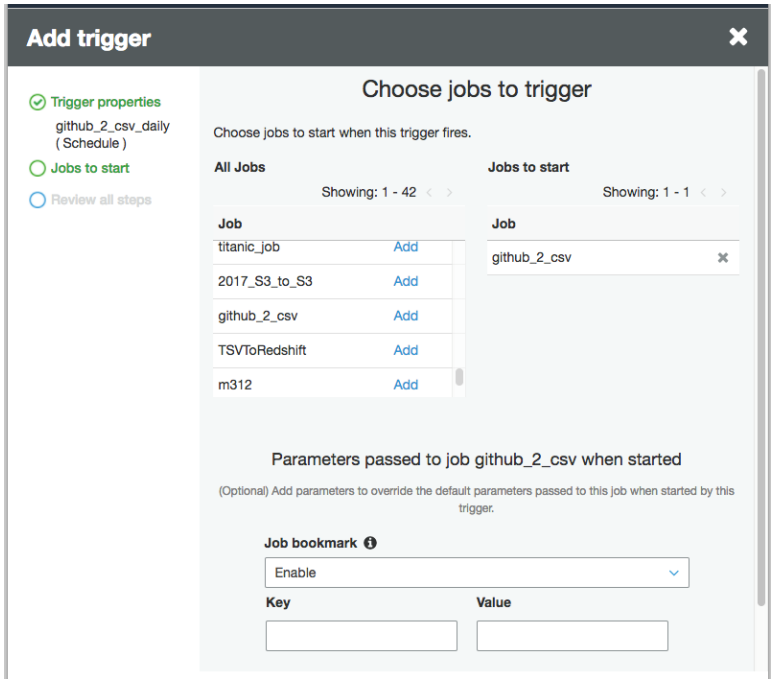

pass parameters

Serverless job execution

No need to provision, configure, or manage servers

Auto-configure VPC & role-based access
security & isolation preserved

Customers can specify job capacity (DPU)

Automatically scale resources

Only pay for the resources you consume
per-second billing (10-minute min)

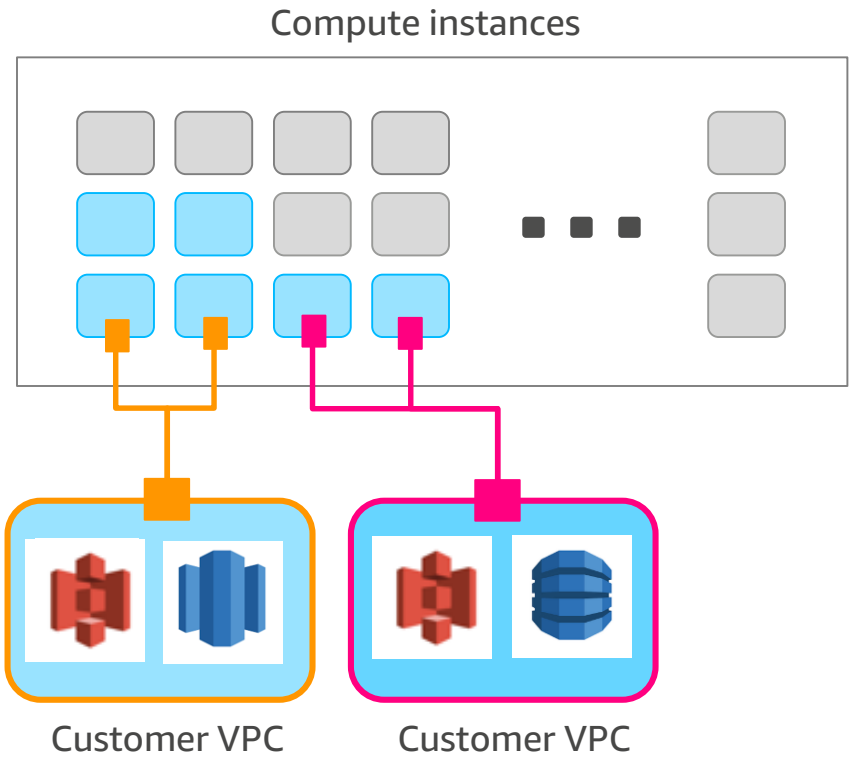

Under the hood: customize AWS Glue scripts

Apache Spark and AWS Glue ETL

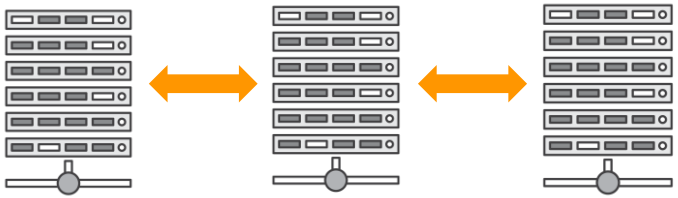

What is Apache Spark?

- Parallel, scale-out data processing engine
- Fault-tolerance built-in
- Flexible interface: Python scripting, SQL
- Rich eco-system: ML, Graph, analytics, …

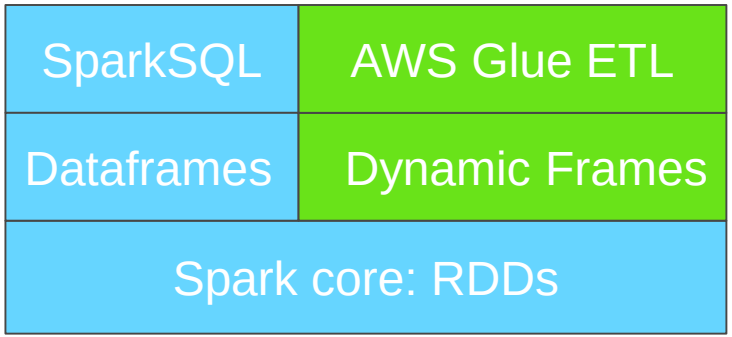

AWS Glue ETL libraries

- Integration: Data Catalog, job orchestration, code-generation, job bookmarks, S3, RDS
- ETL transforms, more connectors & formats
- New data structure: Dynamic Frames

Dataframes and Dynamic Frames

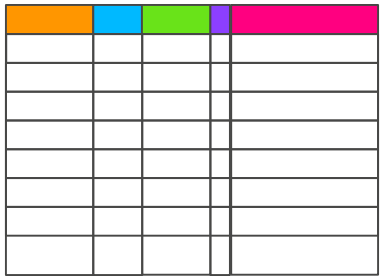

Dataframes

- Core data structure for SparkSQL
- Like structured tables
  - Need schema up-front
  - Each row has same structure
- Suited for SQL-like analytics

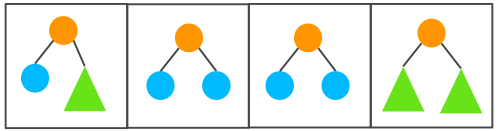

Dynamic Frames

- Like dataframes for ETL
- Designed for processing **semi-structured** data, e.g. JSON, Avro, Apache logs ...

Public GitHub timeline is …

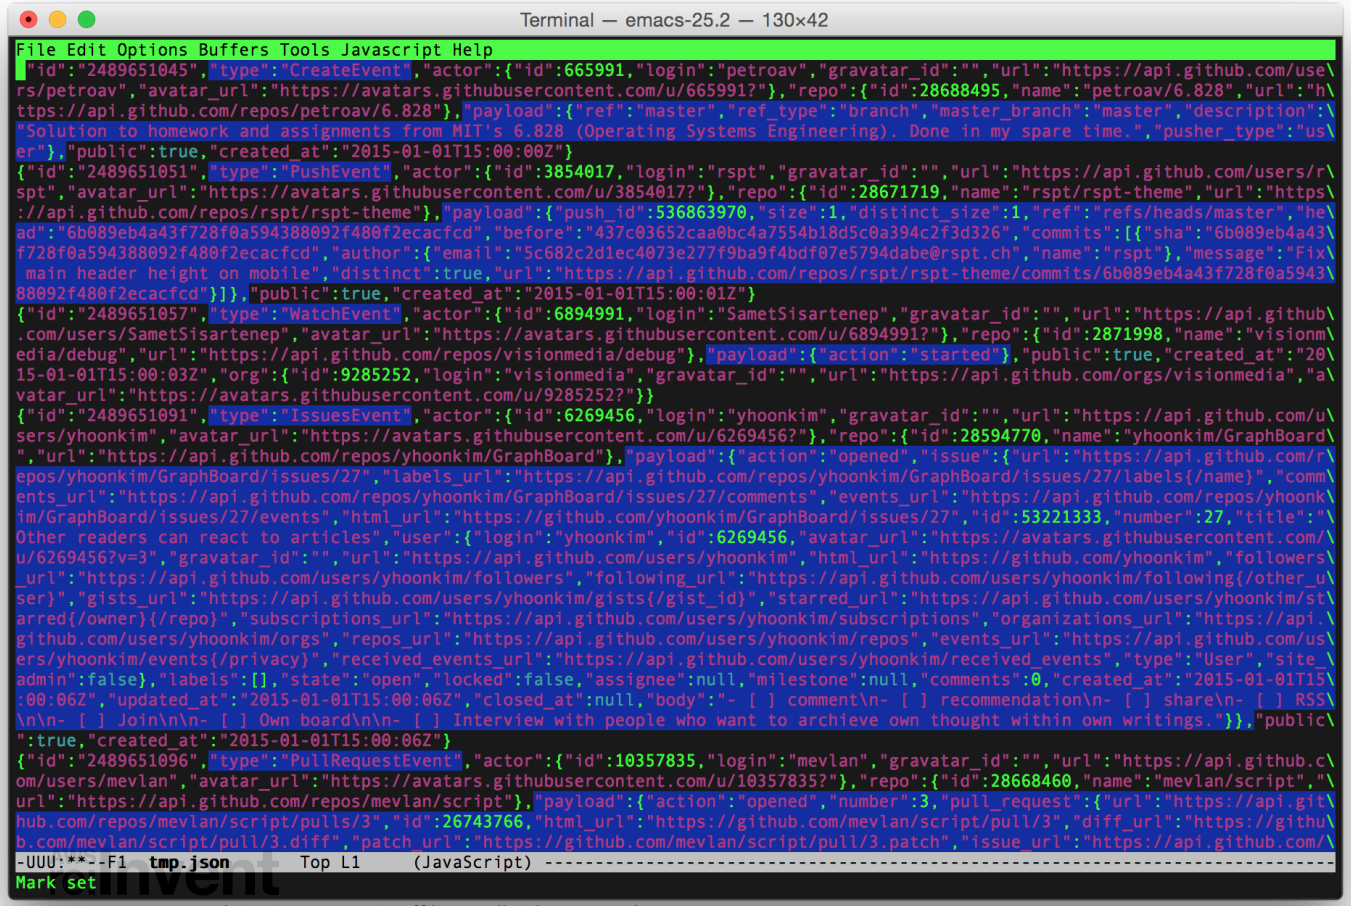

semi-structured

35+ event types

payload structure and size varies by event type

Dynamic Frame internals

Dynamic Records

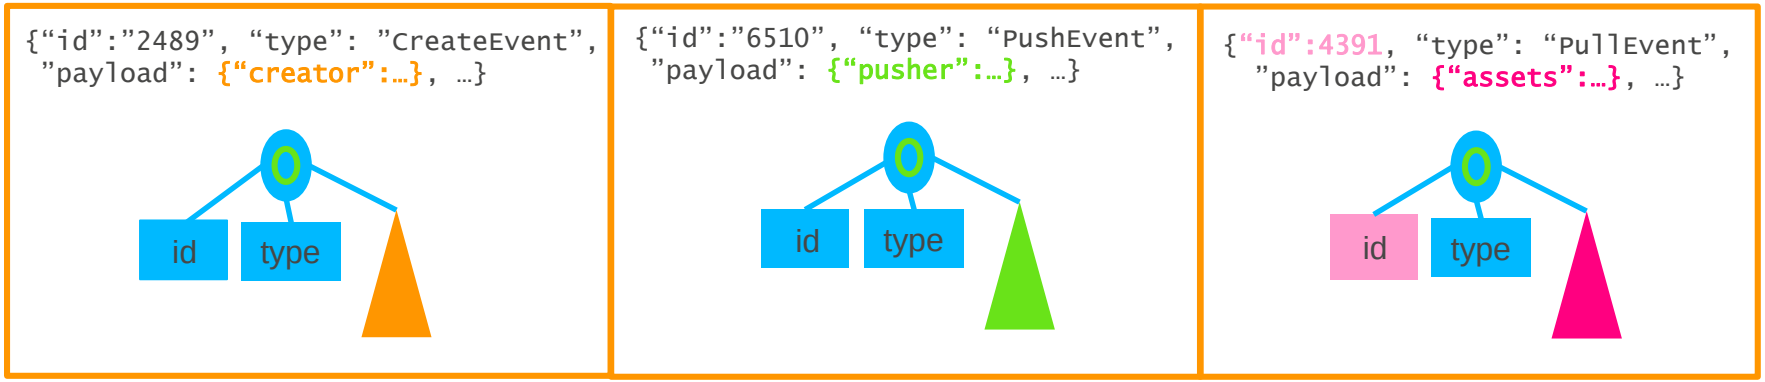

Dynamic Frame Schema

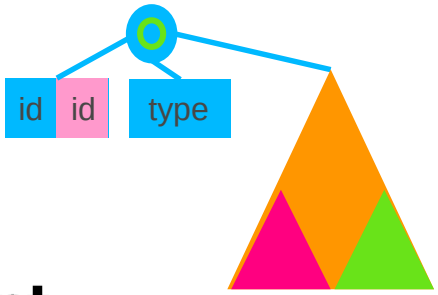

schema per-record, no up-front schema needed

Easy to restructure, tag, modify

Can be more compact than dataframe rows

Many flows can be done in single-pass

Dynamic Frame transforms

15+ transforms out-of-the box

ResolveChoice()

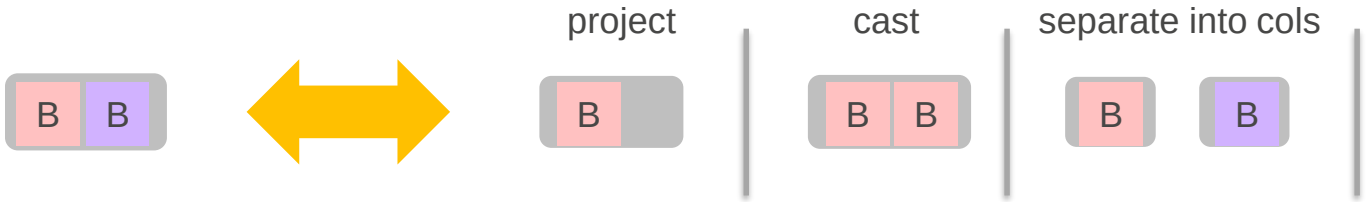

ApplyMapping()

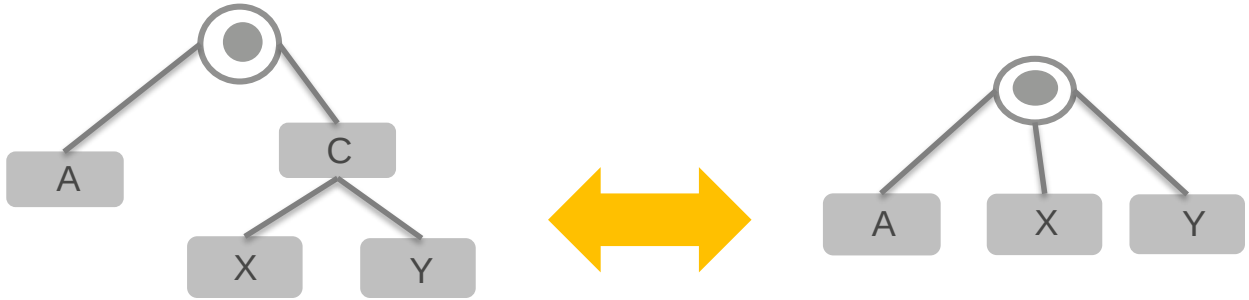

Relationalize() transform

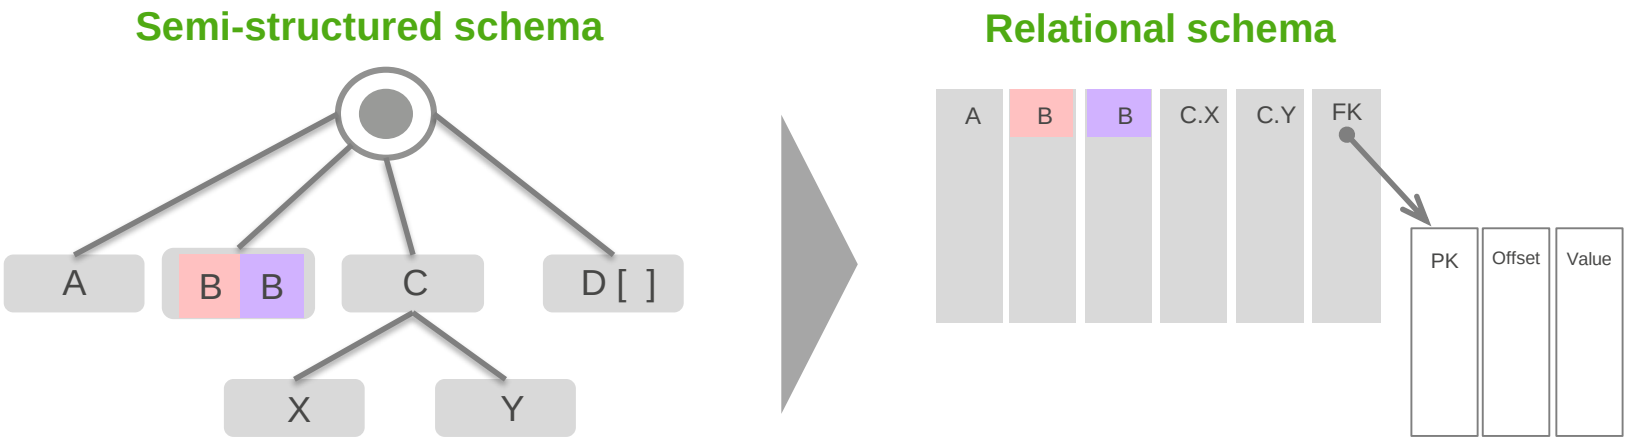

Transforms and **adds new** columns, types, and tables on-the-fly

Tracks **keys** and **foreign keys** across runs

SQL on the relational schema is orders of **magnitude faster** than JSON processing

Useful AWS Glue transforms

toDF(): Convert to a Dataframe

Spigot(): Sample data of any Dynamic Frame to S3

Unbox(): Parse string column as given format into Dynamic Frame

Filter(), Map(): Apply Python UDFs to Dynamic Frames

Join(): Join two Dynamic Frames

And more ….

Performance: AWS Glue ETL

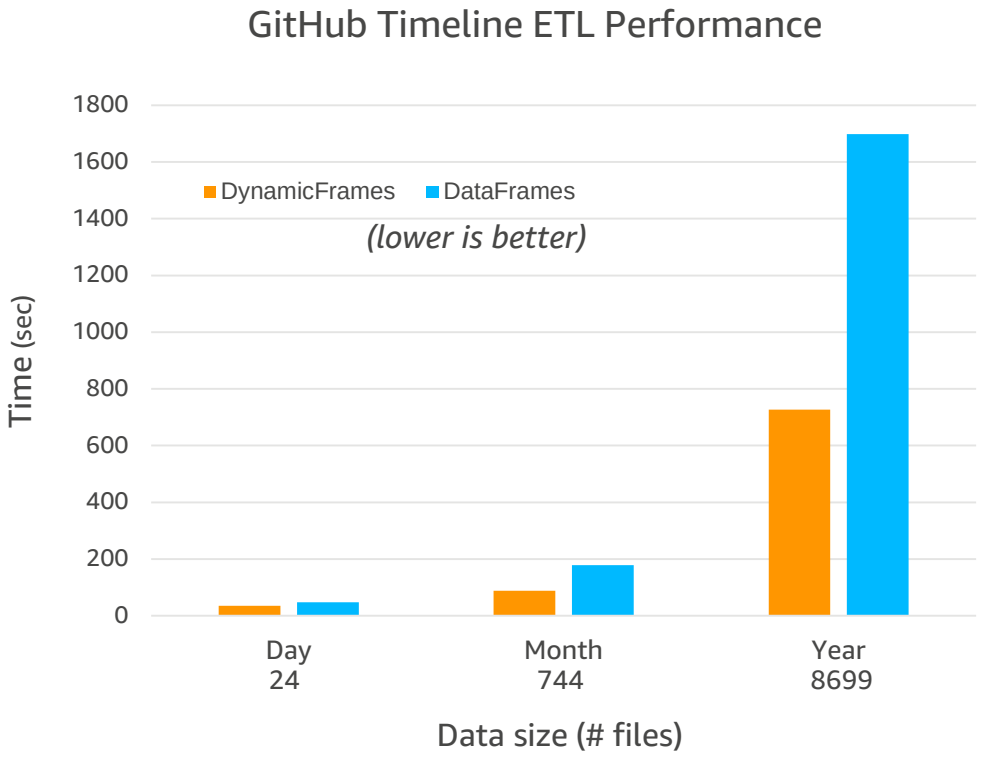

Configuration

- 10 DPUs
- Apache Spark 2.1.1

Workload

- JSON to CSV
- Filter for Pull events

On average: 2x performance improvement

Performance: Lots of small files

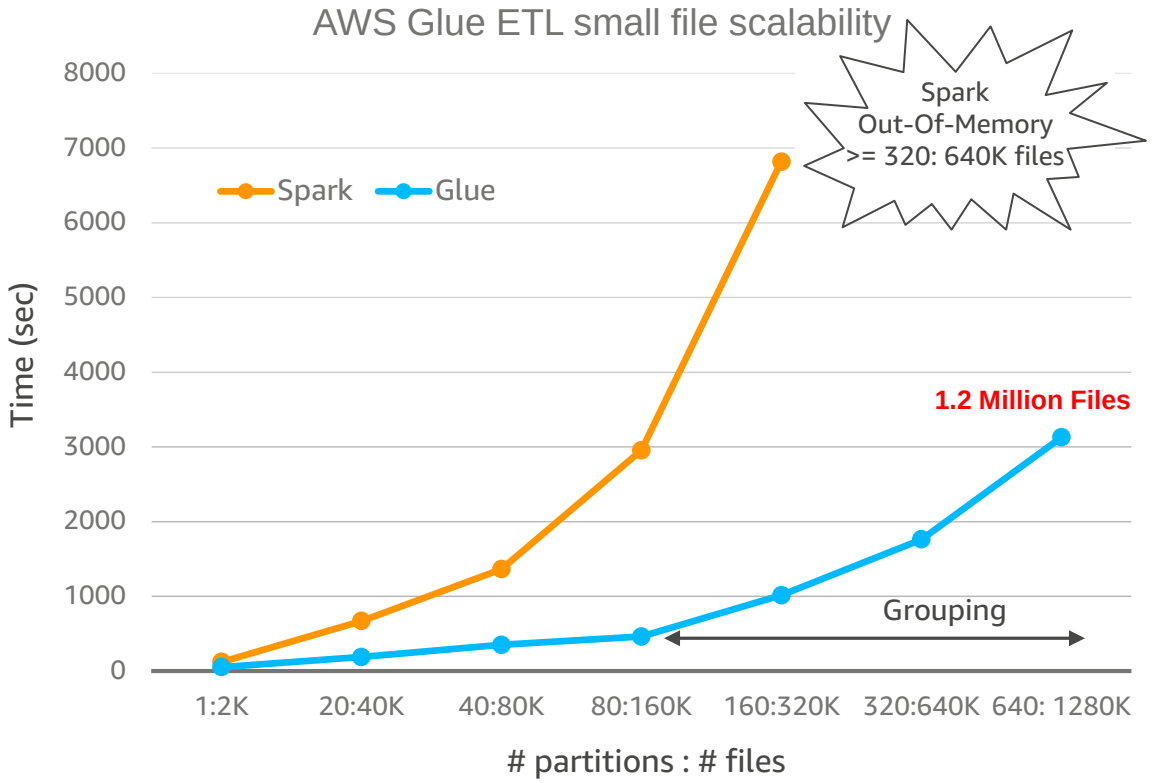

Lots of small files, e.g. Kinesis Firehose

Vanilla Apache Spark (2.1.1) overheads

- Must reconstruct partitions (2-pass)
- Too many tasks: task per file
- Scheduling & memory overheads

AWS Glue Dynamic Frames

- Integration with Data Catalog
- Automatically group files per task
- Rely on crawler statistics

Job bookmark example

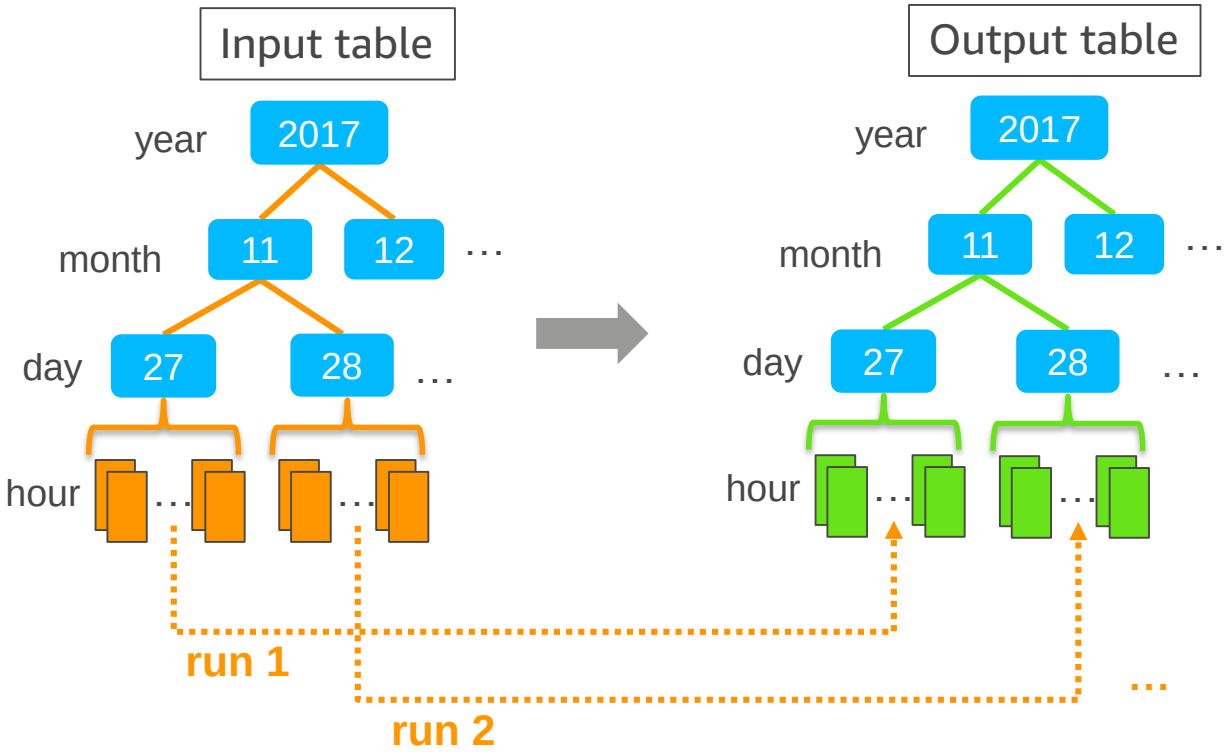

Periodically run a job

avoid reprocessing previous input

avoid generating duplicate output

Job bookmarks

Bookmarks are per-job checkpoints that track the work done in previous runs.

They persist the state of sources, transforms, and sinks on each run.

Examples uses:

Process githubarchive files daily

Process Firehose files hourly

Track timestamps or primary keys in DBs

Track generated foreign keys for normalization

Job bookmark options

Option	Behavior
Enable	Pick up from where you left off
Disable	Ignore and process the entire dataset every time
Pause	Temporarily disable advancing the bookmark

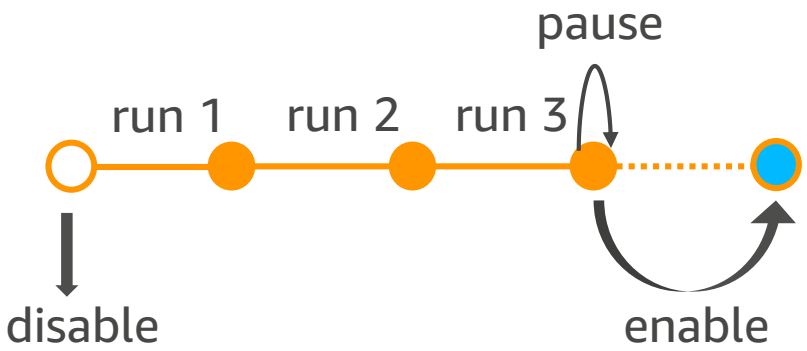

Examples:

Enable: Process the newest githubarchive partition

Disable: Process the entire githubarchive table

Pause: Process the previous githubarchive partition

Job bookmark internals

How do we avoid space blowup?

Use timestamps to filter already processed input

But S3 is eventually consistent?

Maintain exclusion list of files created in *inconsistency window* (size *d*) prior to start.

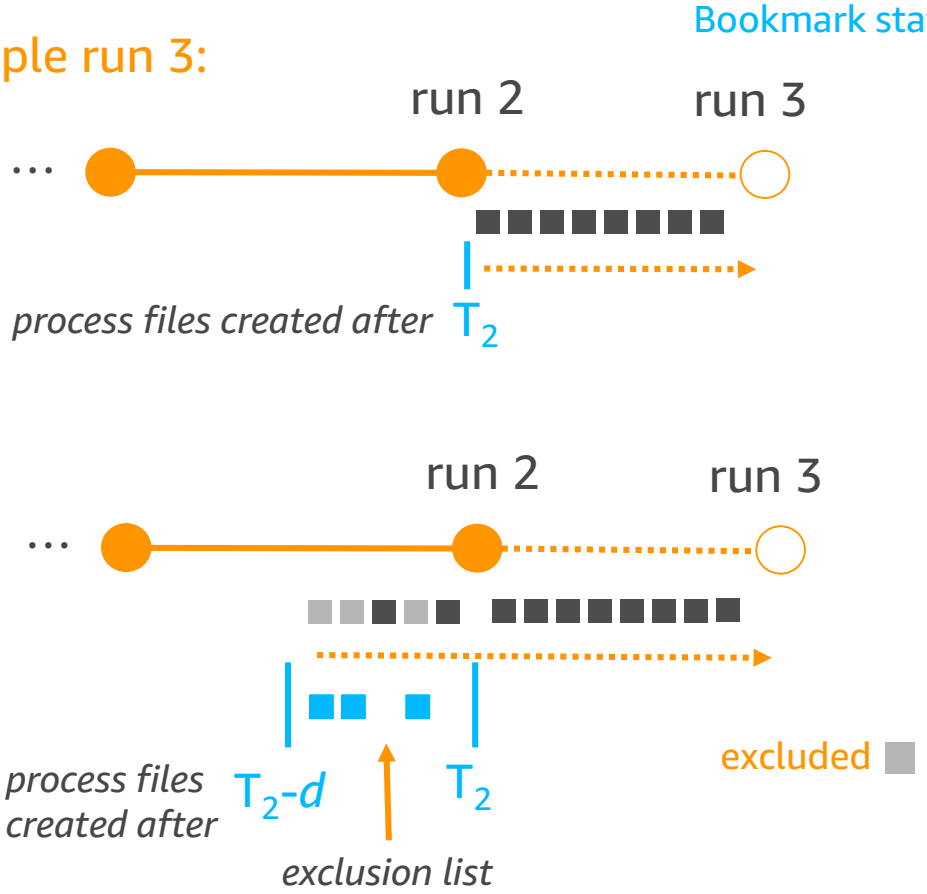

Wrap Up

4 steps to build a production ETL flow

AWS Glue features

- Dynamic frames
- Job bookmarks

AWS Glue Announcements

Scala support

New regions: Asia Pacific (Tokyo) & EU (Ireland)

Merck – customer testimonial

Keith Smola

Global Operations Management, Merck & Co.

AWS re:Invent

aws

FOR MORE THAN A CENTURY, MERCK HAS BEEN INVENTING

TO SOLVE SOME OF THE GREATEST CHALLENGES TO PEOPLE'S HEALTH AND WELL-BEING AROUND THE WORLD.

HEADQUARTERS
Kenilworth, NJ, U.S.A. operating in more than 140 countries

MRK

Merck & Co., Inc. is our legal name and is listed on the New York Stock Exchange under the symbol "MRK."

68,000

EMPLOYEES
approximately 68,000 worldwide (as of 12/31/16)

BUSINESSES
Prescription medicines, Vaccines, Biologic therapies, Animal Health products

2016 REVENUES
$39.8 billion, 54% of sales come from outside the United States

2016 R&D EXPENSE
$10.1 billion; 24 product pipeline programs in late-stage development

MERCK

Problem

We were challenged to quickly align extended data and metrics to a Enterprise Software Delivery Management system. One capability of the system is Environment Management, which is used to provide the knowledge regarding environments, what assets make them up and how they align to different software lifecycle efforts. Other capabilities include deployment management, delivery planning and scope management.

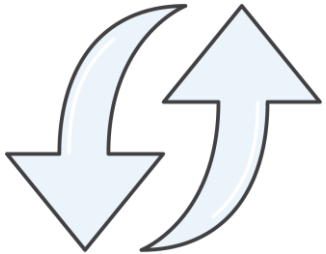

Limited integration with the Enterprise Software Delivery Management System

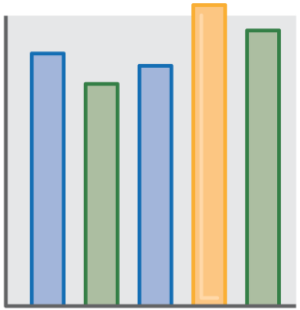

BI tools require a data source with application lifecycle context

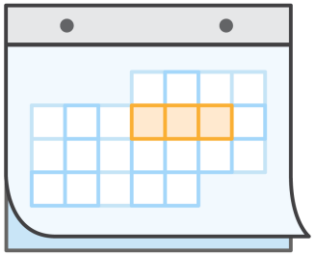

Major projects require fast alignment to Enterprise Software Delivery Management System

Challenges

Short timeline to produce a data layer.

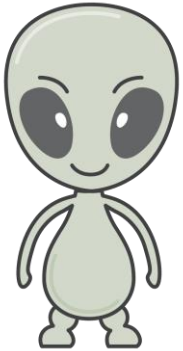

Resources are not data scientists or ETL developers.

Established support does not go beyond AWS.

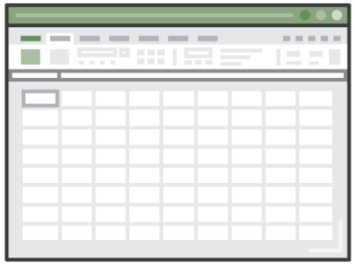

Data is in different spreadsheets or existing databases

Catalog vs. warehousing

Architecture

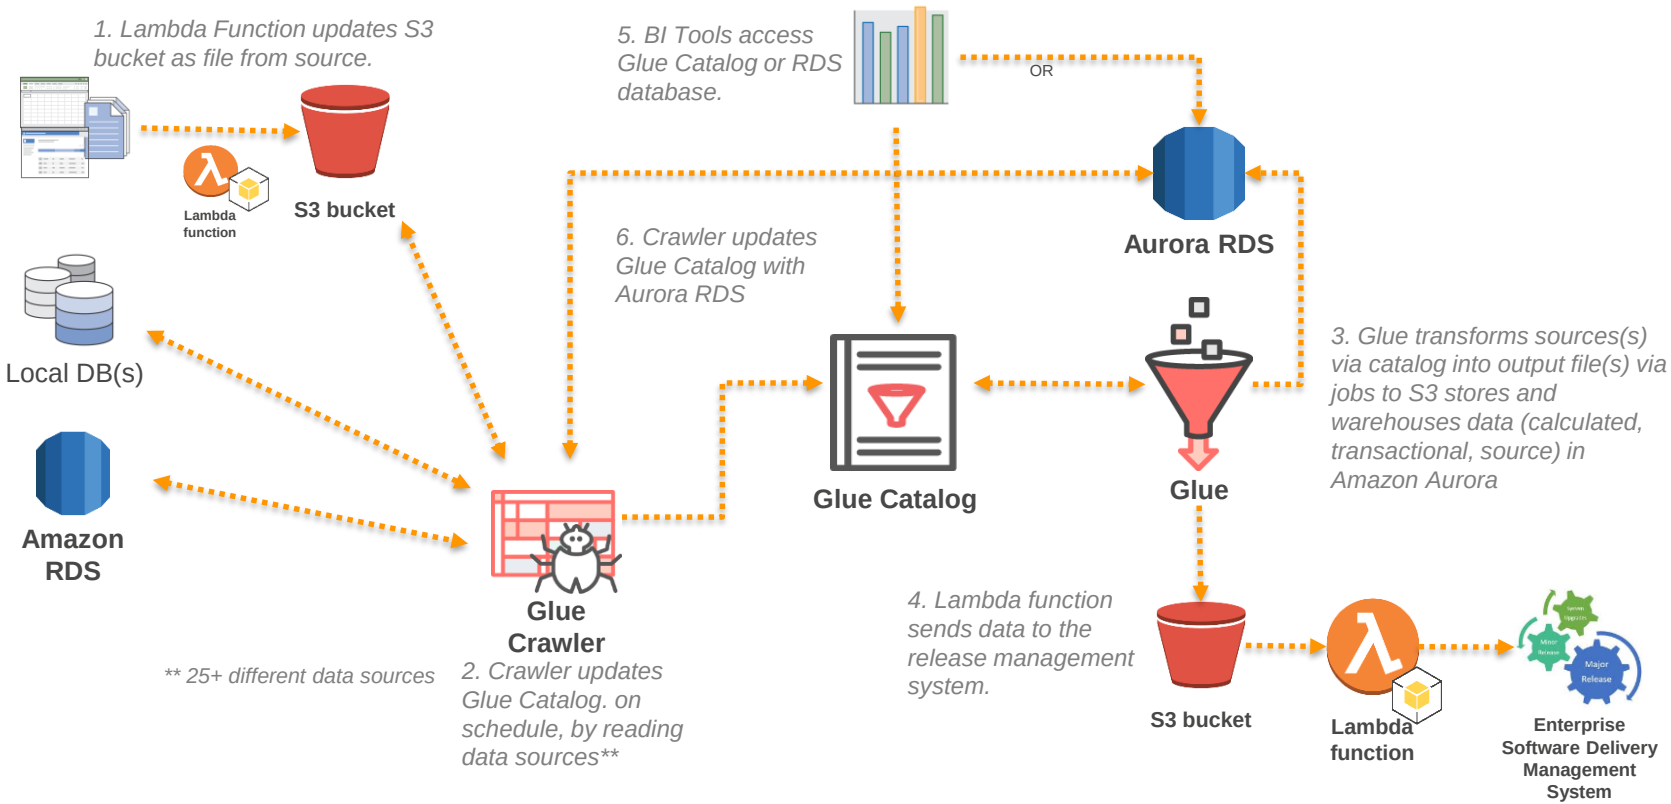

Value

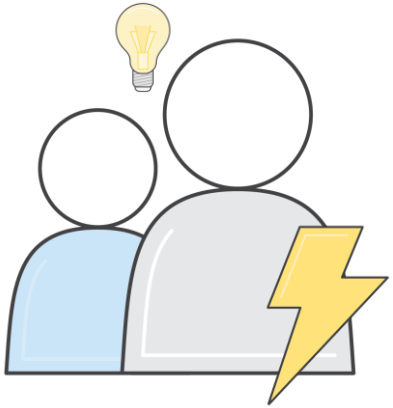

Leverage existing resources.

Glue was quick to learn.

No servers means no additional resources to manage, procure or maintain.

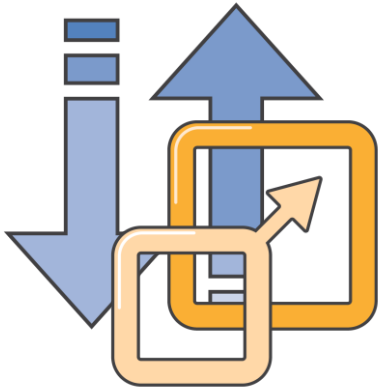

Easy to adapt to new data sources and scale availability within short timeframe.

Provided a single source of data (aligned and calculated).

Enabled a scalable data layer/lake

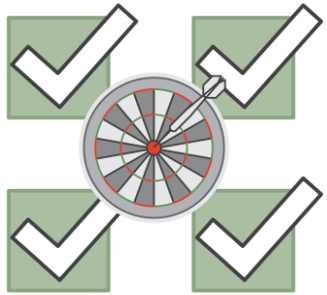

Projects were able to use the data via our Release Management System.

Q&A

AWS re:Invent

THANK YOU!