aws SUMMIT

BDA 308

Deep Dive: Log Analytics With Amazon Elasticsearch Service

Case Study: Amazon ES at Expedia Group

Jon Handler

Principal Search Services SA, AWS

Kuldeep Chowhan

Principal Engineer, Expedia Group

AWS Databases and Analytics

Broad and deep portfolio, purpose-built for builders

Business Intelligence & Machine Learning

- QuickSight
- SageMaker
- Comprehend

Relational Databases

- Aurora
- RDS

Non-Relational Databases

- DynamoDB
- ElastiCache (Redis, Memcached)
- Neptune (Graph)

Analytics

DW | Big Data Processing | Interactive

- Amazon Redshift
- EMR
- Athena

Real-time

- Elasticsearch Service
- Kinesis Data Analytics

Data Lake

- S3/Amazon Glacier
- AWS Glue (ETL & Data Catalog)

Data Movement

Database Migration Service | Snowball | Snowmobile | Kinesis Data Firehose | Kinesis Data Streams

Things with purpose

Database characteristics

Relational

Referential integrity with strong consistency, transactions, and hardened scale

Complex query support via SQL

Key-value

Low-latency key based queries with high throughput and fast ingestion of data

Simple query methods with filters

Document

Indexing and storing documents with support for query on any property

Simple query with filters, projections and aggregates

Graph

Creating and navigating relations between data easily and quickly

Easily express queries in terms of relations

Time Series

Time-stamped data with large range-scans for summarization and processing

Computational support for summarized results

Elasticsearch’s purpose

Text search

Natural language
Boolean queries
Relevance

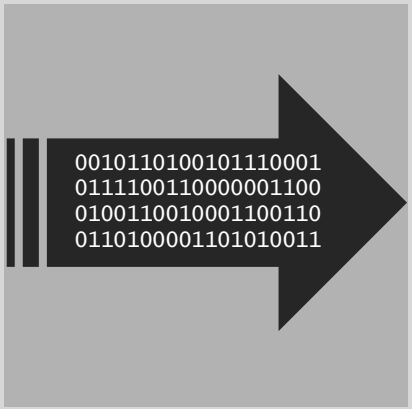

Streaming data

High-volume ingest
Near real time
Distributed storage

Analysis

Time-based visualizations
Nestable statistics
Time series tools

Amazon Elasticsearch Service

Amazon Elasticsearch Service is a **fully managed service** that makes it easy to deploy, manage, and scale Elasticsearch and Kibana in the AWS Cloud

elasticsearch + kibana

aws SUMMIT

Amazon Elasticsearch Service's Storage Layer

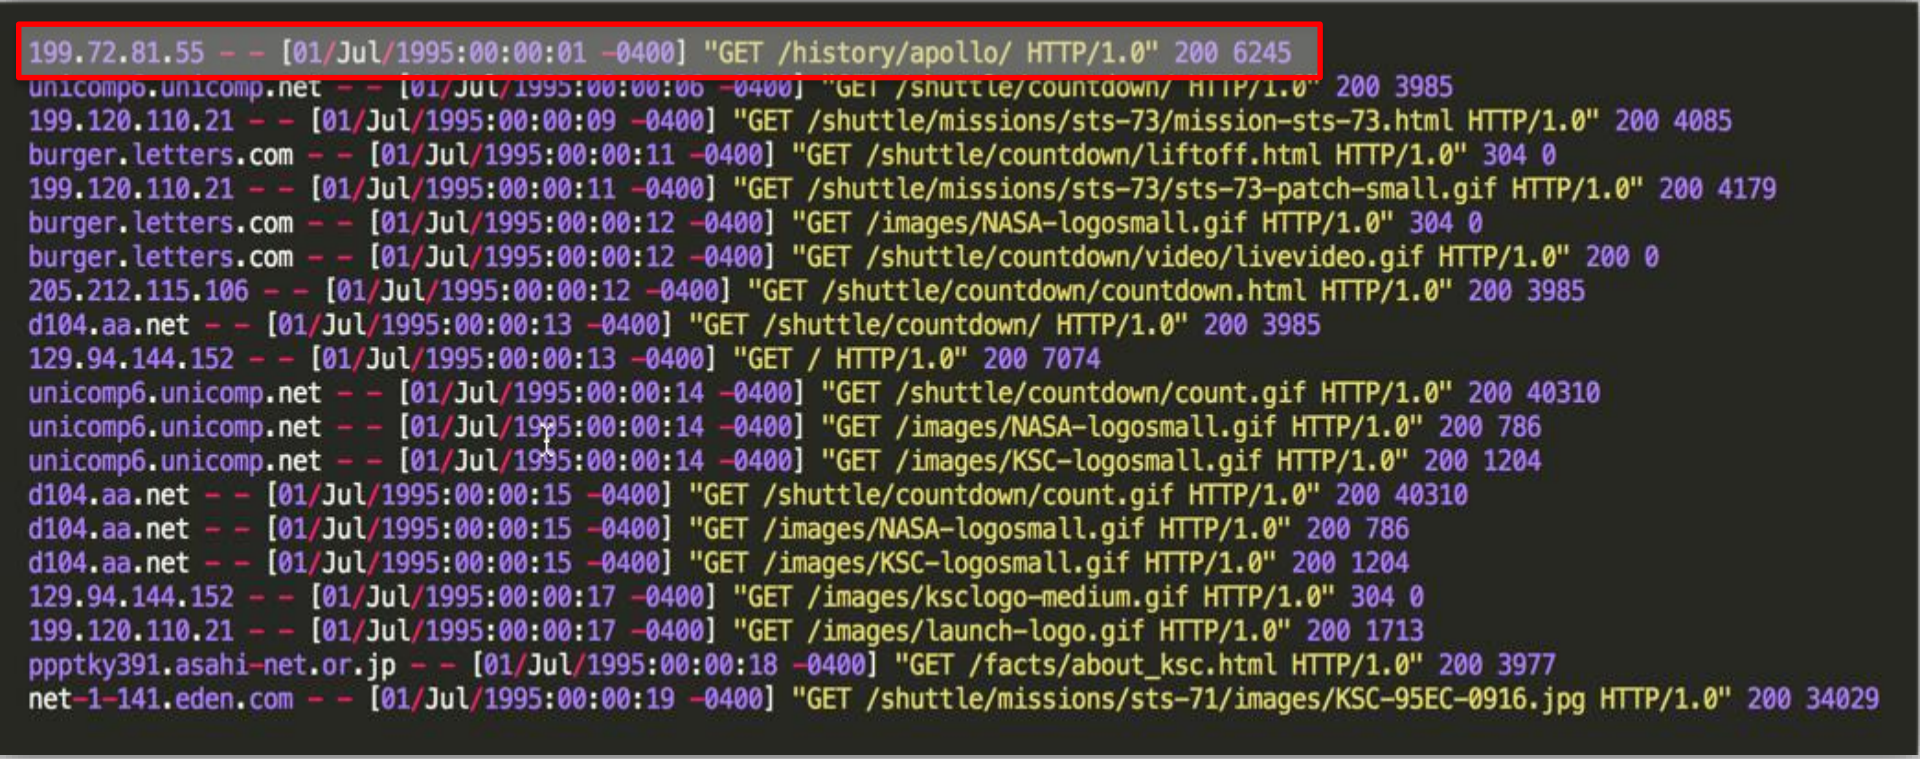

Each log line or other event constitutes a search *document*

```
199.72.81.55 - - [01/Jul/1995:00:00:01 -0400] "GET /history/apollo/ HTTP/1.0" 200 6245
unicomp6.unicomp.net - - [01/Jul/1995:00:00:06 -0400] "GET /shuttle/countdown/ HTTP/1.0" 200 3985
199.120.110.21 - - [01/Jul/1995:00:00:09 -0400] "GET /shuttle/missions/sts-73/mission-sts-73.html HTTP/1.0" 200 4085
burger.letters.com - - [01/Jul/1995:00:00:11 -0400] "GET /shuttle/countdown/liftoff.html HTTP/1.0" 304 0
199.120.110.21 - - [01/Jul/1995:00:00:11 -0400] "GET /shuttle/missions/sts-73/sts-73-patch-small.gif HTTP/1.0" 200 4179
burger.letters.com - - [01/Jul/1995:00:00:12 -0400] "GET /images/NASA-logosmall.gif HTTP/1.0" 304 0
burger.letters.com - - [01/Jul/1995:00:00:12 -0400] "GET /shuttle/countdown/video/livevideo.gif HTTP/1.0" 200 0
205.212.115.106 - - [01/Jul/1995:00:00:12 -0400] "GET /shuttle/countdown/countdown.html HTTP/1.0" 200 3985
d104.aa.net - - [01/Jul/1995:00:00:13 -0400] "GET /shuttle/countdown/ HTTP/1.0" 200 3985
129.94.144.152 - - [01/Jul/1995:00:00:13 -0400] "GET / HTTP/1.0" 200 7074
unicomp6.unicomp.net - - [01/Jul/1995:00:00:14 -0400] "GET /shuttle/countdown/count.gif HTTP/1.0" 200 40310
unicomp6.unicomp.net - - [01/Jul/1995:00:00:14 -0400] "GET /images/NASA-logosmall.gif HTTP/1.0" 200 786
unicomp6.unicomp.net - - [01/Jul/1995:00:00:14 -0400] "GET /images/KSC-logosmall.gif HTTP/1.0" 200 1204
d104.aa.net - - [01/Jul/1995:00:00:15 -0400] "GET /shuttle/countdown/count.gif HTTP/1.0" 200 40310
d104.aa.net - - [01/Jul/1995:00:00:15 -0400] "GET /images/NASA-logosmall.gif HTTP/1.0" 200 786
d104.aa.net - - [01/Jul/1995:00:00:15 -0400] "GET /images/KSC-logosmall.gif HTTP/1.0" 200 1204
129.94.144.152 - - [01/Jul/1995:00:00:17 -0400] "GET /images/ksclogo-medium.gif HTTP/1.0" 304 0
199.120.110.21 - - [01/Jul/1995:00:00:17 -0400] "GET /images/launch-logo.gif HTTP/1.0" 200 1713
ppptky391.asahi-net.or.jp - - [01/Jul/1995:00:00:18 -0400] "GET /facts/about_ksc.html HTTP/1.0" 200 3977
net-1-141.eden.com - - [01/Jul/1995:00:00:19 -0400] "GET /shuttle/missions/sts-71/images/KSC-95EC-0916.jpg HTTP/1.0" 200 34029
```

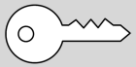

Key Idea

Log lines contain fields

Send JSON to Elasticsearch, with fields and values

```
{
    "host": "199.72.81.55",
    "verb": "GET",
    "request":
        "GET /history/apollo/
         HTTP/1.0",
    "@timestamp":
        "1995-07-01T00:00:01",
    "timezone": "-0400",
    "ident": "-",
    "authuser": "-",
    "response": 200,
    "bytes": 6245
}
```

Lucene creates and stores an index for each field

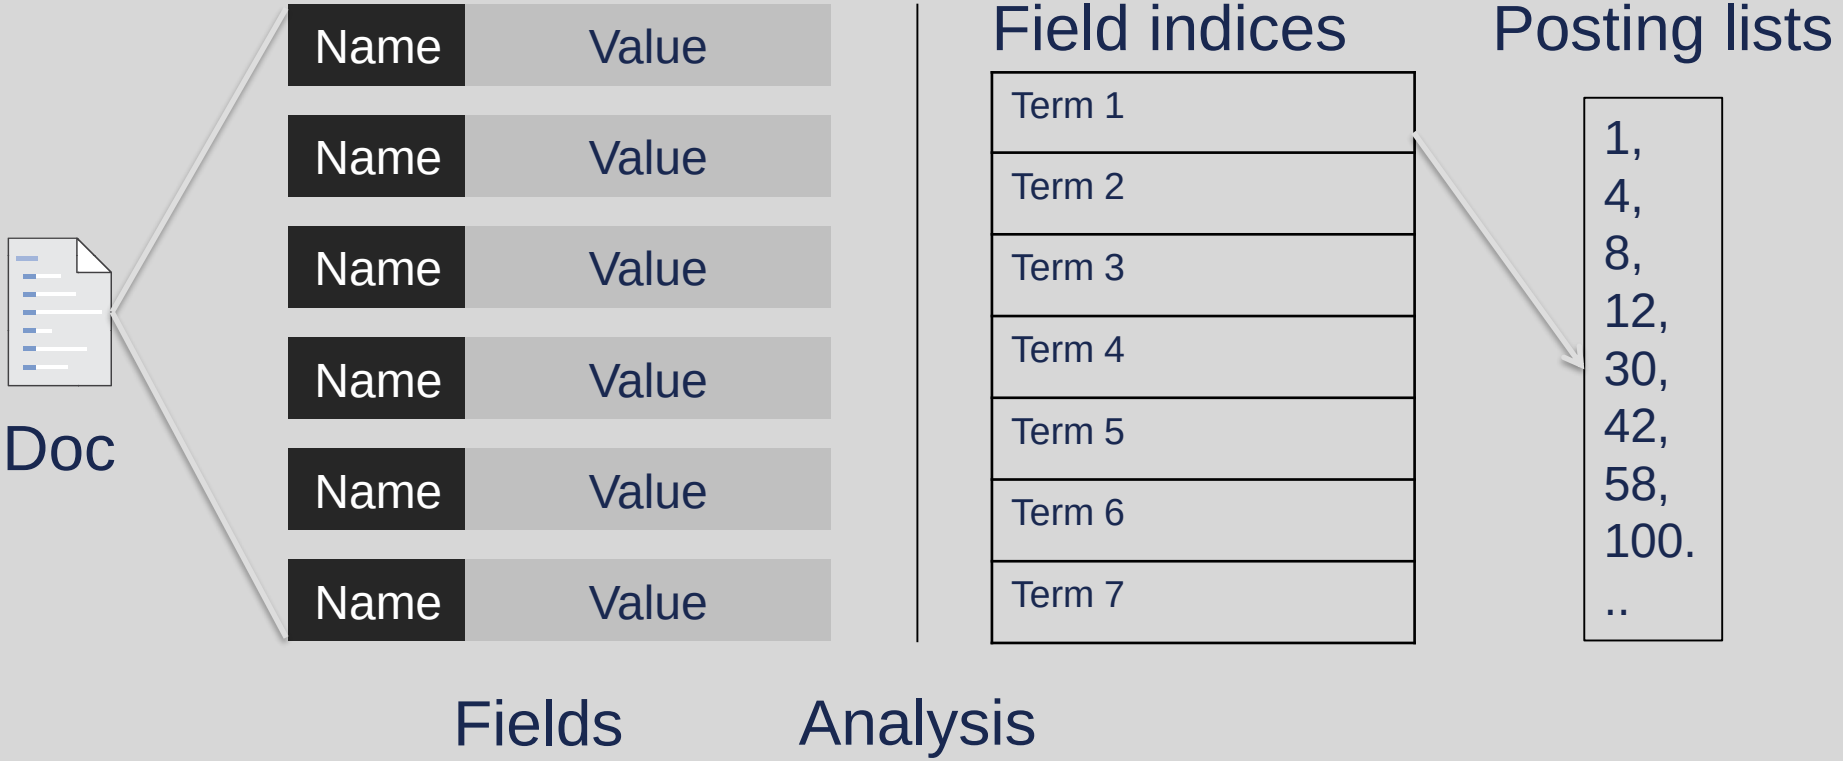

Field indices are managed by shards, organized into API-level indices

Logs index

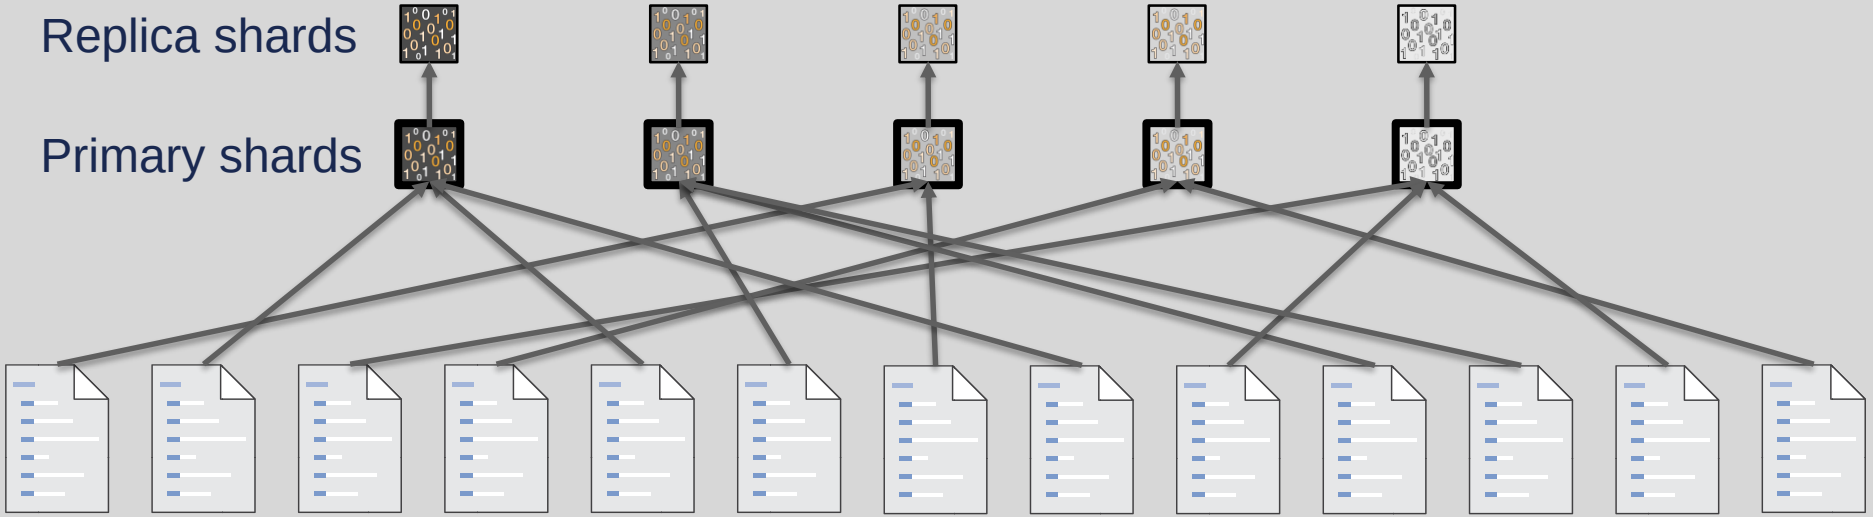

Elasticsearch assigns shards to instances

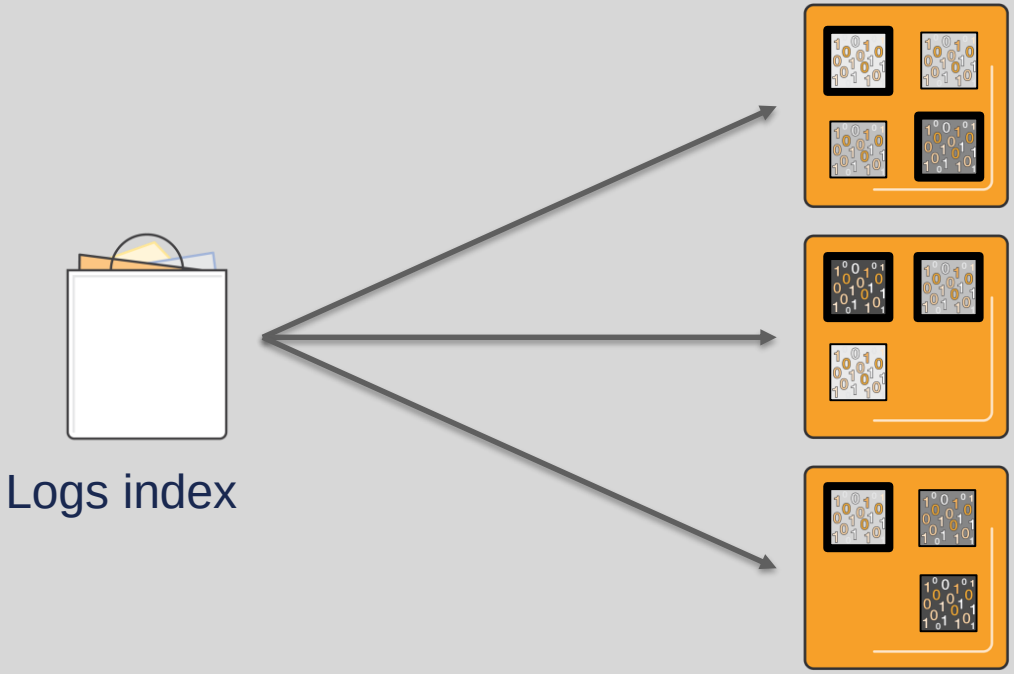

Storage required

- On disk, indices are ~10% larger than source
- Each replica adds an additional 1x storage requirement
- You choose the per-instance storage

Example: a 1 TB corpus will need 2 instances

With one replica and 10% inflation, you need 2.2 TB of storage

Choose 1.5 TB of EBS per instance, and you need 2

Shards as units of storage

- Set primary shard count based on storage, 40 GB per primary (90 GB for I3 instances)
- Always use at least 1 replica in production
- Keep shard sizes as equivalent as possible

Example: Set shard count = 50 for a 2-TB corpus (2 TB / 40 GB = 50 shards)

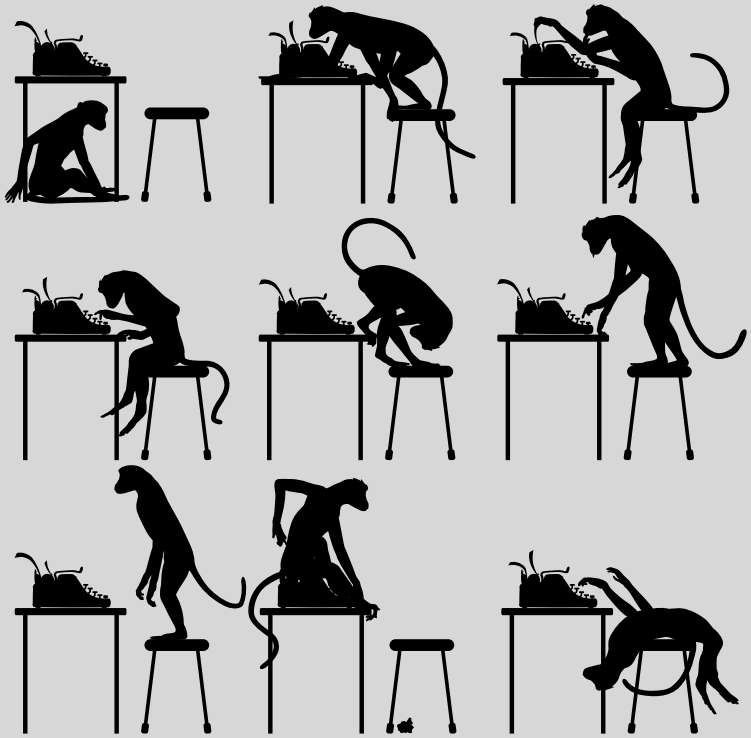

Build an ingest pipeline that completes these tasks

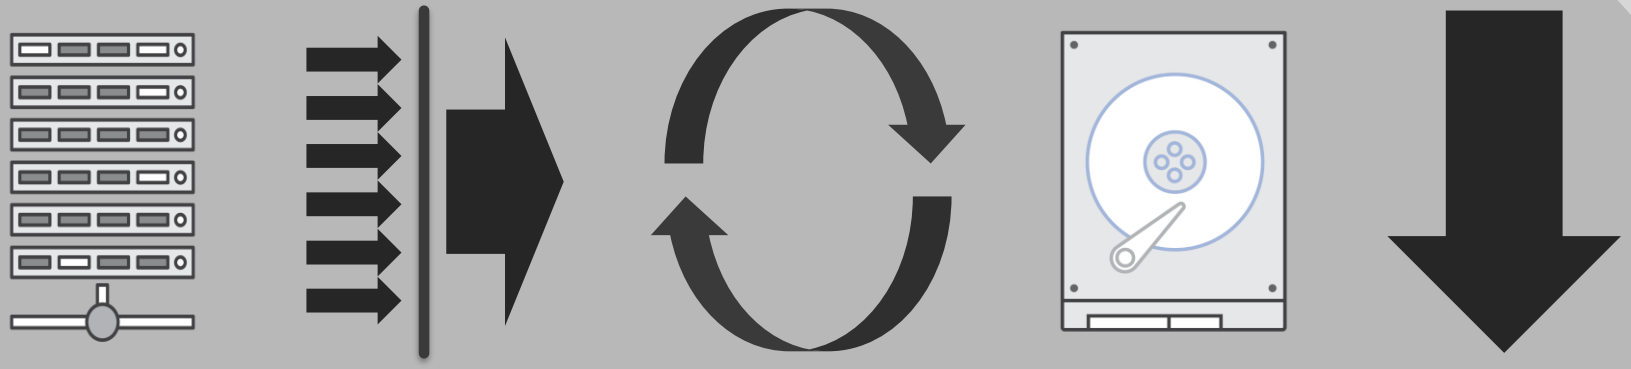

Data source | Collect | Transform | Buffer | Deliver

Organize data in daily indexes

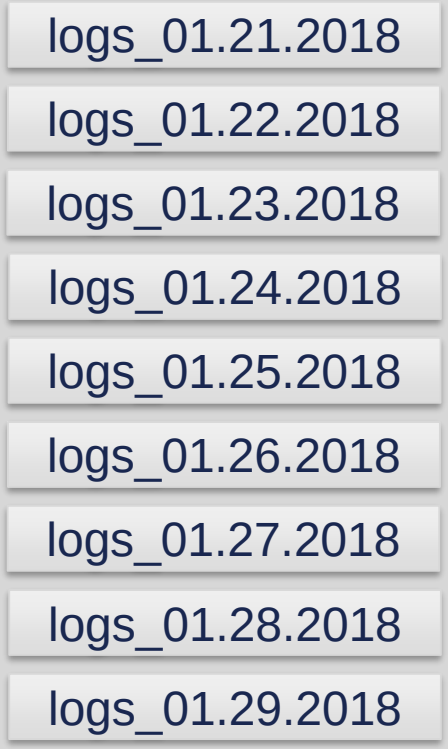

- On ingest, create indexes with a root string, e.g., logs_
- Depending on volume, rotate at regular intervals – normally daily
- Daily indexes simplify index management. Delete the oldest index to create more space on your cluster.

Use templates to set shard count

```
PUT* <endpoint>/_template/template1
{
    "index_patterns": [”movies*"],
    "settings": {
        "number_of_shards": 50,
        "number_of_replicas": 1
    }
}
```

All new indexes that match the index pattern receive the settings

***Note: ES 6.0+ syntax**

aws SUMMIT

Amazon Elasticsearch Service's Query Engine

Query distribution

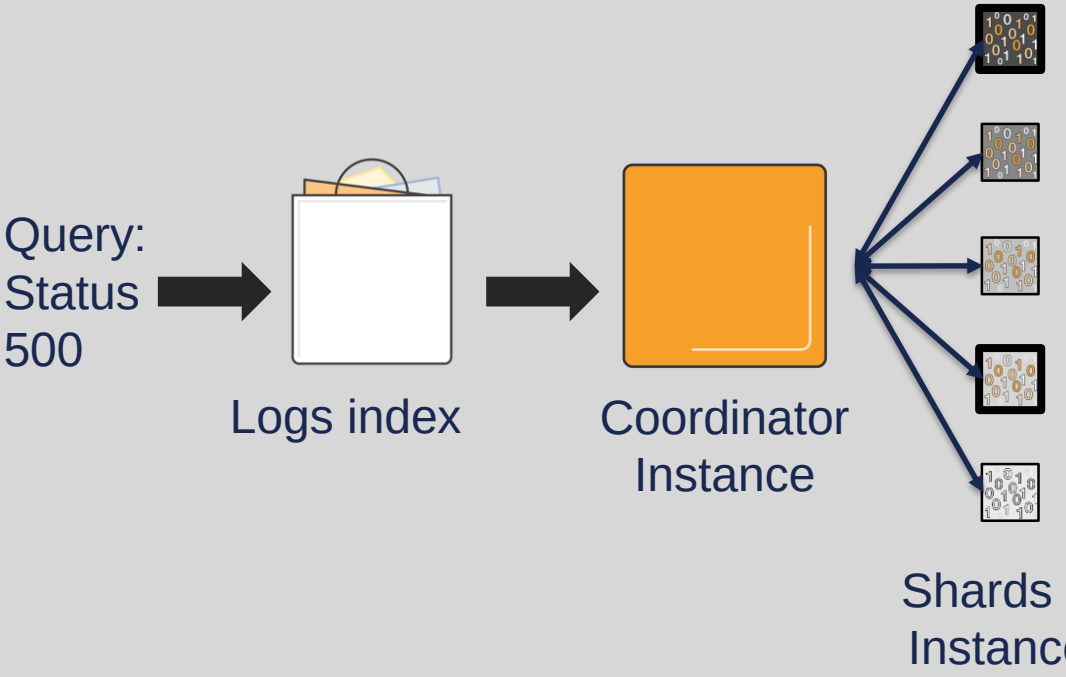

Each shard computes and returns a result to the coordinator, which re-aggregates a final result

Query processing

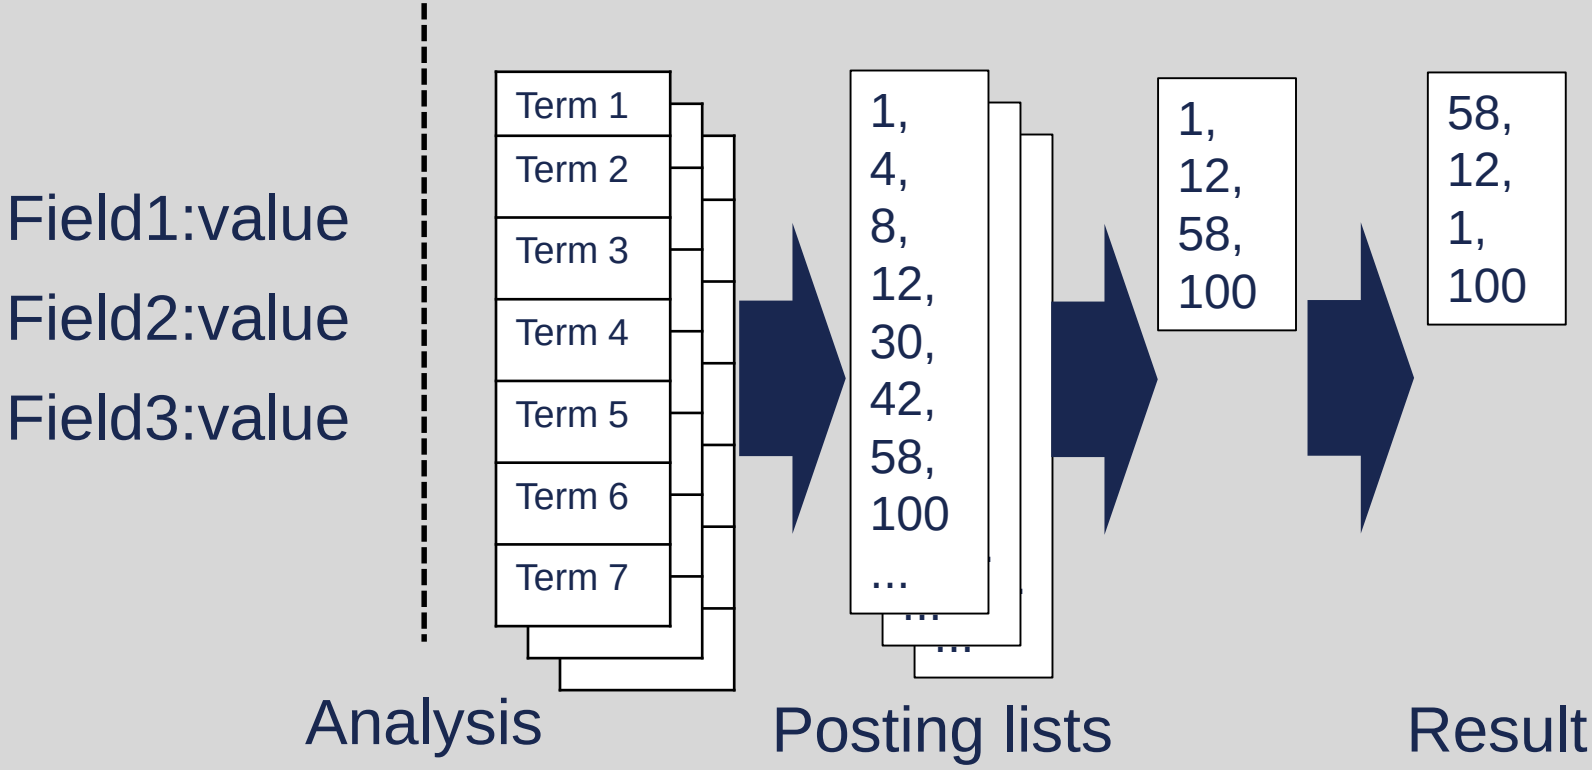

Analyze field values to get statistics and build visualizations

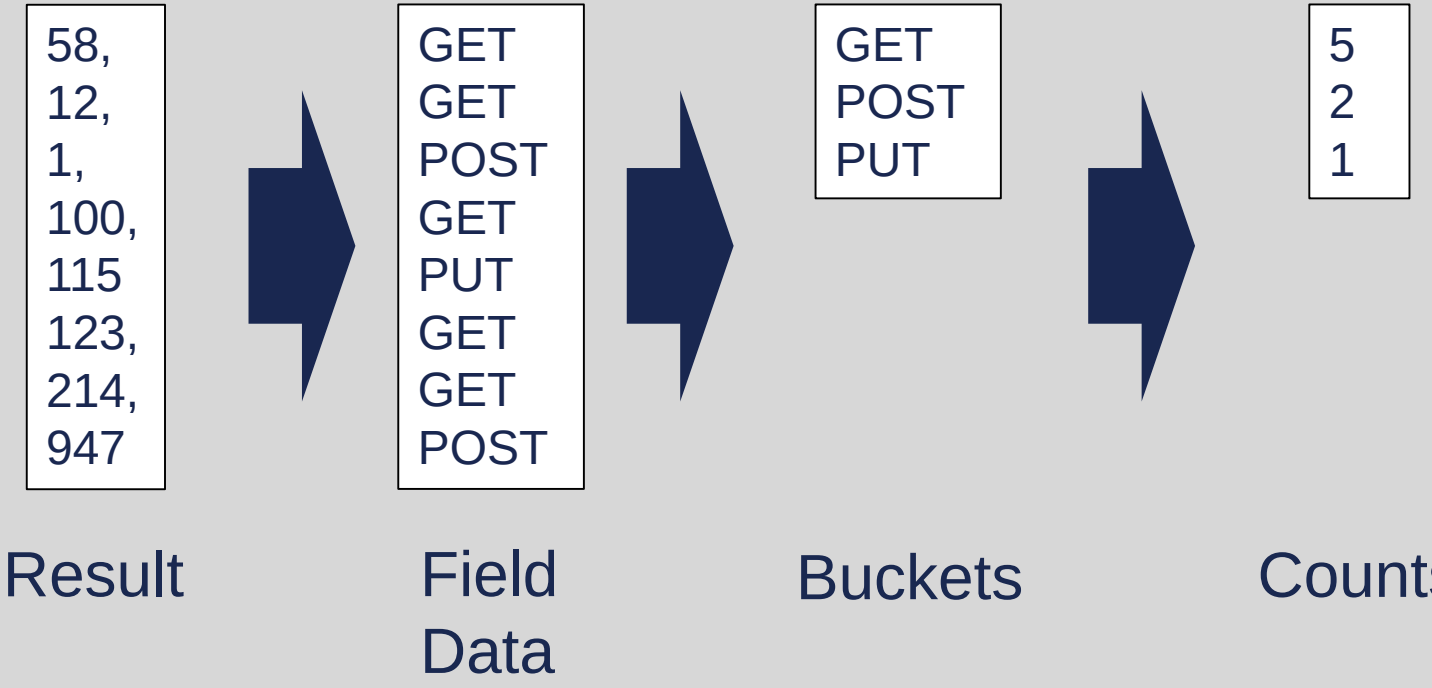

Visualize your data

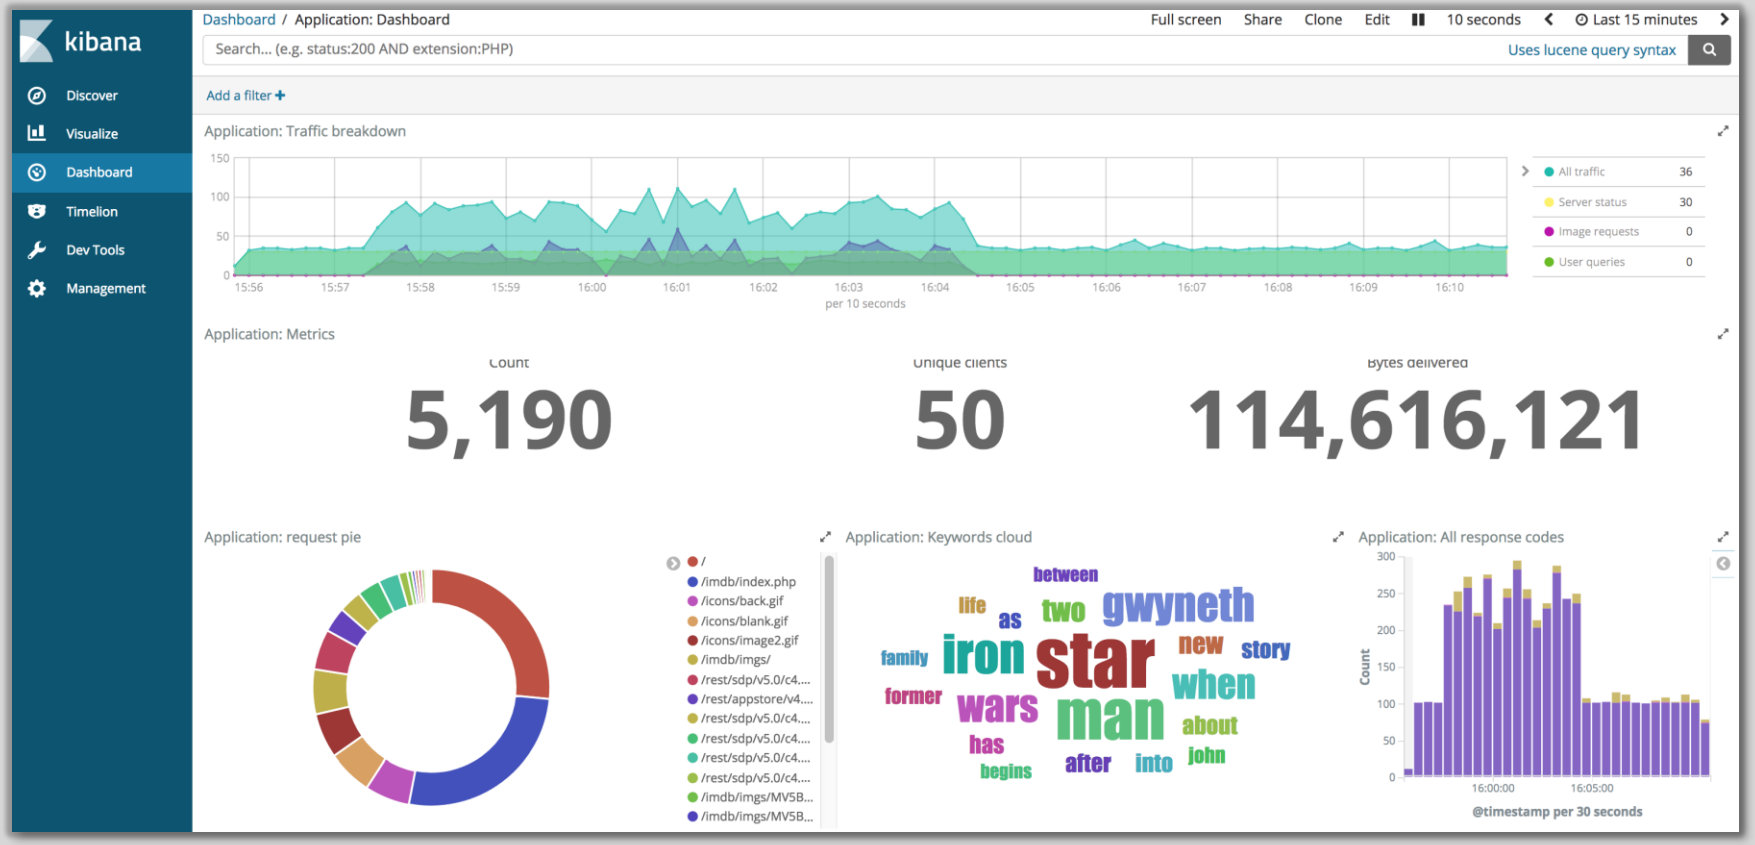

Case study: MirrorWeb

Full text search

MirrorWeb

PROBLEM

Make the UK Government and UK Parliament's web archives searchable

Large scale ingestion scenario: 120 TB of data (1.2 MM 100-MB files), duplicates and bad data, Warc format

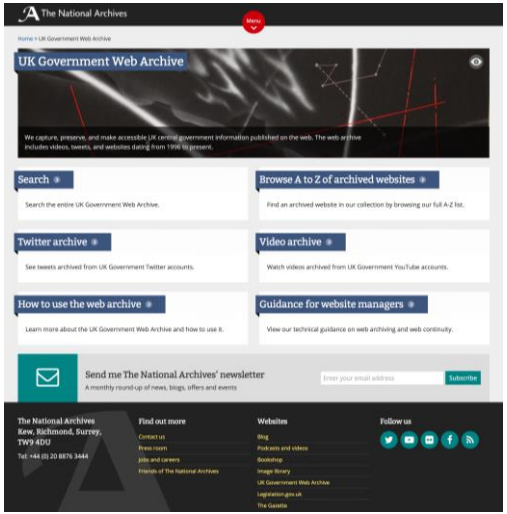

SOLUTION

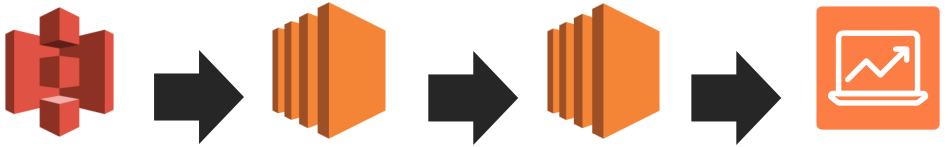

BENEFITS

Scalability: Started on a 9-node, R4.4Xlarge cluster for fast ingest, reduced to 6 R4.Xlarge instances for search. Able to reconfigure the cluster with no down time

Cost effective: Indexed 1.4 billion documents for $337

Fast: 146 MM docs per hour indexed. 14x faster than the previous best for this data set (using Hadoop)

For more on this case, see http://tinyurl.com/ybqwbolq

Case study: Financial Times

Business and Clickstream Analytics

FT FINANCIAL TIMES

PROBLEM

What stories do our readers care about? What's hot?

Required a custom clickstream analytics solution

Need a solution that delivers analytics in real time

Did not have a team to manage analytics infrastructure

SOLUTION

Streaming user data to Amazon ES for analysis. Created their own custom dashboards for editors and journalists – Lantern.

Lantern - "shines a light" on reader activity for the editors and journalists at the FT

Critical tool for making editorial decisions. Daily editorial meetings start by looking at Lantern dashboard

BENEFITS

Reliability : Lantern is used throughout the day by journalists and editors. Relying on Amazon to manage their systems for maximum uptime.

Cost savings: Able to easily tune their cluster to meet their needs with minimal management overhead

aws SUMMIT

Amazon Elasticsearch Service

Benefits of Amazon Elasticsearch Service

Supports open-source APIs and tools

Drop-in replacement with no need to learn new APIs or skills

Easy to use

Deploy a production-ready Elasticsearch cluster in minutes

Scalable

Resize your cluster with a few clicks or a single API call

Secure

Deploy into your VPC and restrict access using security groups and IAM policies

Highly available

Replicate across Availability Zones, with monitoring and automated self-healing

Tightly integrated with other AWS services

Seamless data ingestion, security, auditing and orchestration

Service architecture

AWS SDK

AWS CLI

AWS CloudFormation

Amazon Elasticsearch Service domain

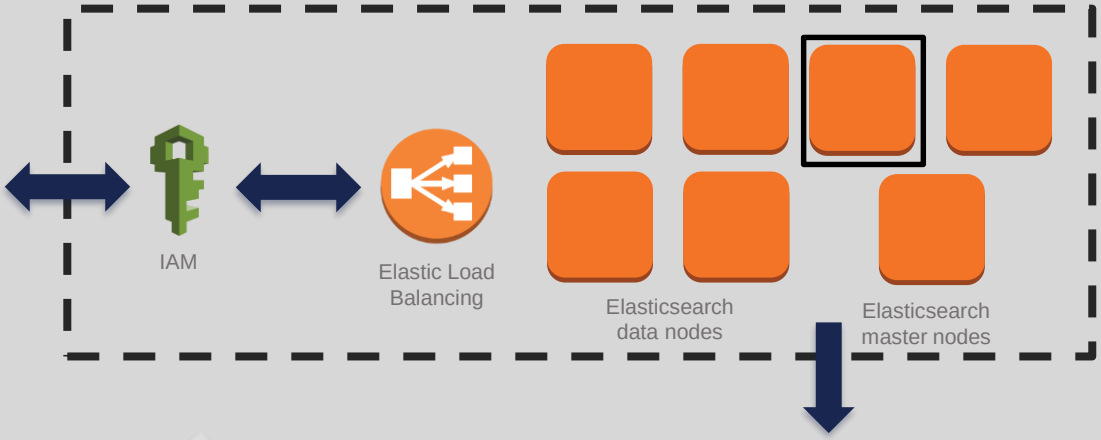

AWS CloudTrail

Amazon CloudWatch

Security

- VPC access (Recommended)
- Public access

Public endpoints – IAM

Private endpoints – IAM and security groups

Encryption

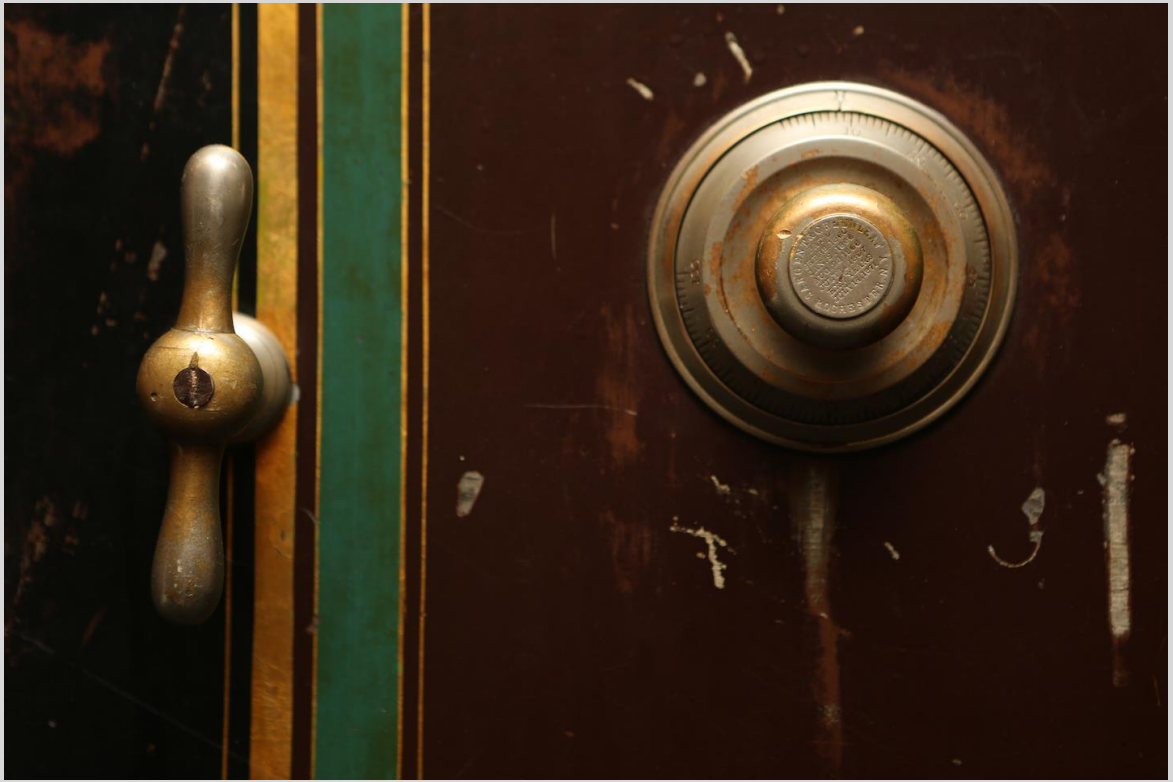

Use three dedicated master instances in production

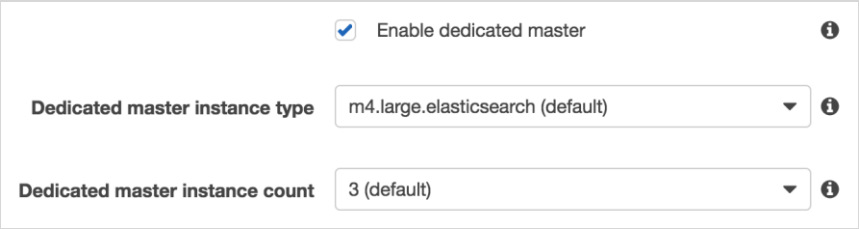

Master instances orchestrate and make your cluster more stable

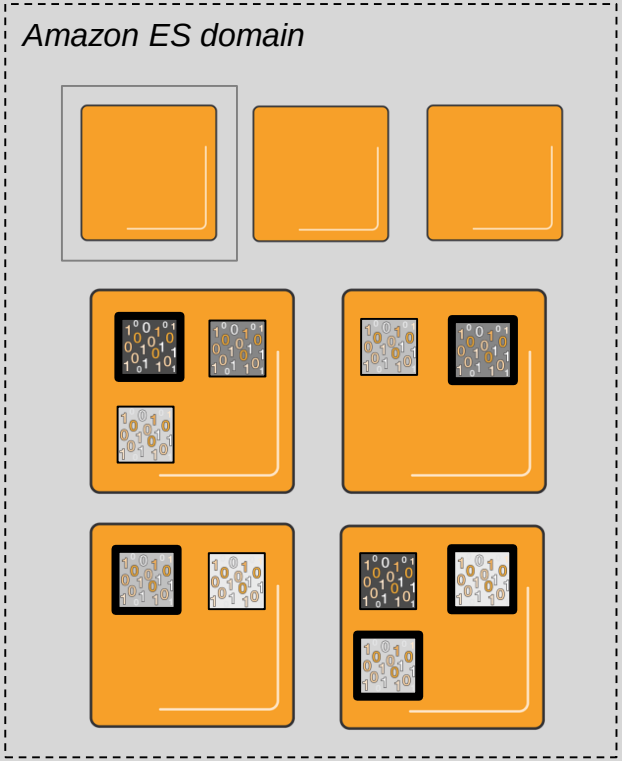

Use zone awareness in production

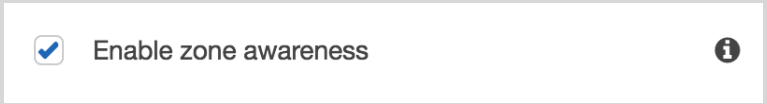

100% data redundancy in two zones makes your cluster more highly available

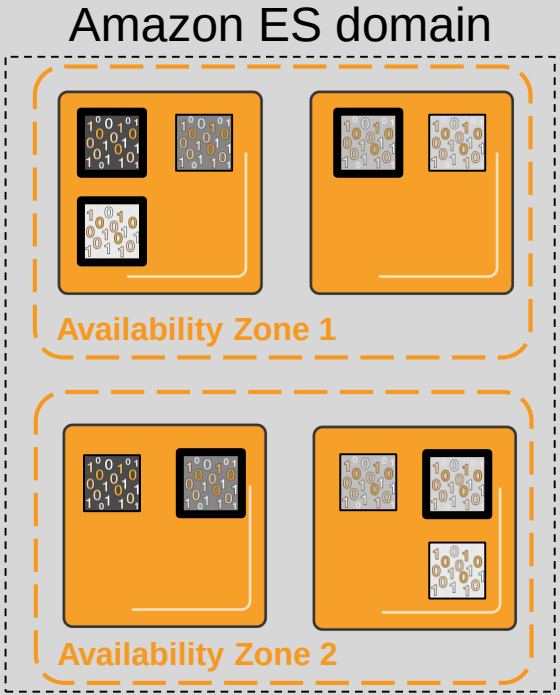

Set CloudWatch metrics and alarms

Name	Metric	Threshold	Periods
ClusterStatus.red	Maximum	>= 1	1
ClusterIndexWritesBlocked	Maximum	>= 1	1
CPUUtilization/MasterCPUUtilization	Average	>= 80%	3
JVMMemoryPressure/Master...	Maximum	>= 80%	3
FreeStorageSpace	Minimum	<= (25% of avail space)	1
AutomatedSnapshotFailure	Maximum	>= 1	1

Monitor Elasticsearch slow logs

- Easy console setup
- Integrated with CloudWatch Logs
- Set thresholds to receive log events corresponding to slow queries and slow indexing

- `index.search.slowlog.threshold.query.warn`
- `index.search.slowlog.threshold.query.info`
- `index.search.slowlog.threshold.query.debug`
- `index.search.slowlog.threshold.query.trace`

- `index.search.slowlog.threshold.fetch.warn`
- `index.search.slowlog.threshold.fetch.info`
- `index.search.slowlog.threshold.fetch.debug`
- `index.search.slowlog.threshold.fetch.trace`

- `index.indexing.slowlog.threshold.index.warn`
- `index.indexing.slowlog.threshold.index.info`
- `index.indexing.slowlog.threshold.index.debug`
- `index.indexing.slowlog.threshold.index.trace`

- `index.indexing.slowlog.level: trace`
- `index.indexing.slowlog.source: 255`

Pay only for what you use

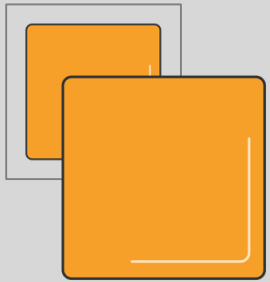

Instance hours

For data and master instances

EBS GB/Mo

For volumes deployed

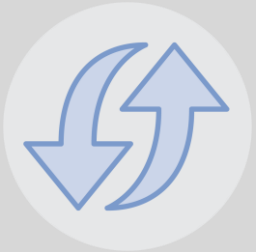

AWS data transfer

For transfer out

aws SUMMIT

Amazon Elasticsearch Service usage @ Expedia

Kuldeep Chowhan

Principal Engineer
@ Expedia Group, Inc.

@this_is_kuldeep

+175 AmazonES clusters

+500 EC2 instances

+40B documents

+35 TB of data

aws SUMMIT

Why did we choose Amazon Elasticsearch Service?

Easy to set up

Set up for high availability

Instance count 12

Instance type i2.2xlarge.elasticsearch

☑ Enable dedicated master

Dedicated master instance type m3.xlarge.elasticsearch

Dedicated master instance count 3 (default)

☑ Enable zone awareness

Storage configuration is Active

Choose a storage type for your data nodes. If you choose the EBS storage type, you will need to specify the EBS volume type and EBS volume size for the cluster. The EBS volume size setting is configured per instance. Multiply the volume size by the number of data nodes in your cluster for the total storage size available in your cluster. Take into account size of indices, shards, and replicas you intend to create in your cluster when configuring storage settings. Storage settings do not apply to any dedicated master nodes in the cluster.

Storage type Instance (default)

Snapshot configuration

Once a day, Amazon ES takes an automated snapshot of your cluster. You can set the start hour for the snapshot. We recommend that you choose a time when traffic on your cluster is low.

Automated snapshot start hour 00:00 UTC (default)

Security

Elasticsearch access policy example

```
{
      "Effect": "Allow",
      "Principal": {
            "AWS": "arn:aws:iam::xxxxx:root"
      },
      "Action": "es:*",
      "Resource": “arn:aws:es:us-west-2:xxxxx:domain/xxxxx/*"
},
{
      "Effect"  "Allow"
      "Principal": {
            "AWS": "*"
      },
      "Action": "es:Http*",
      "Resource": "arn:aws:es:us-west-2:xxxx:domain/xxxxx/*",
      "Condition": {
            "IpAddress": {
                  "aws:SourceIp": [
                        "0.0.0.0/28"
                  ],
```

Monitoring & backups

aws SUMMIT

Different log analytics architectures using Elasticsearch Service @ Expedia

Different Log Analytics Architectures

- Docker startup logs to Elasticsearch
- CloudTrail log analytics using Elasticsearch Service
- Distributed tracing platform using Elasticsearch Service

aws SUMMIT

Docker start-up logs to Elasticsearch

Docker startup logs to Elasticsearch

Amazon
ECS

Docker startup logs to Elasticsearch

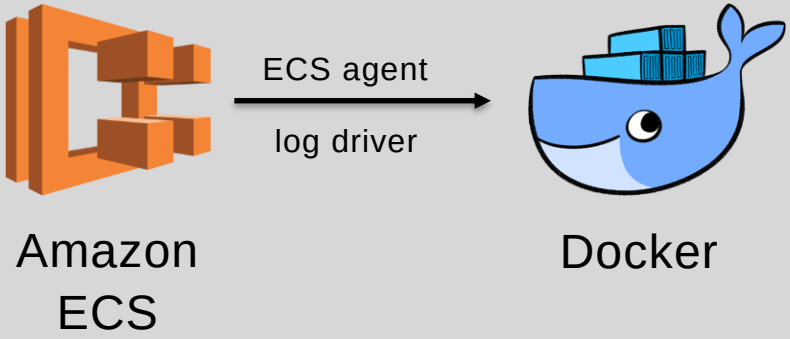

Docker startup logs to Elasticsearch

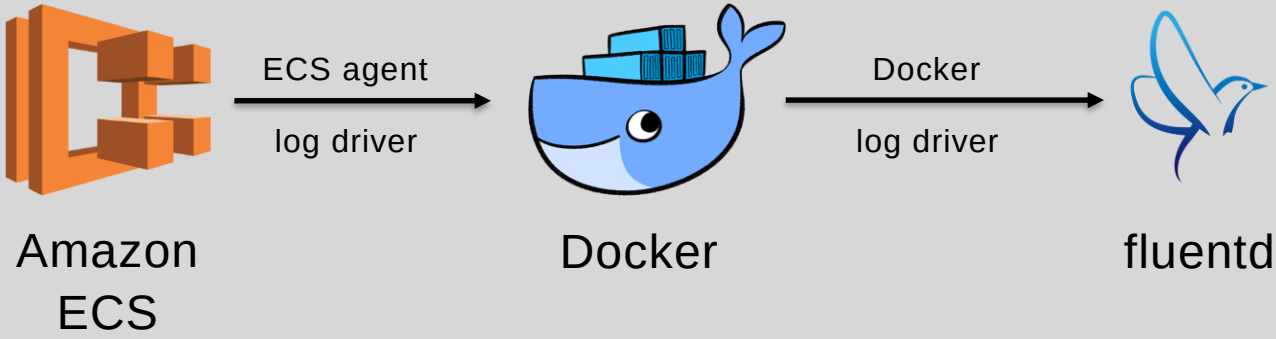

Docker startup logs to Elasticsearch

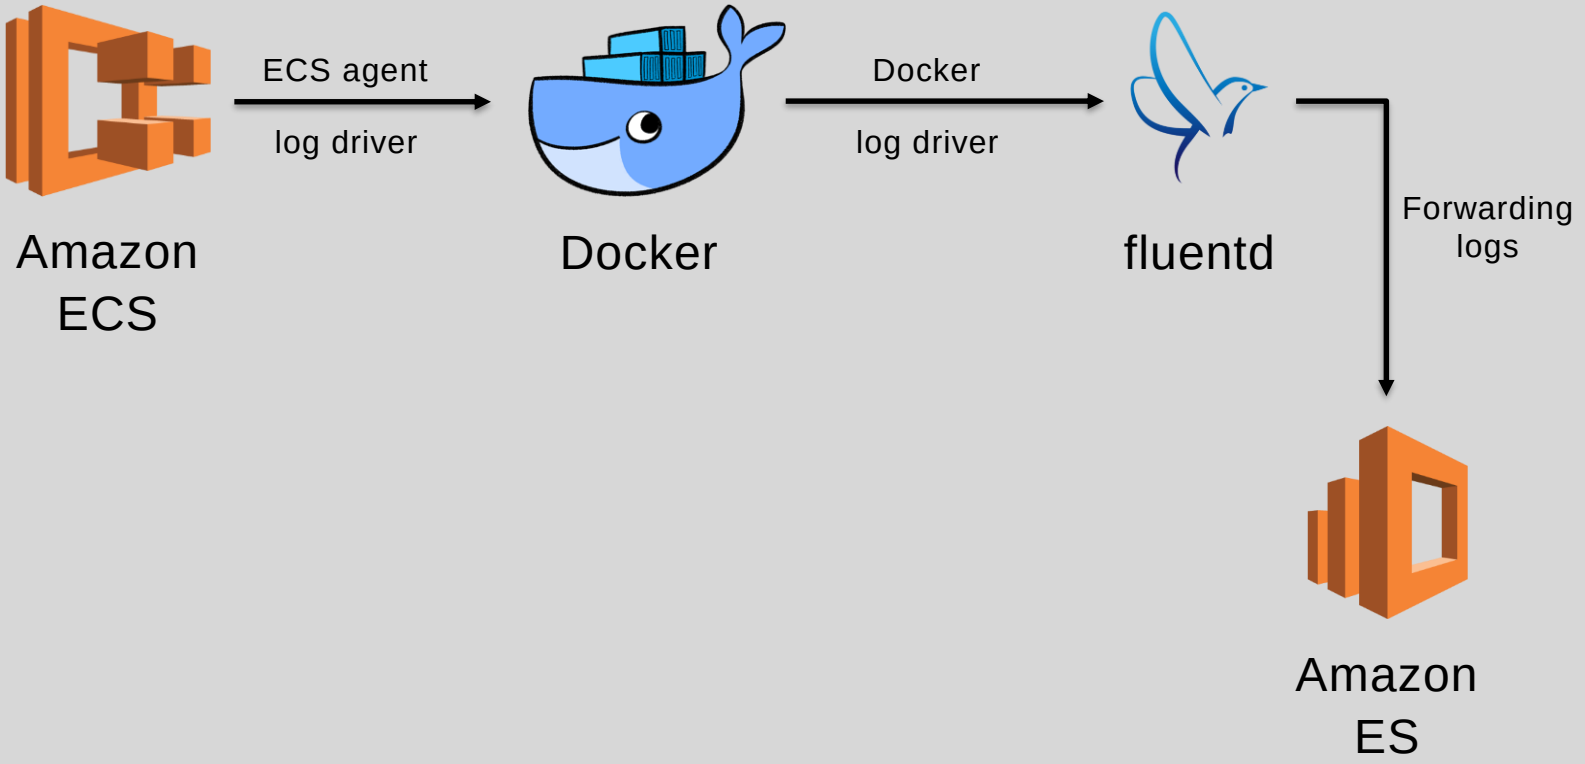

Docker startup logs to Elasticsearch

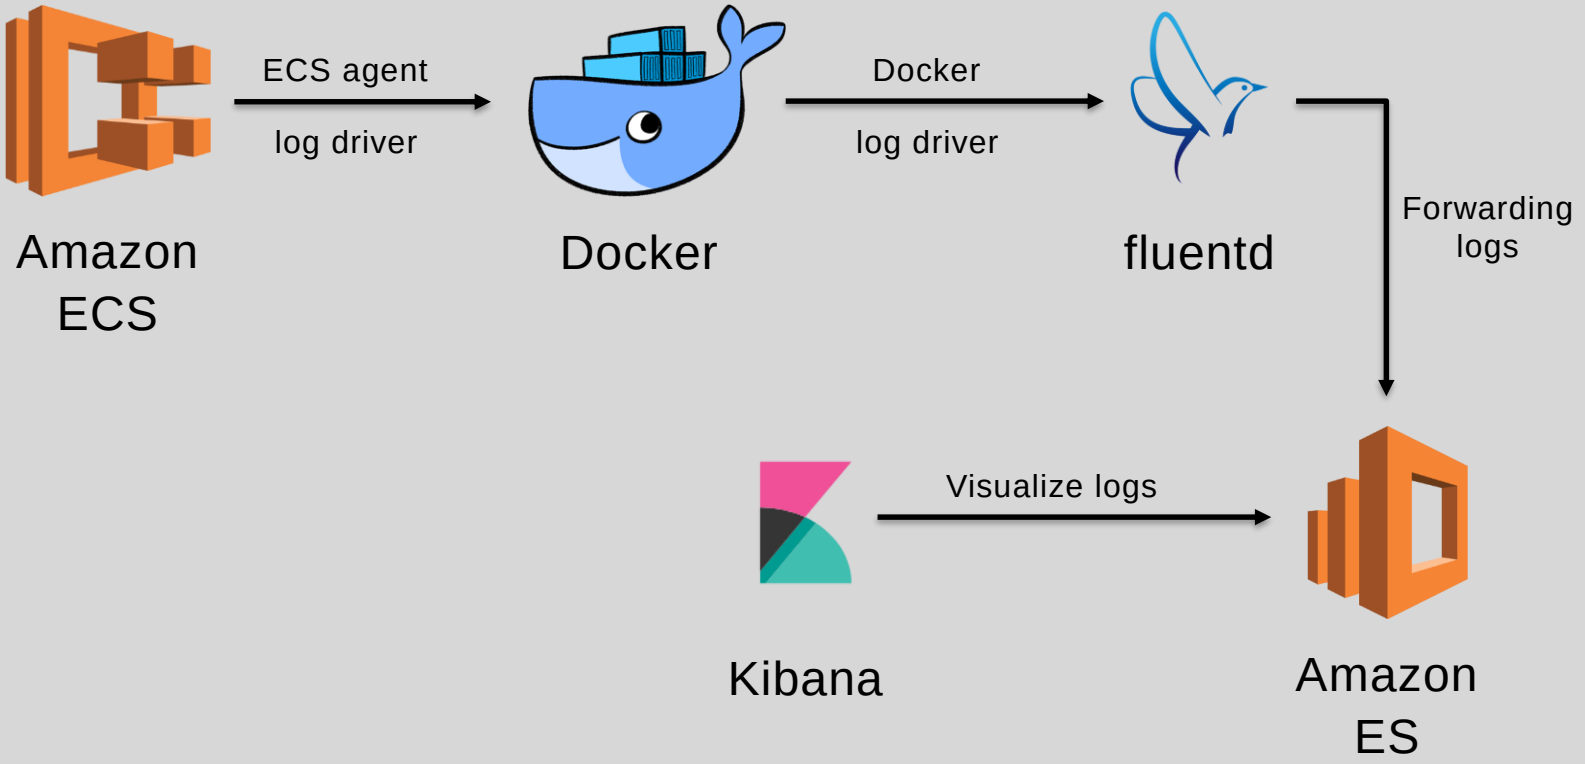

Docker fluentd log_driver configuration

```
{
    "log_driver": "fluentd",
    "options": {
        "fluentd-address": “<fluentd>:24224",
        "tag": "#{ImageName}"
    }
}
```

fluentd configuration to receive Docker logs

```
<source>
   @type forward
   port 24224
   bind 0.0.0.0
</source>
<match *.**>
   @type copy
```

fluentd to ES configuration

```
<match *.**>
    @type copy
    <store>
        @type elasticsearch
        host <elasticsearch domain>
        include_tag_key true
        tag_key @log_name
        flush_interval 1s
    </store>
```

aws SUMMIT

CloudTrail log analytics using Elasticsearch Service

CloudTrail log analytics using Elasticsearch

AWS
CloudTrail

CloudTrail log analytics using Elasticsearch

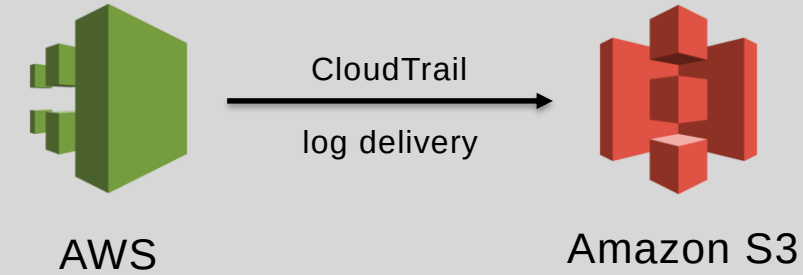

CloudTrail log analytics using Elasticsearch

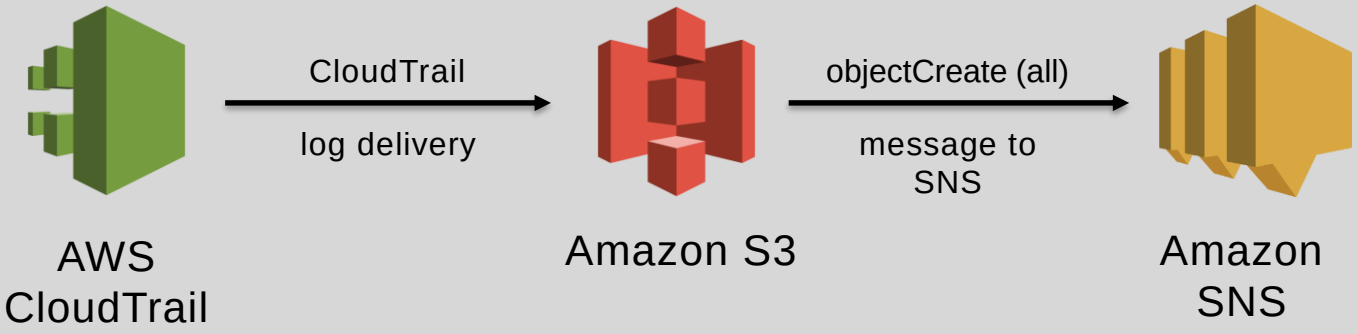

CloudTrail log analytics using Elasticsearch

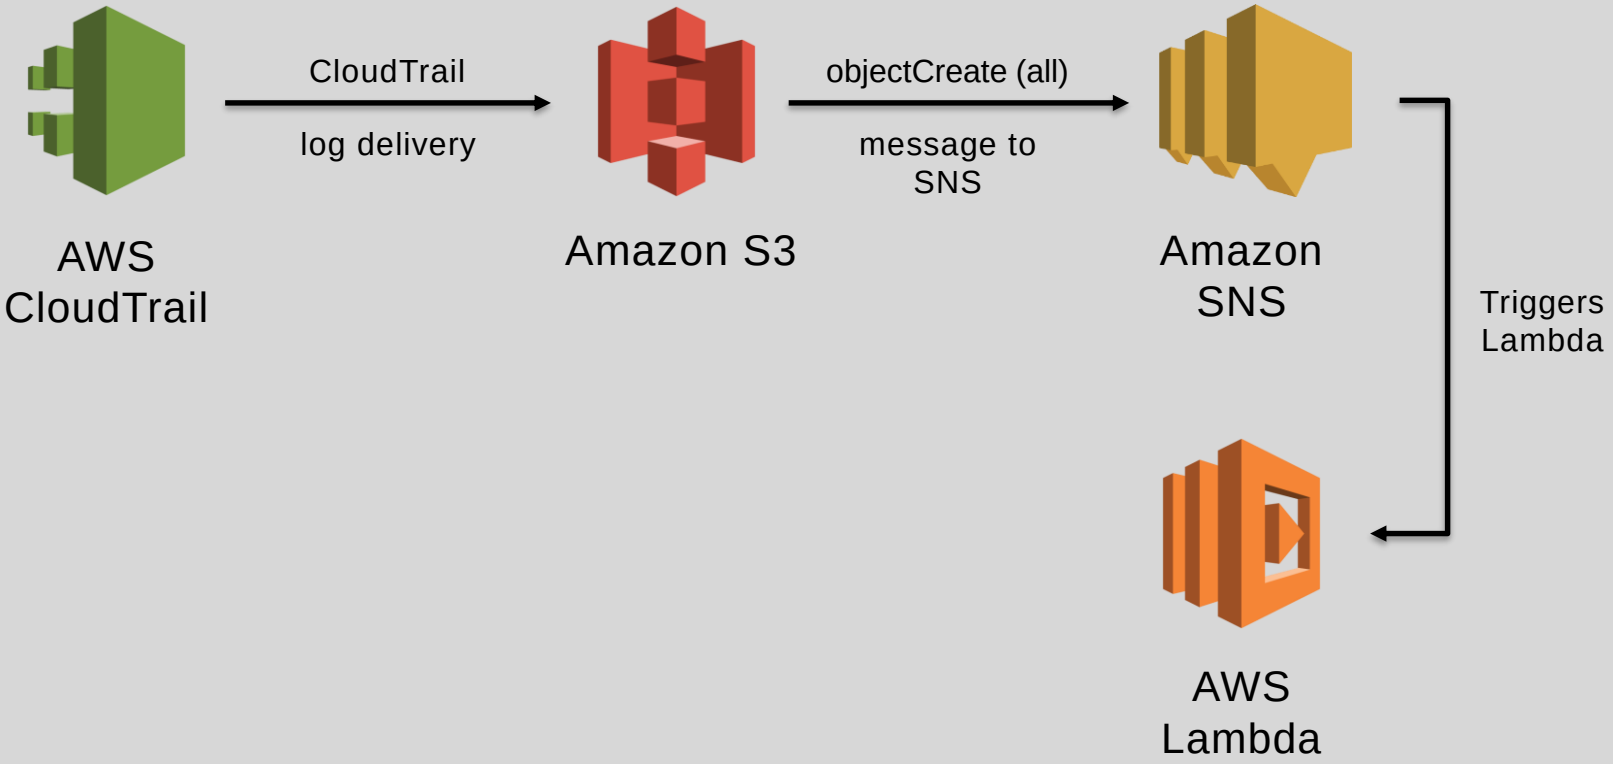

CloudTrail log analytics using Elasticsearch

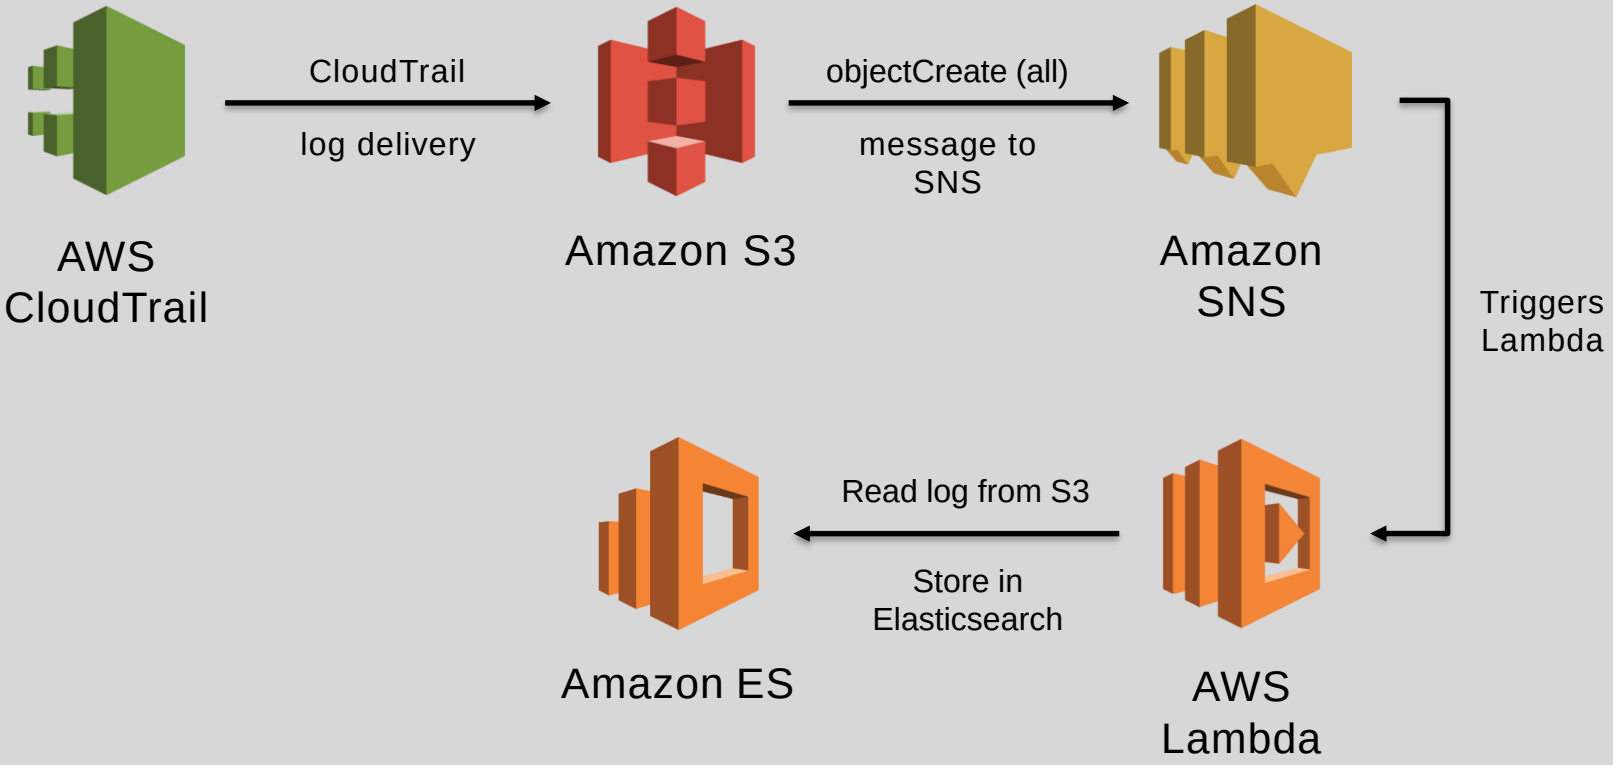

CloudTrail log analytics using Elasticsearch

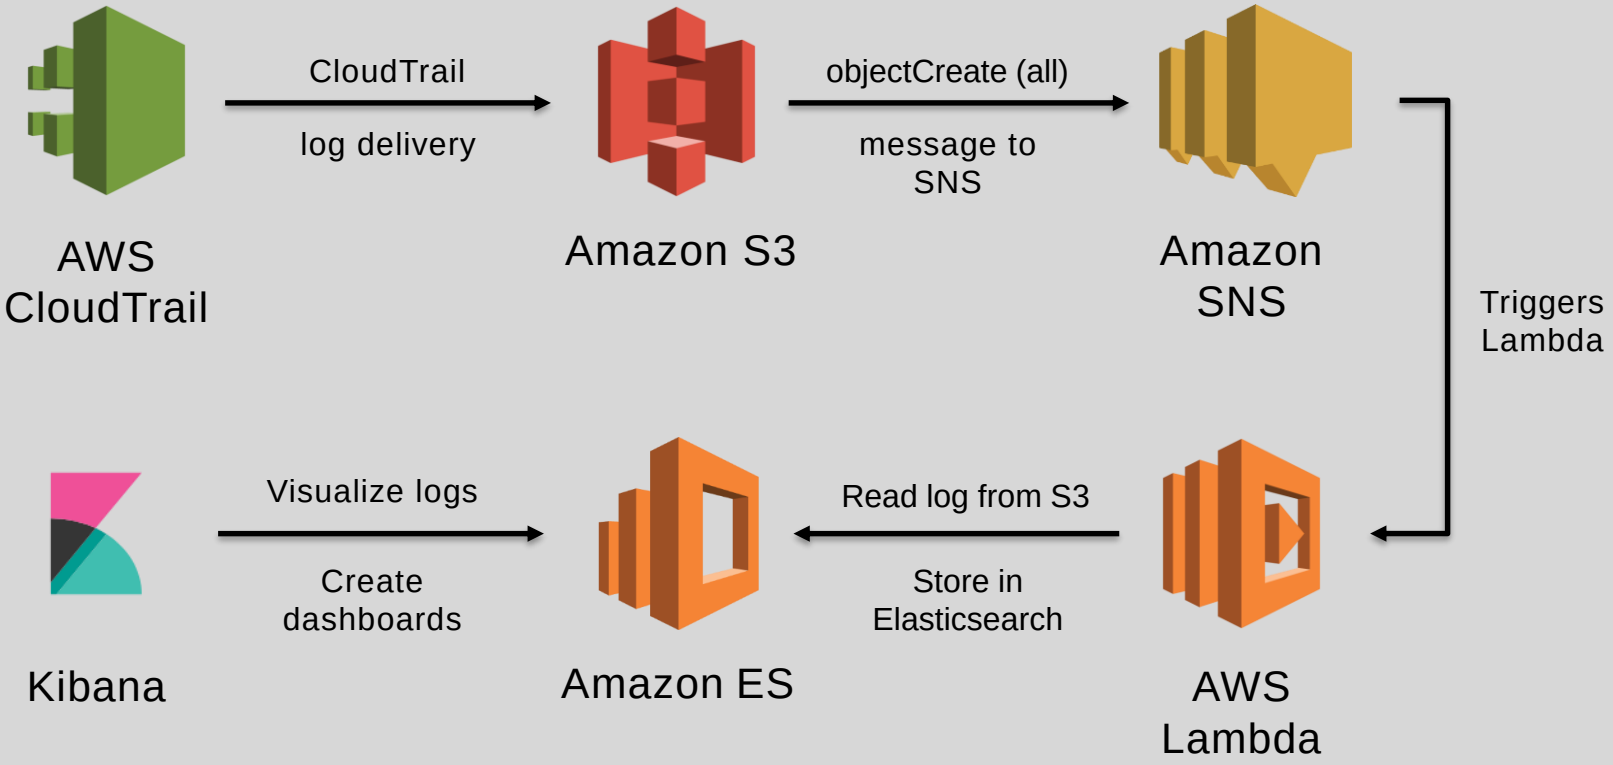

CloudTrail log analytics using Elasticsearch

CloudTrail to S3 and SNS

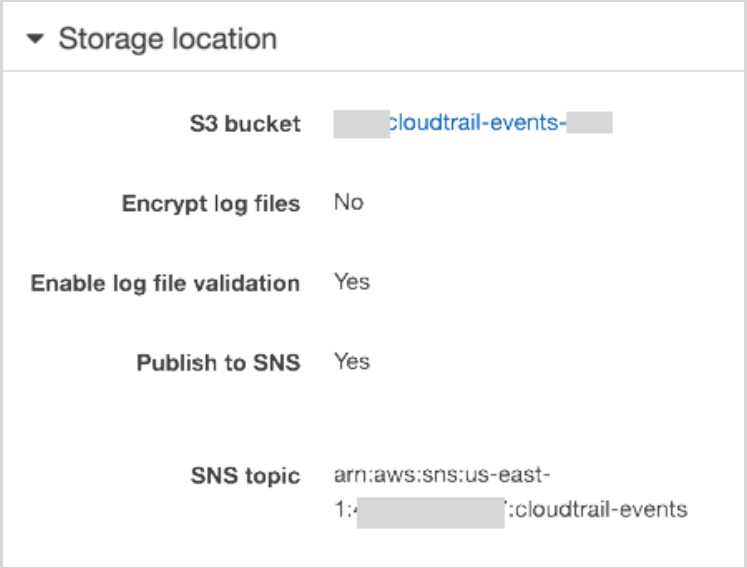

SNS to Lambda trigger

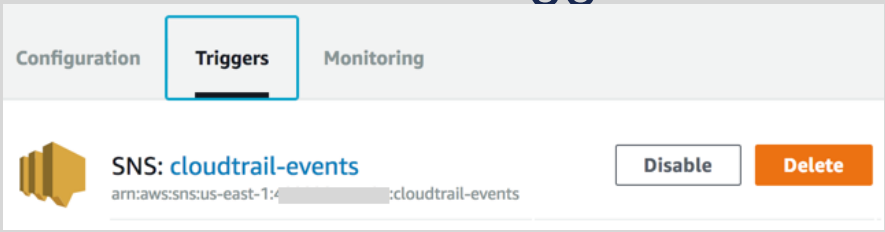

CloudTrail logs from S3 to Elasticsearch

```
try:
    response = s3.get_object(Bucket=s3Bucket, Key=s3ObjectKey)
    content =
gzip.GzipFile(fileobj=StringIO(response['Body'].read())).read()
for record in json.loads(content)['Records']:
    recordJson = json.dumps(record)
    logger.info(recordJson)
    indexName = 'ct-' + datetime.datetime.now().strftime("%Y-%m-%d")
    res = es.index(index=indexName, doc_type='record', id=record['eventID'],
body=recordJson)
    logger.info(res)
return True
```

How did we use this CloudTrail log data that is in Elasticsearch?

Top 10 AWS API calls (per 10 mins) dashboard

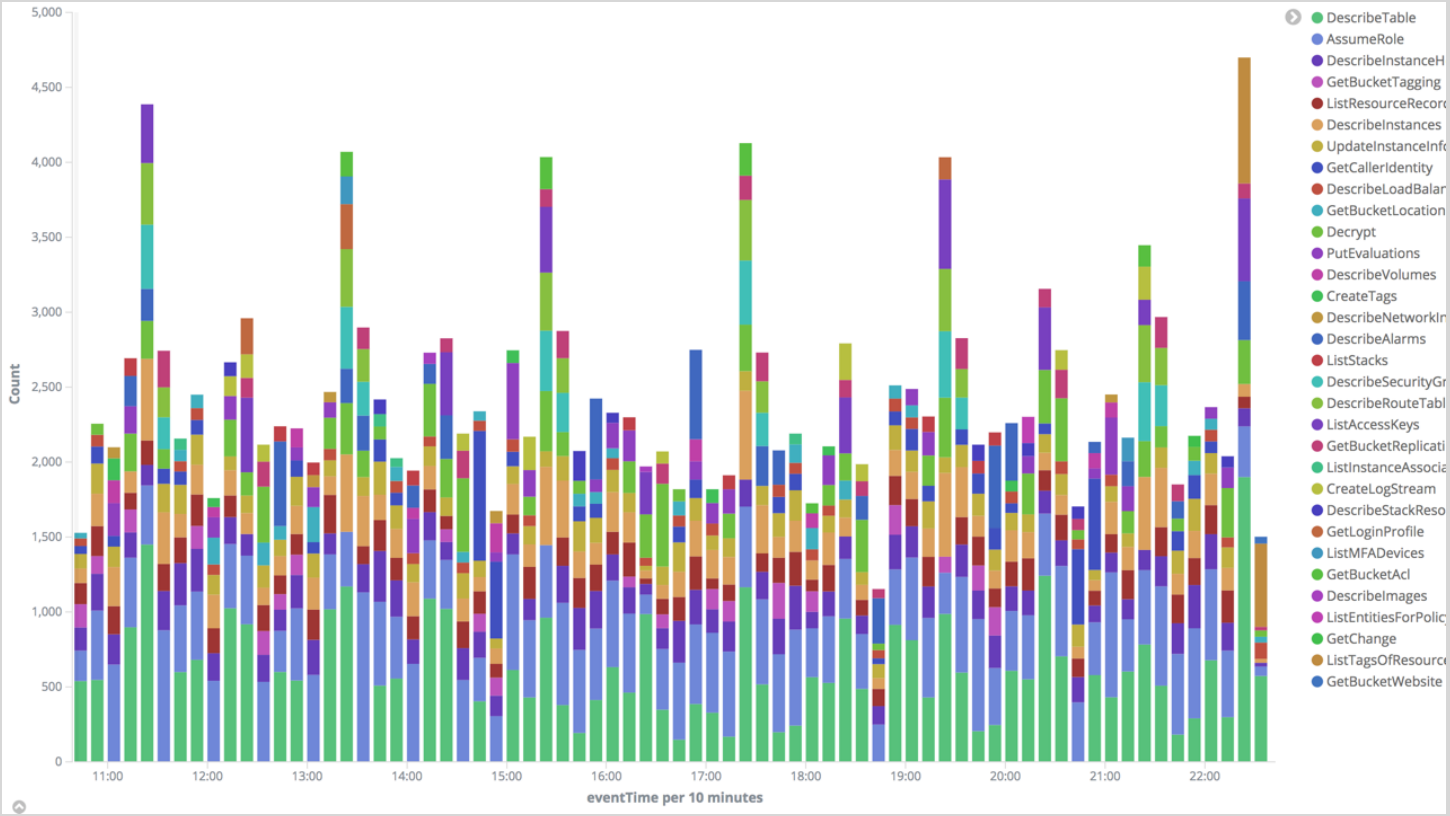

This solution is open sourced at

https://github.com/ExpediaDotCom/cloudtrail-log-analytics

as a serverless application

aws SUMMIT

Distributed tracing platform using Elasticsearch Service

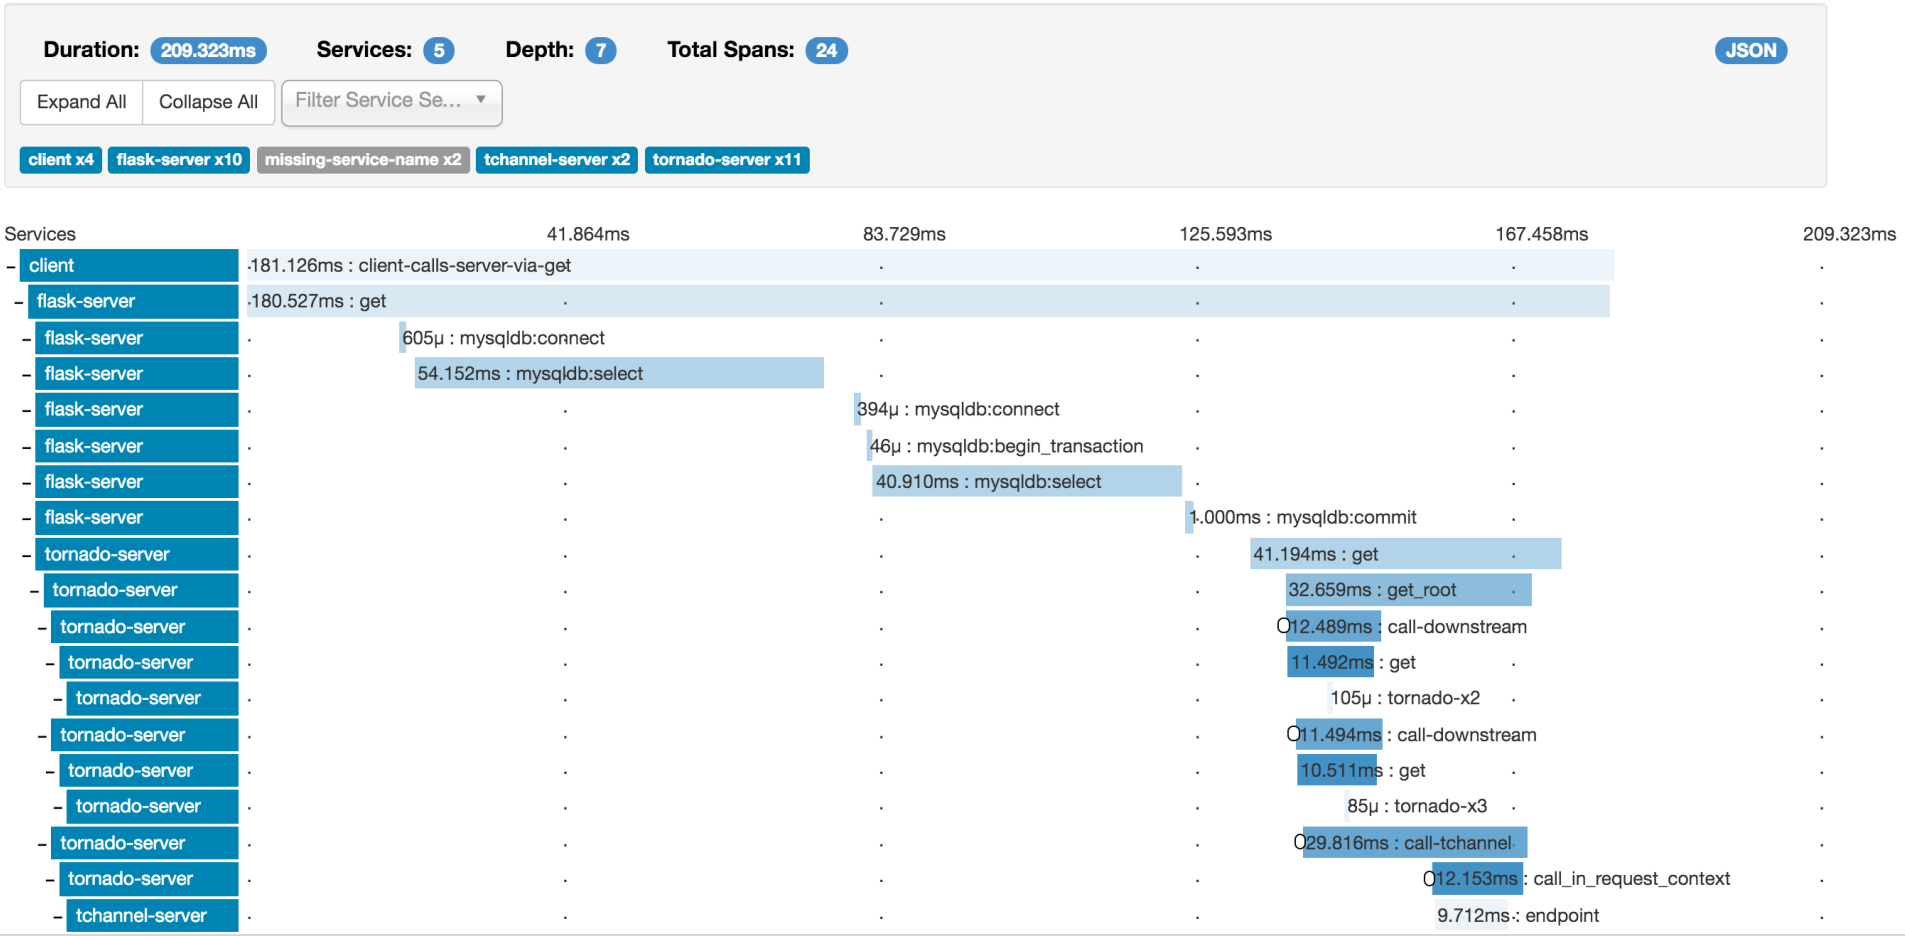

Distributed tracing platform using Elasticsearch

Microservices

Distributed tracing platform using Elasticsearch

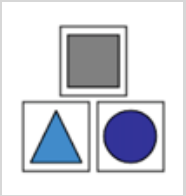

Trace

telemetry

Microservices

Java library

Distributed tracing platform using Elasticsearch

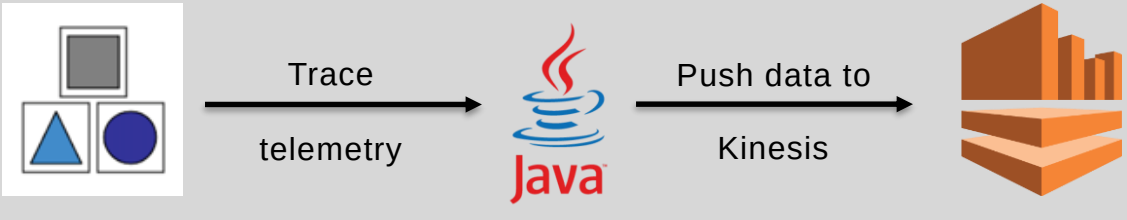

Distributed tracing platform using Elasticsearch

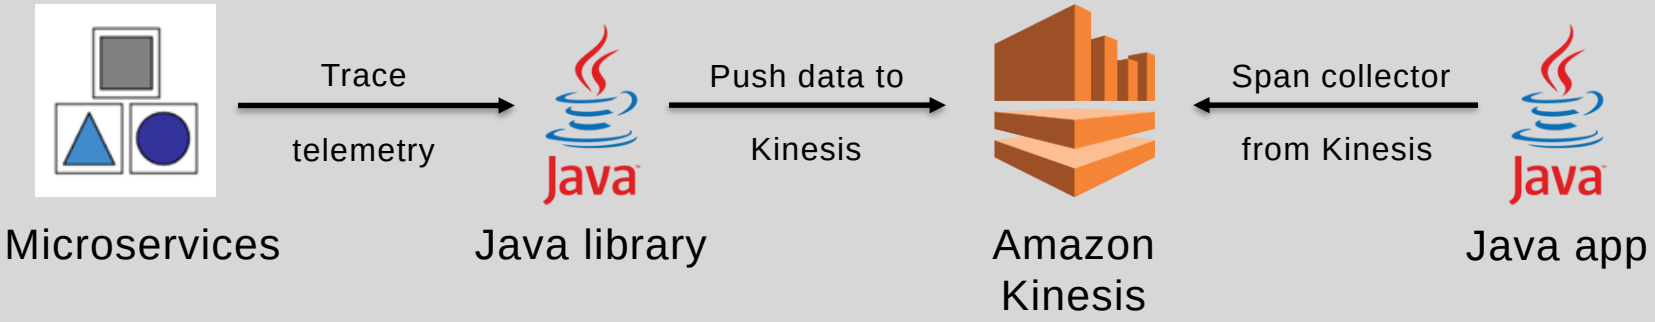

Distributed tracing platform using Elasticsearch

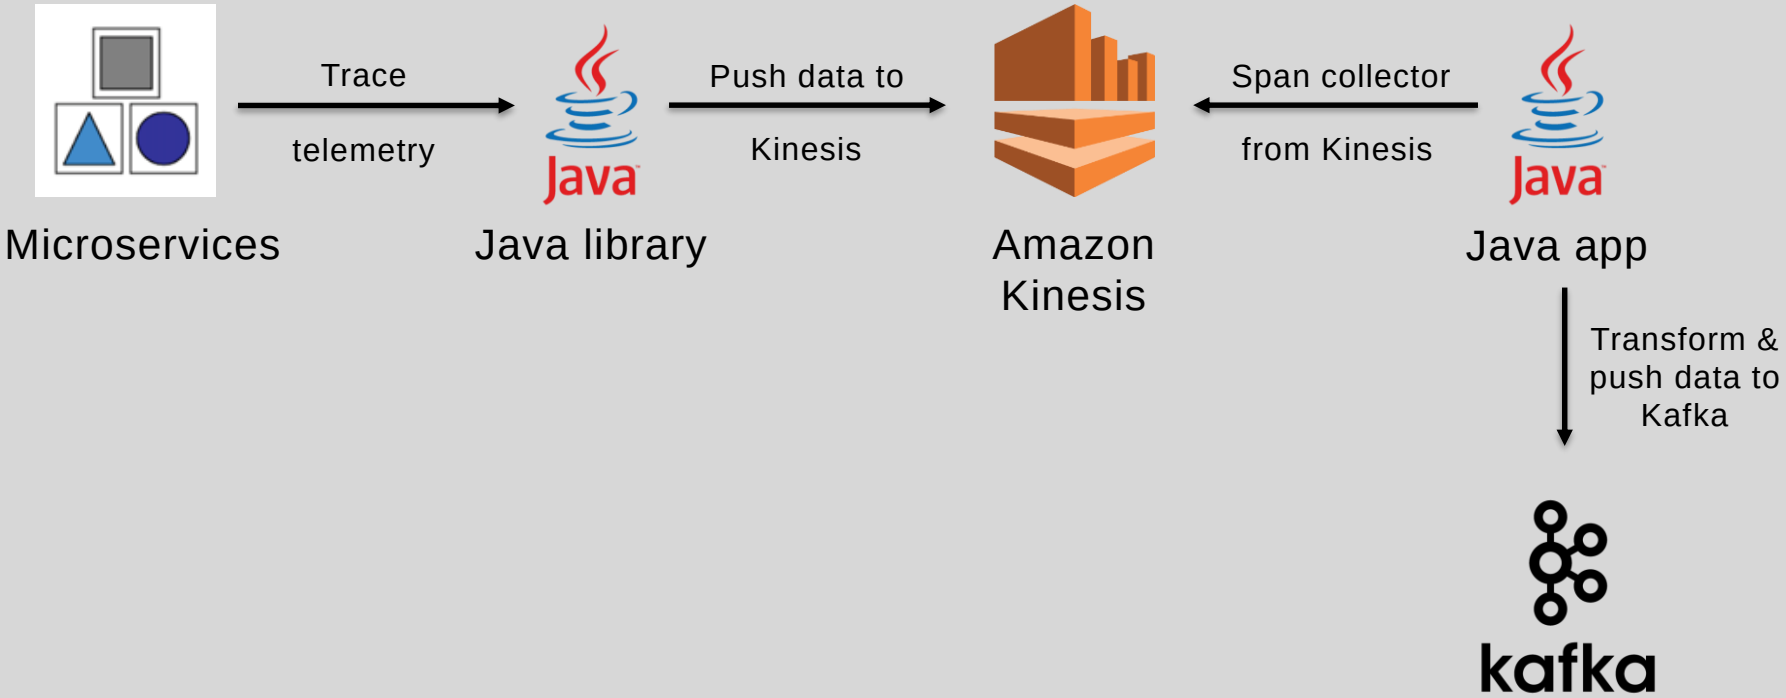

Distributed tracing platform using Elasticsearch

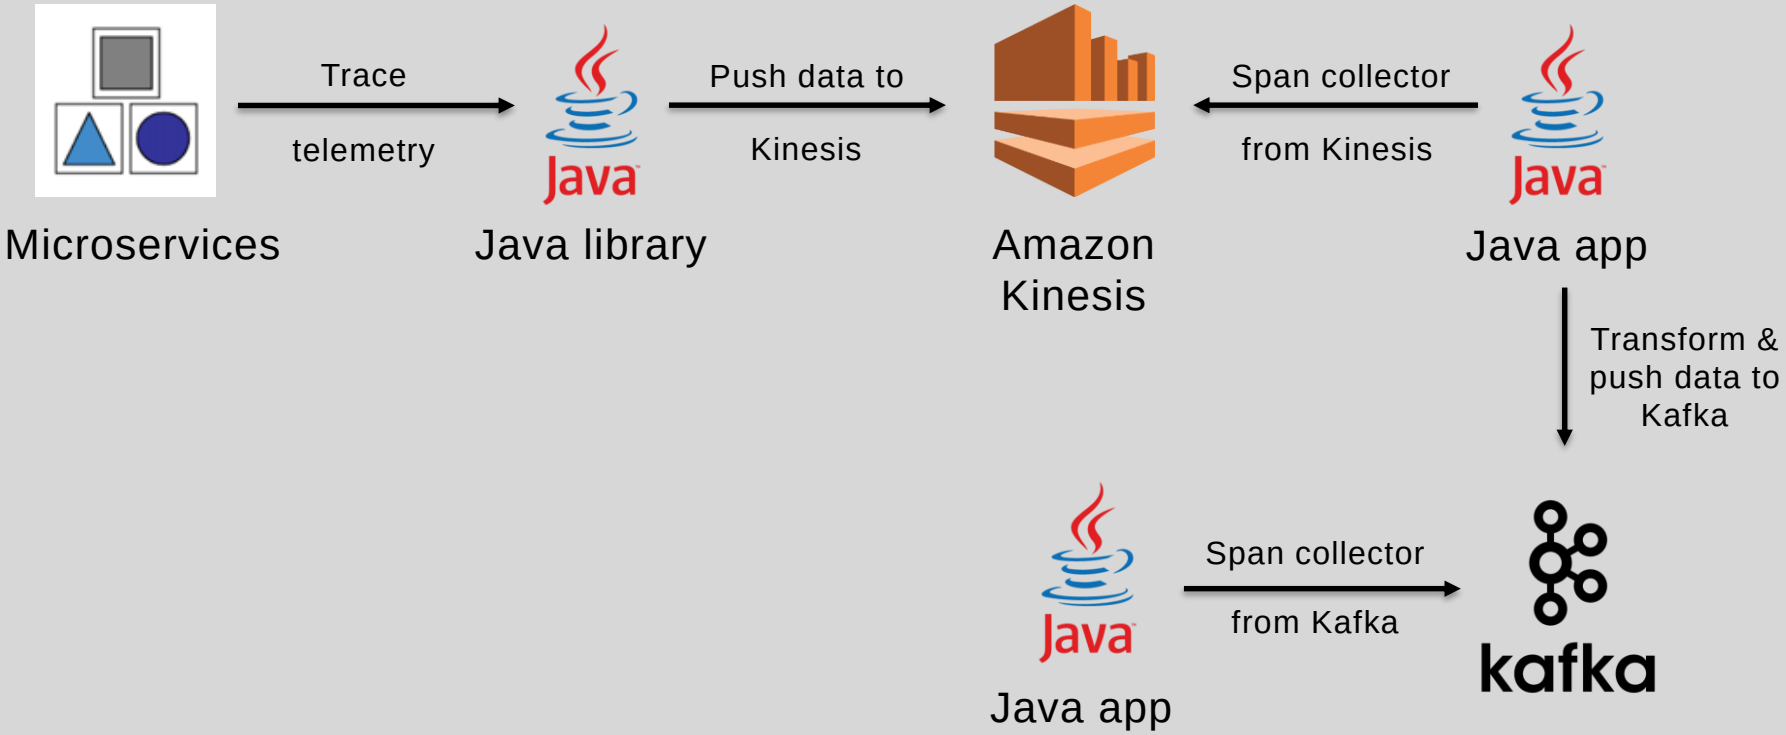

Distributed tracing platform using Elasticsearch

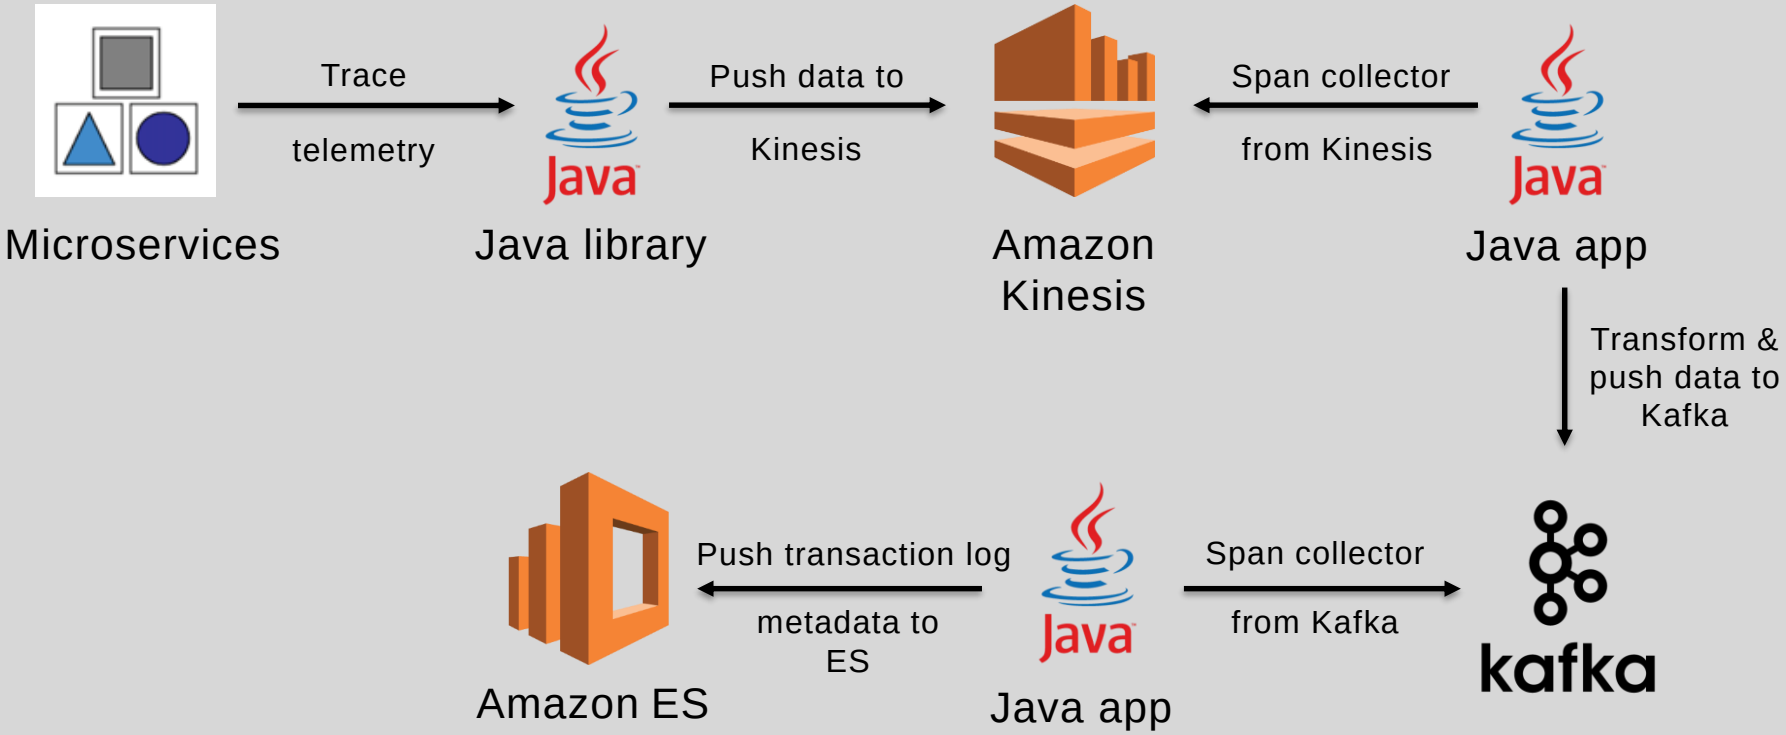

Distributed tracing platform using Elasticsearch

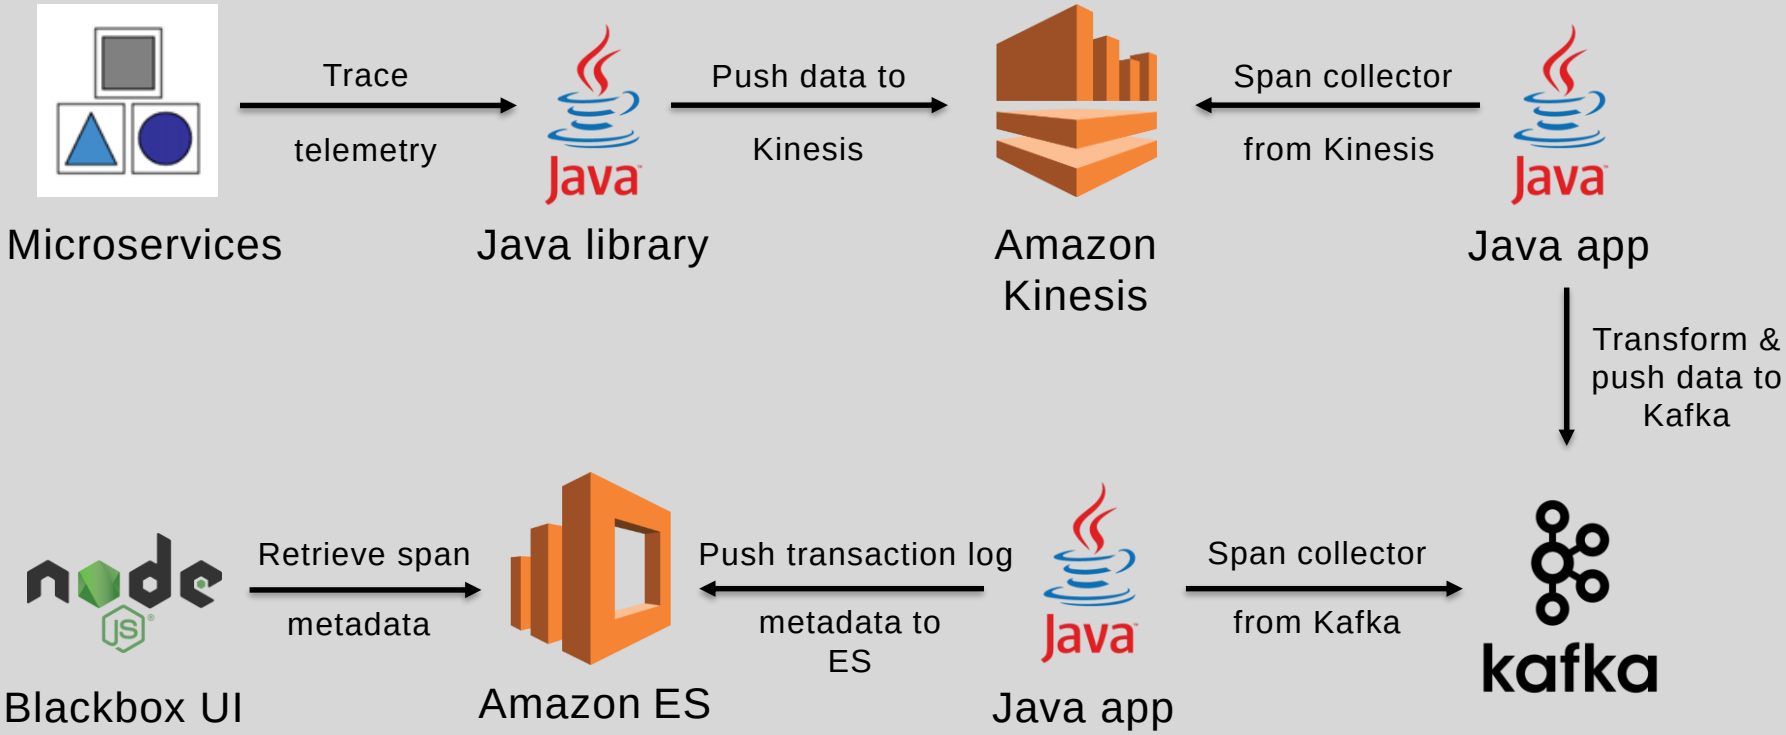

Haystack document in Elasticsearch


```
{
      "_index": "transactions-2017-10-26-04",
      "_type": "transactions_logs",
      "_id": "AV9W9lCq5Jdrc3UjOvCg",
      "_score": null,
      "_source": {
            "transactionid": "5e66cad8-d7ea-49e8-94c8-24d2298d4cdc"
      },
      "fields": {
            "startTime": [
                  1508992503395
            ]
      },
      "sort": [
            1508992503395
      ]
}
```

How is Elasticsearch used in distributed tracing?

Time-based queries for traces

Filtering traces by services

Things to keep in mind

- Scaling of cluster results in a new cluster with the data being synchronized
- Monitor and optimize the cluster yourself
- No upgrade button between Elasticsearch versions (wish list)
- Monitoring doesn’t show how much disk space in use (wish list)

Wrap up

Benefits of using Amazon Elasticsearch Service:

- Easy to set up
- Setup for high availability
- Security
- Monitoring and backup

Please complete the session survey in the summit mobile app.

aws SUMMIT

Thank you!

@this_is_kuldeep

Submit Session Feedback

1. Tap the **Schedule** icon.

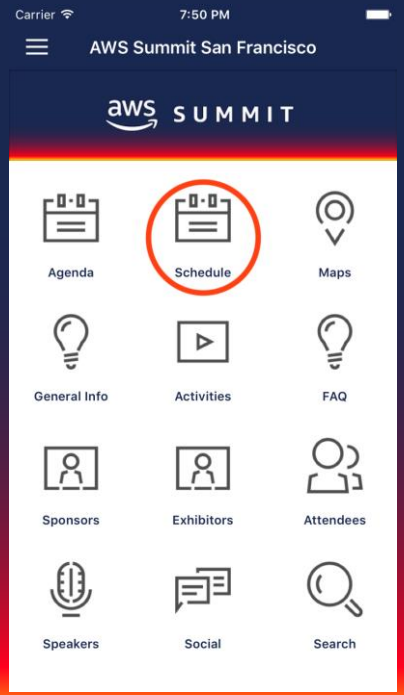

2. Select the session you attended.

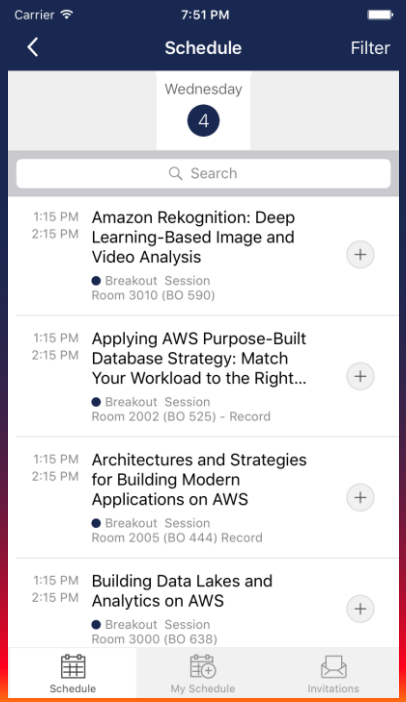

3. Tap **Session Evaluation** to submit your feedback.