

Easy, Scalable, Fault-tolerant Stream Processing with Structured Streaming

Tathagata “TD” Das

 @tathadas

Spark Summit Europe 2017
25th October, Dublin



About Me

Started Spark Streaming project in AMPLab, UC Berkeley

Currently focused on building Structured Streaming

Member of the Apache Spark PMC

Software Engineer at Databricks

building robust
stream processing
apps is hard

Complexities in stream processing

COMPLEX DATA

Diverse data formats
(json, avro, binary, ...)

Data can be dirty,
late, out-of-order

COMPLEX WORKLOADS

Combining streaming with
interactive queries

Machine learning

COMPLEX SYSTEMS

Diverse storage systems
(Kafka, S3, Kinesis, RDBMS, ...)

System failures

Structured Streaming

stream processing on Spark SQL engine

fast, scalable, fault-tolerant

rich, unified, high level APIs

deal with *complex data* and *complex workloads*

rich ecosystem of data sources

integrate with many *storage systems*

you
should not have to
reason about streaming

you
should write simple queries

&

Spark
should continuously update the answer



Streaming word count

Anatomy of a Streaming Word Count

```
spark.readStream  
  .format("kafka")  
  .option("subscribe", "input")  
  .load()
```



Source

- Specify one or more locations to read data from
- Built in support for Files/Kafka/Socket, pluggable.
- Can include multiple sources of different types using `union()`

Anatomy of a Streaming Query

```
spark.readStream
  .format("kafka")
  .option("subscribe", "input")
  .load()
  .groupBy('value.cast("string") as 'key)
  .agg(count("*") as 'value)
```

}

Transformation

- Using DataFrames, Datasets and/or SQL.
- Catalyst figures out how to execute the transformation incrementally.
- Internal processing always exactly-once.


Spark automatically streamifies!

```
input = spark.readStream
    .format("kafka")
    .option("subscribe", "topic")
    .load()


result = input
    .select("device", "signal")
    .where("signal > 15")

result.writeStream
    .format("parquet")
    .start("dest-path")
```


DataFrames,
Datasets, SQL



Logical
Plan



Optimized
Physical Plan



Series of Incremental
Execution Plans

Spark SQL converts batch-like query to a series of incremental execution plans operating on new batches of data

Anatomy of a Streaming Query

```
spark.readStream
  .format("kafka")
  .option("subscribe", "input")
  .load()
  .groupBy('value.cast("string") as 'key)
  .agg(count("*") as 'value)
  .writeStream
  .format("kafka")
  .option("topic", "output")
```



Sink

- Accepts the output of each batch.
- When supported sinks are transactional and exactly once (Files).
- Use foreach to execute arbitrary code.

Anatomy of a Streaming Query

```
spark.readStream
  .format("kafka")
  .option("subscribe", "input")
  .load()
  .groupBy('value.cast("string") as 'key)
  .agg(count("*") as 'value)
  .writeStream
  .format("kafka")
  .option("topic", "output")
  .trigger("1 minute")
  .outputMode("update")
```

Output mode – What's output

- Complete – Output the whole answer every time
- Update – Output changed rows
- Append – Output new rows only



Trigger – When to output

- Specified as a time, eventually supports data size
- No trigger means as fast as possible

Anatomy of a Streaming Query

```
spark.readStream
  .format("kafka")
  .option("subscribe", "input")
  .load()
  .groupBy('value.cast("string") as 'key)
  .agg(count("*") as 'value)
  .writeStream
  .format("kafka")
  .option("topic", "output")
  .trigger("1 minute")
  .outputMode("update")
  .option("checkpointLocation", "...")
  .start()
```

}

Checkpoint

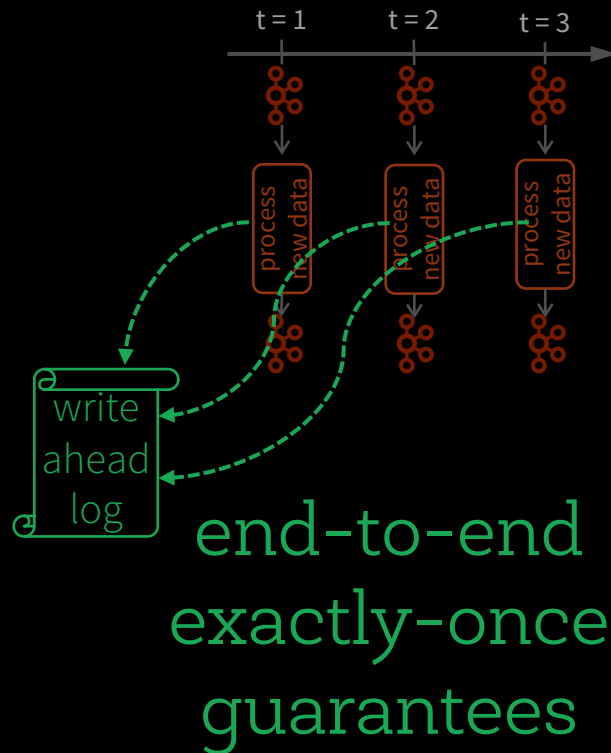
- Tracks the progress of a query in persistent storage
- Can be used to restart the query if there is a failure

Fault-tolerance with Checkpointing

Checkpointing – tracks progress (offsets) of consuming data from the source and intermediate state.


Offsets and metadata saved as JSON

Can resume after changing your streaming transformations



Performance: **YAHOO!** Benchmark

Structured Streaming reuses
the **Spark SQL Optimizer**
and **Tungsten Engine**.



Read more details in our [blog post](#)



Complex Streaming ETL

Traditional ETL



Raw, dirty, un/semi-structured is data dumped as files

Periodic jobs run every few hours to convert raw data to structured data ready for further analytics


Traditional ETL



Hours of delay before taking decisions on latest data

Unacceptable when time is of essence
[intrusion detection, anomaly detection, etc.]

Streaming ETL w/ Structured Streaming



Structured Streaming enables raw data to be available as structured data as soon as possible

Streaming ETL w/ Structured Streaming

Example

Json data being received in Kafka

Parse nested json and flatten it

Store in structured Parquet table

Get end-to-end failure guarantees

```
val rawData = spark.readStream
  .format("kafka")
  .option("kafka.bootstrap.servers", ...)
  .option("subscribe", "topic")
  .load()

val parsedData = rawData
  .selectExpr("cast (value as string) as json")
  .select(from_json("json", schema).as("data"))
  .select("data.*")

val query = parsedData.writeStream
  .option("checkpointLocation", "/checkpoint")
  .partitionBy("date")
  .format("parquet")
  .start("/parquetTable")
```

Reading from Kafka

Specify options to configure

How?

`kafka.bootstrap.servers => broker1,broker2`

```
val rawData = spark.readStream
  .format("kafka")
  .option("kafka.bootstrap.servers",...)
  .option("subscribe", "topic")
  .load()
```

What?

```
subscribe          => topic1,topic2,topic3    // fixed list of topics
subscribePattern    => topic*                  // dynamic list of topics
assign              => {"topicA":[0,1] }        // specific partitions
```


Where?

```
startingOffsets => latest(default) / earliest / {"topicA":{"0":23,"1":345} }
```

Reading from Kafka

rawData dataframe has
the following columns

```
val rawData = spark.readStream  
    .format("kafka")  
    .option("kafka.bootstrap.servers", "...")  
    .option("subscribe", "topic")  
    .load()
```



| key | value | topic | partition | offset | timestamp |
|-----------------|-----------------|----------|-----------|--------|------------|
| <i>[binary]</i> | <i>[binary]</i> | "topicA" | 0 | 345 | 1486087873 |
| <i>[binary]</i> | <i>[binary]</i> | "topicB" | 3 | 2890 | 1486086721 |

Transforming Data

Cast binary *value* to string
Name it column *json*

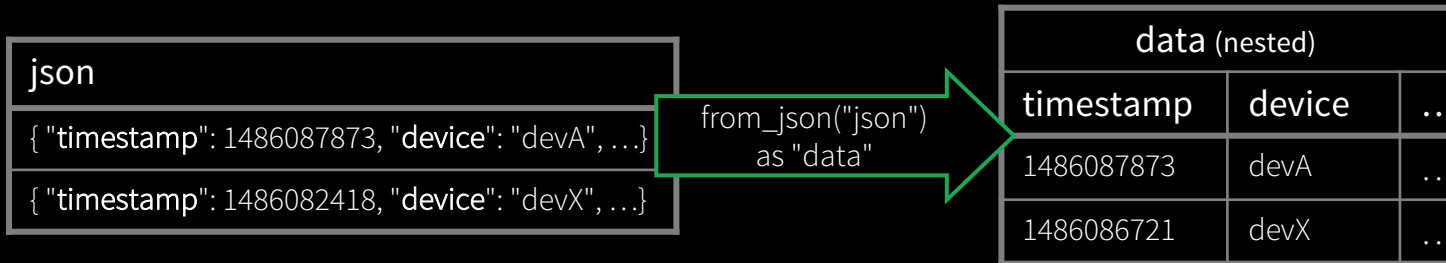
```
val parsedData = rawData
  .selectExpr("cast (value as string) as json")
  .select(from_json("json", schema).as("data"))
  .select("data.*")
```


Transforming Data

Cast binary *value* to string
Name it column *json*

```
val parsedData = rawData
  .selectExpr("cast (value as string) as json")
  .select(from_json("json", schema).as("data"))
  .select("data.*")
```

Parse *json* string and expand into
nested columns, name it *data*



Transforming Data

Cast binary *value* to string
Name it column *json*

```
val parsedData = rawData
  .selectExpr("cast (value as string) as json")
  .select(from_json("json", schema).as("data"))
  .select("data.*")
```

Parse *json* string and expand into
nested columns, name it *data*

Flatten the nested columns

| data (nested) | | |
|---------------|--------|-----|
| timestamp | device | ... |
| 1486087873 | devA | ... |
| 1486086721 | devX | ... |

select("data.*")

| (not nested) | | |
|--------------|--------|-----|
| timestamp | device | ... |
| 1486087873 | devA | ... |
| 1486086721 | devX | ... |

Transforming Data

Cast binary *value* to string
Name it column *json*

```
val parsedData = rawData
  .selectExpr("cast (value as string) as json")
  .select(from_json("json", schema).as("data"))
  .select("data.*")
```

Parse *json* string and expand into
nested columns, name it *data*

Flatten the nested columns

powerful built-in APIs to
perform complex data
transformations

from_json, to_json, explode, ...
100s of functions

(see [our blog post](#))

Writing to Parquet

Save parsed data as Parquet table in the given path

Partition files by date so that future queries on time slices of data is fast

e.g. query on last 48 hours of data

```
val query = parsedData.writeStream
  .option("checkpointLocation", ...)
  .partitionBy("date")
  .format("parquet")
  .start("/parquetTable")
```


Checkpointing

Enable checkpointing by
setting the checkpoint
location to save offset logs

start actually starts a
continuous running
StreamingQuery in the
Spark cluster

```
val query = parsedData.writeStream  
  .option("checkpointLocation", ...)   
  .format("parquet")  
  .partitionBy("date")  
  .start("/parquetTable/")
```

Streaming Query



```
val query = parsedData.writeStream  
  .option("checkpointLocation", ...)   
  .format("parquet")  
  .partitionBy("date")  
  .start("/parquetTable")
```

query is a handle to the continuously running StreamingQuery

Used to monitor and manage the execution

Data Consistency on Ad-hoc Queries



Data available for complex, ad-hoc analytics within seconds

Parquet table is updated atomically, ensures *prefix integrity*

Even if distributed, ad-hoc queries will see either all updates from streaming query or none, read more in our blog

<https://databricks.com/blog/2016/07/28/structured-streaming-in-apache-spark.html>

More Kafka Support

Write out to Kafka

Dataframe must have binary fields
named key and value

```
result.writeStream  
  .format("kafka")  
  .option("topic", "output")  
  .start()
```

Direct, interactive and batch queries on Kafka

Makes Kafka even more powerful
as a storage platform!

```
val df = spark  
  .read           // not readStream  
  .format("kafka")  
  .option("subscribe", "topic")  
  .load()
```

```
df.registerTempTable("topicData")  
spark.sql("select value from topicData")
```

Added to Spark 2.2

Amazon Kinesis

Configure with options (similar to Kafka)
Available with Databricks Runtime

How?

```
region => us-west-2 / us-east-1 / ...  
awsAccessKey (optional) => AKIA...  
awsSecretKey (optional) => ...
```

What?

```
streamName => name-of-the-stream
```

Where?

```
initialPosition => latest(default) / earliest / trim_horizon
```

```
spark.readStream  
  .format("kinesis")  
  .option("streamName", "myStream")  
  .option("region", "us-west-2")  
  .option("awsAccessKey", ...)  
  .option("awsSecretKey", ...)  
  .load()
```



Working With Time

Event Time

Many use cases require aggregate statistics by event time

E.g. what's the #errors in each system in the 1 hour windows?

Many challenges

Extracting event time from data, handling late, out-of-order data

DStream APIs were insufficient for event-time stuff

Event time Aggregations

Windowing is just another type of grouping in Struct.
Streaming

number of records every hour

```
parsedData  
  .groupBy(window("timestamp", "1 hour"))  
  .count()
```

avg signal strength of each
device every 10 mins

```
parsedData  
  .groupBy(  
    "device",  
    window("timestamp", "10 mins"))  
  .avg("signal")
```

Support UDAFs!

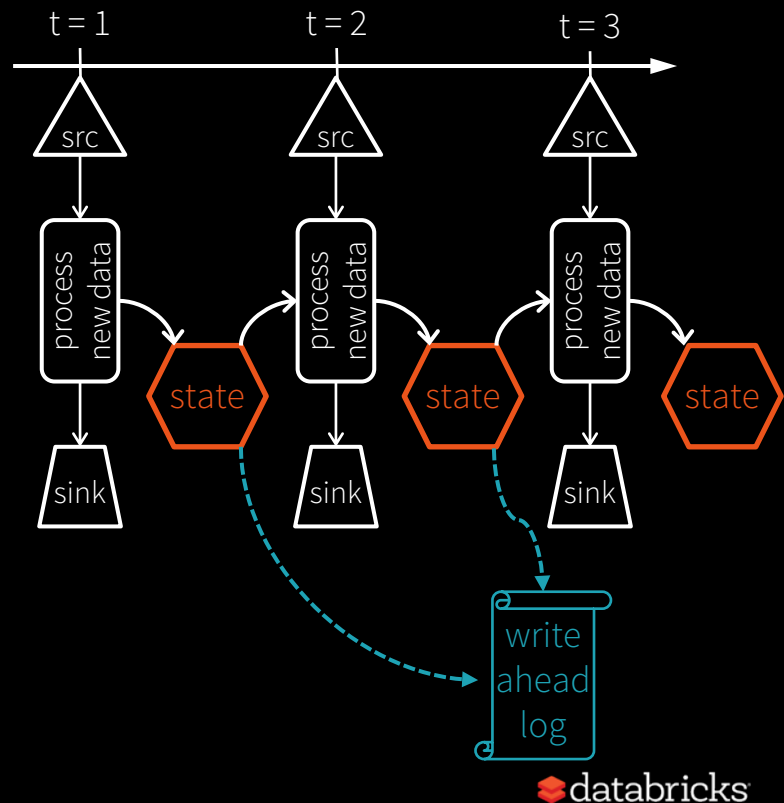
Stateful Processing for Aggregations

Aggregates has to be saved as **distributed state** between triggers

Each trigger reads previous state and writes updated state


State stored in memory, backed by *write ahead log* in HDFS/S3

Fault-tolerant, **exactly-once guarantee!**



Automatically handles Late Data

Keeping state allows late data to update counts of old windows



But size of the state increases indefinitely if old windows are not dropped

red = state updated with late data

Watermarking

Watermark - moving threshold of how late data is expected to be and when to drop old state

Trails behind **max event time** seen by the engine

Watermark delay = trailing gap




Watermarking

Data newer than watermark may be late, but allowed to aggregate

Data older than watermark is "too late" and dropped

Windows older than watermark automatically deleted to limit the amount of intermediate state

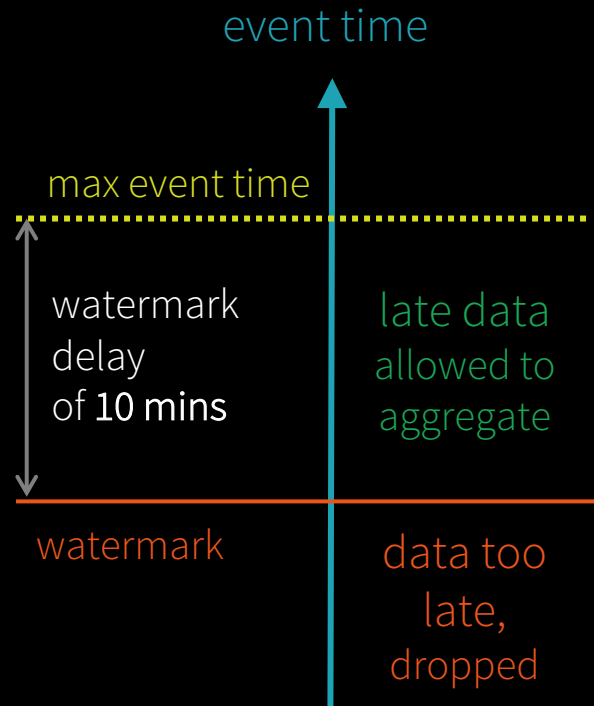


Watermarking


```
parsedData  
  .withWatermark("timestamp", "10 minutes")  
  .groupBy(window("timestamp", "5 minutes"))  
  .count()
```

Useful only in stateful operations

Ignored in non-stateful streaming queries and batch queries



Watermarking



Clean separation of concerns

Query Semantics

separated from

Processing Details

```
parsedData
  .withWatermark("timestamp", "10 minutes")
  .groupBy(window("timestamp", "5 minutes"))
  .count()
  .writeStream
  .trigger("10 seconds")
  .start()
```

Clean separation of concerns

Query Semantics

How to group data by time?
(same for batch & streaming)

```
parsedData
  .withWatermark("timestamp", "10 minutes")
  .groupBy(window("timestamp", "5 minutes"))
  .count()
  .writeStream
  .trigger("10 seconds")
  .start()
```

Processing Details

Clean separation of concerns

Query Semantics

How to group data by time?
(same for batch & streaming)

```
parsedData
  .withWatermark("timestamp", "10 minutes")
  .groupBy(window("timestamp", "5 minutes"))
  .count()
  .writeStream
  .trigger("10 seconds")
  .start()
```

Processing Details

How late can data be?

Clean separation of concerns

Query Semantics

How to group data by time?
(same for batch & streaming)

```
parsedData
  .withWatermark("timestamp", "10 minutes")
  .groupBy(window("timestamp", "5 minutes"))
  .count()
  .writeStream
  .trigger("10 seconds")
  .start()
```

Processing Details

How late can data be?

How often to emit updates?

Other Interesting Operations

Streaming Deduplication

Watermarks to limit state

```
parsedData.dropDuplicates("eventId")
```

Joins

Stream-batch joins supported,
stream-stream joins coming in 2.3

```
parsedData.join(batchData, "device")
```

Arbitrary Stateful Processing

[map|flatMap]GroupsWithState

```
ds.groupByKey(_.id)  
  .mapGroupsWithState  
    (timeoutConf)  
    (mappingWithStateFunc)
```

[See my [other talk](#) at 4:20 PM, today for a deep dive into stateful ops]

Monitoring Streaming Queries

Get last progress of the streaming query

- Current input and processing rates

- Current processed offsets

- Current state metrics

```
streamingQuery.lastProgress()
```

```
{ ...  
  "inputRowsPerSecond" : 10024.225210926405,  
  "processedRowsPerSecond" : 10063.737001006373,  
  "durationMs" : { ... },  
  "sources" : [ ... ],  
  "sink" : { ... }  
  ...  
}
```

Get progress asynchronously through by registering your own StreamingQueryListener

```
new StreamingQueryListener {  
  def onQueryStart(...)   
  def onQueryProgress(...)   
  def onQueryTermination(...)   
}
```


Dropwizard Metrics

Metrics into Ganglia, Graphite, etc.

Enabled using SQL configuration

```
spark.conf.set("spark.sql.streaming.metricsEnabled", "true")
```



Building Complex Continuous Apps

Metric Processing @ databricks®

Events generated by user actions (logins, clicks, spark job updates)



ETL

Clean, normalize and store historical data



Dashboards

Analyze trends in usage as they occur



Alerts

Notify engineers of critical issues



Ad-hoc Analysis

Diagnose issues when they occur

Metric Processing @ databricks®

Difficult with only streaming frameworks



ETL



Dashboards



Alerts



Ad-hoc Analysis



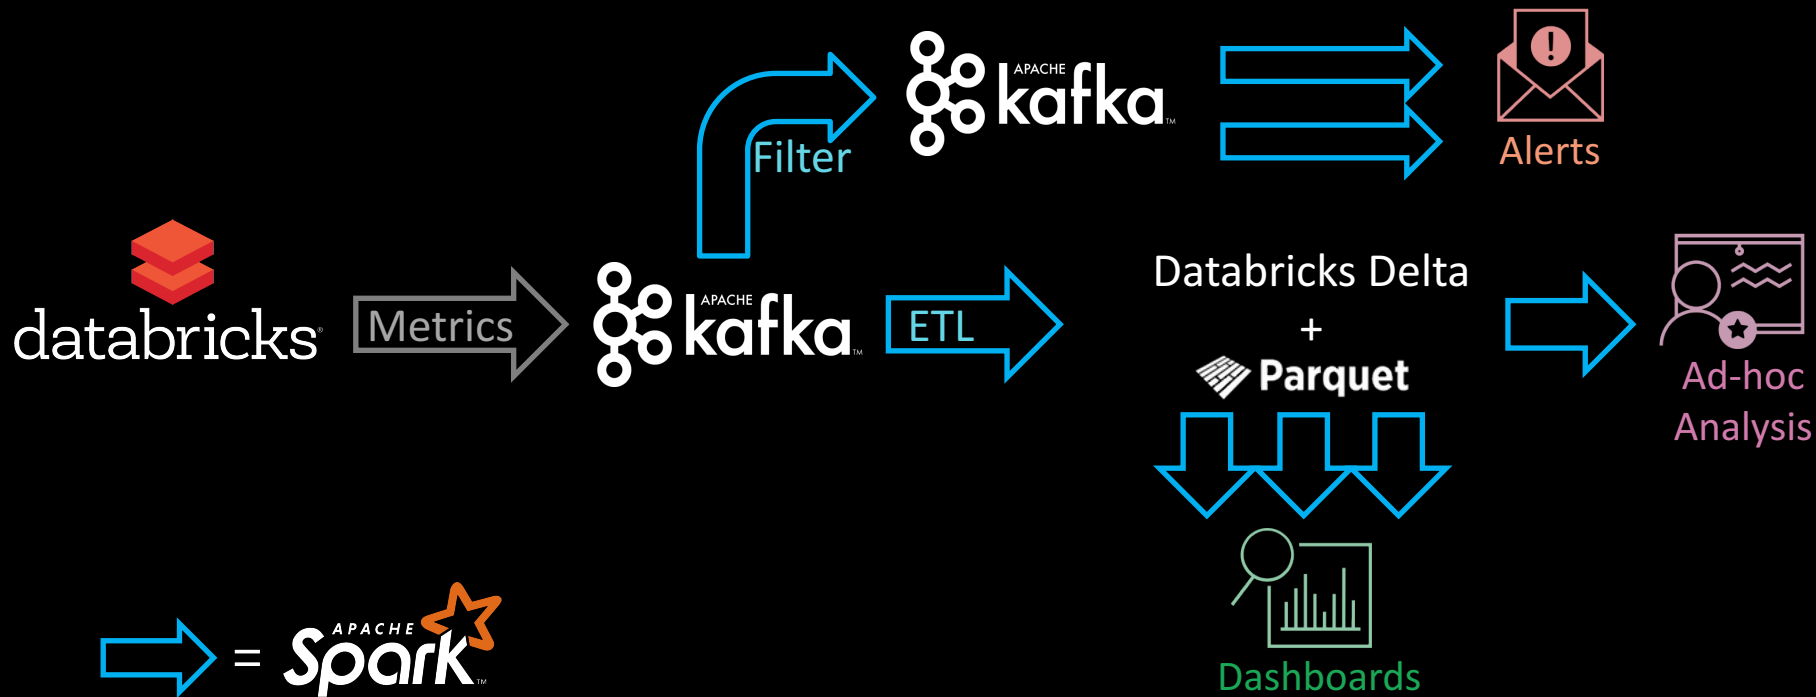
kafka

Limited retention in
streaming storage

Inefficient for ad-hoc queries

Hard for novice users
(limited SQL support)

Metric Processing @ databricks®



Read from kafka




```
rawLogs = spark.readStream  
  .format("kafka")  
  .option("kafka.bootstrap.servers", ...)   
  .option("subscribe", "rawLogs")  
  .load()
```

```
augmentedLogs = rawLogs  
  .withColumn("msg",  
    from_json($"value".cast("string"),  
    schema))  
  .select("timestamp", "msg.*")  
  .join(table("customers"), ["customer_id"])
```

DataFrames can be
reused for multiple
streams

Can build libraries of
useful DataFrames and
share code between
applications

Write to Parquet



Databricks Delta
+
 Parquet

Store augmented stream as efficient columnar data for later processing

Latency: ~1 minute

```
augmentedLogs  
  .repartition(1)  
  .writeStream  
  .format("delta")  
  .option("path", "/data/metrics")  
  .trigger("1 minute")  
  .start()
```

Buffer data and
write one large file
every minute for
efficient reads

Dashboards

Always up-to-date visualizations of important business trends

Latency: ~1 minute to hours (configurable)

```
logins = spark.readStream.parquet("/data/metrics")  
    .where("metric = 'login'")  
    .groupBy(window("timestamp", "1 minute"))  
    .count()
```

```
display(logins)          // visualize in Databricks notebooks
```

Databricks Delta




Dashboards

Filter and write to kafka™

Forward filtered and augmented events back to Kafka

Latency: ~100 ms average



```
filteredLogs = augmentedLogs
  .where("eventType = 'clusterHeartbeat'")
  .selectExpr("to_json(struct("*")) as value")
```

```
filteredLogs.writeStream
  .format("kafka")
  .option("kafka.bootstrap.servers", ...)
  .option("topic", "clusterHeartbeats")
  .start()
```

to_json() to convert columns back into json string, and then save as different Kafka topic

Simple Alerts



Alerts

E.g. Alert when Spark cluster load > threshold

Latency: ~100 ms

```
sparkErrors
  .as[ClusterHeartBeat]
  .filter(_.load > 99)
  .writeStream
  .foreach(new PagerdutySink(credentials))
```

notify PagerDuty



Complex Alerts




Alerts

E.g. Monitor health of Spark clusters
using custom **stateful logic**

Latency: ~10 seconds

```
sparkErrors
  .as[ClusterHeartBeat]
  .groupBy(_.id)
  .flatMapGroupsWithState(Update, EventTimeTimeout) {
    (id: Int, events: Iterator[ClusterHeartBeat], state: GroupState[ClusterState]) =>
    ... // check if cluster non-responsive for a while
  }
```

react if no heartbeat
from cluster for 1 min



Ad-hoc Analysis

Trouble shoot problems as they occur with latest information


Latency: ~1 minute

```
SELECT *  
FROM delta.`/data/metrics`  
WHERE level IN ('WARN', 'ERROR')  
      AND customer = "..."  
      AND timestamp < now() - INTERVAL 1 HOUR
```

Databricks Delta



Parquet



Ad-hoc
Analysis

will read latest data
when query executed

Metric Processing @ databricks®

14+ billion records / hour
with 10 nodes

meet diverse latency requirements
as efficiently as possible



Structured Streaming @ databricks®

100s of customer streaming apps
in production on Databricks

Largest app process 10s of trillions
of records per month

Future Direction: Continuous Processing

Continuous processing mode to run without micro-batches

<1 ms latency (same as per-record streaming systems)

No changes to user code

Proposal in [SPARK-20928](#), expected in Spark 2.3

More Info

Structured Streaming Programming Guide

<http://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>

Databricks blog posts for more focused discussions

<https://databricks.com/blog/2016/07/28/structured-streaming-in-apache-spark.html>

<https://databricks.com/blog/2017/01/19/real-time-streaming-etl-structured-streaming-apache-spark-2-1.html>

<https://databricks.com/blog/2017/02/23/working-complex-data-formats-structured-streaming-apache-spark-2-1.html>

<https://databricks.com/blog/2017/04/26/processing-data-in-apache-kafka-with-structured-streaming-in-apache-spark-2-2.html>

<https://databricks.com/blog/2017/05/08/event-time-aggregation-watermarking-apache-sparks-structured-streaming.html>

<https://databricks.com/blog/2017/10/11/benchmarking-structured-streaming-on-databricks-runtime-against-state-of-the-art-streaming-systems.html>

and more to come, stay tuned!!

Try Apache Spark in Databricks!

UNIFIED ANALYTICS PLATFORM

- Collaborative cloud environment
- Free version (community edition)

DATABRICKS RUNTIME 3.0

- Apache Spark - optimized for the cloud
- Caching and optimization layer - DBIO
- Enterprise security - DBES

Try for free today
databricks.com