

Building Robust, Adaptive Streaming Apps with Spark Streaming

Tathagata “TD” Das

 @tathadas

Spark Summit East 2016



Who am I?

Project Management Committee (PMC) member of Spark

Started Spark Streaming in grad school - AMPLab, UC Berkeley

Current technical lead of Spark Streaming

Software engineer at Databricks

Streaming Apps in the Real World

Building a high volume stream processing system in production has many challenges

Fast and Scalable

Spark Streaming is fast, distributed and scalable by design!


Running in production at many places with large clusters and high data volumes

Streaming Apps in the Real World

Building a high volume stream processing system in production has many challenges

Fast and Scalable

Easy to program



Spark Streaming makes it easy to express complex streaming business logic

Interoperates with Spark RDDs, Spark SQL DataFrames/Datasets and MLlib

Streaming Apps in the Real World

Building a high volume stream processing system in production has many challenges

Fast and Scalable

Easy to program

Fault-tolerant



Spark Streaming is fully fault-tolerant and can provide end-to-end semantic guarantees

See my [previous Spark Summit talk](#) for more details

Streaming Apps in the Real World

Building a high volume stream processing system in production has many challenges

Fast and Scalable

Easy to program

Fault-tolerant

Adaptive



Focus of this talk

Adaptive Streaming Apps

Processing conditions can change dynamically

- Sudden surges in data rates

- Diurnal variations in processing load

- Unexpected slowdowns in downstream data stores

Streaming apps should be able to adapt accordingly

Backpressure

Make apps robust
against data surges

Elastic Scaling

Make apps scale with
load variations

Backpressure

Make apps robust
against data surges

Motivation

Stability condition for any streaming app

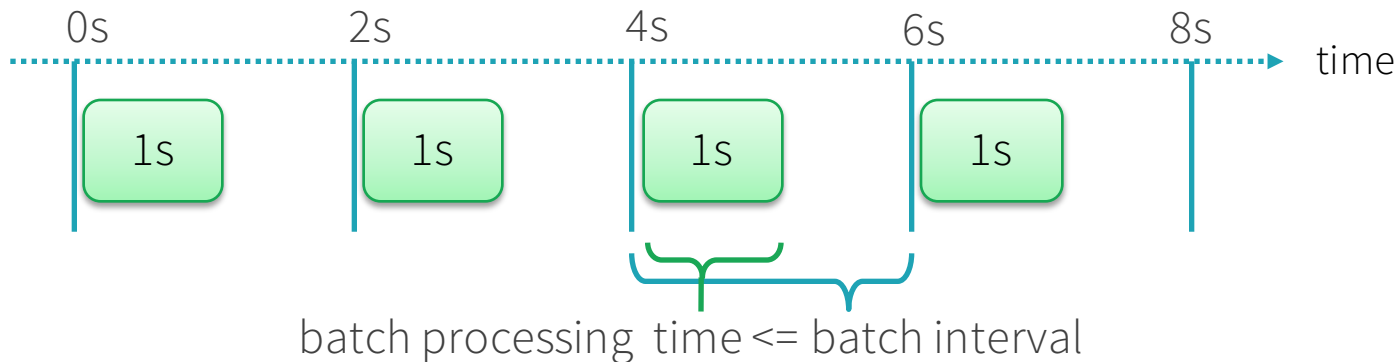
Receive data only as fast as the system can process it

Stability condition for Spark Streaming's "micro-batch" model

Finish processing previous batch before next one arrives

Stable micro-batch operation

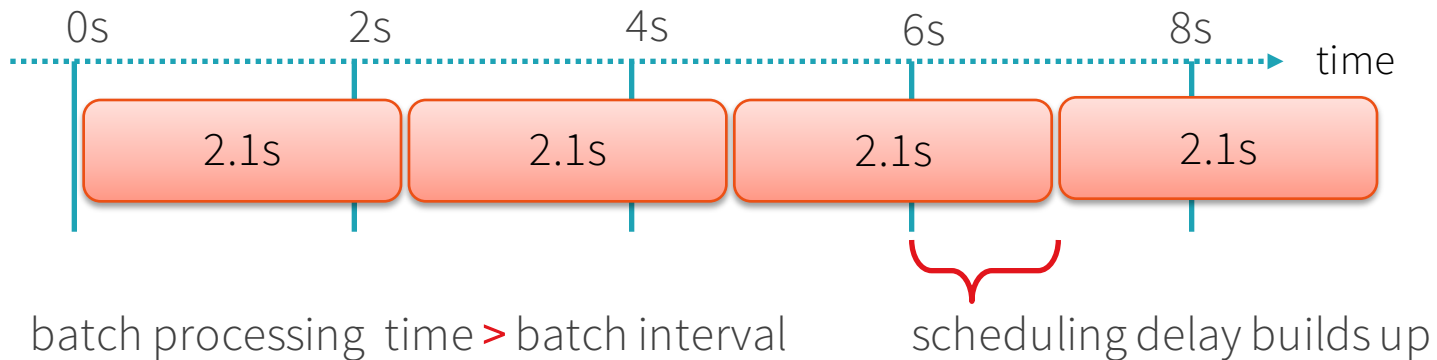
Spark Streaming runs micro-batches at fixed *batch intervals*



Previous batch is processed before next one arrives => **stable**

Unstable micro-batch operation

Spark Streaming runs micro-batches at fixed *batch intervals*




Batches continuously gets delayed and backlogged => **unstable**

Backpressure: Feedback Loop



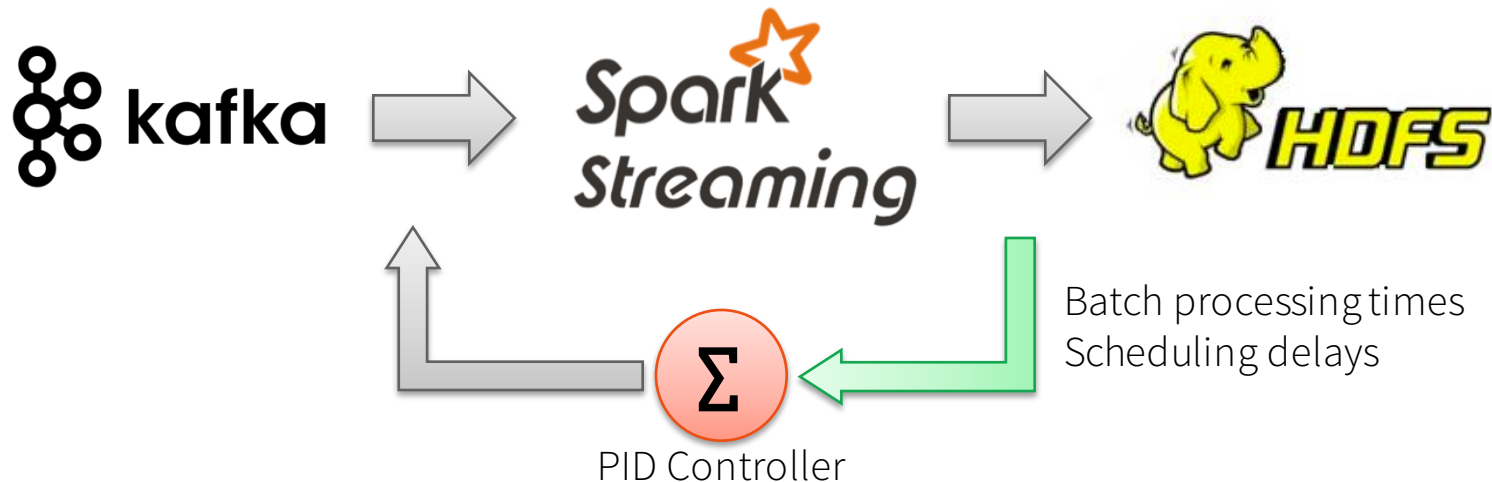
Backpressure introduces a **feedback loop** to dynamically adapt the system and avoid instability

Backpressure: Dynamic rate limiting



Batch processing times and scheduling delays used to continuously estimate current processing rates

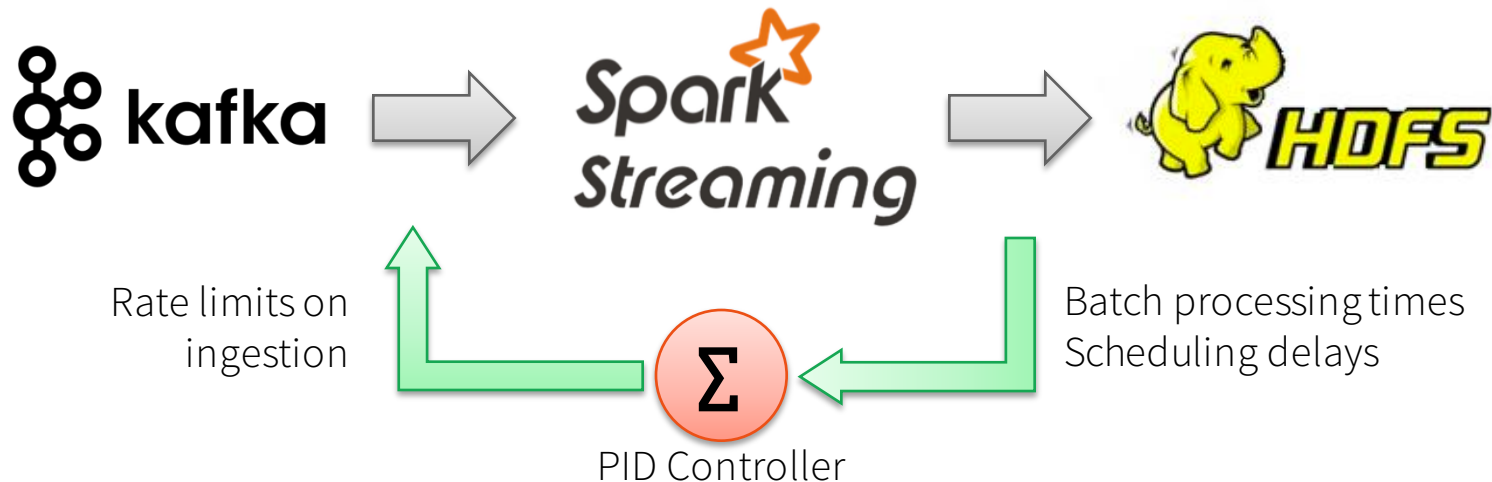
Backpressure: Dynamic rate limiting



Max stable processing rate estimated with **PID Controller theory**

Well known feedback mechanism used in industrial control systems

Backpressure: Dynamic rate limiting



Accordingly, the system dynamically adapts the limits on the data ingestion rates

Backpressure: Dynamic rate limiting



If HDFS ingestion slows down,
processing times increase



SS lowers rate limits to slow
down receiving from Kafka

Data buffered in Kafka, ensures
Spark Streaming stays stable

Backpressure: Configuration

Available since Spark 1.5

Enabled through SparkConf, set

`spark.streaming.backpressure.enabled=true`

Elastic Scaling

Make apps scale with
load variations

Elastic Scaling (aka Dynamic Allocation)

Scaling the number of Spark executors according to the load

Spark already supports Dynamic Allocation for batch jobs


- Scale down if executors are idle

- Scale up if tasks are queueing up


Streaming “micro-batch” jobs need different scaling policy

- No executor is idle for a long time!

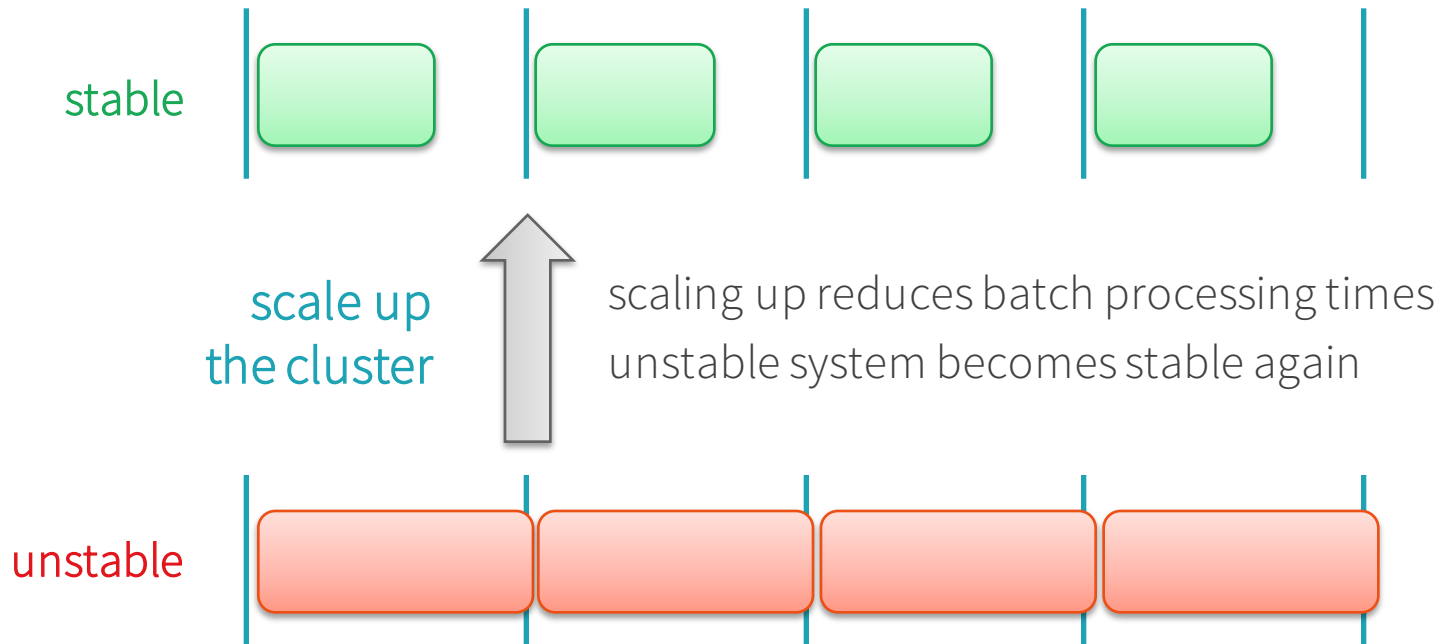
Scaling policy with Streaming



Scaling policy with Streaming



Scaling policy with Streaming



Elastic Scaling



If Kafka gets data faster than
what backpressure allows



SS scales up cluster to
increase processing rate

Data buffered in Kafka starts draining,
allowing app adapt to any data rate

Elastic Scaling: Configuration

Will be available in Spark 2.0

Enabled through SparkConf, set

`spark.streaming.dynamicAllocation.enabled = true`

More parameters will be in the online programming guide

Elastic Scaling: Configuration

Make sure there is enough parallelism to take advantage of max cluster size

- # of partitions in reduce, join, etc.

- # of Kafka partitions

- # of receivers

Gives usual fault-tolerance guarantees with files, Kafka Direct, Kinesis, and receiver-based sources with WAL enabled

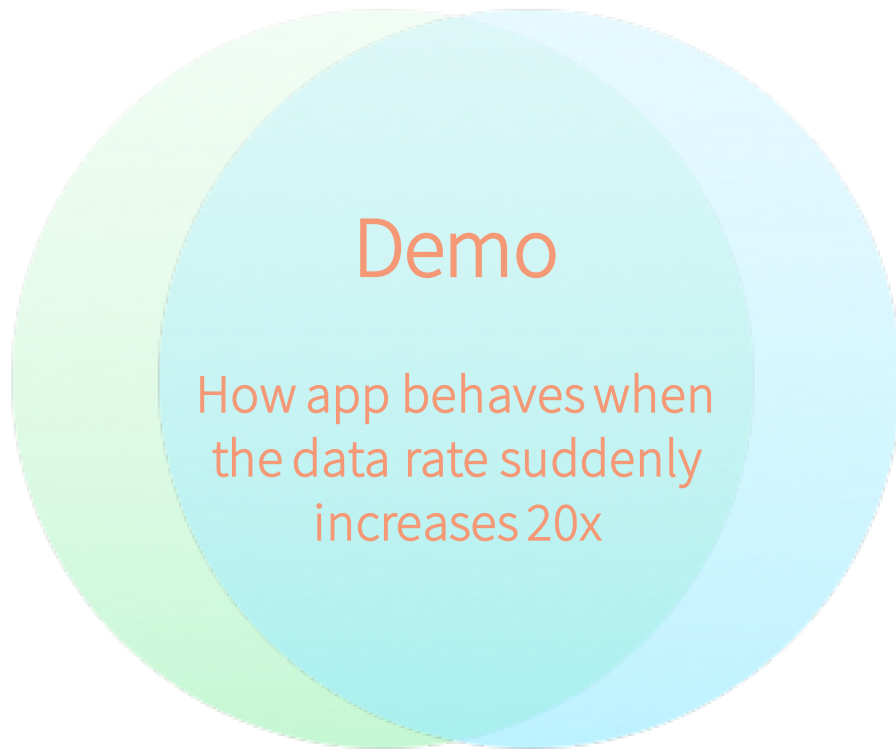


Backpressure



Elastic Scaling






Demo


Processing time increases with data rate, until equal to batch interval

Backpressure limits the ingestion rate lower than 20k recs/sec to keep app stable


Data Ingestion Rate



Processing and Scheduling Delays



Cluster Size




Demo


Elastic Scaling detects heavy load and increases cluster size

Processing times reduces as more resources available


Data Ingestion Rate



Processing and Scheduling Delays



Cluster Size




Demo


Backpressure relaxes limits to allow higher ingestion rate

But still less than 20x as cluster is fully utilized


Data Ingestion Rate



Processing and Scheduling Delays



Cluster Size



Takeaways

Backpressure: makes apps robust to sudden changes

Elastic Scaling: makes apps adapt to slower changes

Backpressure + Elastic Scaling =
Awesome Adaptive Apps