

Magellan – Spark as a Geospatial Analytics Engine

Ram Sriharsha
Hortonworks



SPARK SUMMIT EAST
DATA SCIENCE AND ENGINEERING AT SCALE
FEBRUARY 16-18, 2016 NEW YORK CITY

Who Am I?


- Apache Spark PMC Member
- Hortonworks Architect, Spark + Data Science
- Prior to HWX, Principal Research Scientist @ Yahoo Labs (Large Scale Machine Learning)
 - Login Risk Detection, Sponsored Search Click Prediction, Advertising Effectiveness Models, Online Learning, ...




What is Geospatial Analytics?



How do pickup/ dropoff neighborhood hotspots evolve with time?



Correct GPS errors with more
Accurate landmark measurements



Incorporate location in IR and search
advertising



SPARK SUMMIT EAST
2016


Do we need one more library?

- Spatial Analytics at scale is challenging
 - Simplicity + Scalability = Hard
- Ancient Data Formats
 - metadata, indexing not handled well, inefficient storage
- Geospatial Analytics is not simply BI anymore
 - Statistical + Machine Learning being leveraged in geospatial
- Now is the time to do it!
 - Explosion of mobile data
 - Finer granularity of data collection for geometries
 - Analytics stretching the limits of traditional approaches
 - Spark SQL + Catalyst + Tungsten makes extensible SQL engines easier than ever before!




Introduction to Magellan

Polygon = (
[],
[(0.0, 0.0),(1.0, 0.0),
(2.0, 1.0),(1.0, 1.0),
(1.0, 2.0),(0.0, 2.0),
(0.0, 0.0)
])



Introduction to Magellan




```
Polygon = (  
  [0, 5],  
  [(0.0, 0.0),(1.0, 0.0),  
   (1.0, 2.0),(0.0, 2.0),  
   (0.0, 0.0),  
   (0.3, 0.3),  
   (0.6, 0.3),  
   (0.6, 0.9),  
   (0.3, 0.9),  
   (0.3, 0.3)  
])
```




Introduction to Magellan

polygon	metadata
([0], [(-122.4413024, 7.8066277), ...])	neighborhood -> Marina
([0], [(-122.4111659, 37.8003388), ...])	neighborhood -> North Beach

```
sqlContext.read.format("magellan")  
.load(${neighborhoods.path})
```



```
sqlContext.read.format("magellan")  
.option("type", "geojson")  
.load(${neighborhoods.path})
```



Introduction to Magellan

polygon	metadata
[[0], [[(-122.4413024, 7.8066277), ...]]]	neighborhood -> Marina
[[0], [[(-122.4111659, 37.8003388), ...]]]	neighborhood -> North Beach



polygon	metadata
[[0], [[(-122.4111659, 37.8003388), ...]]]	neighborhood -> North Beach

```
neighborhoods.filter(  
  point(-122.4111659, 37.8003388)  
  within  
  'polygon'  
)  
.show()
```

Shape literal

Boolean Expression



Introduction to Magellan

polygon	metadata
([0], [(-122.4111659, 37.8003388), ...])	neighborhood-> North Beach
([0], [(-122.4413024, 7.8066277), ...])	neighborhood-> Marina

point
(-122.4111659, 37.8003388)
(-122.4343576, 37.8068007)

`points.join(neighborhoods).
where('point within 'polygon).
show()`

point	polygon	metadata
(-122.4343576, 37.8068007)	([0], [(-122.4111659, 37.8003388), ...])	neighborhood-> North Beach



Introduction to Magellan

polygon	metadata
([0], [(-122.4111659, 37.8003388), ...])	neighborhood-> North Beach
([0], [(-122.4413024, 7.8066277), ...])	neighborhood-> Marina




point	polygon	metadata
(-122.4343576, 37.8068007)	([0], [(-122.4111659, 37.8003388), ...])	neighborhood-> North Beach

```
neighborhoods.filter(  
  point(-122.4111659, 37.8003388).buffer(0.1)  
  intersects  
  'polygon'  
)show()
```



'point within 'polygon



the join

- Inherits all join optimizations from Spark SQL
 - if neighborhoods table is small, Broadcast Cartesian Join
 - else Cartesian Join



Status

- Magellan 1.0.3 available as Spark Package.
- Scala
- Spark 1.4
- Spark Package: Magellan
- Github: <https://github.com/harsha2010/magellan>
- Blog: <http://hortonworks.com/blog/magellan-geospatial-analytics-in-spark/>
- Notebook example: <http://bit.ly/1GwLyrV>
- Input Formats: ESRI Shapefile, GeoJSON, OSM-XML
- Please try it out and give feedback!



What is next?


- Magellan 1.0.4
 - Spark 1.6
 - Python
 - Spatial Join
 - Persistent Indices
 - Better leverage Tungsten via codegen + memory layout optimizations
 - More operators: buffer, distance, area etc.



the join revisited

- What is the time complexity?
 - m points, n polygons (assume average k edges)
 - l partitions
 - $O(mn/l)$ computations of 'point within 'polygon
 - $O(ml)$ communication cost
 - Each 'point within 'polygon costs $O(k)$
 - Total cost = $O(ml) + O(mnk/l)$
 - $O(m\sqrt{n}\sqrt{k})$ cost, with $O(\sqrt{n}\sqrt{k})$ partitions





SPARK SUMMIT EAST
2016

Optimization?

- *Do we need to send every point to every partition?*
- *Do we need to compute ‘point in neighborhood for each neighborhood within a given partition?*



2d indices

- Quad Tree
- R Tree
- Dimensional Reduction
 - Hashing
 - PCA
 - Space Filling Curves




dimensional reduction


- What does a good dimensional reduction look like?
 - (Approximately) preserve nearness
 - enable range queries
 - No (little) collision




row order curve




snake order curve




z order curve



Binary Representation



Binary Representation



properties

- Locality: Two points differing in a small # of bits are near each other
 - converse not necessarily true!
- Containment
- Efficient construction
- Nice bounds on precision



geohash

- *Z order curve with base 32 encoding*
- *Start with world boundary $(-90,90) \times (-180, 180)$ and recursively subdivide based on precision*



encode (-74.009, 40.7401)


- 40.7401 is in $[0, 90)$ \Rightarrow first bit = 1
- 40.7401 is in $[0, 45)$ \Rightarrow second bit = 0
- 40.7401 is in $[22.5, 45)$ \Rightarrow third bit = 1
- ...
- do same for longitude

answer = dr5rgb



decode dr5rgb

- Decode from Base 32 -> Binary
 - 01100 10111 00101 01111 01010
- lat = 101110001111, long = 0100101111000
- Now decode binary -> decimal.
 - latitude starts with 1 => between 0 - 90
 - second bit = 0 => between 0 - 45
 - third bit = 1 => between 22.5 - 45
 - ...



An algorithm to scale join?

- Preprocess points
 - *For each point compute geohash of precision p covering point*
- Preprocess neighborhoods
 - *For each neighborhood compute geohashes of precision p that intersect neighborhood.*
- Inner join on geohash
- Filter out edge cases



Implementation in Spark SQL

- Override Strategy to define SpatialJoinStrategy
 - Logic to decide when to trigger this join
 - Only trigger if geospatial queries
 - Only trigger if join is complex: if $n \sim O(1)$ then broadcast join is good enough
 - Override BinaryNode to handle the physical execution plan ourselves
 - Override execute(): RDD to execute join and return results
 - Stitch it up using Experimental Strategies in SQLContext



**JOIN IMPLEMENTATION
TOO SLOW**



**LEVERAGED SPATIAL INDICES
AND CATALYST TO FIX IT**

imgflip.com



SPARK SUMMIT EAST
2016

Persistent Indices

- Often, the geometry dataset does not change... eg. neighborhoods
- Index the dataset once and for all?




Persistent Indices

- Use Magellan to generate spatial indices
 - *Think of geometry as document, list of geohashes as words!*
- Persist indices to Elastic Search
- Use Magellan Data Source to query indexed ES data
- Pushdown geometric predicates where possible
 - *Predicate rewritten to IR query*



Overall architecture



That's all Folks!



SPARK SUMMIT EAST
2016

THANK YOU.

Twitter: @halfabrane, Github: @harsha2010



SPARK SUMMIT EAST
DATA SCIENCE AND ENGINEERING AT SCALE
FEBRUARY 16-18, 2016 NEW YORK CITY