

# BUILDING A GRAPH OF ALL U.S. BUSINESSES USING SPARK TECHNOLOGIES.

Alexis Roos @alexisroos

Engineering manager / Lead Data Science Engineer

Radius Intelligence



**SPARK SUMMIT EAST**  
DATA SCIENCE AND ENGINEERING AT SCALE  
FEBRUARY 16-18, 2016 NEW YORK CITY

## AGENDA

- Radius Intelligence
- Objectives and terminology
- Business graph data pipeline
- Lessons learned
- Q&As

# RADIUS INTELLIGENCE

## Predictive Marketing software

Radius transforms the way B2B marketers **discover** markets, **acquire** customers, and **measure** success.

Our software is powered by the *Radius Intelligence Cloud* a proprietary data science engine which provides predictive analytics, powerful segmentation, and seamless integrations.



# RADIUS INTELLIGENCE

- Founded in **2009**
- **125+ employees** (and growing)
- Headquartered in **San Francisco, CA**
- **\$125 million** in funding to date
- Cutting edge in **Data Science** and **Data Engineering Technologies**



# BUSINESS GRAPH OBJECTIVE

Produce an **accurate** and **comprehensive** **Business Graph** from **records** coming from dozens of sources comprising hundreds of data points such as:



...



Latest News

## BUSINESS GRAPH BENEFITS

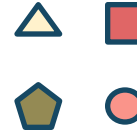
Answer complex connected questions such as:

- **Employees:** Find **engineering executives** at a given company focusing on Big Data or security
- **News:** Show **News** related to companies in my marketing segment such as Funding announcements or mergers and acquisitions
- **Intent:** Show companies with a buying **Intent** for my product or engineering managers looking for security product
- **Organization:** Show all locations or Employees for a given business  
etc



# BUSINESS GRAPH TERMINOLOGY

- **Business:** a company or organization  
ex: Blue Bottle company
- **Location:** a business at a particular [address](#)  
ex: Blue Bottle at Sansome St, San Francisco
- **Attributes:** information associated with a business and/or location  
ex: address, website, phone number, etc.
- **Additional Data types:**  
ex: Employees, Intent, News, Technologies, etc.



# BUSINESS GRAPH DATA PIPELINE

## Data acquisition

- Crawling
- Licensing



Dozens sources  
7B+ records  
50B+ Data Points



## Raw Records





# BUSINESS GRAPH DATA PIPELINE



- Standardization
- Validation
- Normalization

## Raw Records

## Normalized Records



# BUSINESS GRAPH DATA PIPELINE

Data acquisition

Data preparation

**Clustering**

Construction

Cluster records on:

- **Address**
- Name of the business
- Websites
- Phone numbers
- Employees
- Description
- Category codes
- Etc.

Normalized Records



Clusters of records by address/business



**Spark** **Scala**

**GraphX** **databricks**

# BUSINESS GRAPH DATA PIPELINE

Data acquisition

Data preparation

Clustering

Construction

- Locations  
- Businesses  
& more

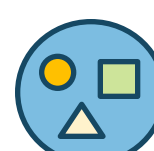
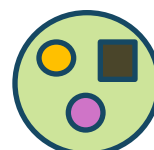
## Construct Locations:

- Filter bad clusters
- Use all records in cluster to create Location with attributes:
  - Select best value(s)
  - Score values: confidence based on specific record features including historical
  - Impute missing values
- Etc.

Clusters



Locations



Spark

Scala

MLLib

databricks



SPARK SUMMIT EAST  
2016

# BUSINESS GRAPH DATA PIPELINE

Data acquisition

Data preparation

Clustering

Construction  
Businesses  
& more

Construct business & more:

- Create business and link locations
- Add business level info: news, employees, intent, etc.
- Correct/ impute fields at business level
- Etc.

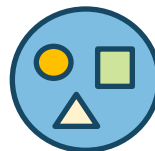
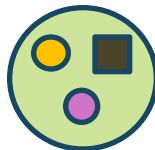


GraphX

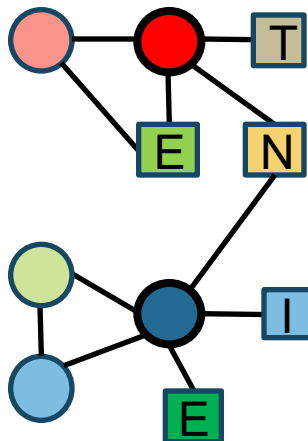
MLLib



Locations



Business graph



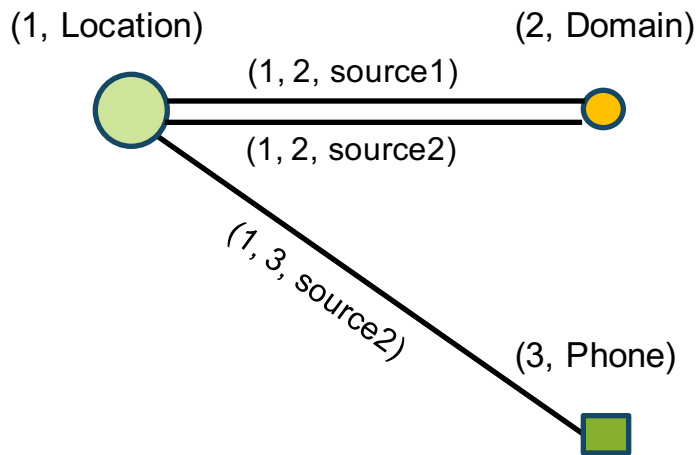
SPARK SUMMIT EAST  
2016

# SPARK / GRAPHX

GraphX models a “property graph” which is an multi directed, **attributed** graph.

VD - Vertex data attribute type

ED - Edge data attribute type



```
class Graph[VD, ED] {  
  val vertices: VertexRDD[VD]  
  val edges: EdgeRDD[ED]  
  val triplets: RDD[EdgeTriplet[VD, ED]]  
}
```

VertexRDD[VD] extends RDD[(VertexID, VD)]  
EdgeRDD[ED] extends RDD[Edge[ED]]

# SPARK / GRAPHX

Graph operations are like RDD operations: functional and lazily evaluated

## Operations from Graph and GraphOps:

- **Information:** *numEdges, numVertices, inDegrees, outDegrees, degrees*
- **Collections:** *vertices, edges, triplets*
- **Transformations and modifications:** *mapVertices, mapEdges, mapTriplets, filter, reverse, subgraph, mask, groupEdges, convertToCanonicalEdges*
- **Join:** *joinVertices, outerJoinVertices*
- **Aggregations:** *collectNeighbor(Id)s, aggregateMessages, sendMsg, mergeMsg, tripletFields*
- **Caching and partitioning:** *persist, cache, checkpoint, unpersist(Vertices), partitionBy*
- **Graph algorithms:** *pregel, pageRank, connectedComponents, triangleCount*

## And data operations on VertexRDD and EdgeRDD:

- **VertexRDD:** *filter, mapValues, minus, diff, innerJoin, leftJoin, aggregateUsingIndex*
- **EdgeRDD:** *mapValues, reverse, innerJoin*



# GRAPH DATA PIPELINE CONSTRUCTION ZOOM IN

Data acquisition

Data preparation

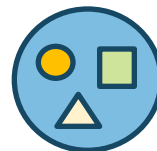
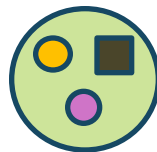
Clustering

**Construction**  
**Businesses**  
**& more**

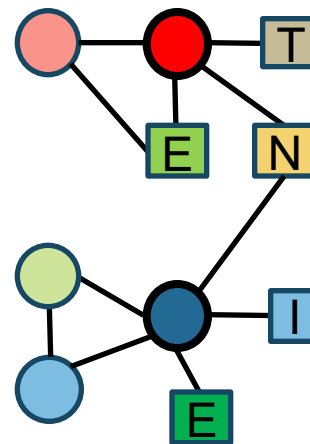
Construct business graph:

- Create business and link locations
- Add business level info: news, employees, intent, etc.
- Correct/ impute fields at business level
- Etc.

Locations

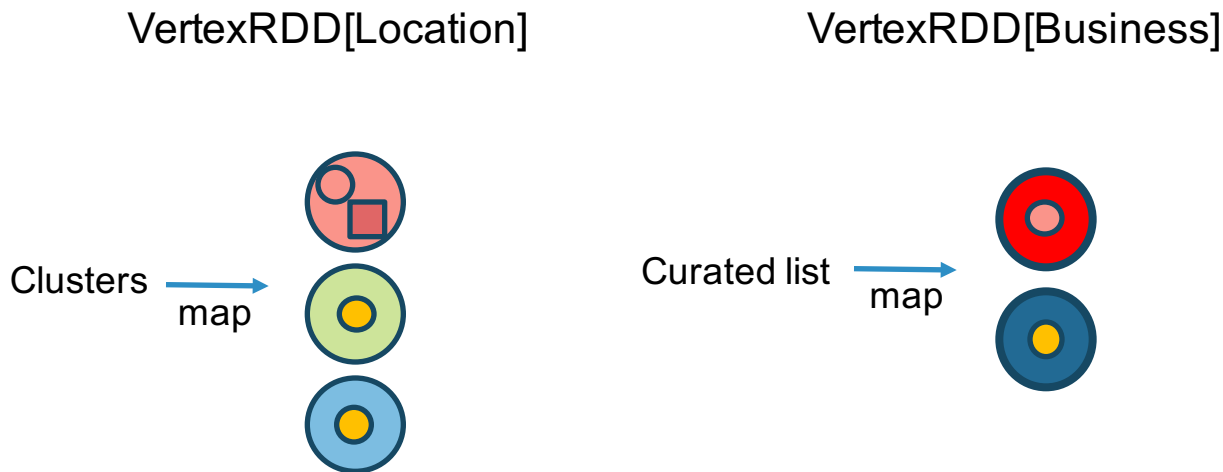


Business graph



# GRAPH DATA PIPELINE CONSTRUCTION ZOOM IN

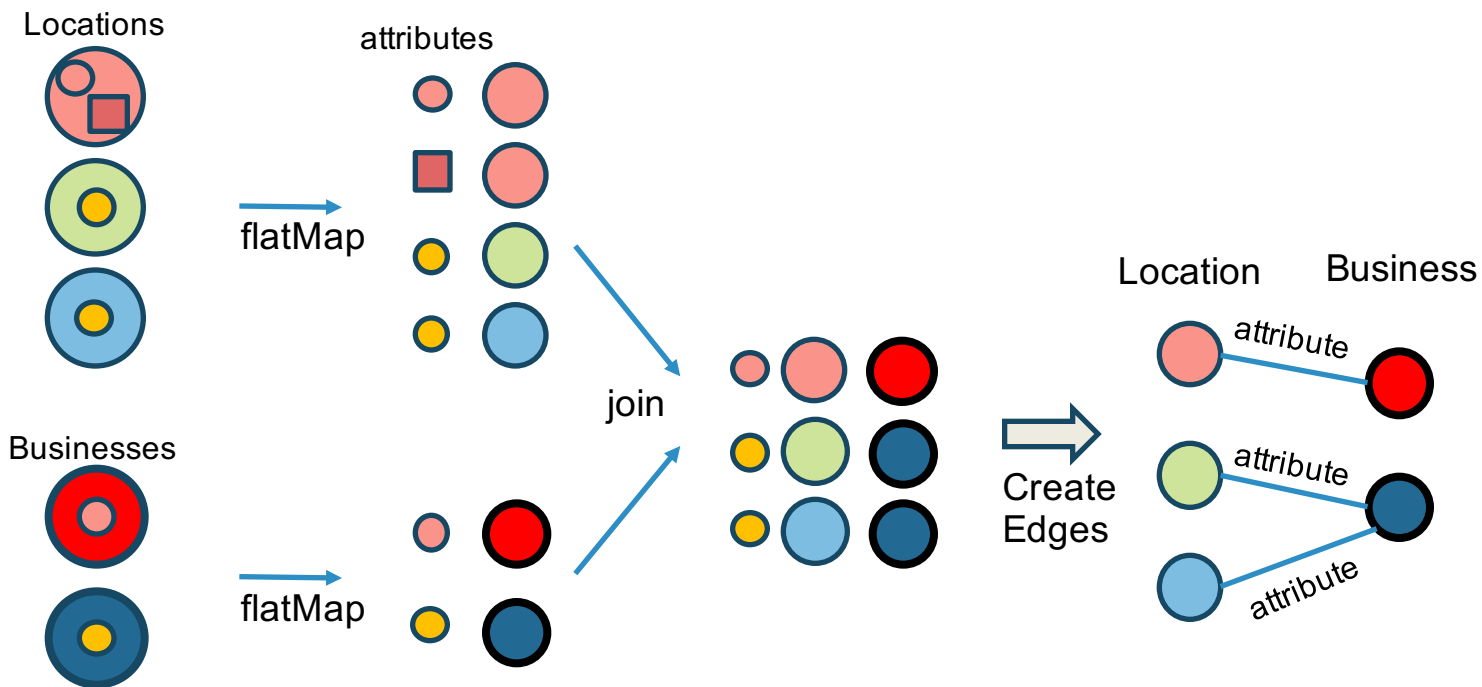
1. Create vertex for each location and each known business





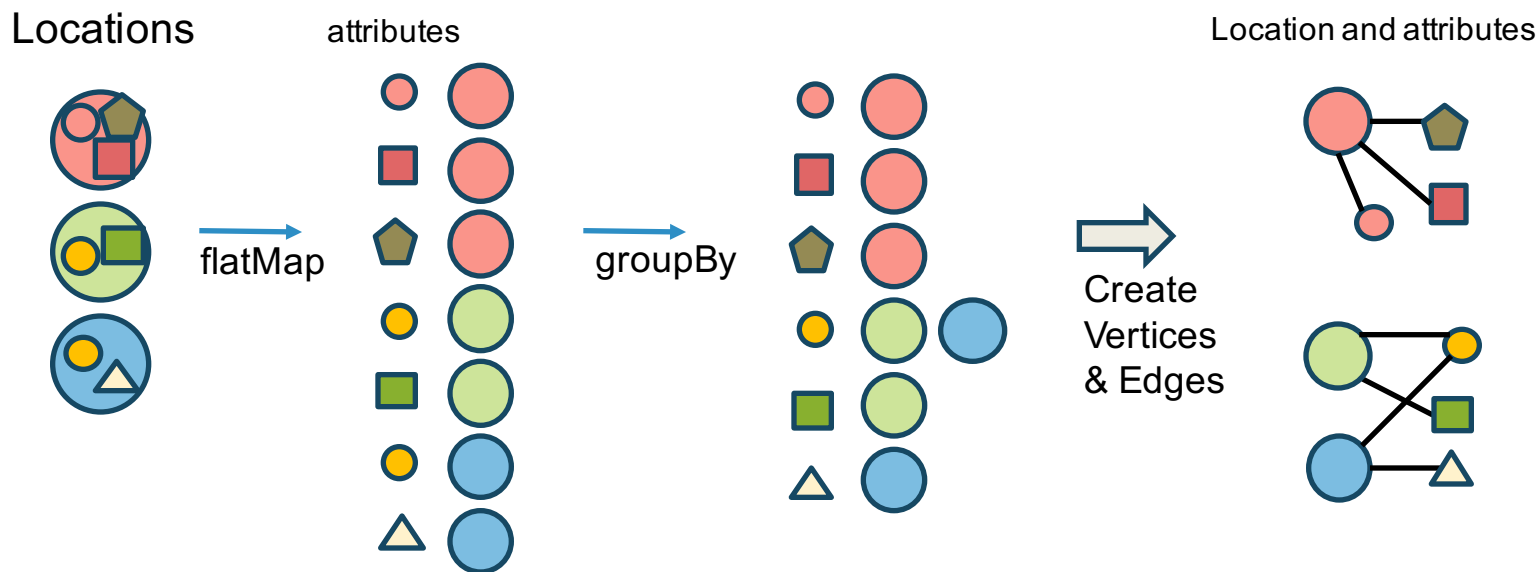
# GRAPH DATA PIPELINE CONSTRUCTION ZOOM IN

## 2. Link Locations to Known Business per relationships (such as domain)



## GRAPH DATA PIPELINE CONSTRUCTION ZOOM IN

3. Create vertex for unique top attributes such as address, phone, domain, etc. and link locations to these attributes

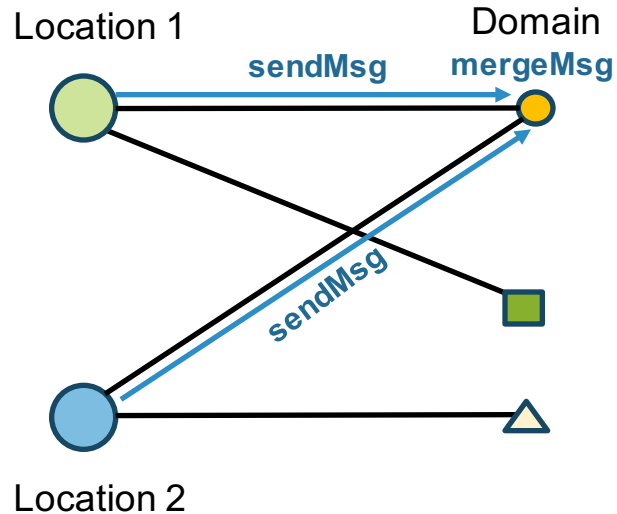


# SPARK / GRAPHX

Aggregate messages:  
sendMsg, mergeMsg

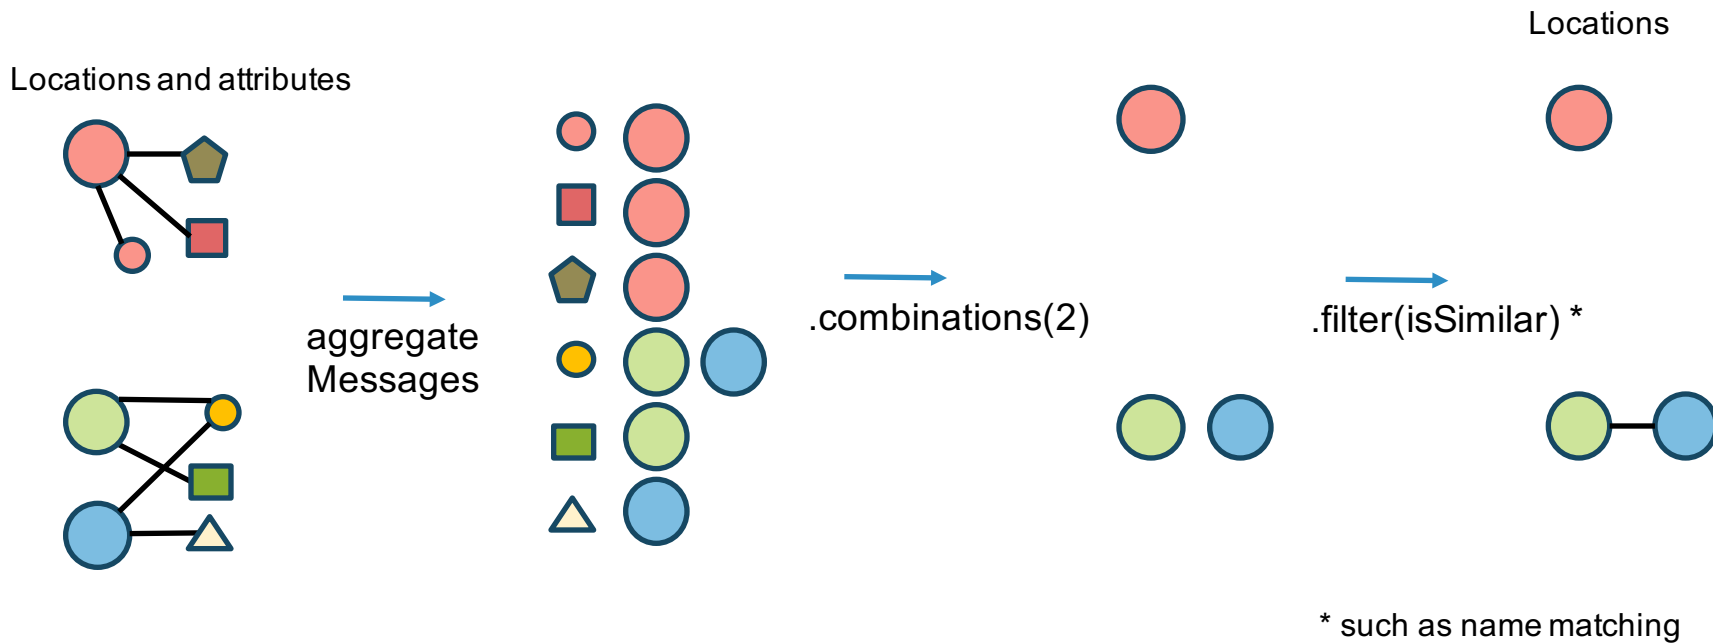
sendMsg: (srcID, name)

mergeMsg: x ++ y



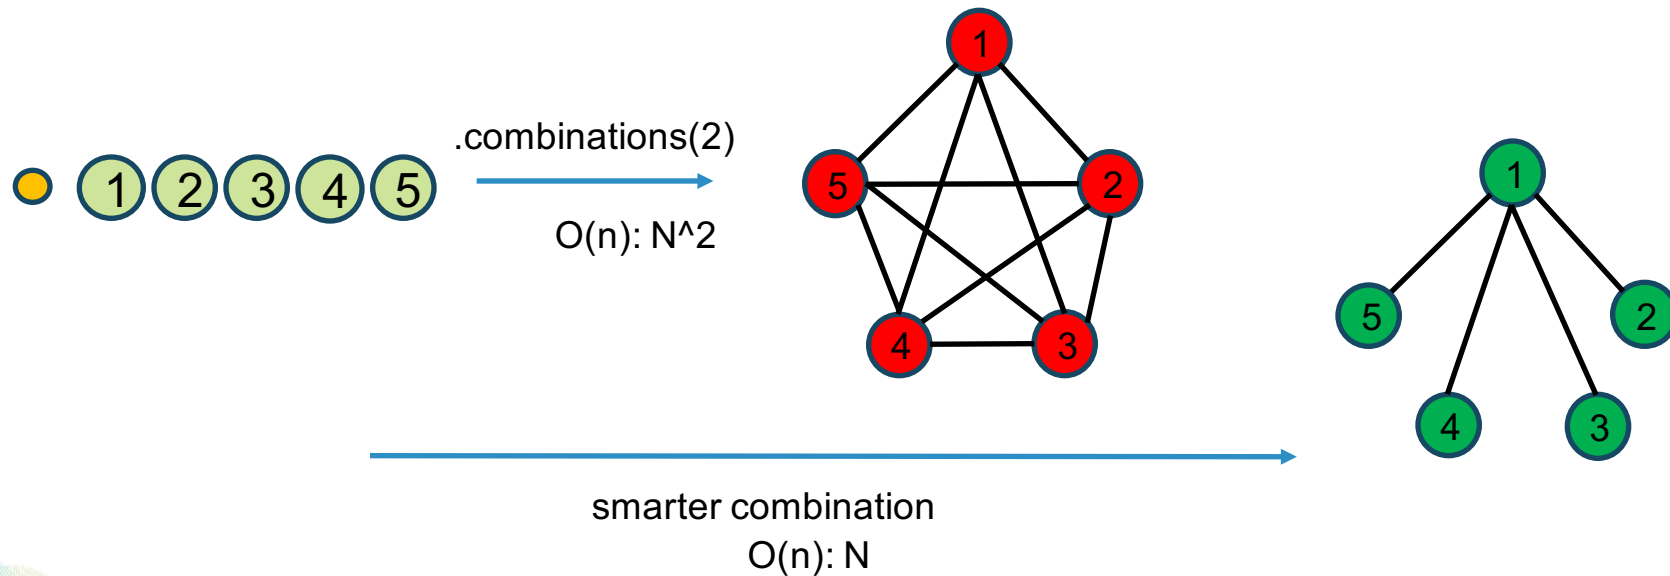
# GRAPH DATA PIPELINE CONSTRUCTION ZOOM IN

## 4. Create location to location edges using attributes relationship and graph cutting



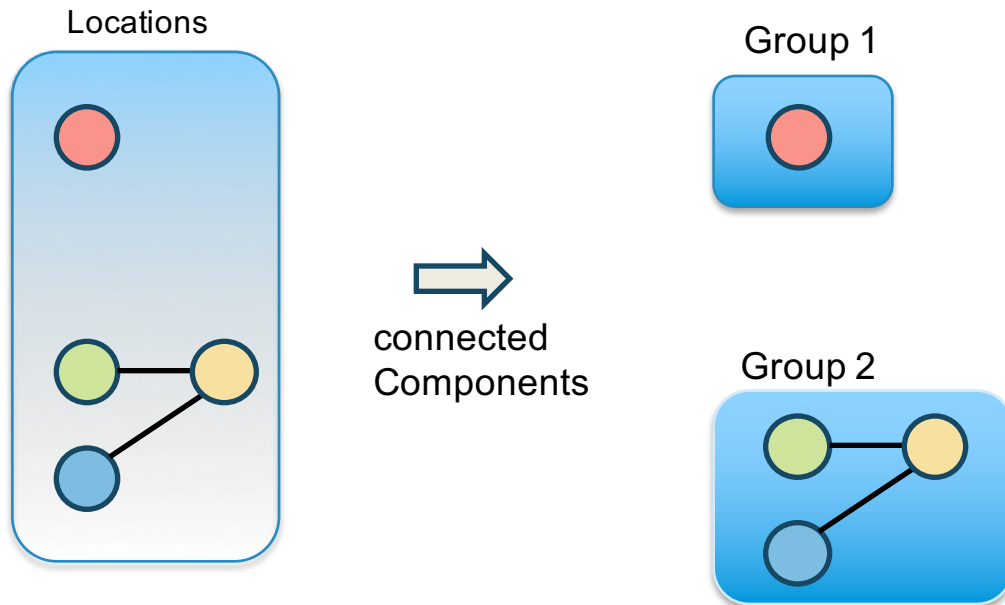
## PERFORMANCE TIP

Note: `.combinations(2)` and bottleneck



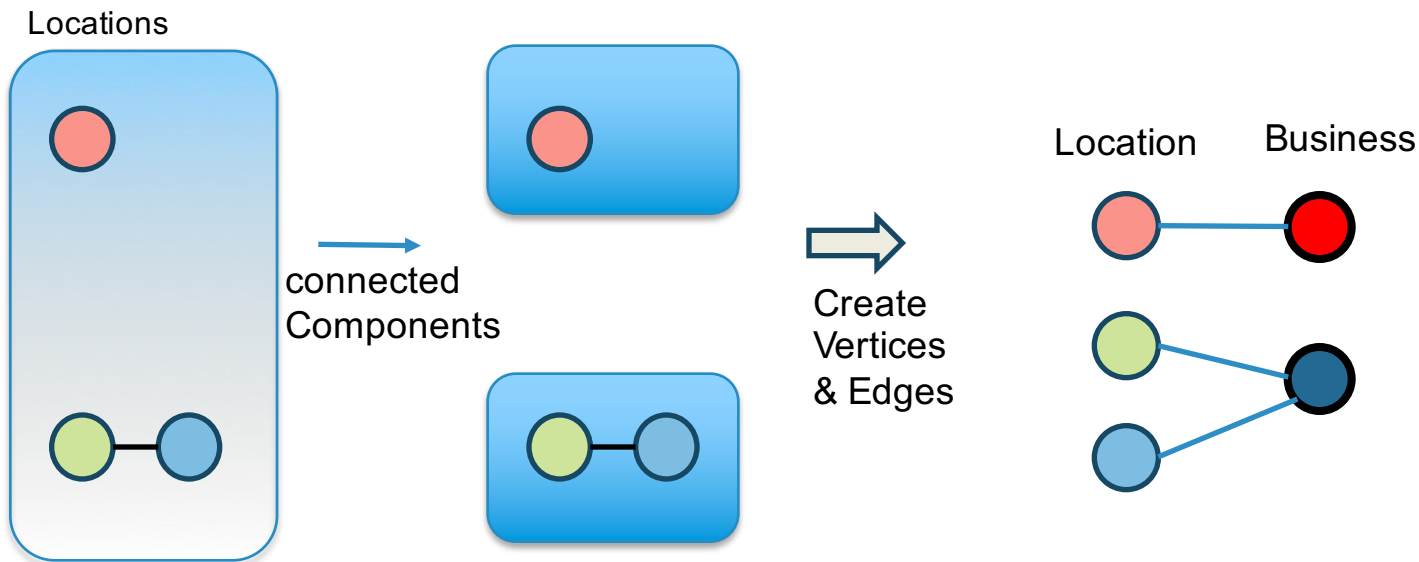
# SPARK / GRAPHX

## Connected components



# GRAPH DATA PIPELINE CONSTRUCTION ZOOM IN

## 5. Call connected components on linked locations and create businesses



## PERFORMANCE TIP

GraphX Connected components implementation already performs optimizations:

- Incremental Updates to Mirror Caches
- Join Elimination
- Index Scanning for Active Sets
- Local Vertex and Edge Indices
- Index and Routing Table Reuse

But algorithm is iterative so optimization matters:

- Vertices and edges should be in memory: more efficient to subset / keep the graph minimal or filter and rejoin later
- Check-pointing improves performance (less lineage)



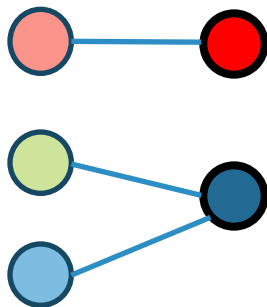


# GRAPH DATA PIPELINE CONSTRUCTION ZOOM IN

## 6. More steps:

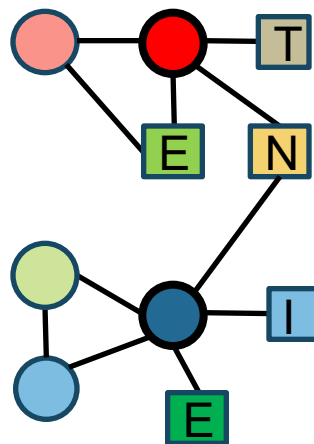
- Assign headquarter
- Join additional information: Employees, News, Intent, etc.
- Further corrections / imputations: attributes, revenues, etc.

Location      Business



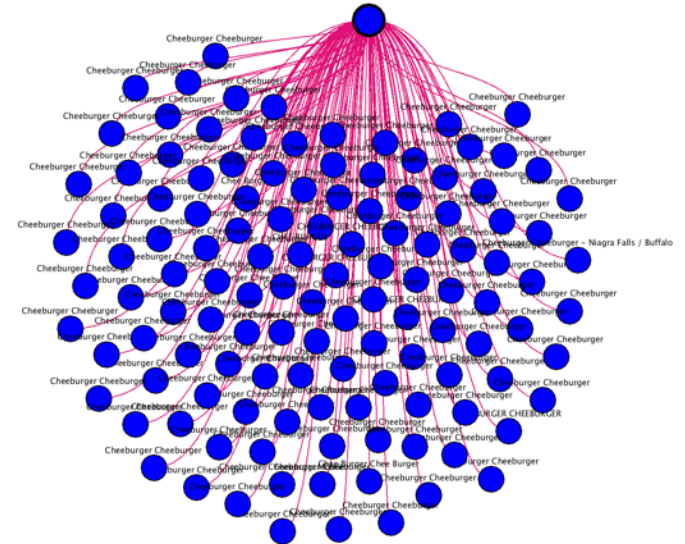
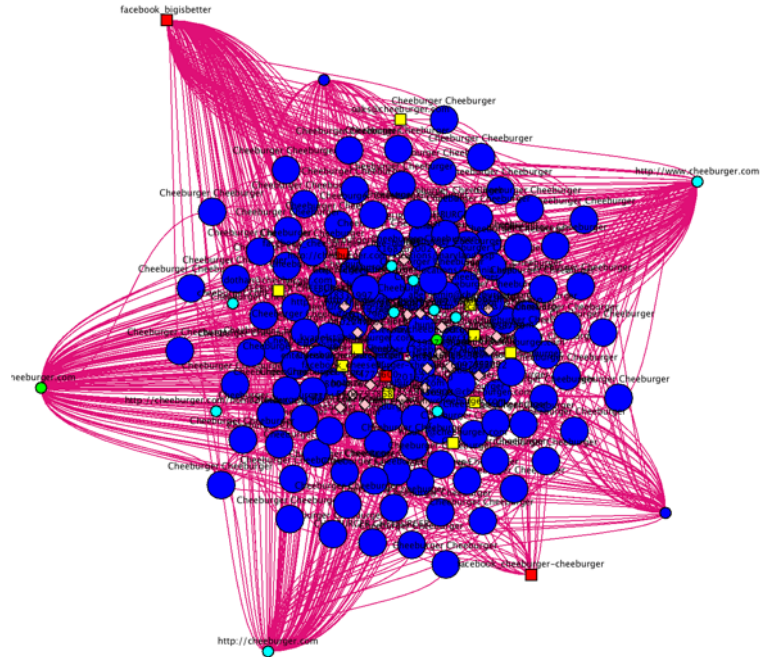
➔  
more  
operations

Business graph



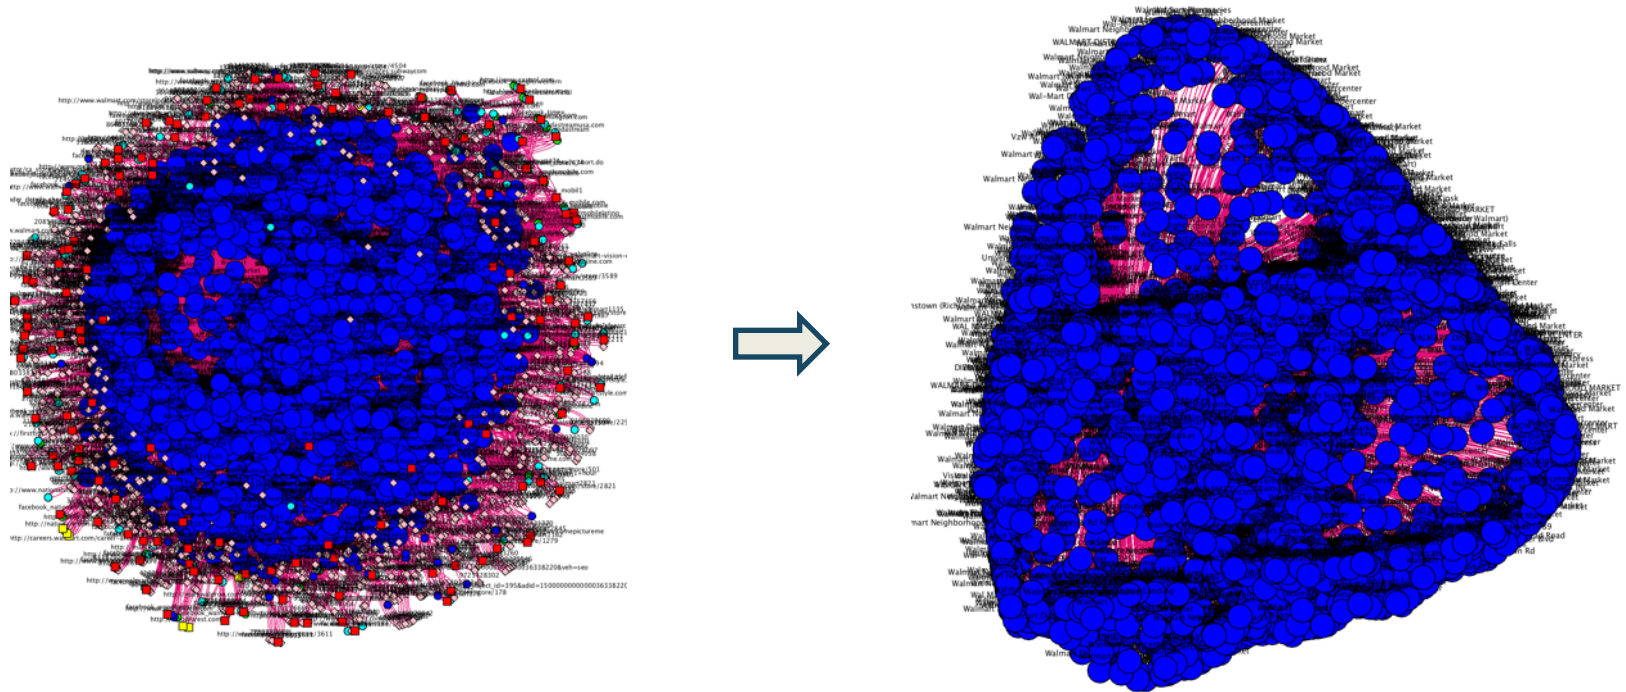
# GRAPHICAL VIEW

From Locations with attributes to Locations with Business



# GRAPHICAL VIEW

From Locations with attributes to Locations with Business



## LESSONS LEARNED

- Using a Graph allows us to naturally model business relationships and append new data types over time. Also allowed us to debug data!
- Creating and updating the graph is easy using RDD operations
- GraphX scales: 250M vertices and 1b edges in clustering
- Edge cutting is an art as much as a science
- Connected Components is an expensive operation



# THANK YOU.

WE ARE HIRING!

<http://radius.com/jobs/>

Alexis Roos

[alexis@radius.com](mailto:alexis@radius.com)

[@alexisroos](#)



**SPARK SUMMIT EAST**

DATA SCIENCE AND ENGINEERING AT SCALE  
FEBRUARY 16-18, 2016 NEW YORK CITY