

# Generalized Linear Models in Spark MLlib and SparkR

Xiangrui Meng

joint with Joseph Bradley, Eric Liang, Yanbo Liang (MiningLamp), DB Tsai (Netflix), et al.

2016/02/17 - Spark Summit East



# About me


- Software Engineer at Databricks
- Spark PMC member and MLlib/PySpark maintainer
- Ph.D. from Stanford on randomized algorithms for large-scale linear regression problems

# Outline

- Generalized linear models (GLMs)
  - linear regression / logistic regression / general form
  - accelerated failure time (AFT) model for survival analysis
  - intercept / regularization / weights
- GLMs in MLlib and SparkR
  - demo: R formula in Spark
- Implementing GLMs
  - gradient descent / L-BFGS / OWL-QN
  - weighted least squares / iteratively re-weighted least squares (IRLS)
  - performance tips

# Generalized linear models

# Linear regression



inference / prediction

# Linear least squares

- m observations:  $(x_1, y_1), (x_2, y_2), \dots, (x_m, y_m)$
- x: explanatory variables, y: dependent variable
- assumes linear relationship between x and y

$$y = x^T \beta + \varepsilon$$


- minimizes the sum of the squares of the errors

$$\text{minimize}_{\beta \in \mathbb{R}^n} \quad \frac{1}{2} \sum_{i=1}^m \|y_i - x_i^T \beta\|_2^2$$

# Linear least squares

- the oldest linear model, trace back to Gauss
- the simplest and the most studied linear model
- has analytic solutions
- easy to solve
- easy to inspect
  
- sensitive to outliers

# Logistic regression



[https://en.wikipedia.org/wiki/Logistic\\_regression](https://en.wikipedia.org/wiki/Logistic_regression)



# Logistic regression

- classification with binary response:  $y \in \{1, -1\}$ 
  - true/false, clicked/not clicked, liked/disliked
- uses logistic function to indicate the likelihood

$$P(y = 1) = \frac{1}{1 + e^{-x^T \beta}}$$

- maximizes the sum of the log-likelihoods, i.e.,

$$\text{minimize}_{\beta} \quad \sum_{i=1}^m \log(1 + e^{-y_i \cdot x_i^T \beta})$$

# Logistic regression

- one of the simplest binary classification models
- widely used in industry
- relatively easy to solve
- easy to interpret

# Multinomial logistic regression

- classification with multiclass response:  $y \in \{1, 2, \dots, K\}$
- uses softmax function to indicate likelihood

$$P(y = k) = e^{x^T \beta_k} / Z, \text{ where } Z = \sum_{l=1}^K e^{x^T \beta_l}$$

- maximizes the sum of log-likelihoods

$$\text{maximize}_{\beta_1, \dots, \beta_K} \sum_{i=1}^m \sum_{l=1}^K I_{y_i=l} \log \left( e^{x_i^T \beta_l} / Z_i \right)$$

# Generalized linear models (GLMs)

- Both linear least squares and logistic regression are special cases of generalized linear models.
- A GLM is specified by the following:
  - a distribution of the response (from the exponential family),
  - a link function  $g$  such that  $\mathbf{E}(y) = g^{-1}(x^T \beta)$
- maximizes the sum of log-likelihoods

$$\text{maximize}_{\beta} \quad \sum_{i=1}^m \log p(y_i | x_i; \beta)$$

# Distributions and link functions

Model	Distribution	Link
linear least squares	normal	identity
logistic regression	binomial	logit
multinomial logic regression	multinomial	generalized logit
Poisson regression	Poisson	log
gamma regression	gamma	inverse

# Accelerated failure time (AFT) model

- m observations:  $(x_1, y_1, c_1), \dots, (x_m, y_m, c_m)$
- y: survival time, c: censor variable (alive or dead)
- assumes the effect of an explanatory variable is to accelerate or decelerate the life time by some constant
- uses maximum likelihood estimation while treating censored and uncensored observations differently

# AFT model for survival analysis

- one popular parametric model for survival analysis
- widely used for lifetime estimation and churn analysis
- could be solved under the same framework as GLMs

# Intercept, regularization, and weights

In practice, a linear model is often more complex

$$\text{maximize}_{\beta} \quad \sum_{i=1}^m w_i \cdot \log p(y_i | x_i^T \beta + \beta_0) + \lambda \cdot \sigma(\beta)$$

where  $w$  describes instance weights,  $\beta_0$  is the intercept term to adjust bias, and  $\sigma$  regularized  $\beta$  with a constant  $\lambda > 0$  to avoid overfitting.



# Types of regularization

- Ridge (L2):  $\frac{1}{2}\|\beta\|_2^2$ 
  - easy to solve (strongly convex)
- Lasso (L1):  $\|\beta\|_1$ 
  - enhance model sparsity
  - harder to solve (though still convex)
- Elastic-Net:  $\alpha\|\beta\|_1 + \frac{1-\alpha}{2}\|\beta\|_2^2, \quad \alpha \in [0, 1]$
- Others: group lasso, nonconvex, etc

# GLMs in MLlib and SparkR

# GLMs in Spark MLlib

Linear models in MLlib are implemented as ML pipeline estimators. They accept the following params:

- **featuresCol**: a vector column containing features ( $x$ )
- **labelCol**: a double column containing responses ( $y$ )
- **weightCol**: a double column containing weights ( $w$ )
- **regType**: regularization type, “none”, “l1”, “l2”, “elastic-net”
- **regParam**: regularization constant
- **fitIntercept**: whether to fit an intercept term
- ...

# Fit a linear model in MLlib

```
from pyspark.ml.classification import LogisticRegression

# Load training data
training = sqlContext.read.parquet("path/to/training")

lr = LogisticRegression(
    weightCol="weight", fitIntercept=False, maxIter=10,
    regParam=0.3, elasticNetParam=0.8)

# Fit the model
model = lr.fit(training)
```

# Make predictions and evaluate models

```
from pyspark.ml.evaluation import BinaryClassificationEvaluator

test = sqlContext.read.parquet("path/to/test")

# make predictions by calling transform
predictions = model.transform(test)

# create a binary classification evaluator
evaluator = BinaryClassificationEvaluator(
    metricName="areaUnderROC")
evaluator.evaluate(predictions)
```

# GLMs in SparkR

In Python/Scala/Java, we keep the APIs about the same for consistency. But in SparkR, we make the APIs similar to existing ones in R (or R packages).

```
# Create the DataFrame
df <- read.df(sqlContext, "path/to/training")

# Fit a Gaussian GLM model
model <- glm(y ~ x1 + x2, data = df, family = "gaussian")
```

# R formula in SparkR

- R provides model formula to express linear models.
- We support the following R formula operators in SparkR:
  - ``~`` separate target and terms
  - ``+`` concat terms, "+ 0" means removing intercept
  - ``-`` remove a term, "- 1" means removing intercept
  - ``:`` interaction (multiplication for numeric values, or binarized categorical values)
  - ``.`` all columns except target
- For example, "y ~ x + z + x:z -1" means using x, z, and their interaction (x:z) to predict y without intercept (-1).

# Demo: GLMs in Spark

... using [Databricks Community Edition](#)!



# Implementing GLMs

# Row-based distributed storage

w	x	y

partition 1

w	x	y

partition 2

# Gradient descent methods

- Stochastic gradient descent (SGD):  $\beta := \beta - \mu \cdot g(\beta; x_i, y_i)$ 
  - trade-offs on the merge scheme and convergence
- Mini-batch SGD:  $\beta := \beta - \mu \cdot \sum_{i \in \mathcal{B}_j} g(\beta; x_i, y_i)$ 
  - hard to sample mini-batches efficiently
  - communication overhead on merging gradients
- Batch gradient descent:  $\beta := \beta - \mu \cdot \sum_{i=1}^m g(\beta; x_i, y_i)$ 
  - slow convergence

# Quasi-Newton methods

- Newton's method converges much faster than GD, but it requires second-order information:  $\beta := \beta - H^{-1}g$
- L-BFGS works for smooth objectives. It approximates the inverse Hessian using only first-order information.
- OWL-QN works for objectives with L1 regularization.
- MLlib calls L-BFGS/OWL-QN implemented in breeze.

# Direct methods for linear least squares

- Linear least squares has an analytic solution:

$$\beta = (X^T X)^{-1} X^T y$$

- The solution could be computed directly or through QR factorization, both of which are implemented in Spark.
- requires only a single pass
- efficient when the number of features is small (<4000)
- provides R-like model summary statistics

# Iteratively re-weighted least squares (IRLS)

- Generalized linear models with exponential family can be solved via iteratively re-weighted least squares (IRLS).
  - linearizes the objective at the current solution
  - solves the weighted linear least squares problem
  - repeat above steps until convergence
- efficient when the number of features is small (<4000)
- provides R-like model summary statistics
- This is the implementation in R.

# Verification using R

Besides normal tests, we also verify our implementation using R.

```
/*  
  df <- as.data.frame(cbind(A, b))  
  for (formula in c(b ~ . -1, b ~ .)) {  
    model <- lm(formula, data=df, weights=w)  
    print(as.vector(coef(model)))  
  }  
  
  [1] -3.727121  3.009983  
  [1] 18.08  6.08 -0.60  
*/  
val expected = Seq(Vectors.dense(0.0, -3.727121, 3.009983),  
                   Vectors.dense(18.08, 6.08, -0.60))
```

# Standardization

To match the result in both R and glmnet, the most popular R package for GLMs, we provide options to standardize features and labels before training:

$$\sigma(\beta) = \frac{1}{2\delta} \sum_{j=1}^n (\sigma_j \beta_j)^2$$

where delta is the stddev of labels, and sigma\_j is the stddev of the j-th feature column.



# Performance tips

- Utilize sparsity.
- Use tree aggregation and torrent broadcast.
- Watch numerical issues, e.g.,  $\log(1+\exp(x))$ .
- Do not change input data. Scaling could be applied after each iteration and intercept could be derived later.

# Future directions

- easy handling of categorical features and labels
- better R formula support
- more model summary statistics
- feature parity in different languages
- model parallelism
  - vector-free L-BFGS with 2D partitioning (WIP)
- using matrix kernels

# Other GLM implementations on Spark

- [CoCoA+](#): communication-efficient optimization
- [LIBLINEAR for Spark](#): a Spark port of LIBLINEAR
- [sparkGLM](#): an R-like GLM package for Spark
- [TFOCS for Spark](#): first-order conic solvers for Spark
- General-purpose packages that implement GLMs
  - [aerosolve](#), [DistML](#), [sparkling-water](#), [thunder](#), [zen](#), etc
- ... and more on [Spark Packages](#)

# Thank you.

- MLlib [user guide](#) and [roadmap](#) for Spark 2.0
- [GLMs on Wikipedia](#)
- Databricks Community Edition, blog posts, and careers

