

OPTIMIZING TERASCALE MACHINE
LEARNING PIPELINES WITH

KeystoneML

Apache

Evan R. Sparks, UC Berkeley AMPLab

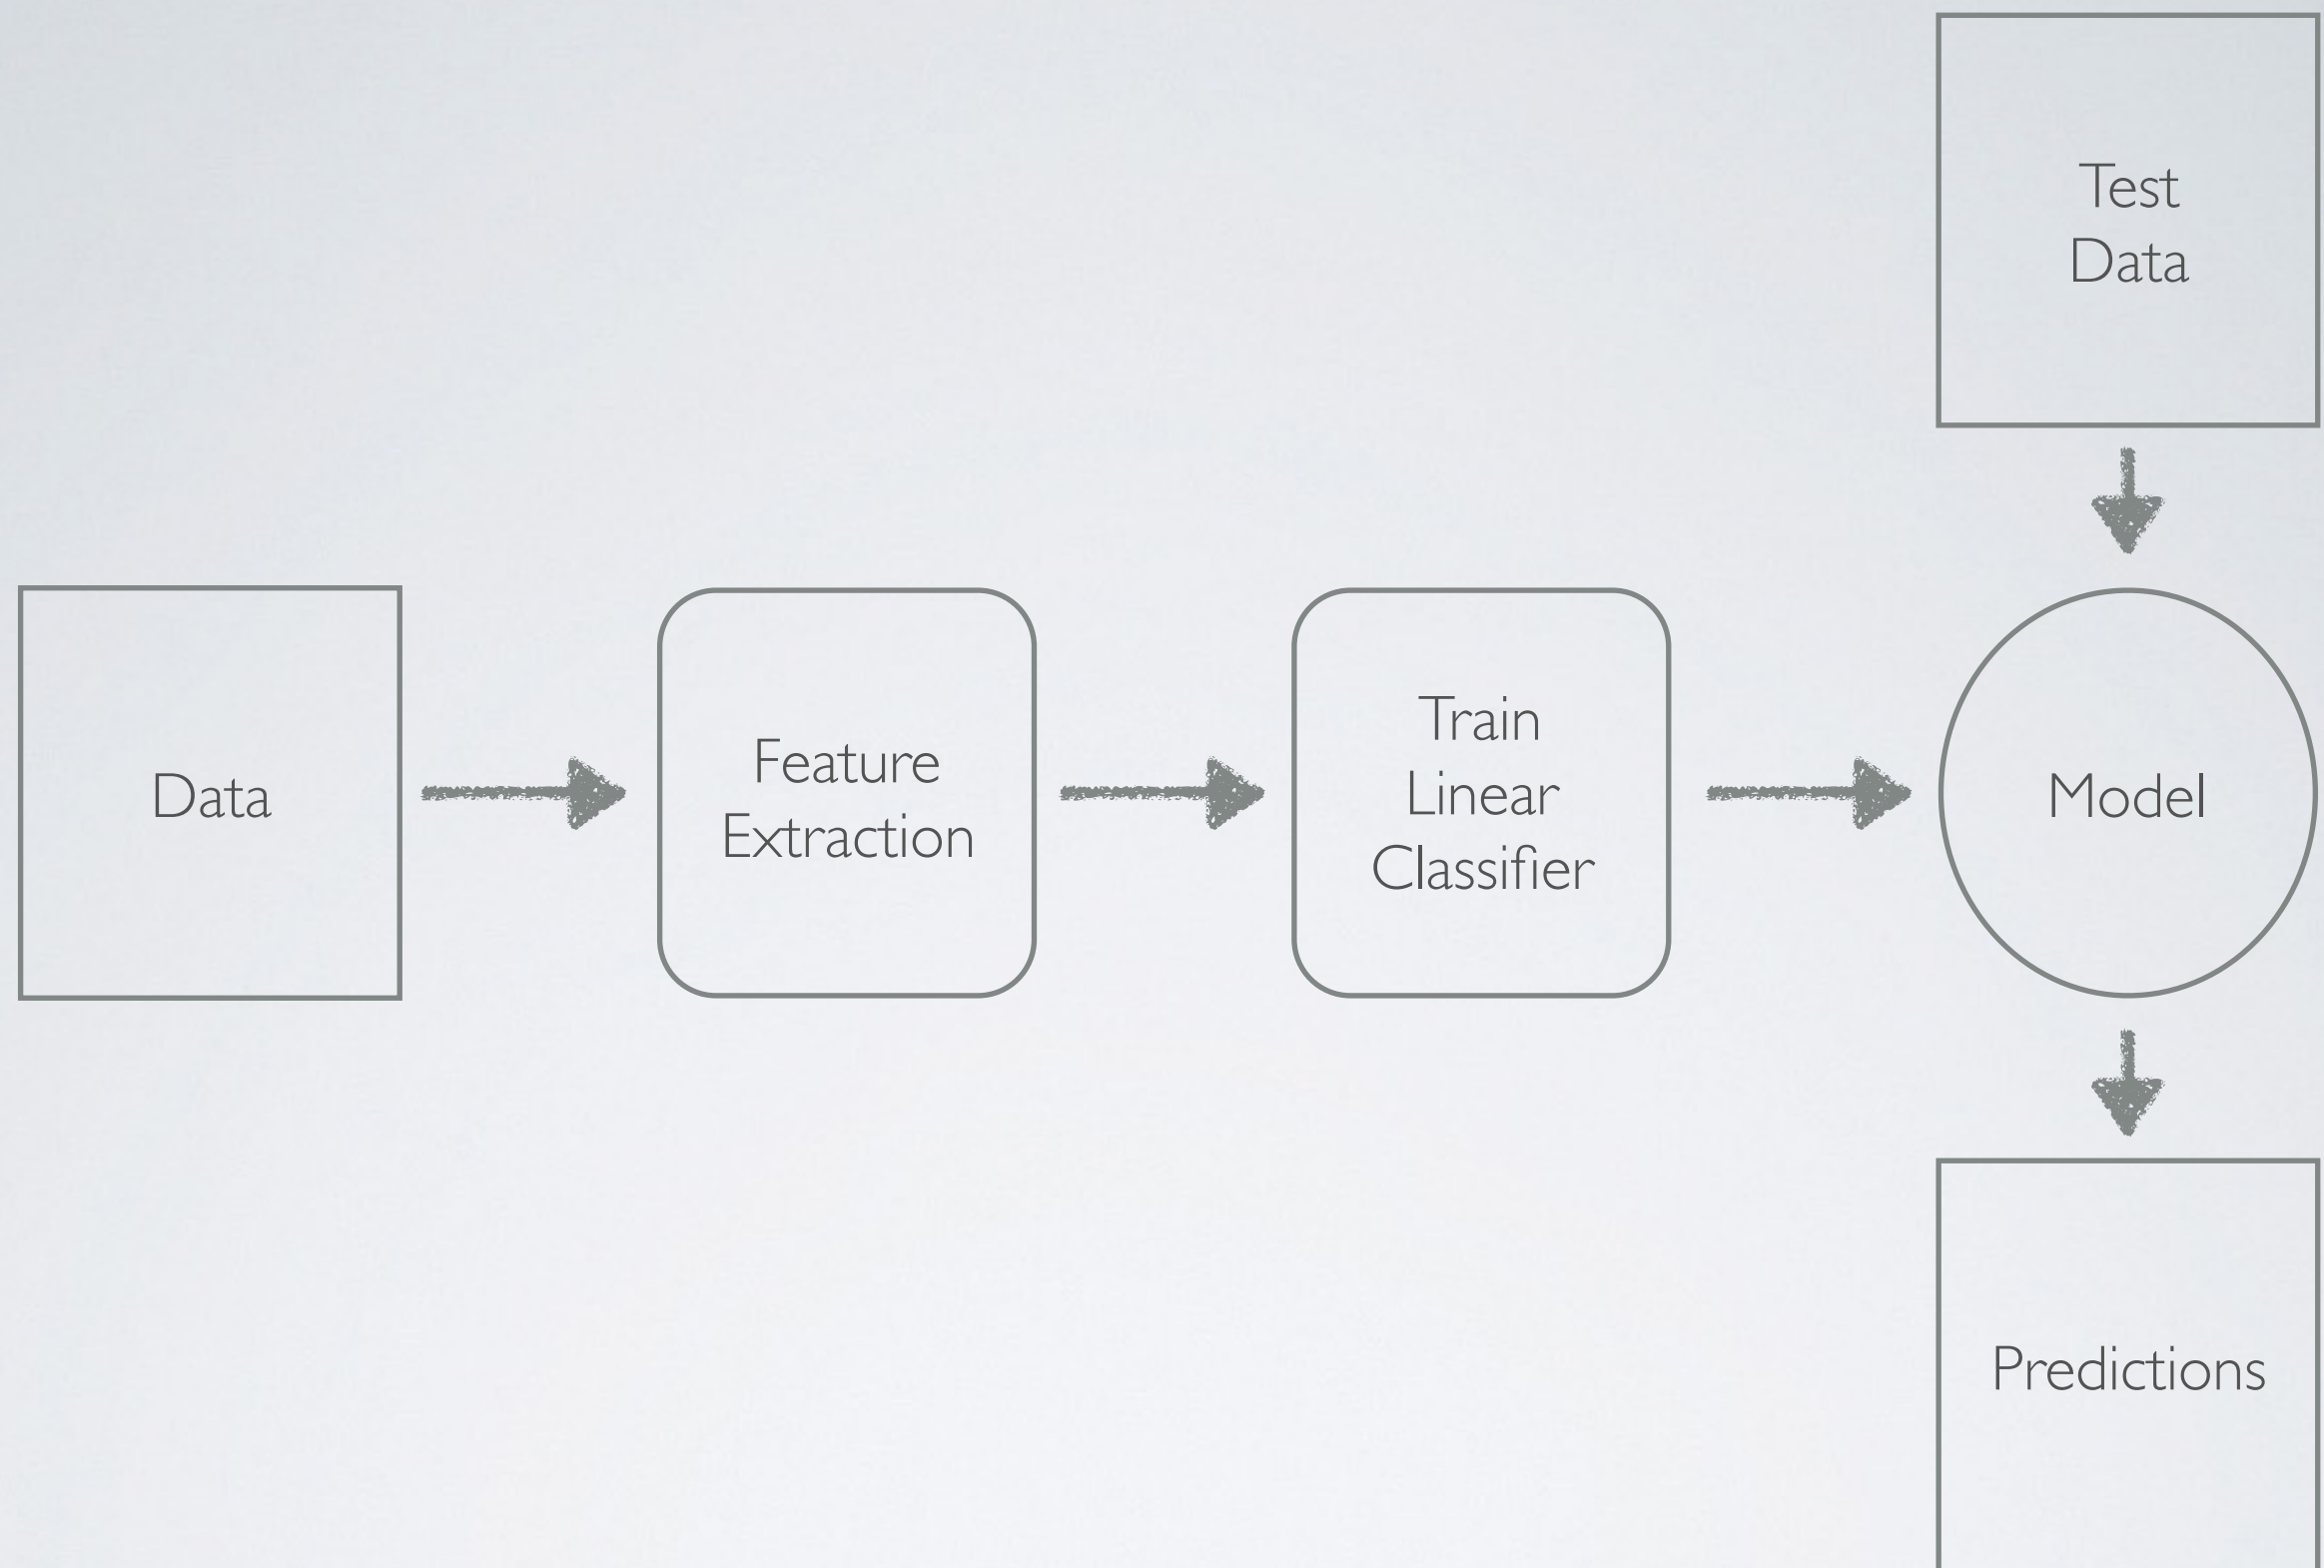
with Shivaram Venkataraman, Tomer Kaftan, Michael Franklin, Benjamin Recht

WHAT'S A MACHINE
LEARNING PIPELINE?



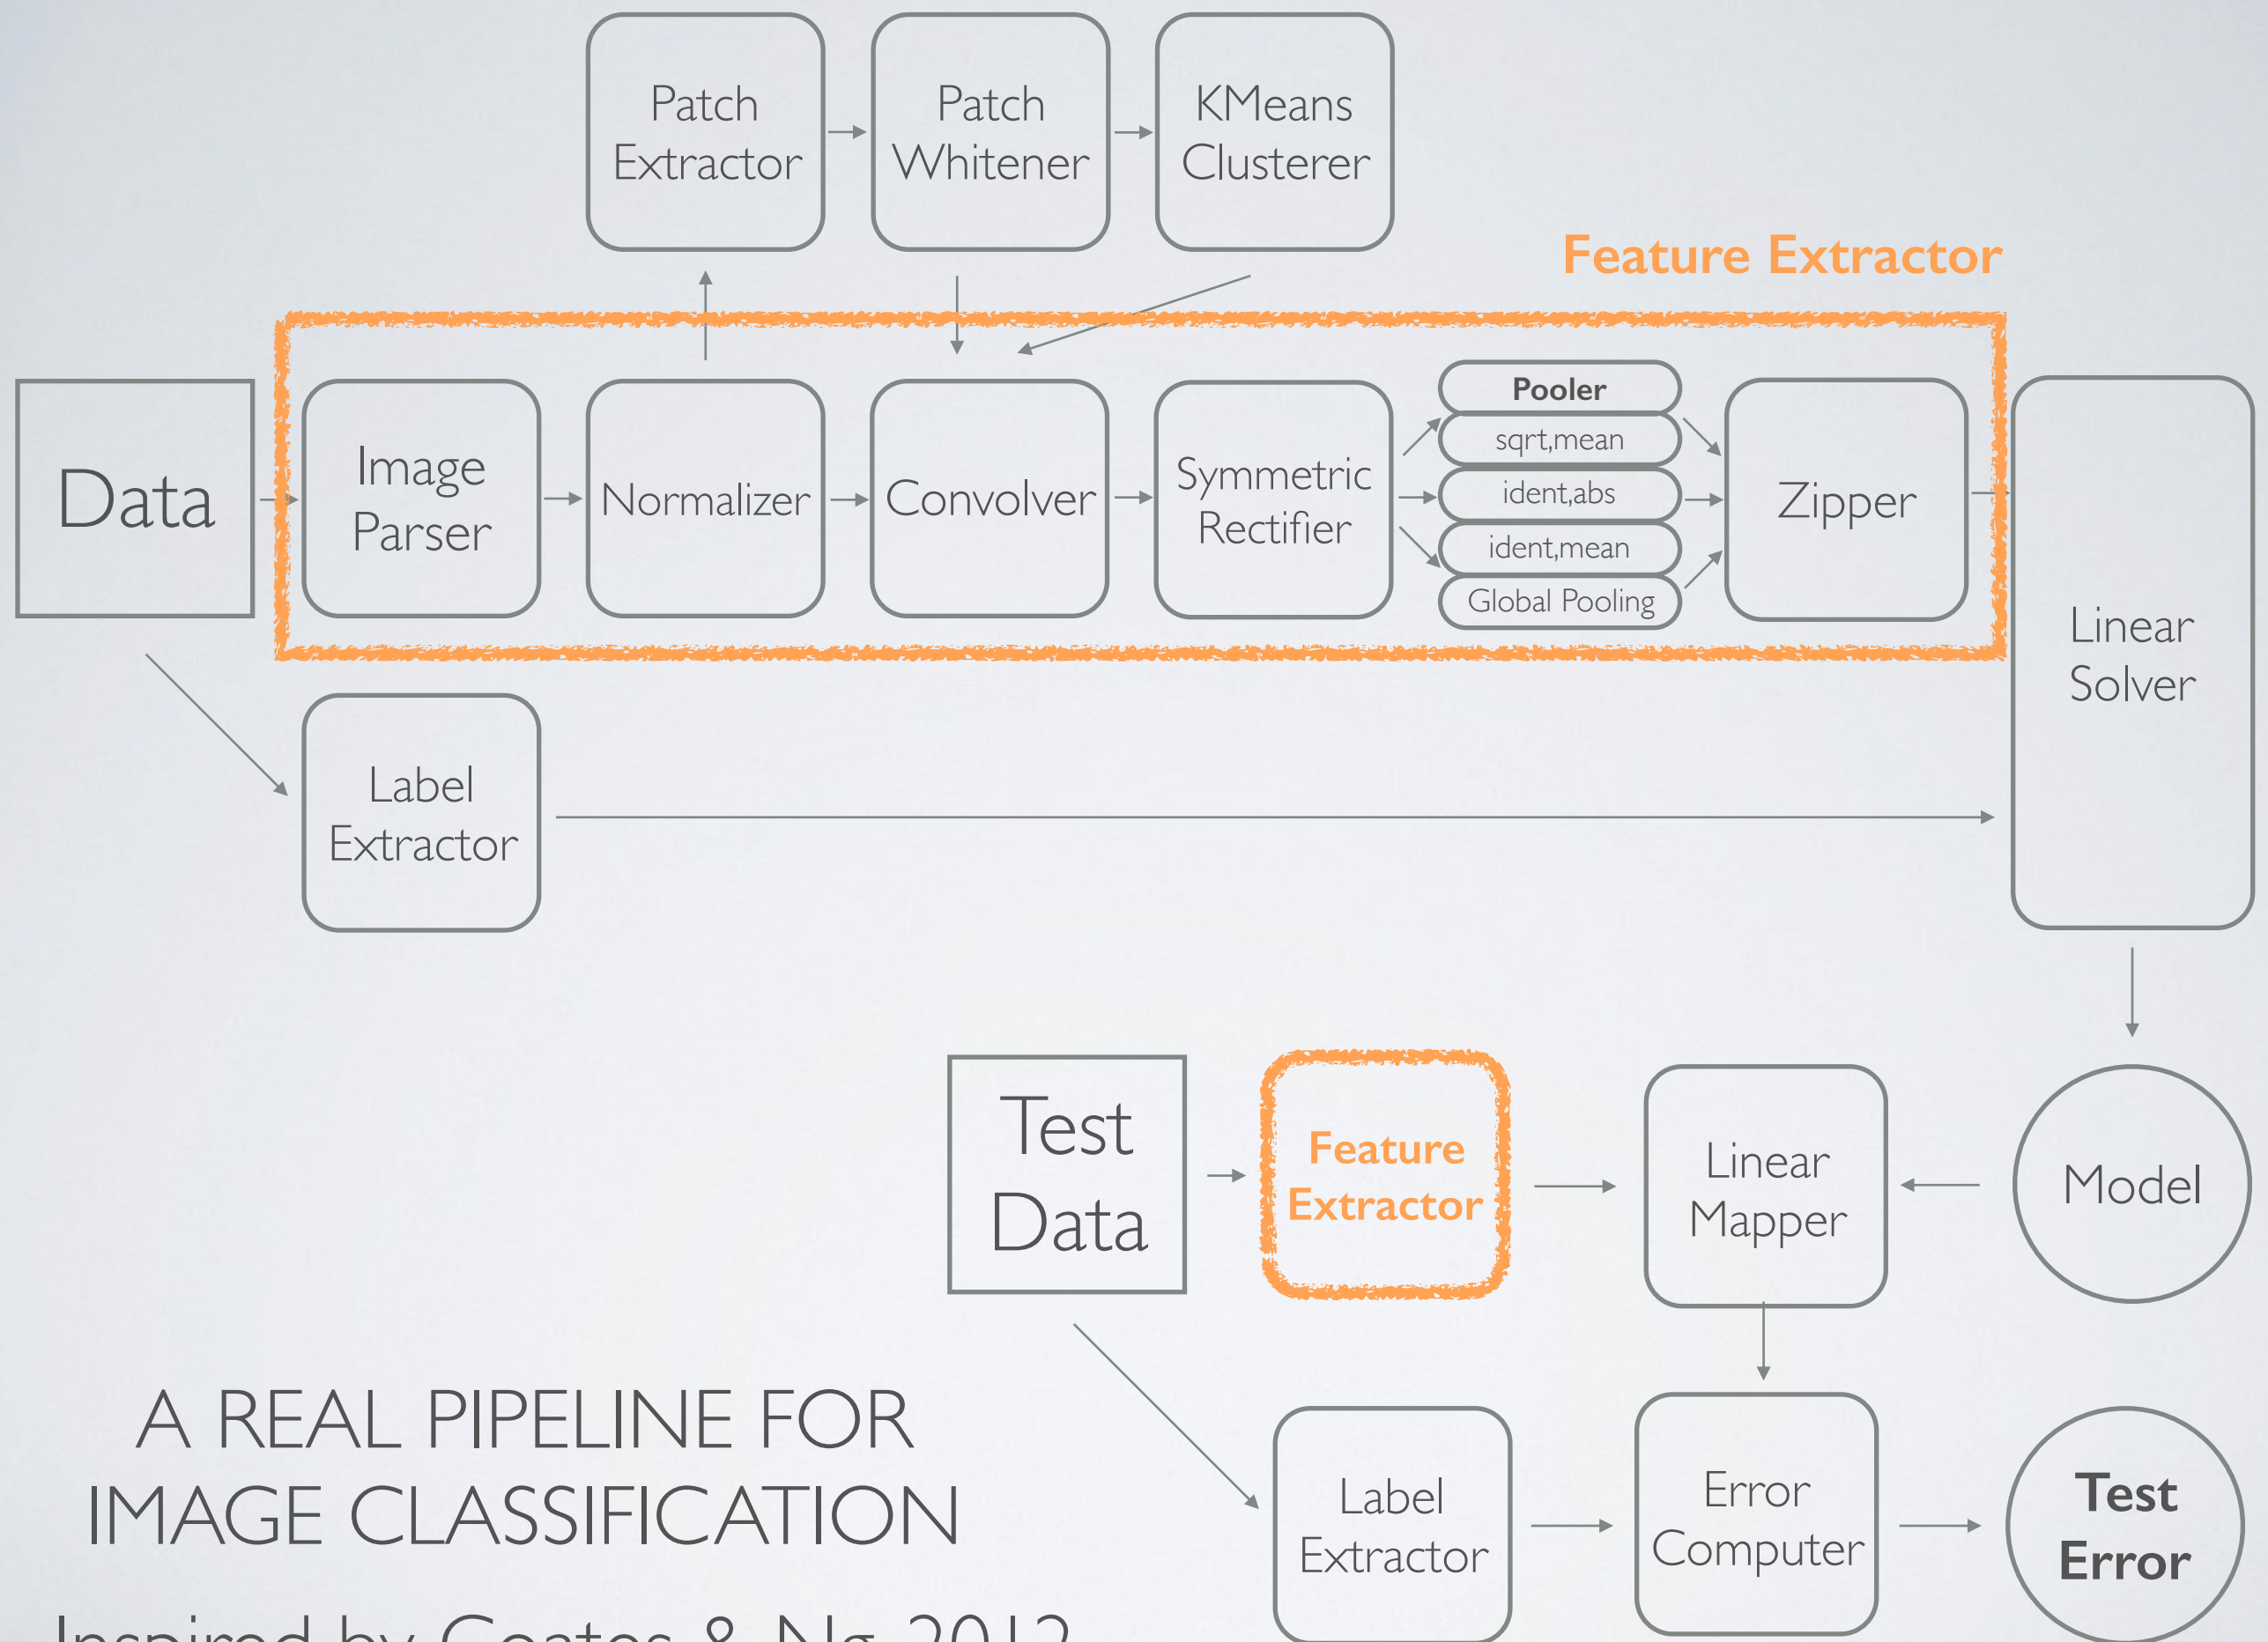
A STANDARD MACHINE LEARNING PIPELINE

Right?

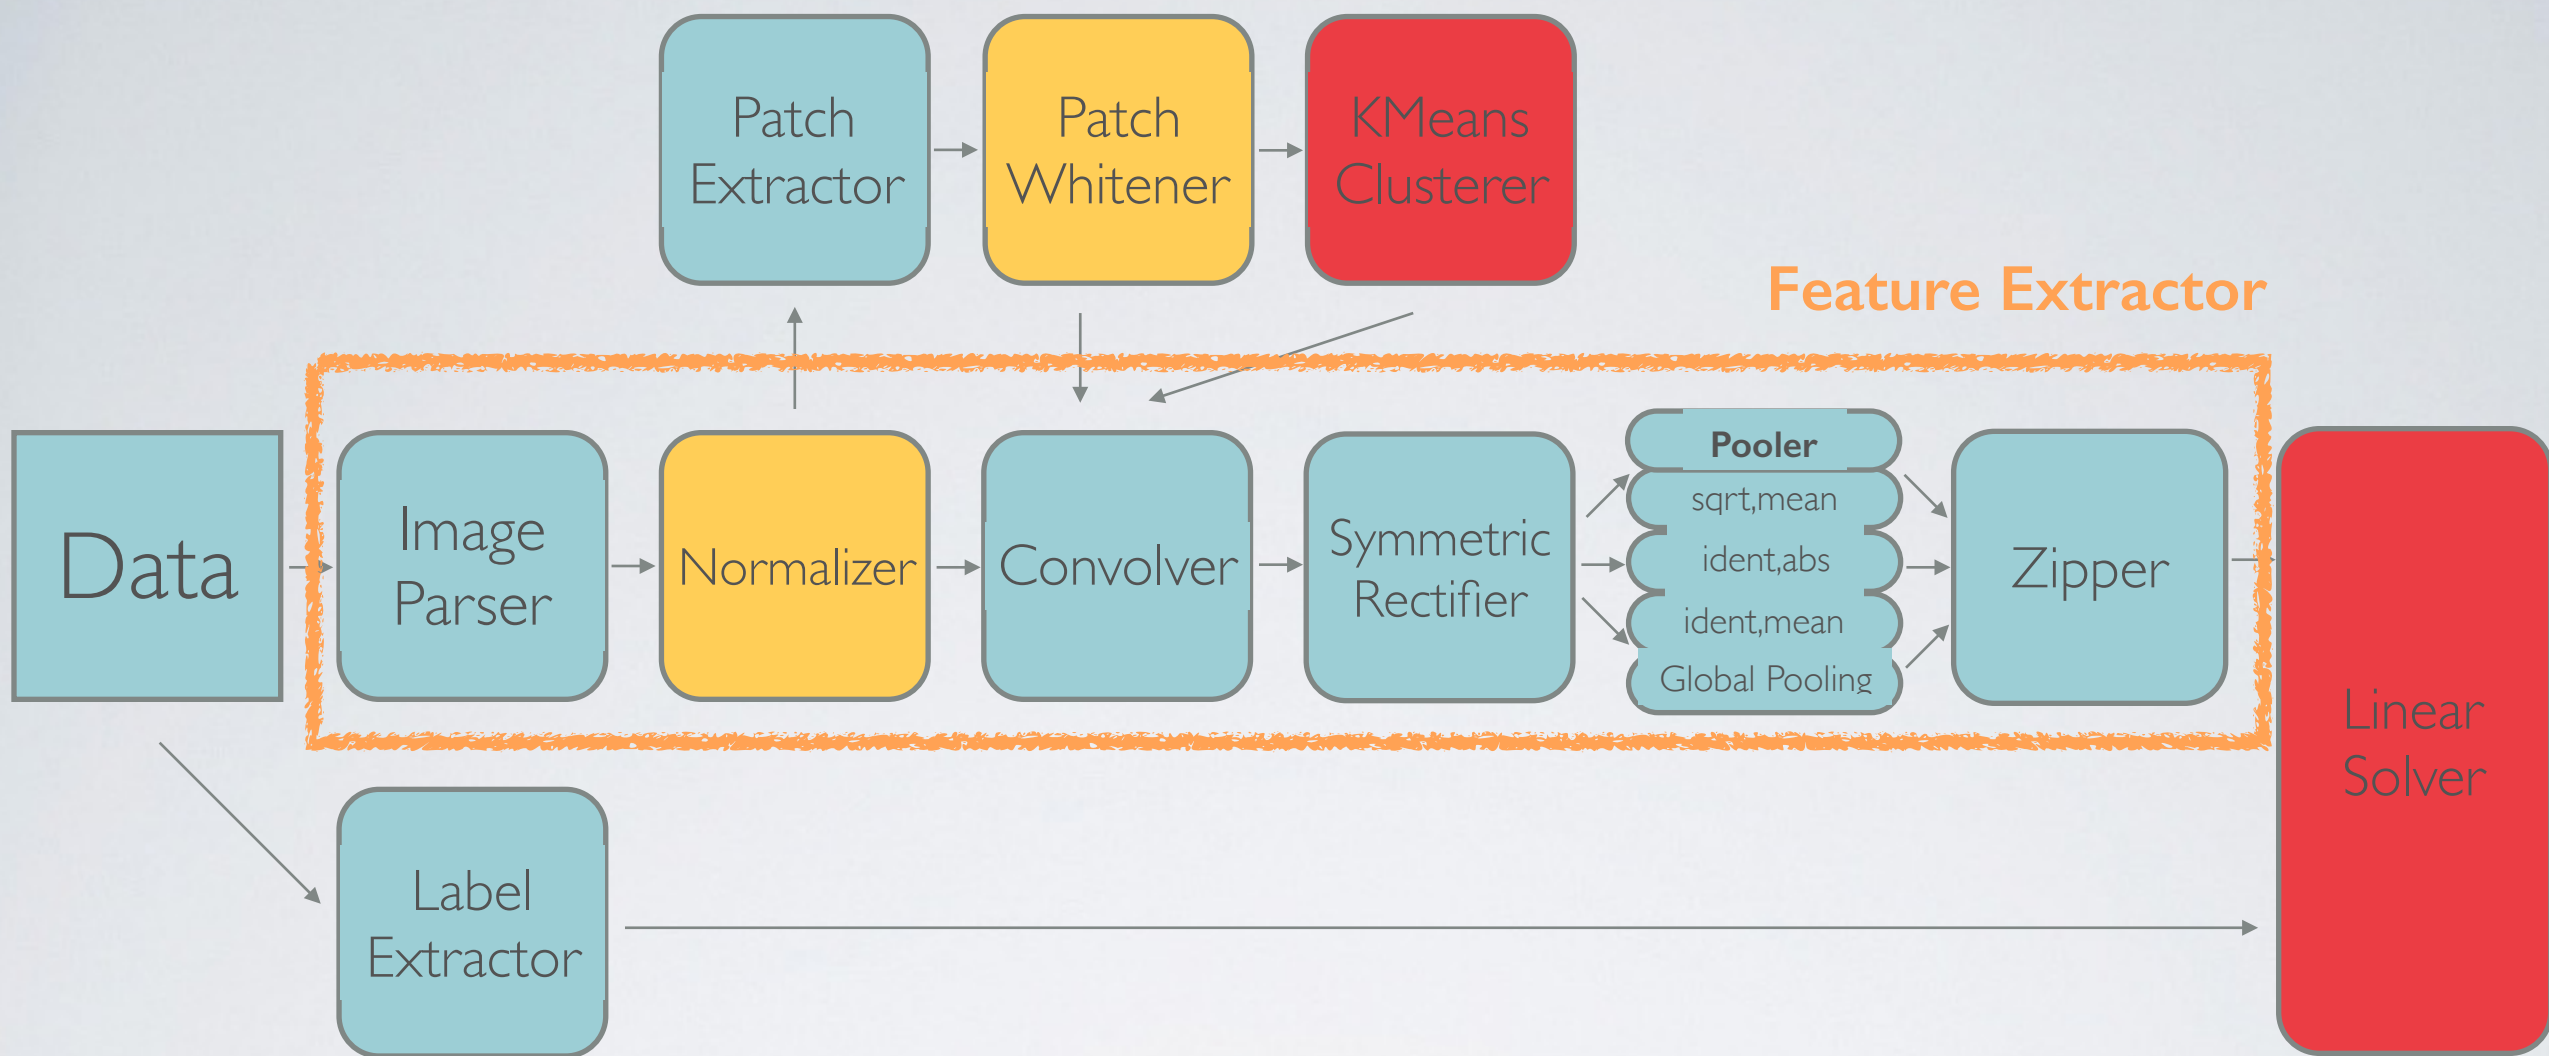


A STANDARD MACHINE LEARNING PIPELINE

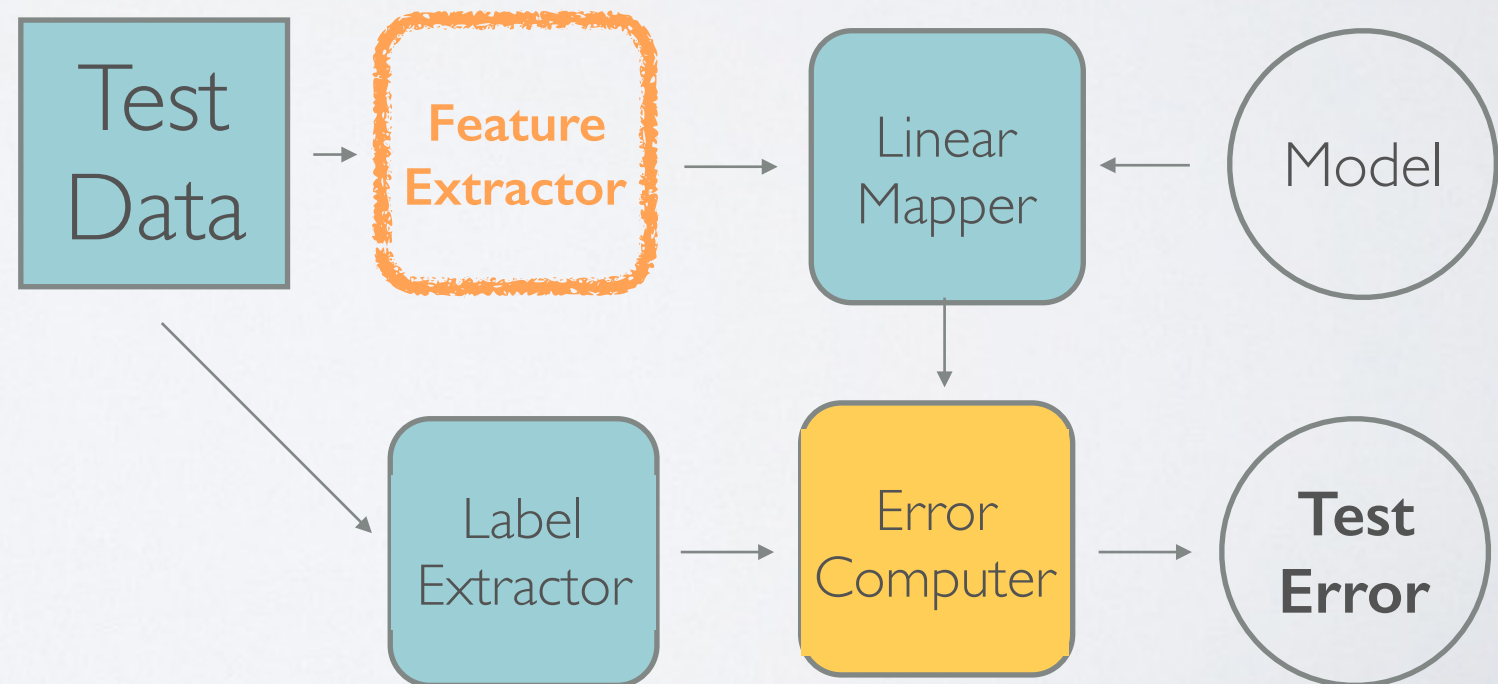
That's more like it!



A REAL PIPELINE FOR
IMAGE CLASSIFICATION
Inspired by Coates & Ng, 2012



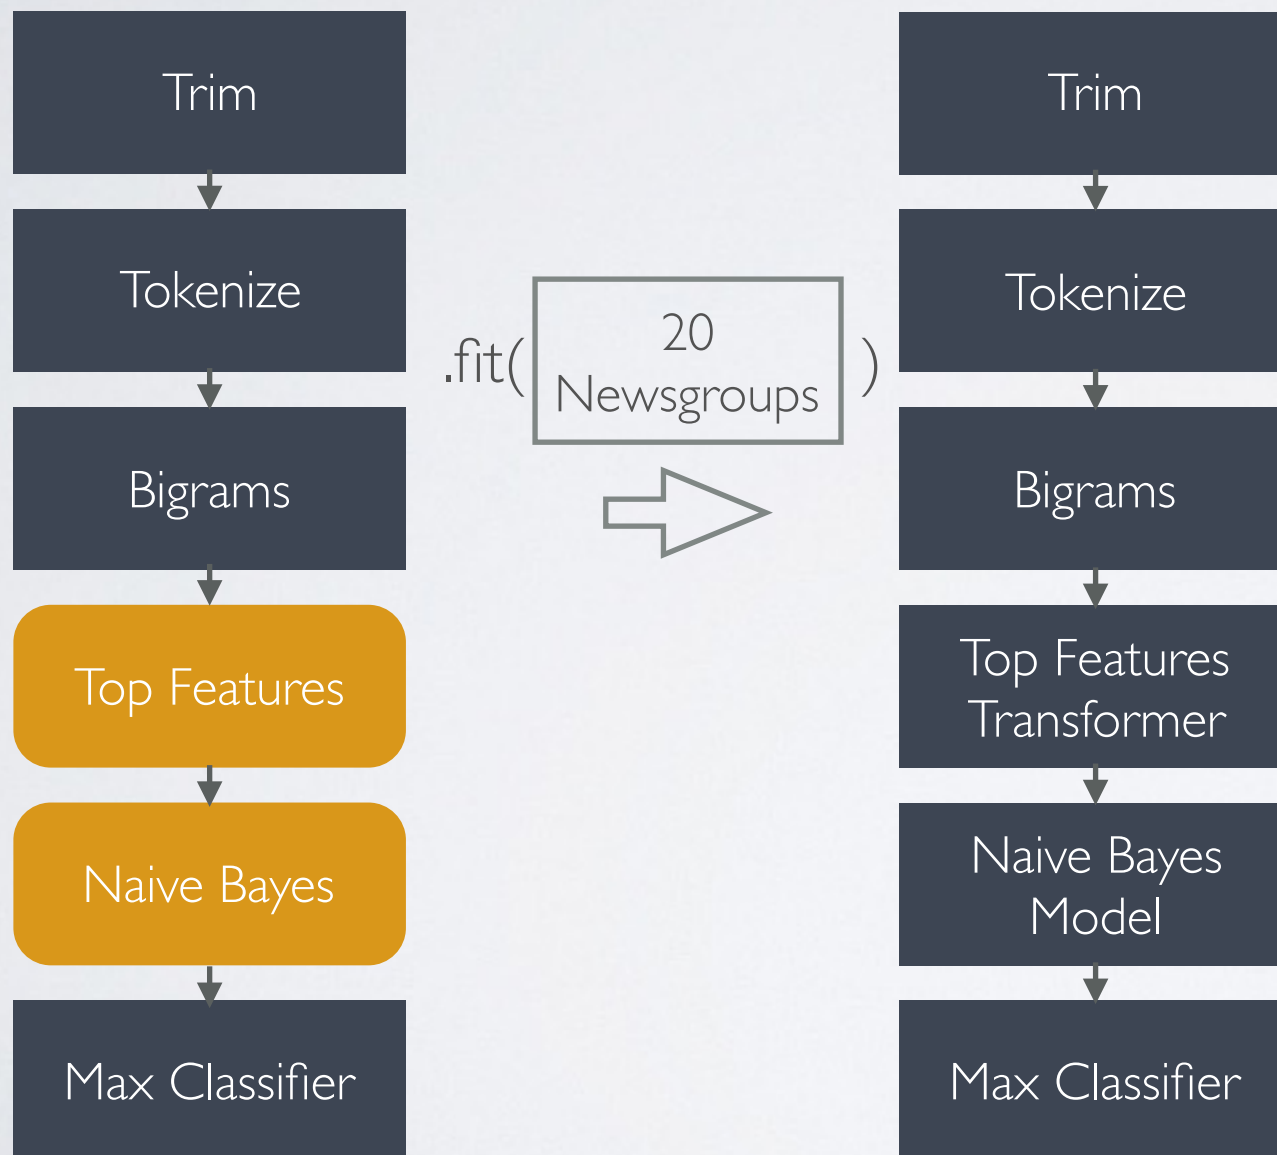
- Embarrassingly Parallel
- Requires Coordination
- Tricky to Scale



ABOUT KEYSTONEML

- Software framework for building **scalable end-to-end** machine learning pipelines on **Apache Spark**.
- Helps us understand what it means to build systems for **robust, scalable**, end-to-end **advanced analytics** workloads and the **patterns** that emerge.
- Example pipelines that achieve **state-of-the-art** results on **large scale datasets** in computer vision, NLP, and speech - **fast**.
- Open source software, available at: <http://keystone-ml.org/>

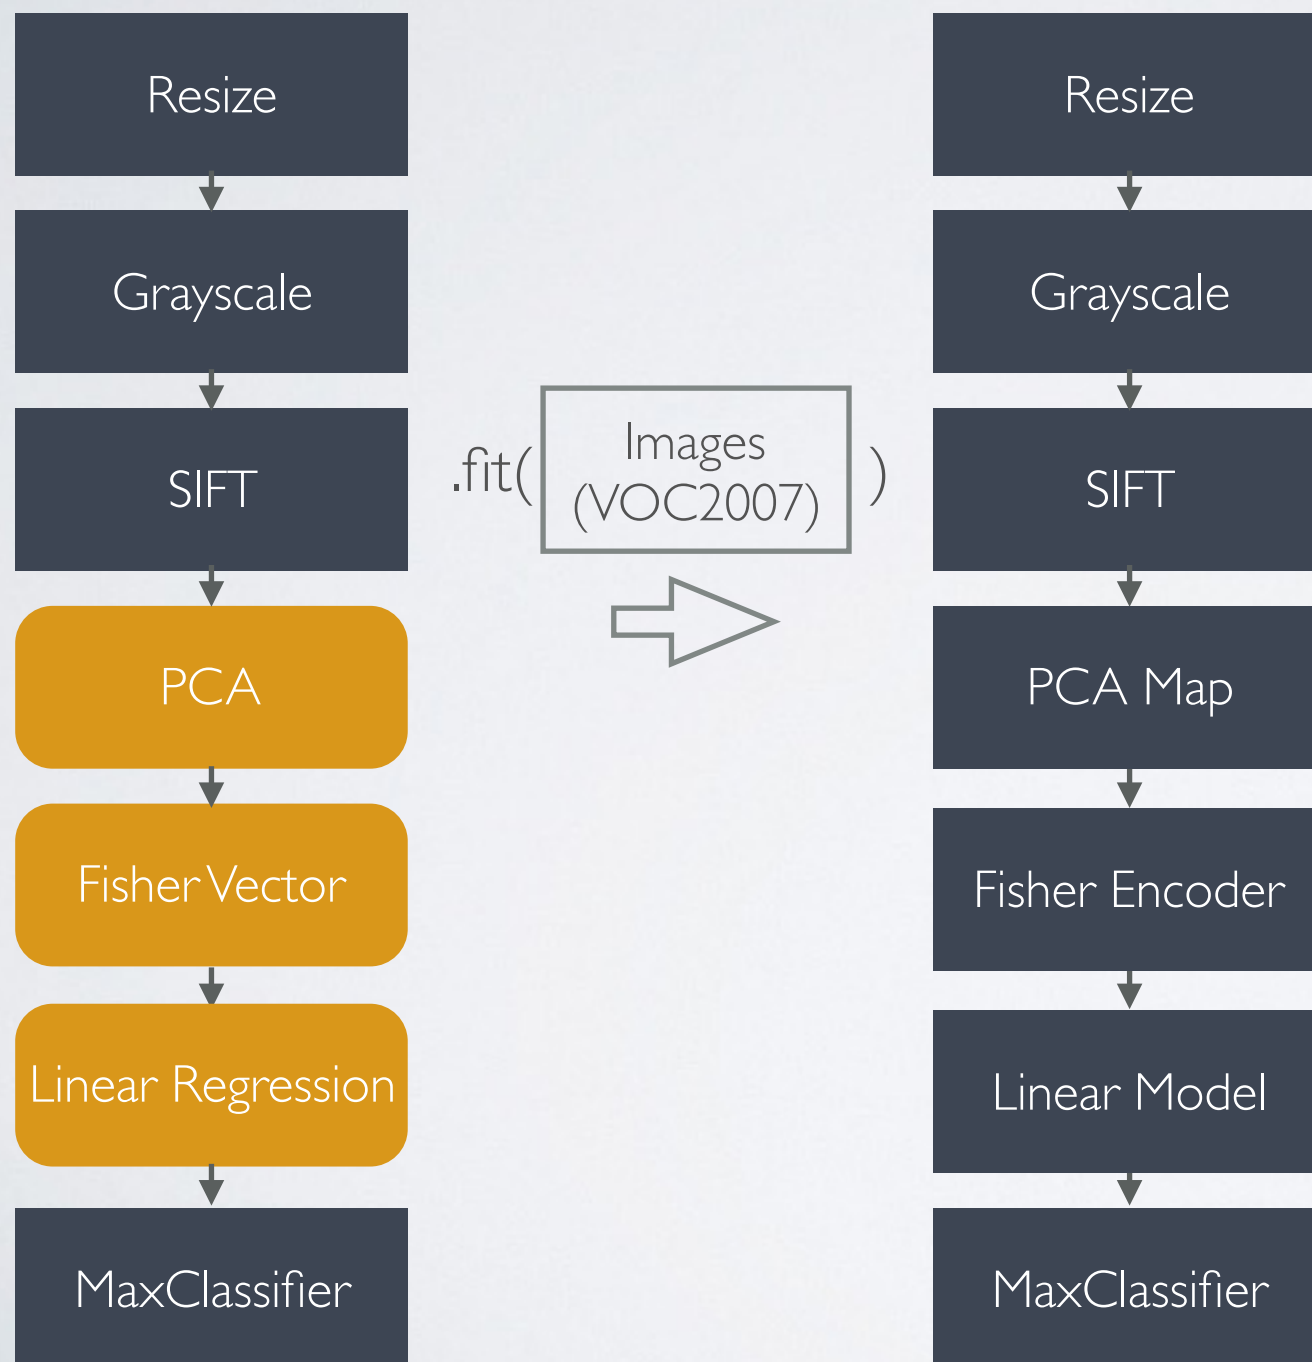
SIMPLE EXAMPLE: TEXT CLASSIFICATION



```
val predictorPipeline = Trim andThen LowerCase() andThen  
Tokenizer() andThen  
NGramsFeaturizer(1 to conf.nGrams) andThen  
TermFrequency(x => 1) andThen  
(CommonSparseFeatures(conf.commonFeatures), training) andThen  
(NaiveBayesEstimator(2), training, labels) andThen  
MaxClassifier
```

Once estimated - apply
these steps to your
production data in an
online or batch fashion.

NOT SO SIMPLE EXAMPLE: IMAGE CLASSIFICATION



```
val predictor = PixelScaler andThen
  GrayScaler andThen
  new Cacher andThen
  new SIFTExtractor(scaleStep = conf.scaleStep) andThen
  new BatchPCATransformer(pca) andThen
  new FisherVector(gmm) andThen
  FloatToDouble andThen
  MatrixVectorizer andThen
  NormalizeRows andThen
  SignedHellingerMapper andThen
  (new LeastSquaresEstimator, trainingData, trainingLabels)
```

5,000 examples, 40,000
features, 20 classes

Pleasantly parallel
featurization and evaluation.

7 minutes on a modest cluster.

Achieves performance
of Chatfield et. al., 2011

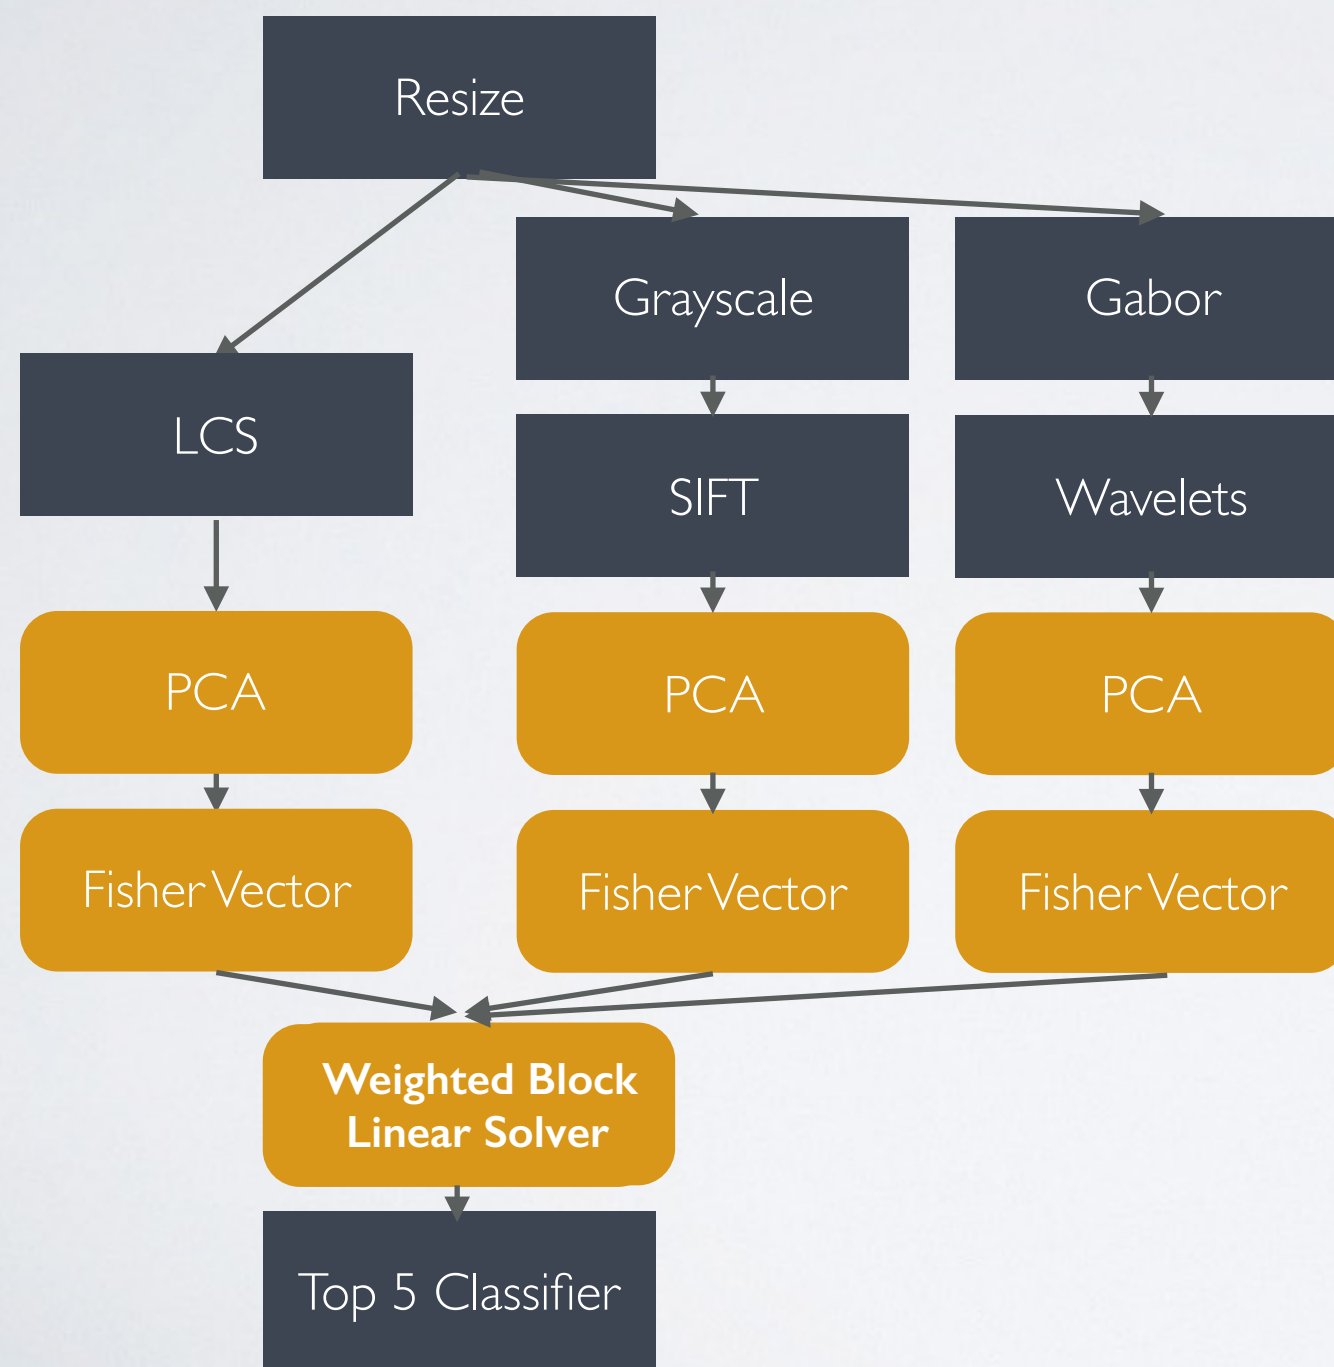
EVEN LESS SIMPLE: IMAGENET

Color

Edges

Texture

< 100 SLOC



1000 class classification.
1,200,000 examples
64,000 features.

90 minutes on 100 nodes.

Upgrading the solver
for higher precision
means changing 1 LOC.

Adding 100,000 more
texture features is easy.

OPTIMIZING KEYSTONEML PIPELINES

High-level API enables rich space of optimizations

Automated ML operator selection.

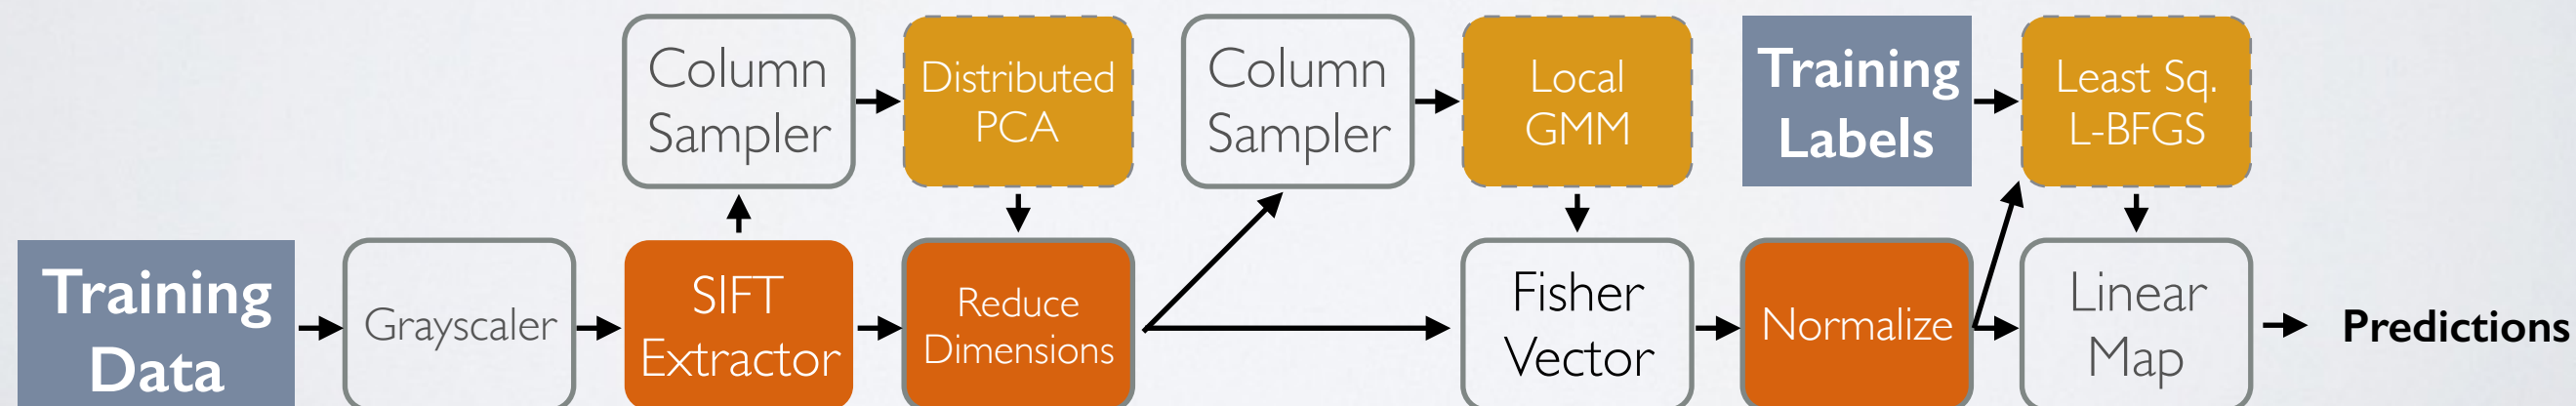
Linear
Solver

Direct
Solver

L-BFGS

Iterative
SGD

Auto-caching for iterative workloads.



KEYSTONEML OPTIMIZER

- Sampling-based **cost model** projects resource usage
 - CPU, Memory, Network
- Utilization tracked through pipeline.
 - Decisions made to minimize total cost of execution.
- Catalyst-based optimizer does the heavy lifting.

Stage	n	d	size (GB)
Input	5000	1m pixel IPEG	0.4
Resize	5000	260k pixels	3.6
Grayscale	5000	260k pixels	1.2
SIFT	5000	65000x128	309
PCA	5000	65000x80	154
FV	5000	256x64x2	1.2
Linear Regression	5000	20	0.0007
Max Classifier	5000	1	0.00009

CHOOSING A SOLVER

- Datasets have a number of interesting degrees of freedom.
 - Problem size (n, d, k)
 - sparsity (nnz)
 - condition number
- Platform has degrees of freedom:
 - Memory, CPU, Network, Nodes
- Solvers are predictable!

Objective:

$$\min_X \|AX - B\|_2^2 + \lambda \|X\|_2^2$$

Where:


$$A \in \mathbb{R}^{n \times d}$$

$$X \in \mathbb{R}^{d \times k}$$

$$B \in \mathbb{R}^{n \times k}$$





CHOOSING A SOLVER

- Three Solvers
 - Exact, Block, LBFGS
- Two datasets
 - Amazon - >99% sparse, $n=65m$
 - TIMIT - dense, $n=2m$
- Exact solve works well for small # features.
- Use LBFGS for sparse problems.
- Block solver scales well to big dense problems.
 - Hundreds of thousands of features.



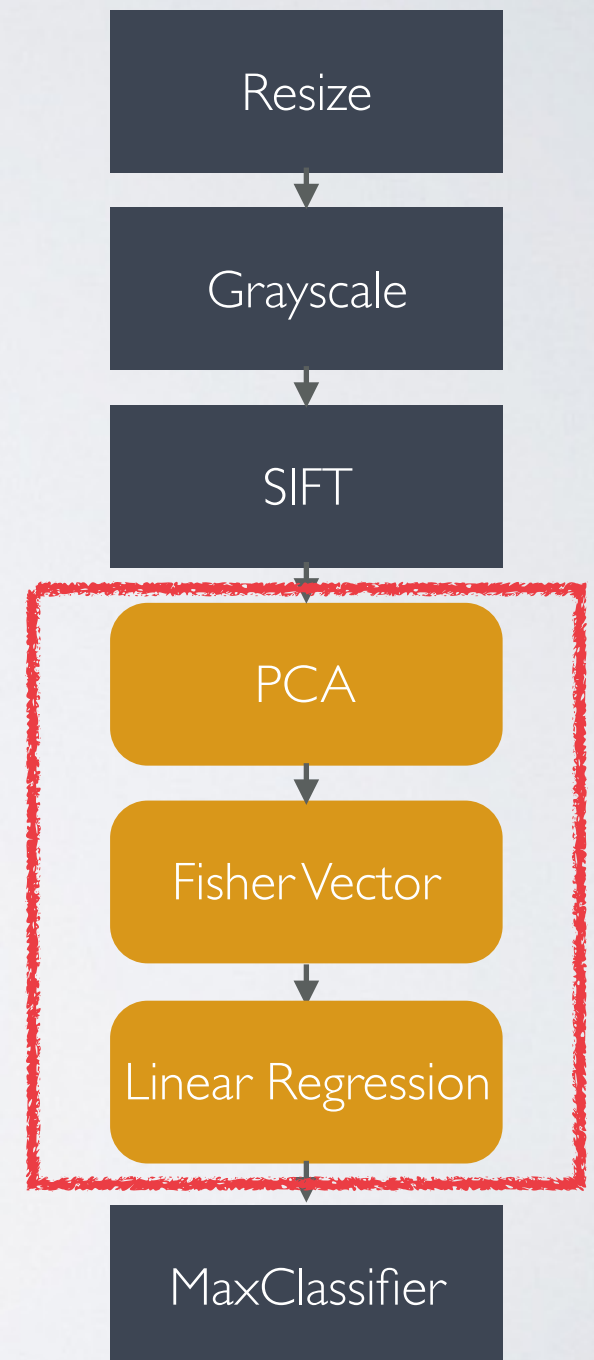
SOLVER PERFORMANCE

- Compared KeystoneML with:
 - VowpalWabbit - specialized system for large, sparse problems.
 - SystemML - general purpose, optimizing ML system.
- Two problems:
 - Amazon - Sparse text features.
 - Binary TIMIT - Dense phoneme data.
- High Order Bit:
 - KeystoneML *pipelines featurization* and *adapts* to workload changes.




DECIDING WHAT TO SAVE

- Pipelines Generate Lots of intermediate state.
 - E.g. SIFT features blow up a 0.42GB VOC dataset to 300GB.
- Iterative algorithms —> state needed many times.
- How do we determine what to save for later and what to reuse, given fixed resource budget?
- Can we **adapt** to workload changes?



CACHING PROBLEM

- Output is computed via depth-first execution of DAG.
 - Caching “truncates” a path after first visit.
- Want to minimize execution time.
 - Subject to memory constraints.
- Picking optimal set is hard!



Cache set	Time	Memory
ABCDE	140s	340g
B	140s	200g
A	180s	50g
{}	240s	0g

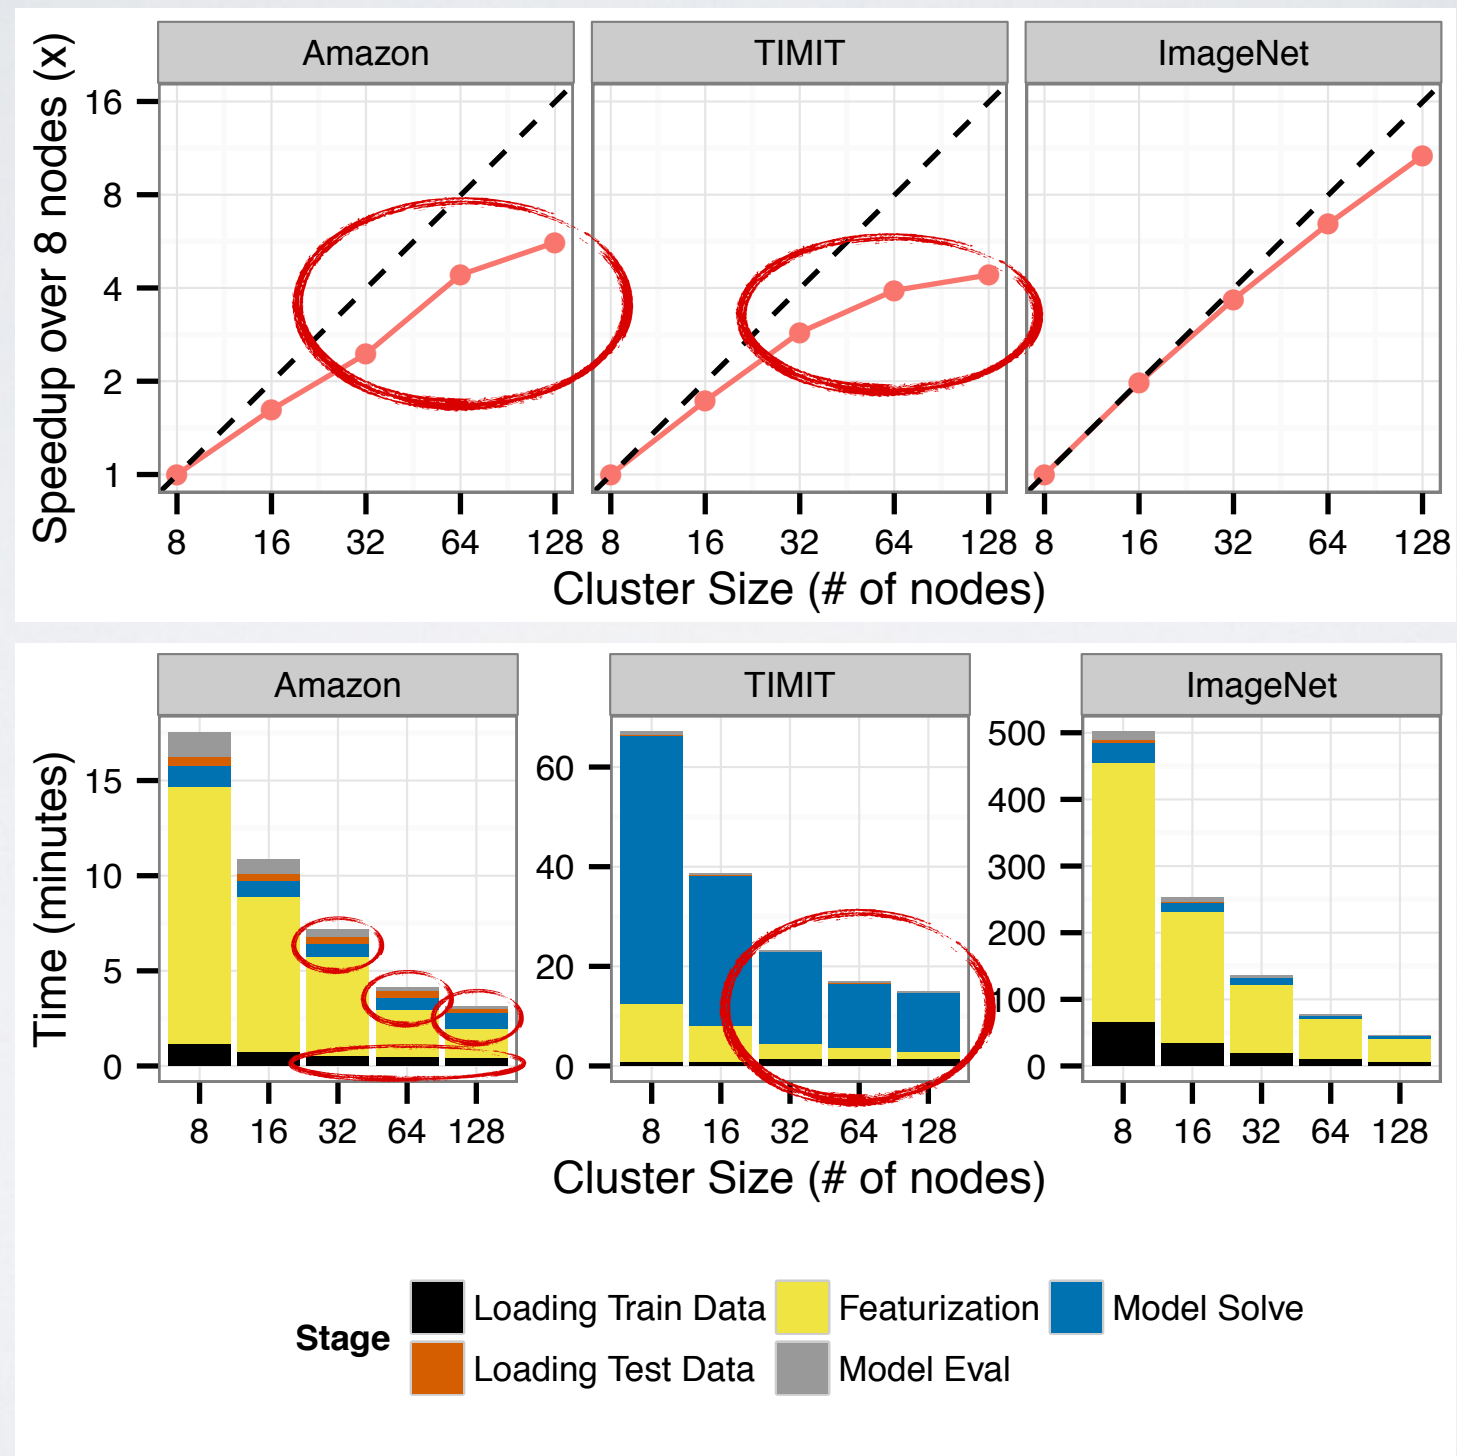
END-TO-END PERFORMANCE

Dataset	Training Examples	Features	Raw Size (GB)	Feature Size (GB)
Amazon	65 million	100k (sparse)	14	89
TIMIT	2.25 million	528k	7.5	8800
ImageNet	1.28 million	262k	74	2500
VOC	5000	40k	0.43	1.5

END-TO-END PERFORMANCE

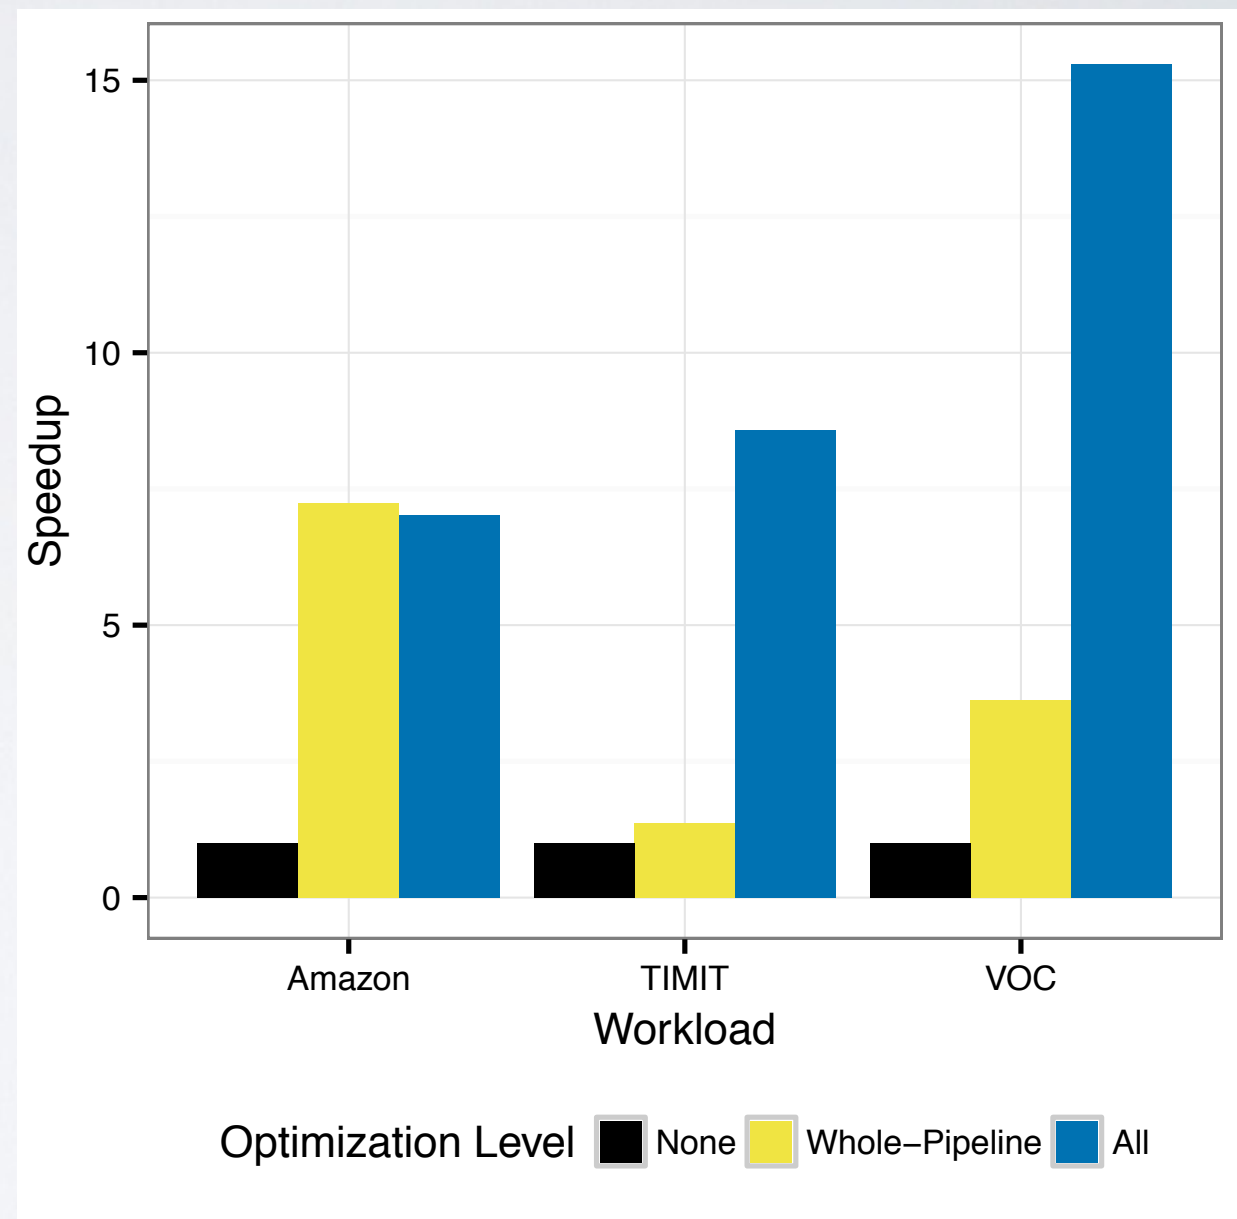
Dataset	KeystoneML Accuracy	Reported Accuracy	KeystoneML Time (m)	Reported Time (m)	Speedup over Reported
Amazon	91.6%	N/A	3.3	N/A	N/A
TIMIT	66.1%	66.3%	138	120	0.87x
ImageNet	67.4%	66.6%	270	5760	21x
VOC	57.2%	59.2%	7	87	12x

END-TO-END PERFORMANCE



END-TO-END PERFORMANCE

- Tested three levels of optimization
 - None
 - Auto-caching only
 - Auto-caching and operator-selection.
- 7x to 15x speedup



QUESTIONS?

Project Page

<http://keystone-ml.org/>

Code

<http://github.com/amplab/keystone>

Training

<http://goo.gl/axbkkc>

BACKUP SLIDES

SOFTWARE FEATURES


- Data Loaders
 - CSV, CIFAR, ImageNet, VOC, TIMIT, 20 Newsgroups
- Transformers
 - NLP - Tokenization, n-gram parsing*
 - Images - Convolution, Gradient Descent, FisherVector*, Pooling, Word Embeddings, CIFAR, VOC, ImageNet
 - Speech - MFCCs*
 - Stats - Random Features, Normalization, Scaling*, Signed Hellinger Mapping, FFT
 - Utility/misc - Caching, Top-K classifier, indicator label mapping, sparse/dense encoding transformers.
- Estimators
 - Learning - Block linear models, Linear Discriminant Analysis, PCA, ZCA Whitening, Naive Bayes*, GMM*
- Example Pipelines
 - NLP - Amazon Product Review Classification, 20 Newsgroups, Wikipedia
- Evaluation Metrics
 - Binary Classification
 - Multiclass Classification
 - Multilabel Classification

Just 11k Lines of Code,
5k of which are Tests or JavaDoc.

* - Links to external library

KEY API CONCEPTS

TRANSFORMERS



```
abstract class Transformer[In, Out] {  
  def apply(in: In): Out  
  def apply(in: RDD[In]): RDD[Out] = in.map(apply)  
  ...  
}
```


TYPE SAFETY HELPS ENSURE ROBUSTNESS

ESTIMATORS




```
abstract class Estimator[In, Out] {  
  def fit(in: RDD[In]): Transformer[In, Out]  
  ...  
}
```

CHAINING

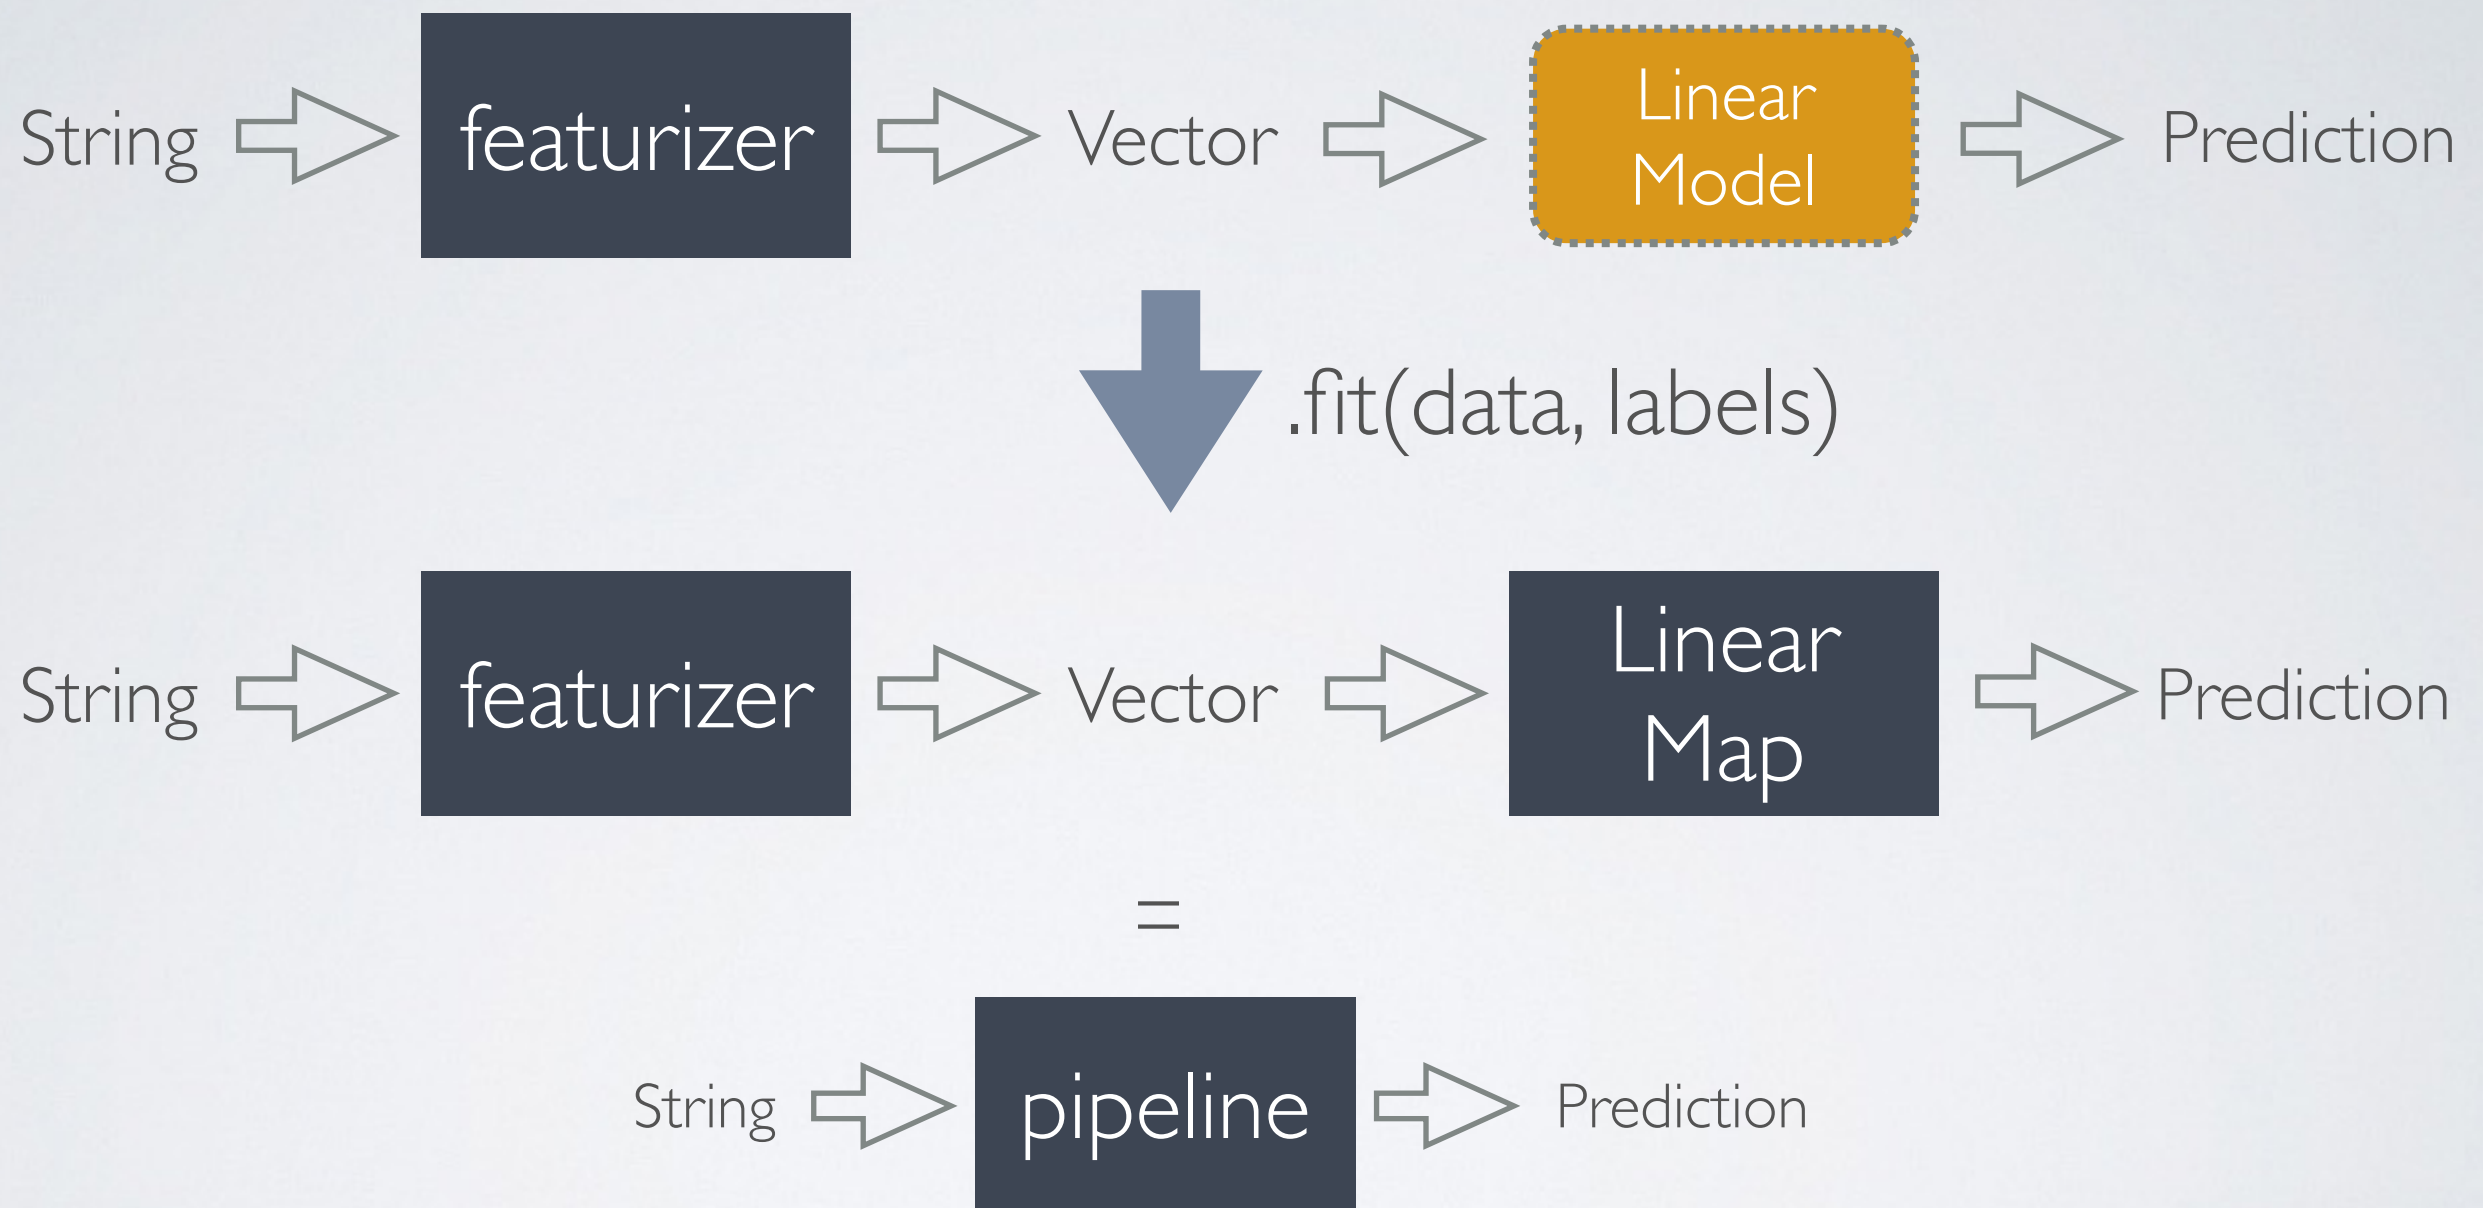


=



```
val featurizer:Transformer[String,Vector] = NGrams(2) then Vectorizer
```


COMPLEX PIPELINES



```
val pipeline = (featurizer thenLabelEstimator LinearModel).fit(data, labels)
```