

# Operational Tips for Deploying Apache® Spark™

Miklos Christine  
Solutions Architect  
Databricks



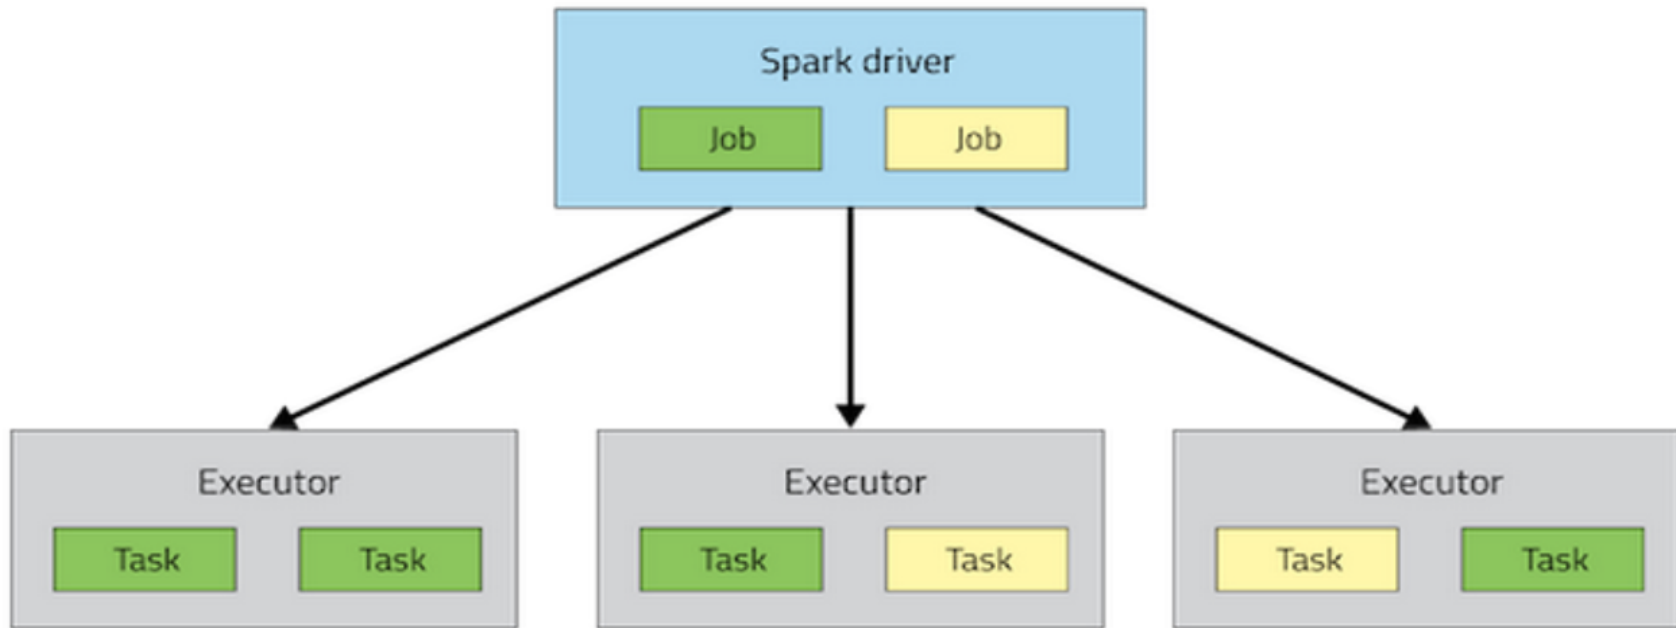
# \$ whoami

- Previously Systems Engineer @ Cloudera
- Deep Knowledge of Big Data Stack
- Apache Spark Expert
- Solutions Architect @ Databricks!

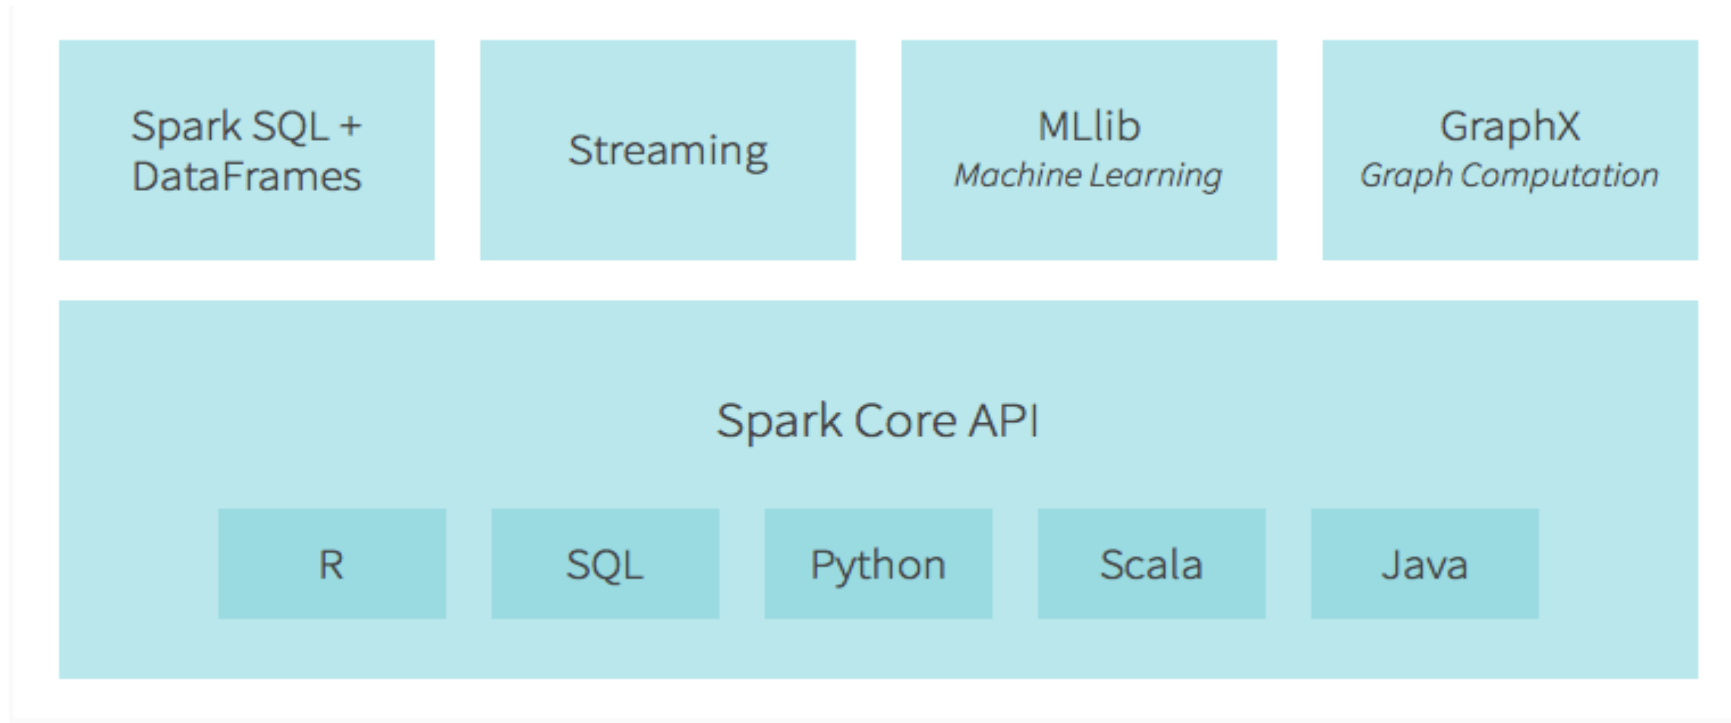
# What Will I Learn?

- Quick Apache Spark Overview
- Configuration Systems
- Pipeline Design Best Practices
- Debugging Techniques

# Apache Spark



# Apache Spark Configuration



# Apache Spark Core Configurations

- **Command Line:**  
`spark-defaults.conf`  
`spark-env.sh`  
  
`// Print SparkConfig`  
`sc.getConf.toDebugString`  
  
`// Print Hadoop Config`
- **Programmatically:**  
`SparkConf()`  
  
`val hdConf =`  
`sc.hadoopConfiguration.iterator()`  
`while (hdConf.hasNext){`  
`println(hdConf.next().toString())`  
`}`
- **Hadoop Configs:**  
`core-site.xml`  
`hdfs-site.xml`

# SparkSQL Configurations

- Set SQL Configs Through SQL Interface

```
SET key=value;
```

```
sqlContext.sql("SET spark.sql.shuffle.partitions=10;")
```

- Tools to see current configurations

```
// View SparkSQL Config Properties
```

```
val sqlConf = sqlContext.getAllConfs
```

```
sqlConf.foreach(x => println(x._1 + " : " + x._2))
```

# Apache Spark Pipeline Design



- File Formats
- Compression Codecs
- Apache Spark APIs
- Job Profiles



# File Formats

- Text File Formats
  - CSV
  - JSON
- Avro Row Format
- Parquet Columnar Format

A	B	C
A1	B1	C1
A2	B2	C2
A3	B3	C3



User Story:

260GB CSV Data Converted to 23GB Parquet

# Compression


- Choose and Analyze Compression Codecs
  - Snappy, Gzip, LZ0
- Configuration Parameters
  - `io.compression.codecs`
  - `spark.sql.parquet.compression.codec`

# Small Files Problem

- Small files problem still exists
- Metadata loading
- Use `coalesce()`

```
display(df)
```

▶ (2) Spark Jobs




Ref:

<http://spark.apache.org/docs/latest/api/python/pyspark.sql.html#pyspark.sql.DataFrame>

# What are Partitions?

- 2 Types of Partitioning
  - Spark
  - Table Level

```
# Get Number of Spark  
> df.rdd.getNumPartitions()  
40
```



# Partition Controls

- **Apache Spark APIs**

- `repartition()`
- `coalesce()`

```
# Re-partition a DataFrame  
> df_10 = df.repartition(10)
```

```
df = sqlContext.read.\  
    jdbc(url=jdbcUrl, \  
          table='employees', \  
          column='emp_no', \  
          lowerBound=1, \  
          upperBound=100000, \  
          numPartitions=100)  
  
df.repartition(20).write.\  
    parquet('/mnt/mwc/jdbc_part/')
```

# Table Partitions

- Partition by a column value within the table

```
> df.write.\
  partitionBy("colName").\
  saveAsTable("tableName")
```

```
ls -l /dbfs/user/hive/warehouse/tableName/
```

```
total 0
-rw-r--r-- 1 root root 0 Jan  1  1970 _SUCCESS
drwxr-xr-x 1 root root 0 Jan  1  1970 year=2014
drwxr-xr-x 1 root root 0 Jan  1  1970 year=2015
```

```
Command took 0.38s
```

# Those Other Partitions

- **SparkSQL Shuffle Partitions**

`spark.sql.shuffle.partitions`

```
sqlCtx.sql("set spark.sql.shuffle.partitions=600")  
sqlCtx.sql("select a1.name, a2.name from adult a1 \  
    join adult a2 \  
    where a1.age = a2.age")
```

```
sqlCtx.sql("select count(distinct(name)) from adult")
```

# How Does This Help?

- Q: Will increasing my cluster size help with my job?
- A: It depends.

```
> %sql
-- 1.6
set spark.sql.shuffle.partitions=200;
select count(distinct(author)) from reddit_all where year = 2015;
```

▼ (1) Spark Jobs

▼ Job 13 [View](#) (Stages: 3/3)

Stage 13: 240/240 ⓘ

Stage 14: 200/200 ⓘ

Stage 15: 1/1 ⓘ

\_c0

5006186



Command took 28.17s

```
> %sql
-- 1.6
set spark.sql.shuffle.partitions=40;
select count(distinct(author)) from reddit_all where year = 2015;
```

▼ (1) Spark Jobs

▼ Job 14 [View](#) (Stages: 3/3)

Stage 16: 240/240 ⓘ

Stage 17: 40/40 ⓘ

Stage 18: 1/1 ⓘ

\_c0

5006186



Command took 19.53s



# Apache Spark Job Profiles

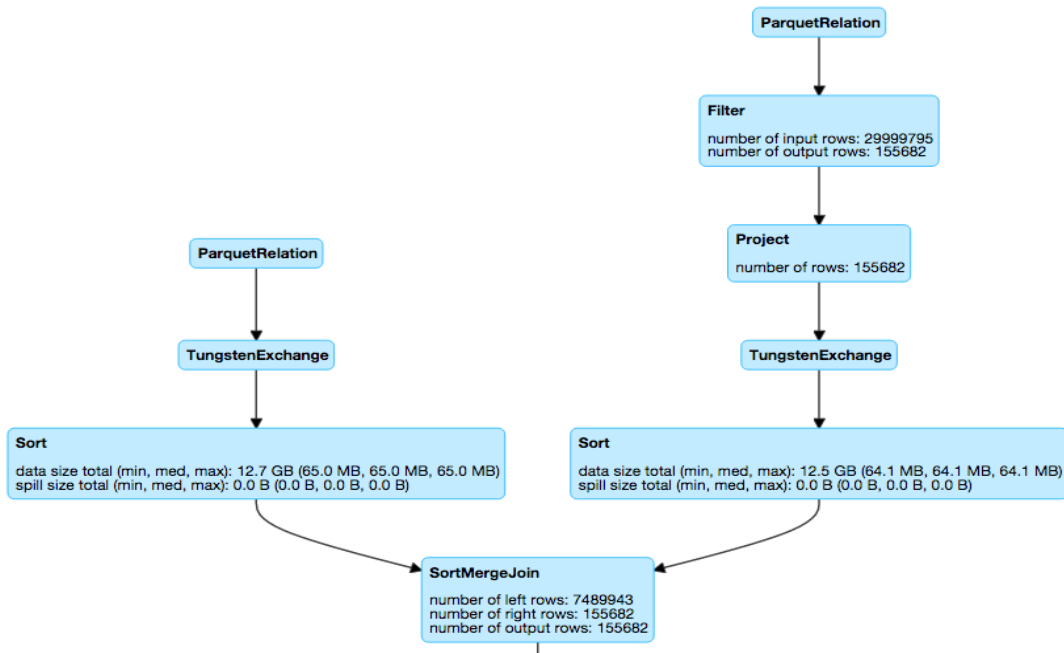
- Leverage Spark UI
  - SQL
  - Streaming

## Details for Query 49

Submitted Time: 2016/02/10 02:06:25

Duration: 12 s

Succeeded Jobs: 85



# Apache Spark Job Profiles

## Aggregated Metrics by Executor

Executor ID ▲	Address	Task Time	Total Tasks	Failed Tasks	Succeeded Tasks	Input Size / Records	Shuffle Write Size / Records
0	ip-10-0-197-52.us-west-2.compute.internal:45222	39 min	291	0	291	10.0 GB / 82423	11.9 KB / 291
1	ip-10-0-238-218.us-west-2.compute.internal:41578	39 min	292	0	292	10.0 GB / 82795	12.0 KB / 292
2	ip-10-0-229-111.us-west-2.compute.internal:43837	39 min	294	0	294	10.1 GB / 82721	12.1 KB / 294
3	ip-10-0-176-36.us-west-2.compute.internal:47072	39 min	294	0	294	10.2 GB / 84025	12.1 KB / 294
4	ip-10-0-207-37.us-west-2.compute.internal:41034	39 min	293	0	293	10.1 GB / 82656	12.0 KB / 293
5	ip-10-0-156-34.us-west-2.compute.internal:44449	39 min	298	0	298	10.2 GB / 83917	12.2 KB / 298
6	ip-10-0-156-35.us-west-2.compute.internal:58803	39 min	298	0	298	10.2 GB / 84277	12.2 KB / 298
7	ip-10-0-232-3.us-west-2.compute.internal:55489	39 min	297	0	297	10.3 GB / 84628	12.2 KB / 297
8	ip-10-0-232-2.us-west-2.compute.internal:38085	39 min	296	0	296	10.2 GB / 83582	12.1 KB / 296
9	ip-10-0-207-36.us-west-2.compute.internal:55661	39 min	288	0	288	9.9 GB / 81804	11.8 KB / 288
10	ip-10-0-178-163.us-west-2.compute.internal:59616	39 min	294	0	294	10.1 GB / 83350	12.1 KB / 294
11	ip-10-0-178-162.us-west-2.compute.internal:41800	39 min	295	0	295	10.2 GB / 83822	12.1 KB / 295

# Apache Spark Job Profiles

## Tasks

Index ▲	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	GC Time	Input Size / Records	Errors
0	16607	0	SUCCESS	PROCESS_LOCAL	18 / ip-10-0-171-33.us-west-2.compute.internal	2016/02/09 19:12:40	24 s	0.2 s	46.8 MB (hadoop) / 614400	
1	16608	0	SUCCESS	PROCESS_LOCAL	19 / ip-10-0-214-114.us-west-2.compute.internal	2016/02/09 19:12:40	23 s		46.8 MB (hadoop) / 614400	
2	16609	0	SUCCESS	PROCESS_LOCAL	18 / ip-10-0-171-33.us-west-2.compute.internal	2016/02/09 19:12:40	4 s		6.5 MB (hadoop) / 85355	
3	16610	0	SUCCESS	PROCESS_LOCAL	19 / ip-10-0-214-114.us-west-2.compute.internal	2016/02/09 19:12:40	23 s		46.8 MB (hadoop) / 614400	
4	16611	0	SUCCESS	PROCESS_LOCAL	18 / ip-10-0-171-33.us-west-2.compute.internal	2016/02/09 19:12:40	24 s	0.2 s	46.8 MB (hadoop) / 614400	

# Apache Spark Job Profile: Metrics

- Monitoring & Metrics
  - Spark
  - Servers
- Toolset
  - Ganglia
  - Graphite



Ref:

<http://www.hammerlab.org/2015/02/27/monitoring-spark-with-graphite-and-grafana/>

# Debugging Apache Spark

- Analyze the Driver's stacktrace.
- Analyze the executors stacktraces
  - Find the initial executor's failure.
- Review metrics
  - Memory
  - Disk
  - Networking

# Debugging Apache Spark

- Know your tools: JDBC vs ODBC. How to test? What can I test?
  - RedShift / Mysql / Tableau to Apache Spark ,etc.

- Json SparkSQL for corrupt records

```
sqlCtx.read.json("/jsonFiles/").registerTempTable("jsonTable")  
sqlCtx.sql("SELECT _corrupt_record \  
            FROM jsonTable \  
            WHERE _corrupt_record IS NOT NULL")
```

- Steps to debug SQL issues
  - Where's the data, what's the DDL?

# Top Support Issues

- OutOfMemoryErrors
  - Driver
  - Executors
- Out of Disk Space Issues
- Long GC Pauses
- API Usage

# Top Support Issues

- Use builtin functions instead of custom UDFs
  - import [pyspark.sql.functions](#)
  - import org.apache.spark.sql.functions
- Examples:
  - to\_date()
  - get\_json\_object()
  - monotonically\_increasing\_id()
  - hour() / minute()

Ref: <http://spark.apache.org/docs/latest/api/python/pyspark.sql.html#module-pyspark.sql.functions>



# Top Support Issues

- SQL query not returning new data
  - `REFRESH TABLE <table_name>`
- Exported Parquet from External Systems
  - `spark.sql.parquet.binaryAsString`
- Tune number of Shuffle Partitions
  - `spark.sql.shuffle.partitions`

# Try Apache Spark with Databricks

- Download notebook for this talk at: [dbricks.co/xyz](https://dbricks.co/xyz)
- Try latest version of Apache Spark and preview of Spark 2.0

<http://databricks.com/try>

Create Cluster

New Cluster

Cluster Name

Spark Version

- ✓ Spark 2.0 (apache/branch-2.0 preview)
- Spark 1.3.0 (Hadoop 1)
- Spark 1.4.1 (Hadoop 1)
- Spark 1.5.2 (Hadoop 1)
- Spark 1.6.0 (Hadoop 1)
- Spark 1.6.1 (Hadoop 1)
- Spark 1.6.1 (Hadoop 2)

... automatically t  
... de your Databri

# Thank you.

[mwc@databricks.com](mailto:mwc@databricks.com)

<https://www.linkedin.com/in/mrchristine>

@Miklos\_C

