

What no one tells you about writing a streaming app

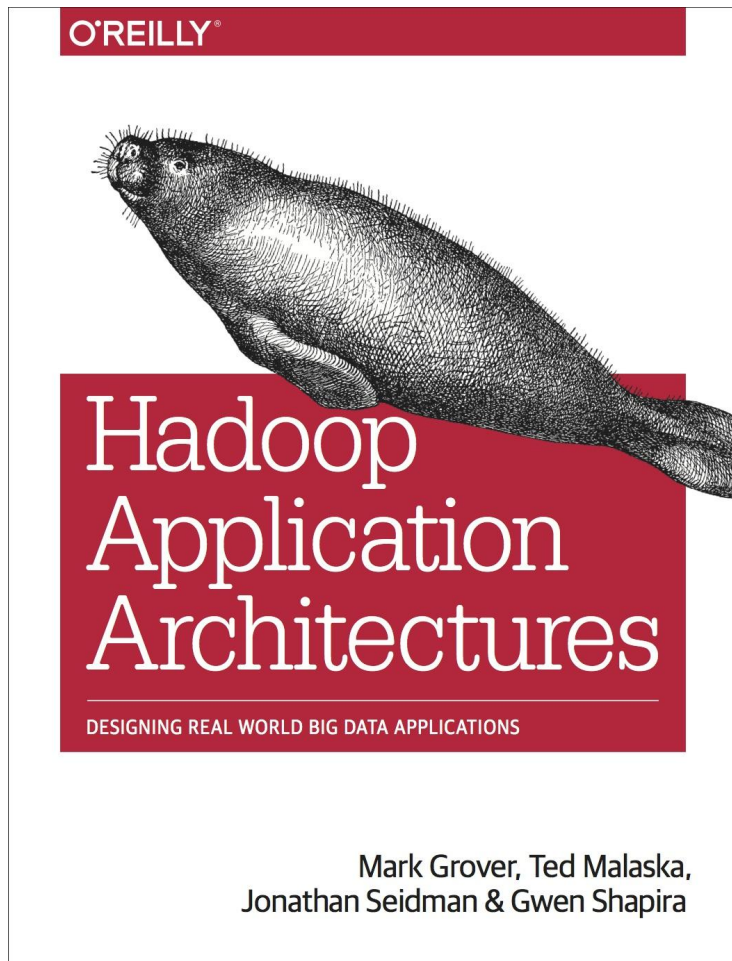
Mark Grover (@mark_grover) - Software Engineer, Cloudera
Ted Malaska (@TedMalaska) - Group Technical Architect, Blizzard

tiny.cloudera.com/streaming-app



Book

- Book signing, 6:30pm
 - Cloudera booth
- hadooparchitecturebook.com
- [@hadooparchbook](https://twitter.com/hadooparchbook)
- github.com/hadooparchitecturebook
- slideshare.com/hadooparchbook



Agenda

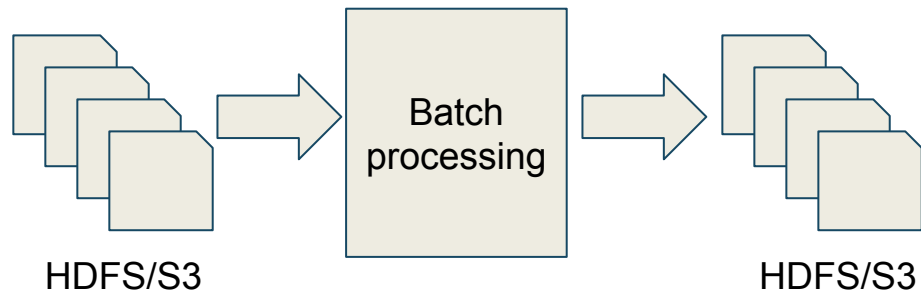
- 5 things no one tells you about writing a streaming app

#5 - Monitoring and managing jobs

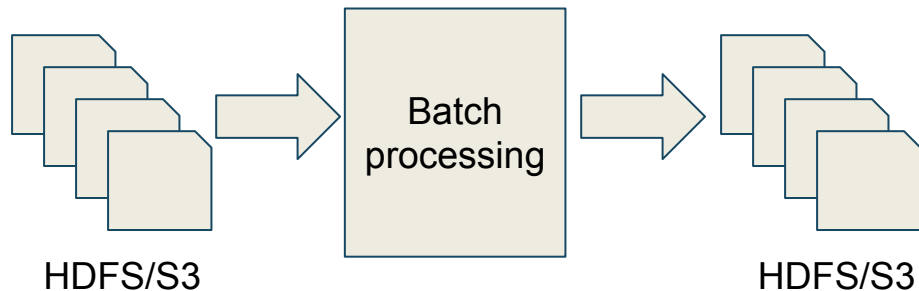


**SPARK
SUMMIT**
EAST 2017

Batch systems

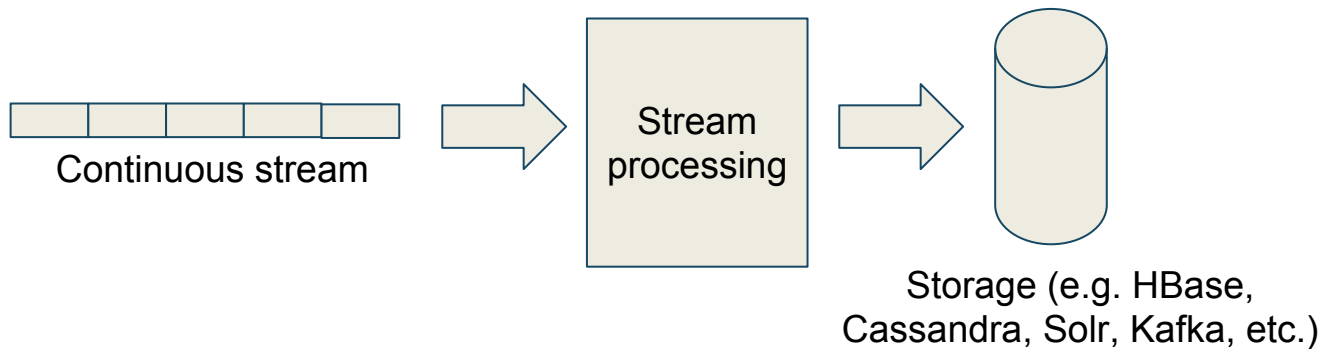


Managing and monitoring batch systems



- Use Cron/Oozie/Azkaban/Luigi for orchestration
- Validation logic in the job (e.g. if input dir is empty)
- Logs aggregated at end of the job
 - Track times against previous runs, etc.

Streaming systems systems



- “Automatic” orchestration (micro-batching)
- Long running driver process

Not originally built for streaming

- YARN
 - Doesn't aggregate logs until job finishes
- Spark
 - Checkpoints can't survive app or Spark upgrades
 - Need to clear checkpoint directory during upgrade

Big questions remain unanswered

1. Management

- a. Where to run the driver?
- b. How to restart the driver automatically if it fails?
- c. How to pause the job?

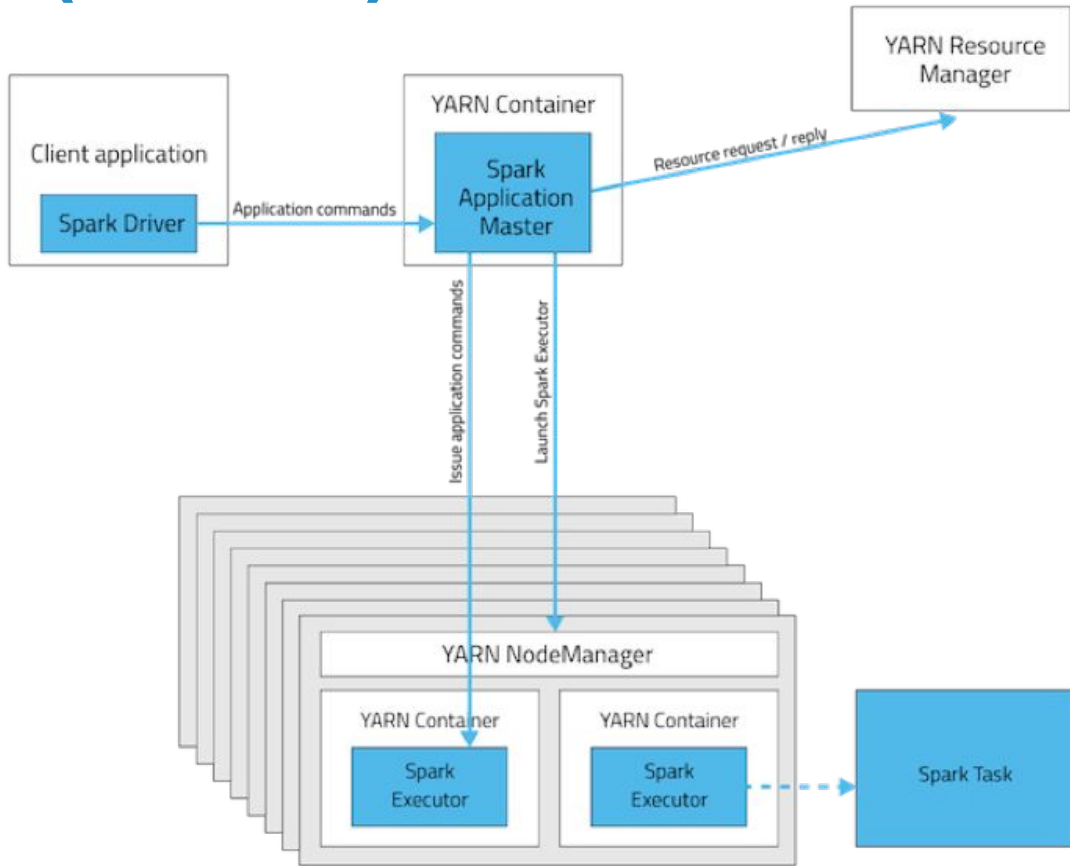
2. Monitoring

- a. How to prevent backlog i.e. make sure processing time $<$ batch interval?
- b. How to keep tabs on health of the long running driver process, etc.?

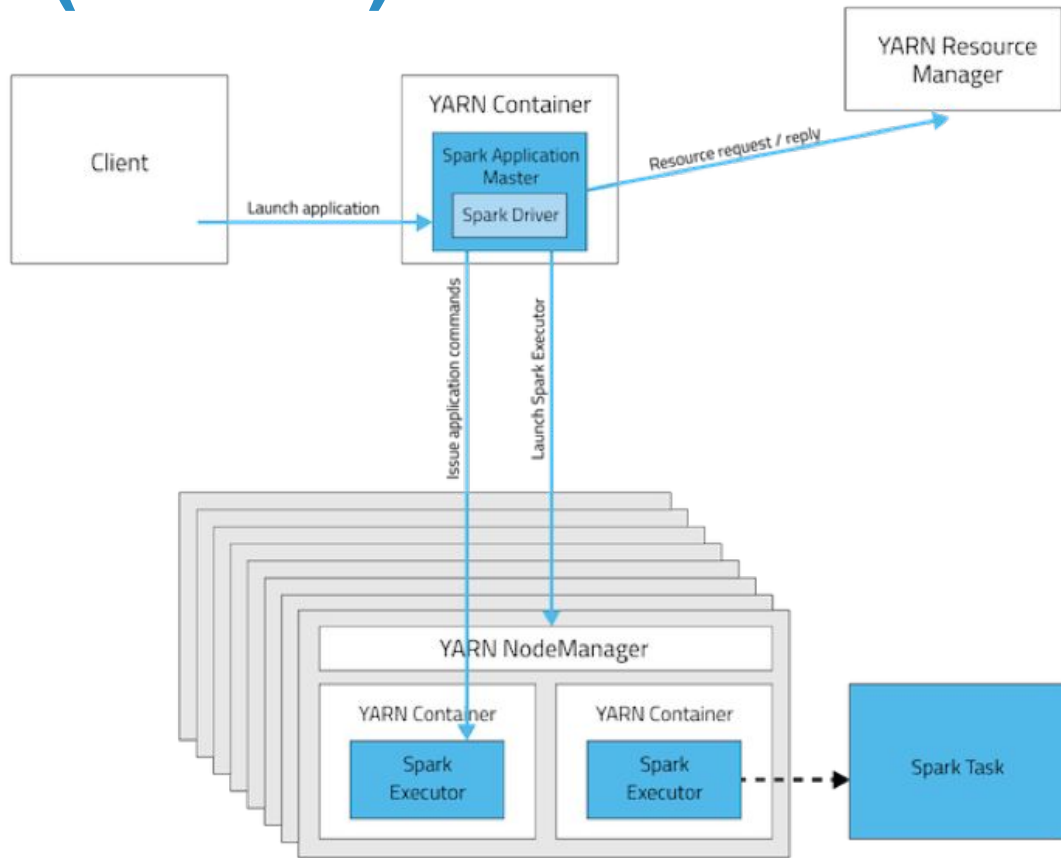
Disclaimer

- Most discussion that follows corresponds to YARN but can be applied to other cluster managers as well.

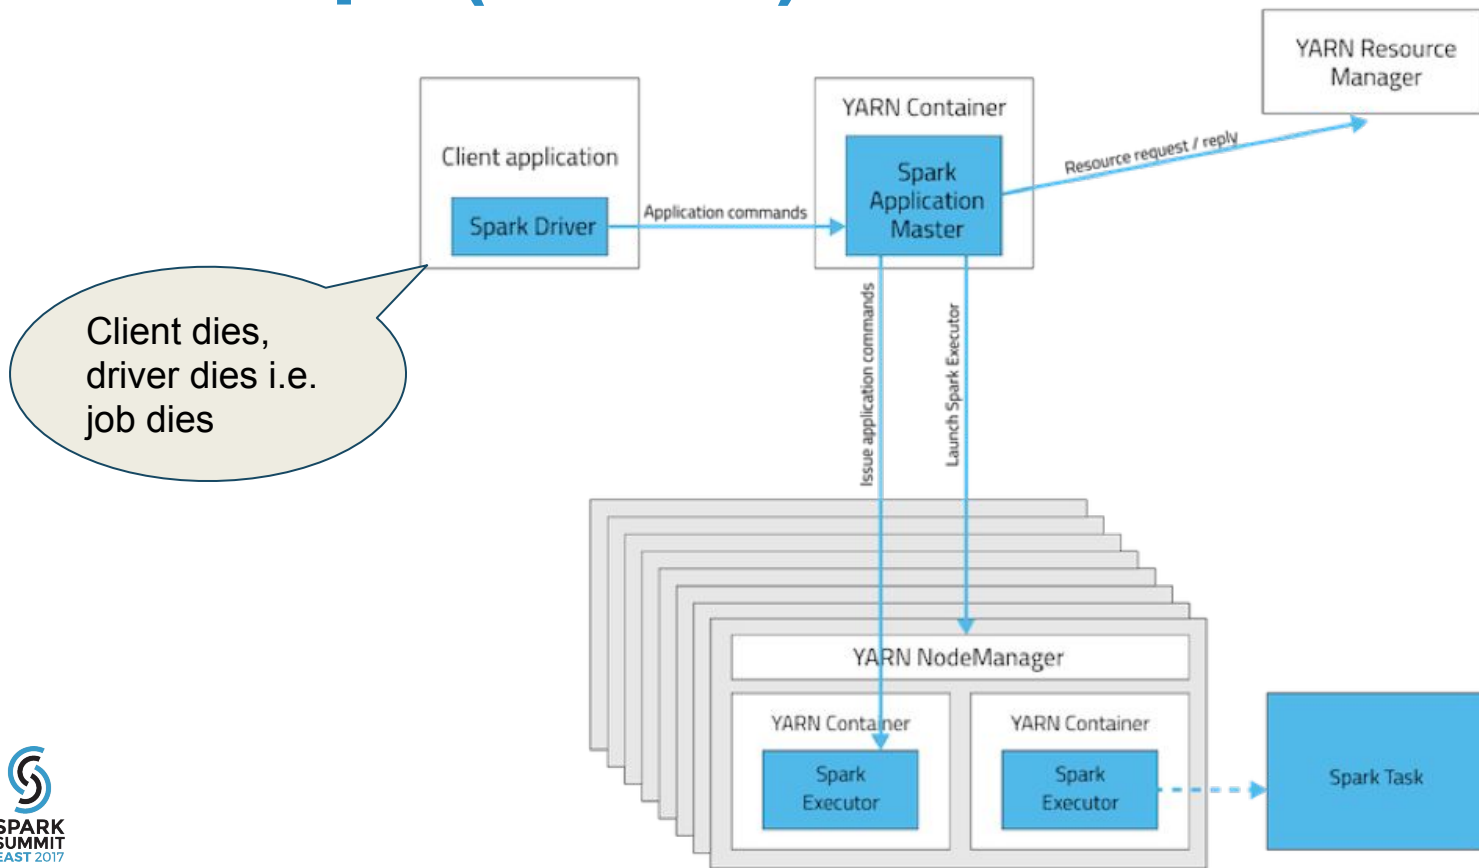
Recap: (YARN) Client mode



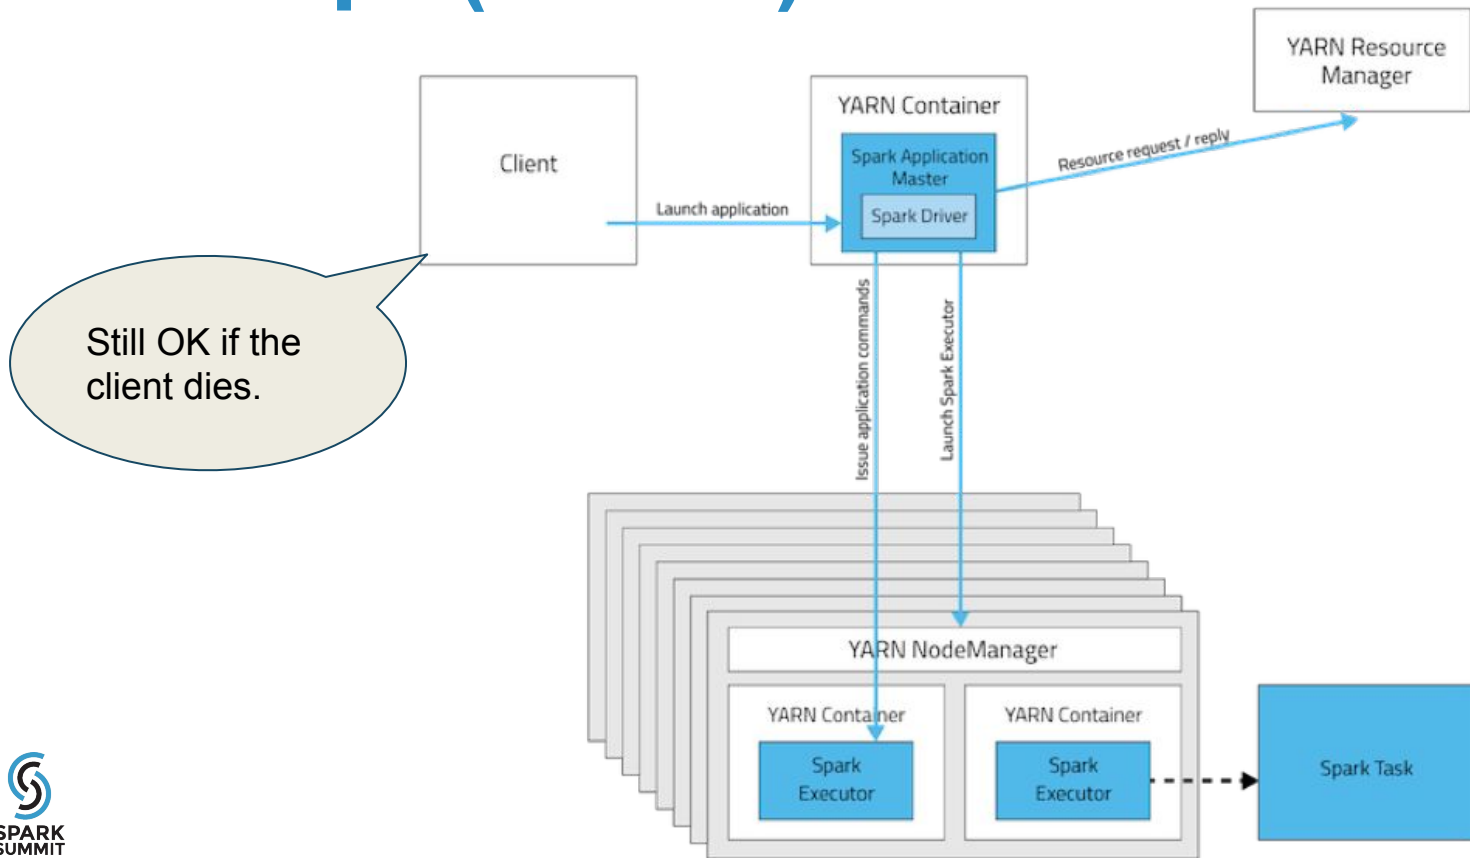
Recap: (YARN) Cluster mode



Recap: (YARN) Client mode



Recap: (YARN) Cluster mode



1a. Where to run the driver?

- Run on YARN Cluster mode
 - Driver will continue running when the client machine goes down

1b. How to restart driver?

- Set up automatic restart

<code>spark.yarn.maxAppAttempts</code>	<code>yarn.resourcemanager.am.max-attempts</code> in YARN	The maximum number of attempts that will be made to submit the application. It should be no larger than the global number of max attempts in the YARN configuration.
<code>spark.yarn.am.attemptFailuresValidityInterval</code>	(none)	Defines the validity interval for AM failure tracking. If the AM has been running for at least the defined interval, the AM failure count will be reset. This feature is not enabled if not configured, and only supported in Hadoop 2.6+.

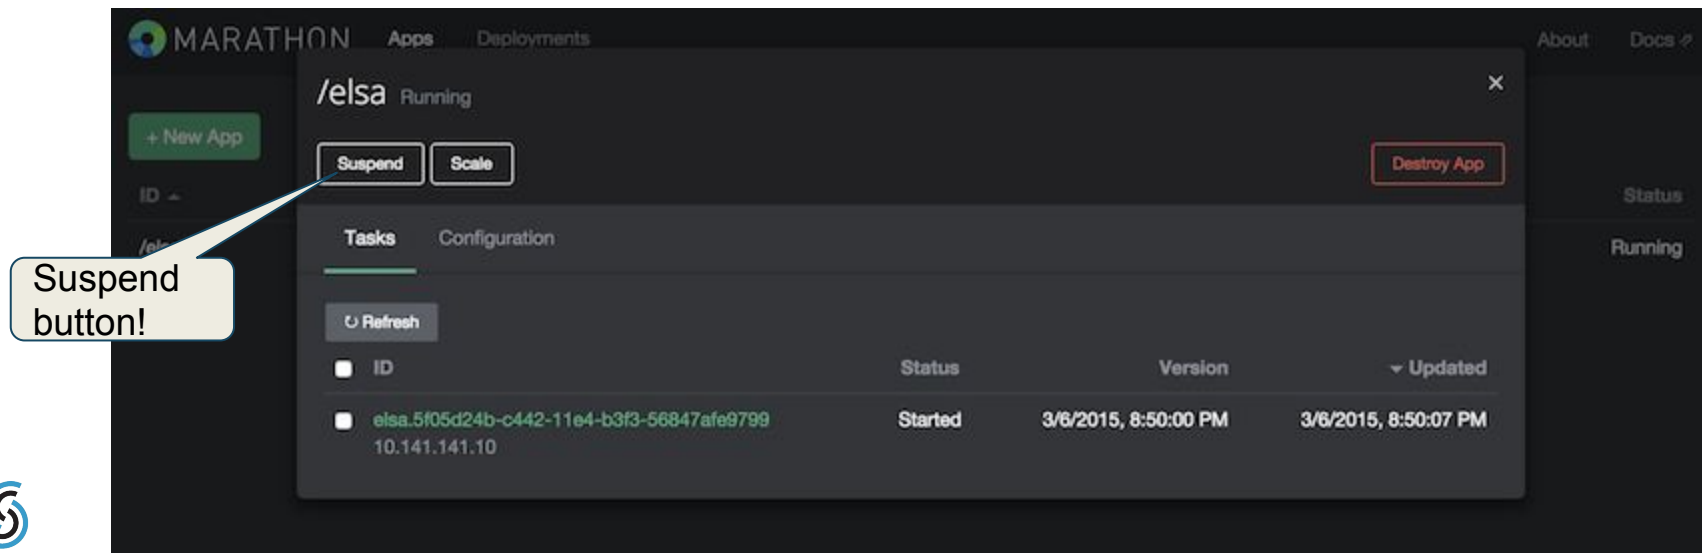
In spark configuration (e.g. `spark-defaults.conf`):

```
spark.yarn.maxAppAttempts=2
```

```
spark.yarn.am.attemptFailuresValidityInterval=1h
```


1c. How to pause a job?

- If running on Mesos, use Marathon:
- See “Graceful Shutdown” later for YARN, etc.



2. Monitoring - Spark Streaming UI

Streaming

Started at: Wed Apr 09 22:47:31 PDT 2014

Time since start: 50 seconds 97 ms

Network receivers: 5

Batch interval: 1 second

Processed batches: 50

Waiting batches: 0

Statistics over last 50 processed batches

Network Input Statistics

Receiver	Location	Records in last batch	Minimum rate [records/sec]	25th percentile rate [records/sec]	Median rate [records/sec]	75th percentile rate [records/sec]	Maximum rate [records/sec]
SocketReceiver-0	localhost	10.7 K	1263.0	9.4 K	10.4 K	11.0 K	12.4 K
SocketReceiver-1	localhost	10.6 K	1177.0	9.3 K	10.3 K	10.9 K	12.4 K
SocketReceiver-2	localhost	10.7 K	1258.0	9.3 K	10.4 K	10.9 K	12.5 K
SocketReceiver-3	localhost	10.6 K	1217.0	9.4 K	10.4 K	10.9 K	12.4 K
SocketReceiver-4	localhost	10.7 K	1245.0	9.3 K	10.4 K	11.0 K	12.5 K

Batch Processing Statistics

Metric	Last batch	Minimum	25th percentile	Median	75th percentile	Maximum
Processing Time	100 ms	96 ms	113 ms	122 ms	145 ms	534 ms
Scheduling Delay	0 ms	0 ms	0 ms	0 ms	0 ms	5 ms
Total Delay	100 ms	96 ms	113 ms	122 ms	146 ms	539 ms

Processing time,
scheduling delay, total
delay, etc.

Base image from: <http://i.imgur.com/1ooDGhm.png>

2. Monitoring - But I want more!

- Spark has a configurable metrics system
 - Based on Dropwizard Metrics Library
- Use Graphite/Grafana to dashboard metrics like:
 - Number of records processed
 - Memory/GC usage on driver and executors
 - Total delay
 - Utilization of the cluster

Summary - Managing and monitoring

- Afterthought
- But it's possible
- In Structured Streaming, use `StreamingQueryListener` (starting Apache Spark 2.1)

References

- Spark Streaming on YARN by Marcin Kuthan

<http://mkuthan.github.io/blog/2016/09/30/spark-streaming-on-yarn/>

- Marathon

<https://mesosphere.github.io/marathon/>

#4 - Prevent data loss

Prevent data loss

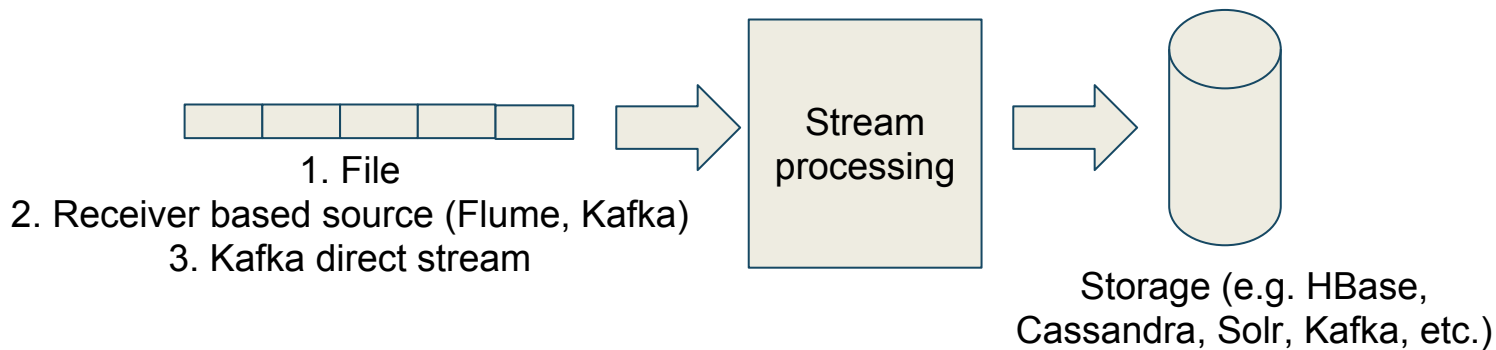
- Ok, my driver is automatically restarting
- Can I lose data in between driver restarts?

No data loss!

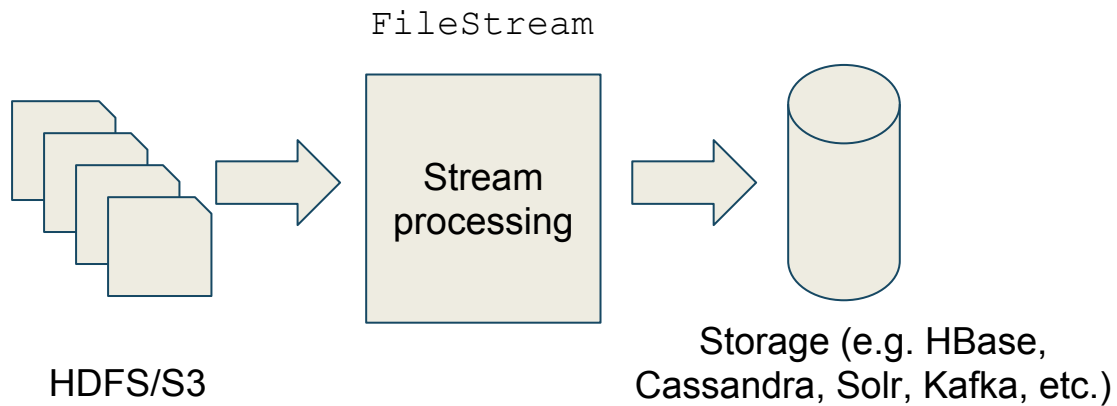
As long as you do things the right way



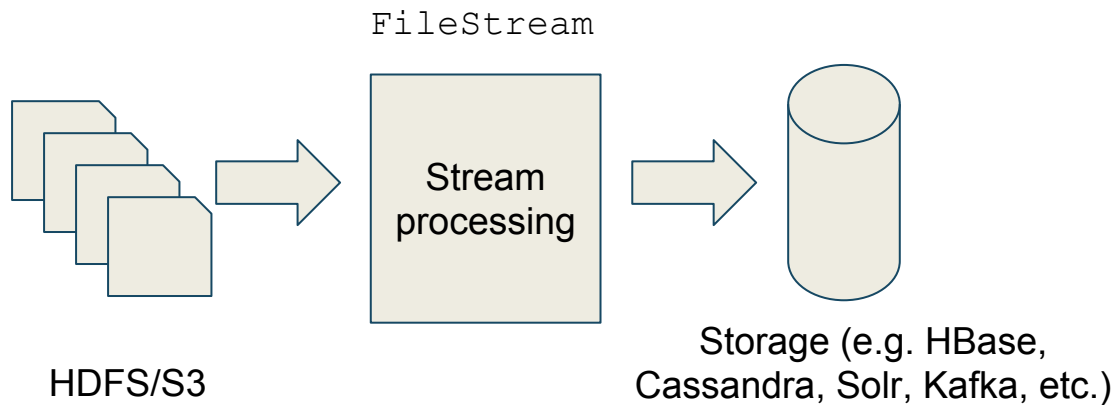
How to do things the right way



1. File sources



1. File sources



- Use checkpointing!

Checkpointing

```
// new context
```

```
val ssc = new StreamingContext (...)
```

```
...
```

```
// set checkpoint directory
```

```
ssc.checkpoint(checkpointDirectory)
```

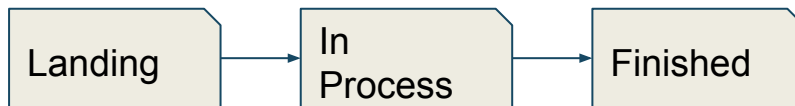
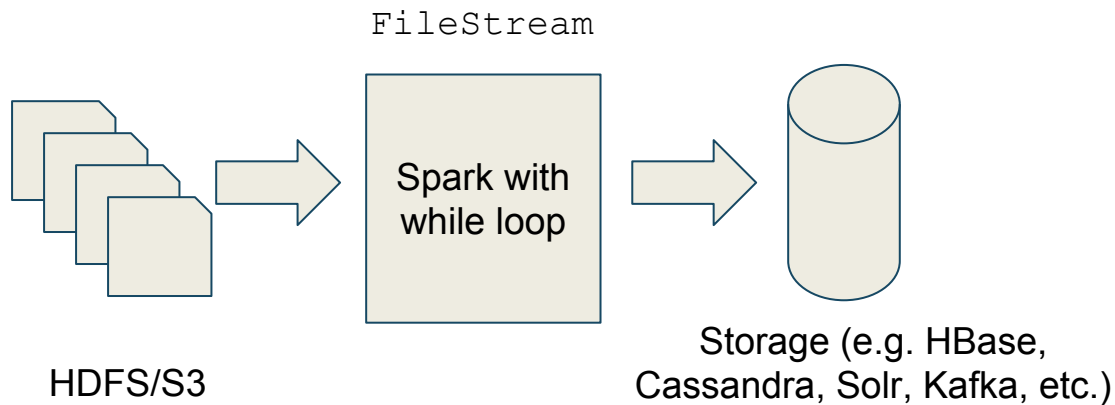
What is checkpointing

- Metadata checkpointing (Configuration, Incomplete batches, etc.)
 - Recover from driver failures
- Data checkpointing
 - Trim lineage when doing stateful transformations (e.g. `updateStateByKey`)
- Spark streaming checkpoints both data and metadata

Checkpointing gotchas

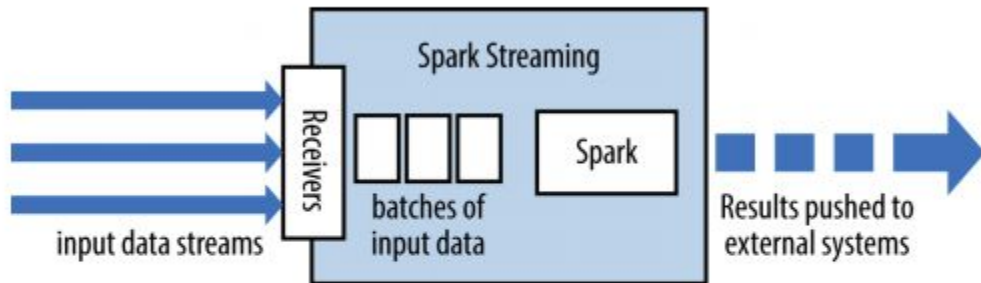
- Checkpoints don't work across app or Spark upgrades
- Clear out (or change) checkpointing directory across upgrades

Ted's Rant



2. Receiver based sources

High-level architecture of Spark Streaming



DStream as a continuous series of RDDs

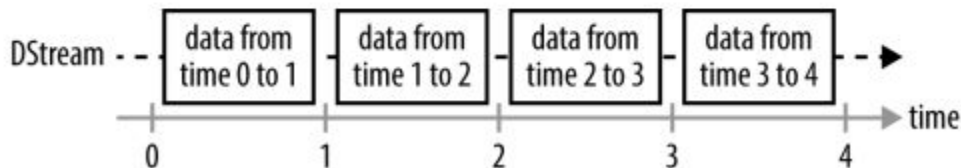


Image source: http://static.oschina.net/uploads/space/2015/0618/110032_9Fvp_1410765.png

Receiver based sources

- Enable checkpointing, AND
- Enable Write Ahead Log (WAL)
 - Set the following property to `true`
`spark.streaming.receiver.writeAheadLog.enable`
 - Default is `false`!

Why do we need a WAL?

- Data on Receiver is stored in executor memory
- Without WAL, a failure in the middle of operation can lead to data loss
- With WAL, data is written to durable storage (e.g. HDFS, S3) before being acknowledged back to source.

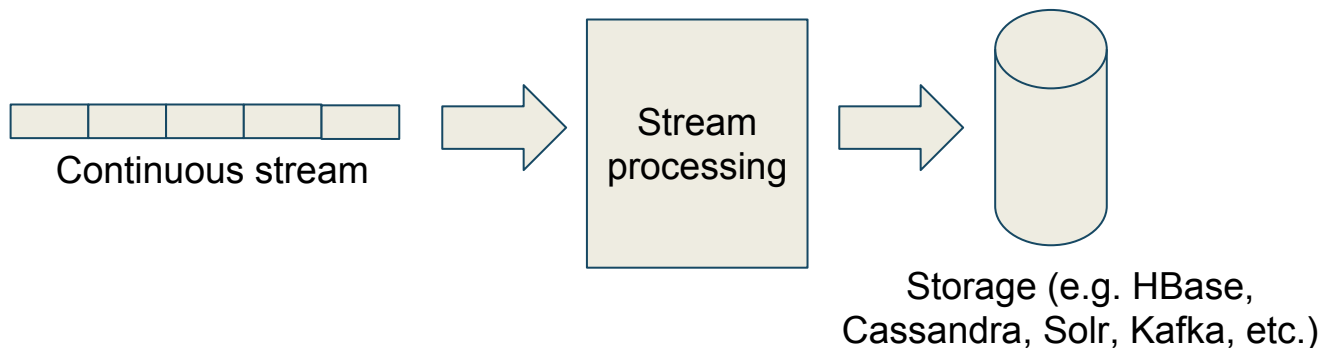
WAL gotchas!

- Makes a copy of all data on disk
- Use `StorageLevel.MEMORY_AND_DISK_SER` storage level for your `DStream`
 - No need to replicate the data in memory across Spark executors
- For S3
`spark.streaming.receiver.writeAheadLog.closeFileAfterWrite` to `true` (similar for driver)

But what about Kafka?

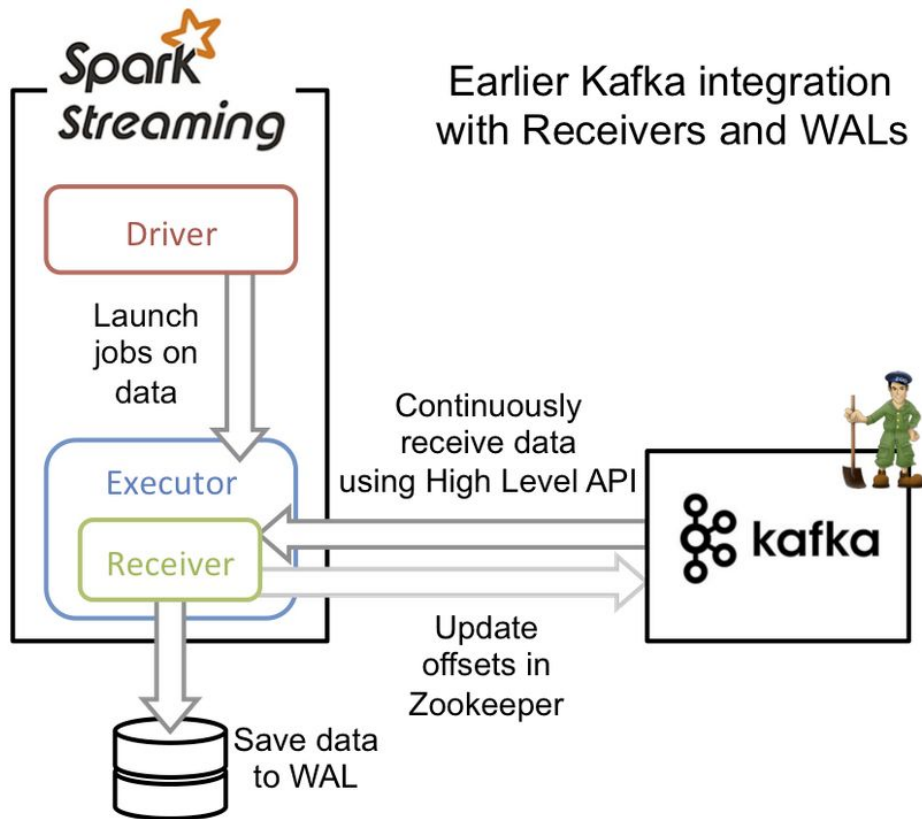
- Kafka already stores a copy of the data
- Why store another copy in WAL?

3. Spark Streaming with Kafka

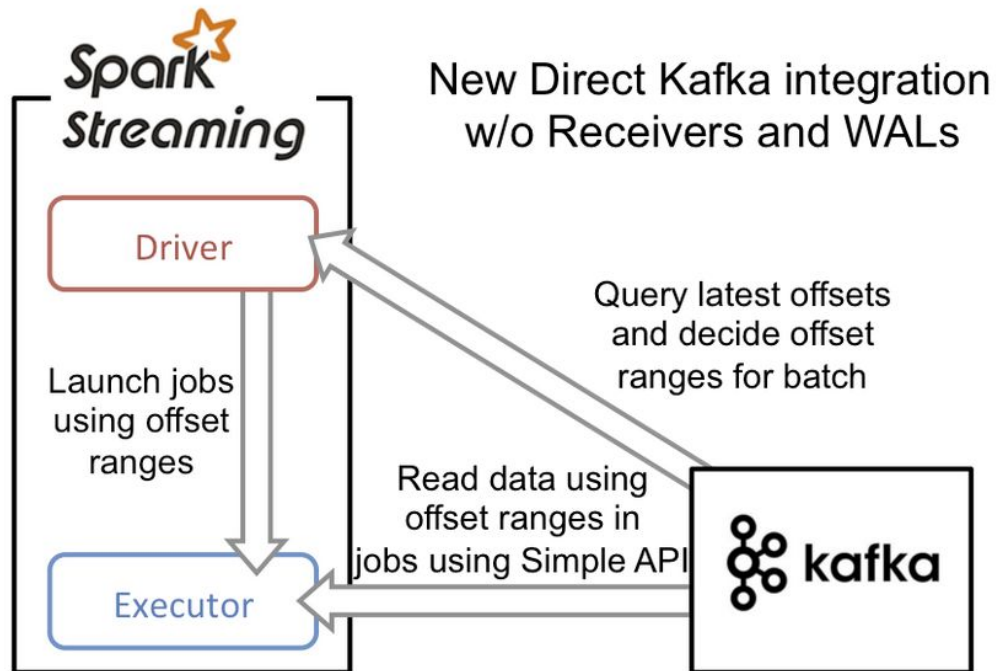


- Use “direct” connector
- No need for a WAL with the direct connector!

Kafka with WAL



Kafka direct connector (without WAL)



Why no WAL?

- No receiver process creating blocks
- Data is stored in Kafka, can be directly recovered from there

Direct connector gotchas

- Need to track offsets for driver recovery
- Checkpoints?
 - No! Not recoverable across upgrades
- Track them yourself
 - In ZooKeeper, HDFS, or a database
- For accuracy
 - Processing needs to be idempotent, OR
 - Update offsets in the same transaction when updating results

Summary - prevent data loss

- Different for different sources
- Preventable if done right
- In Structured Streaming, state is stored in memory (backed by HDFS/S3 WAL), starting Spark 2.1

References

- Improved fault-tolerance and zero data loss
<https://databricks.com/blog/2015/01/15/improved-driver-fault-tolerance-and-zero-data-loss-in-spark-streaming.html>
- Tracking offsets when using direct connector
<http://blog.cloudera.com/blog/2015/03/exactly-once-spark-streaming-from-apache-kafka/>

#3 - Do I really need to use Spark Streaming?



When do you use streaming?

When do you use streaming?

ALWAYS!

Do I need to use Spark Streaming?

Think about your goals

Types of goals

- Atomic Enrichment
- Notifications
- Joining
- Partitioning
- Ingestion
- Counting
- Incremental machine learning
- Progressive Analysis

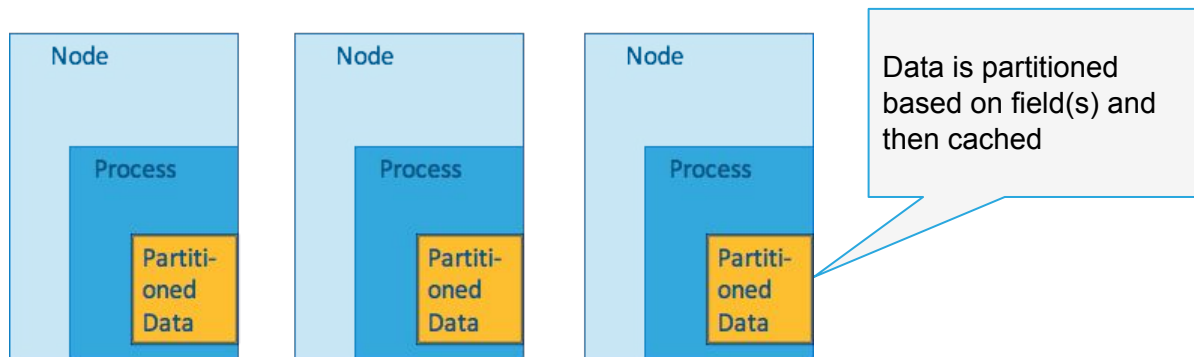
Types of goals

- Atomic enrichment
 - No cache/context needed

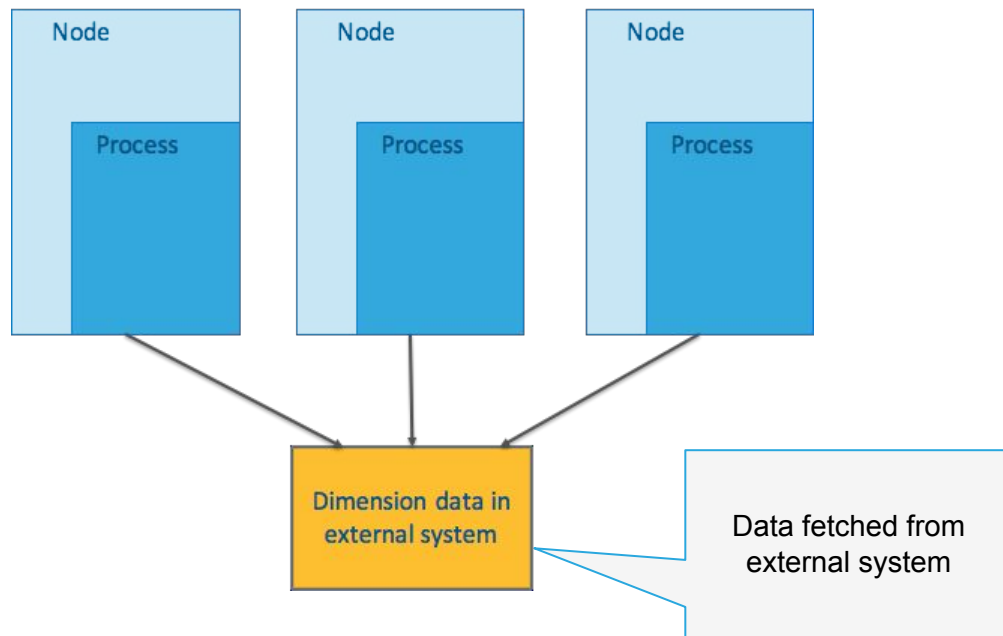
Types of goals

- Notifications
 - NRT
 - Atomic
 - With Context
 - With Global Summary

#2 - Partitioned cache data



#3 - External fetch



Types of goals

- Joining
 - Big Joins
 - Broadcast Joins
 - Streaming Joins

Types of goals

- !!Ted to add a diagram!!
- Partitioning
 - Fan out
 - Fan in
 - Shuffling for key isolation

Types of goals

- Ingestion
 - File Systems or Block Stores
 - No SQLs
 - Lucene
 - Kafka

Types of goals

- Counting
 - Streaming counting
 - Count on query
 - Aggregation on query

Types of goals

- Machine Learning

Types of goals

- Progressive Analysis
 - Reading the stream
 - SQL on the stream

Summary - Do I need to use streaming?

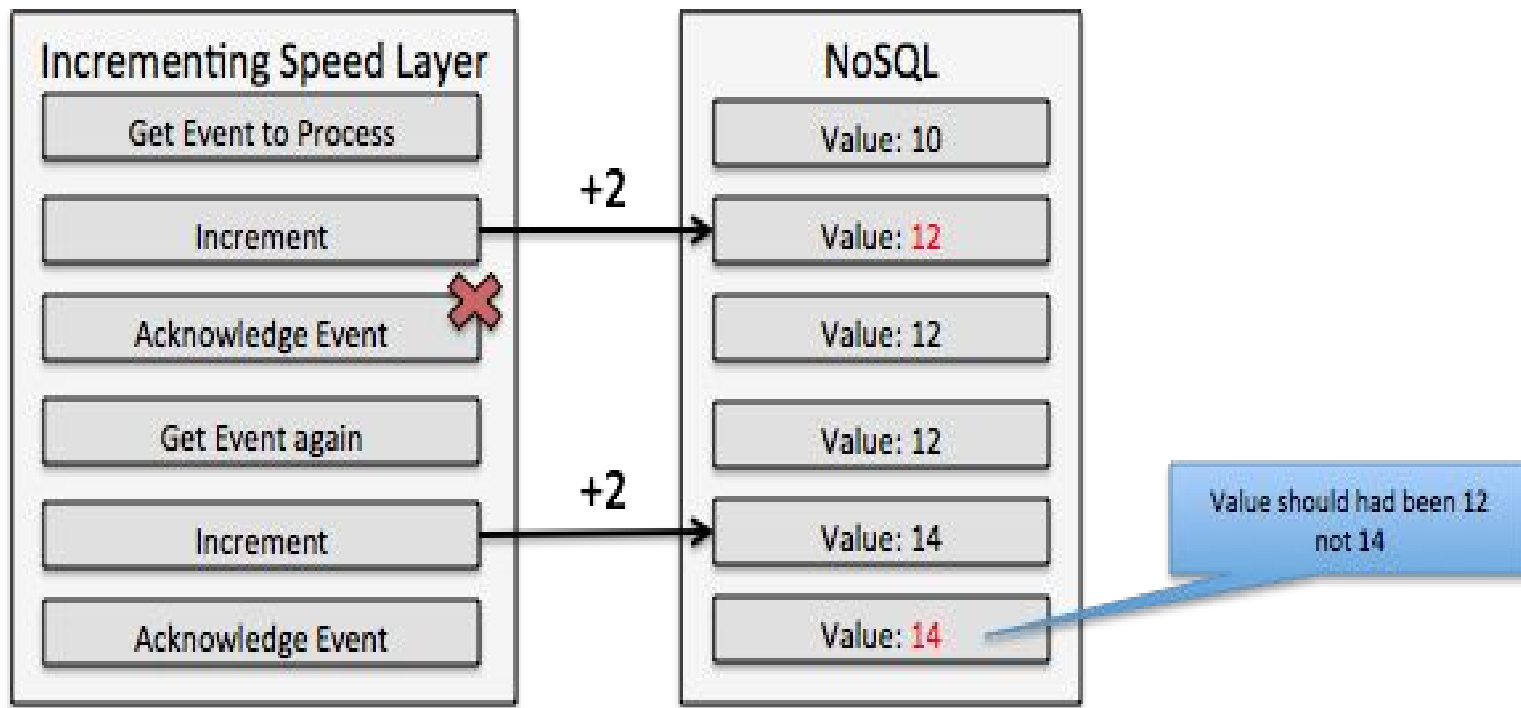
- Spark Streaming is great for
 - Accurate counts
 - Windowing aggregations
 - Progressive Analysis
 - Continuous Machine Learning

#2 - Exactly once semantics

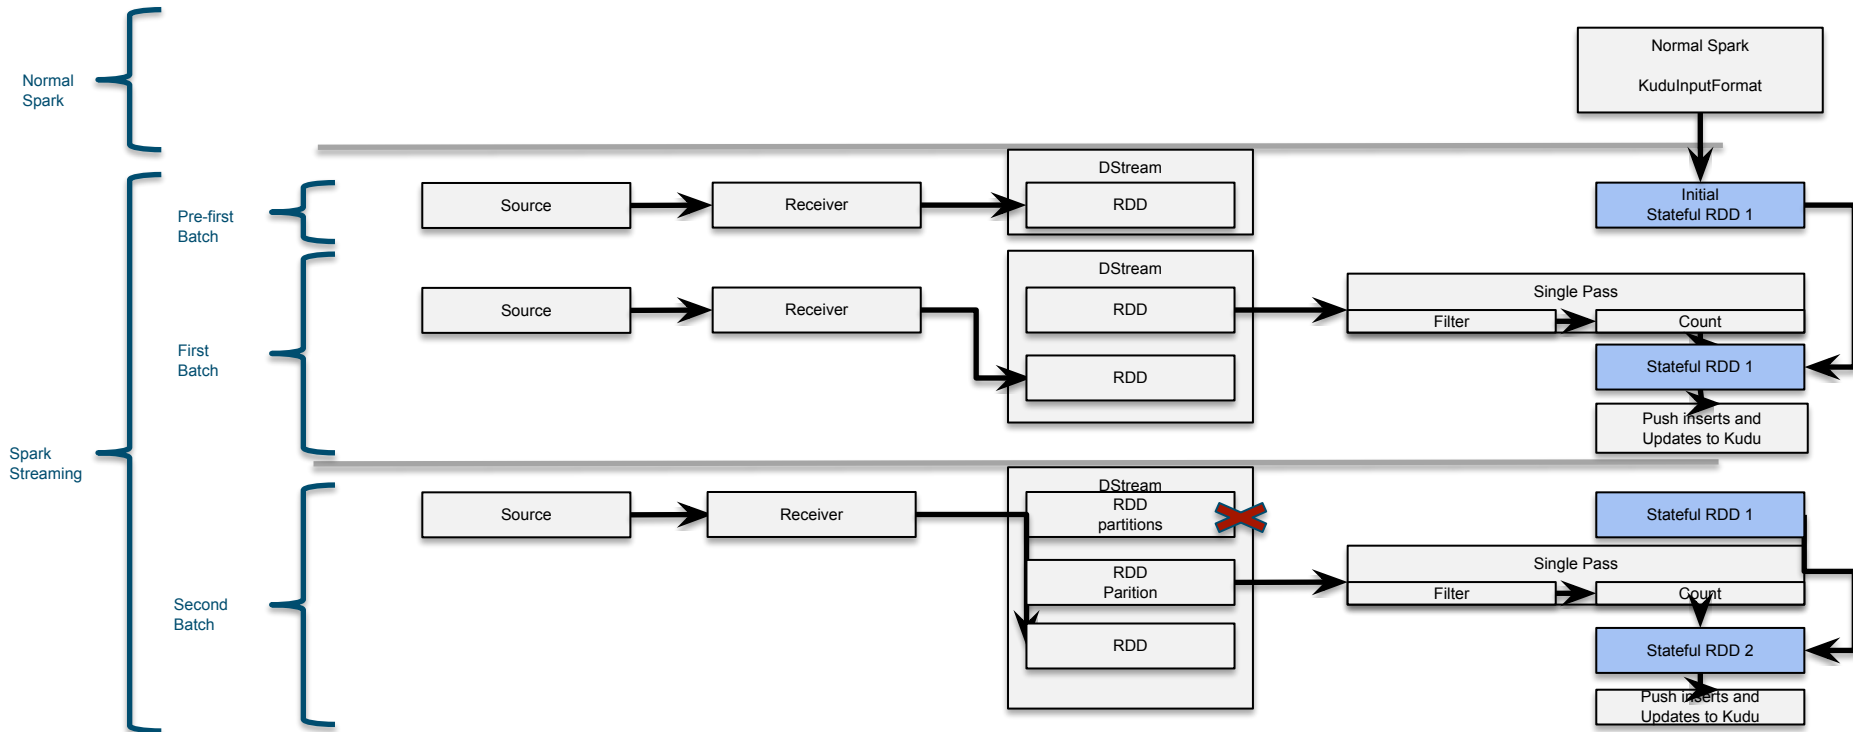
Exactly once semantics

- No duplicate records
- Perfect Counting

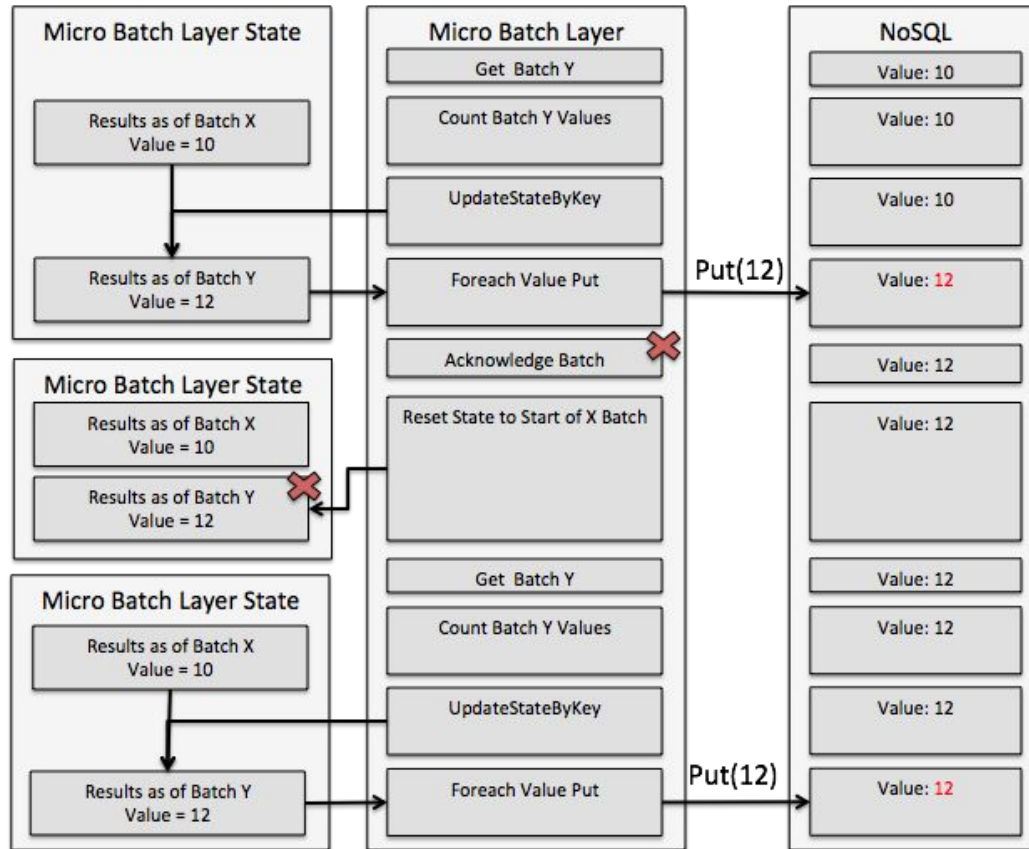
In the days of Lambda



Spark Streaming State



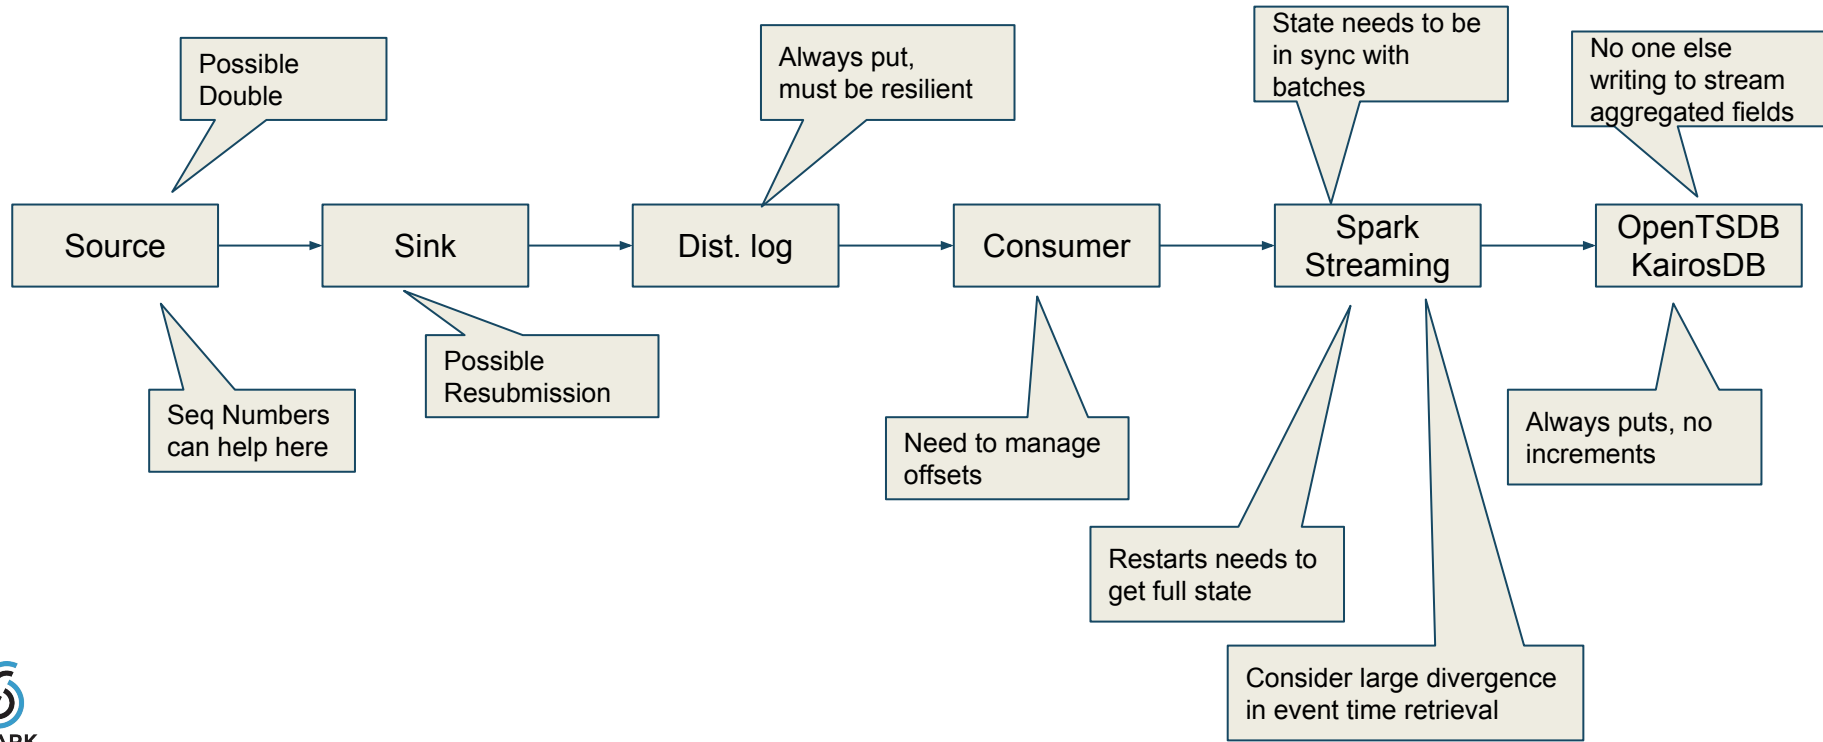
State within



Things can go wrong in many places



Things can go wrong in many places



Summary - Exactly once semantics

- Seq number tied to sources
- Puts for external storage system
- Consider large divergence in event time retrieval
 - Increase divergence window
 - Retrieve state from external storage system *
 - Ignore
 - Process off line
 - Source aggregation (pre-distributed log)

#1 - Graceful shutting down your streaming app



**SPARK
SUMMIT**
EAST 2017

Graceful shutdown

Can we define graceful?

- Offsets known
- State stored externally
- Stopping at the right place (i.e. batch finished)

How to be graceful?

- Thread hooks
 - Check for an external flag every N seconds

```
/**  
 * Stop the execution of the streams, with option of ensuring all received data  
 * has been processed.  
 *  
 * @param stopSparkContext if true, stops the associated SparkContext. The underlying SparkContext  
 * will be stopped regardless of whether this StreamingContext has been  
 * started.  
 * @param stopGracefully if true, stops gracefully by waiting for the processing of all  
 * received data to be completed  
 */  
def stop(stopSparkContext: Boolean, stopGracefully: Boolean): Unit = {
```

Under the hood of grace

```
receiverTracker.stop(processAllReceivedData) //default is to wait 10 second, grace waits until done
jobGenerator.stop(processAllReceivedData) // Will use spark.streaming.gracefulStopTimeout
jobExecutor.shutdown()
val terminated = if (processAllReceivedData) {
    jobExecutor.awaitTermination(1, TimeUnit.HOURS) // just a very large period of time
} else {
    jobExecutor.awaitTermination(2, TimeUnit.SECONDS)
}
if (!terminated) {
    jobExecutor.shutdownNow()
}
```

How to be graceful?

- cmd line

- \$SPARK_HOME_DIR/bin/spark-submit --master \$MASTER_REST_URL --kill \$DRIVER_ID
- spark.streaming.stopGracefullyOnShutdown=true

```
private def stopOnShutdown(): Unit = {  
  val stopGracefully = conf.getBoolean("spark.streaming.stopGracefullyOnShutdown", false)  
  logInfo(s"Invoking stop(stopGracefully=$stopGracefully) from shutdown hook")  
  // Do not stop SparkContext, let its own shutdown hook stop it  
  stop(stopSparkContext = false, stopGracefully = stopGracefully)  
}
```


How to be graceful?

- By marker file
 - Touch a file when starting the app on HDFS
 - Remove the file when you want to stop
 - Separate thread in Spark app, calls
`streamingContext.stop(stopSparkContext = true, stopGracefully = true)`

Storing offsets

- Externally in ZK, HDFS, HBase, Database
- Recover on restart

Summary - Graceful shutdown

- Use provided command line, or
- Background thread with a marker file

Conclusion

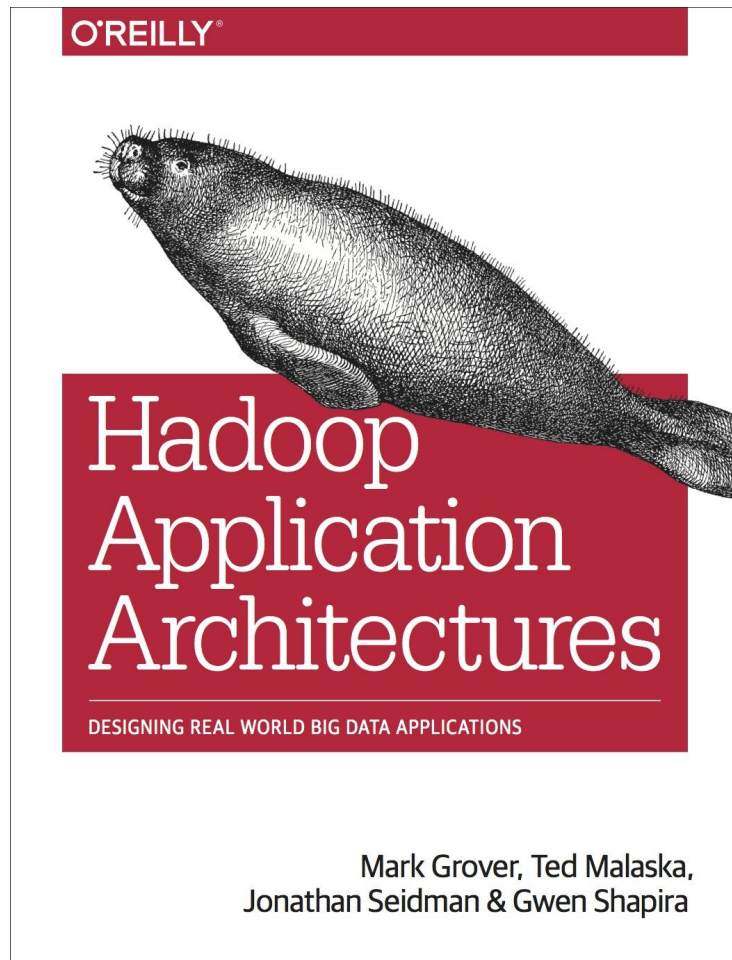


Conclusion

- #5 – How to monitor and manage jobs?
- #4 – How to prevent data loss?
- #3 – Do I need to use Spark Streaming?
- #2 – How to achieve exactly/effectively once semantics?
- #1 – How to gracefully shutdown your app?

Book signing

- Cloudera booth @ 6:30pm



Thank You.

tiny.cloudera.com/streaming-app

