

# Building realtime BI Systems with Kafka, Spark and Kudu

Ruhollah Farchtchi  
Zoomdata



# Drivers for Streaming Data


Data Freshness



Time to Analytic



Business Context



# Streaming Data @ Zoomdata

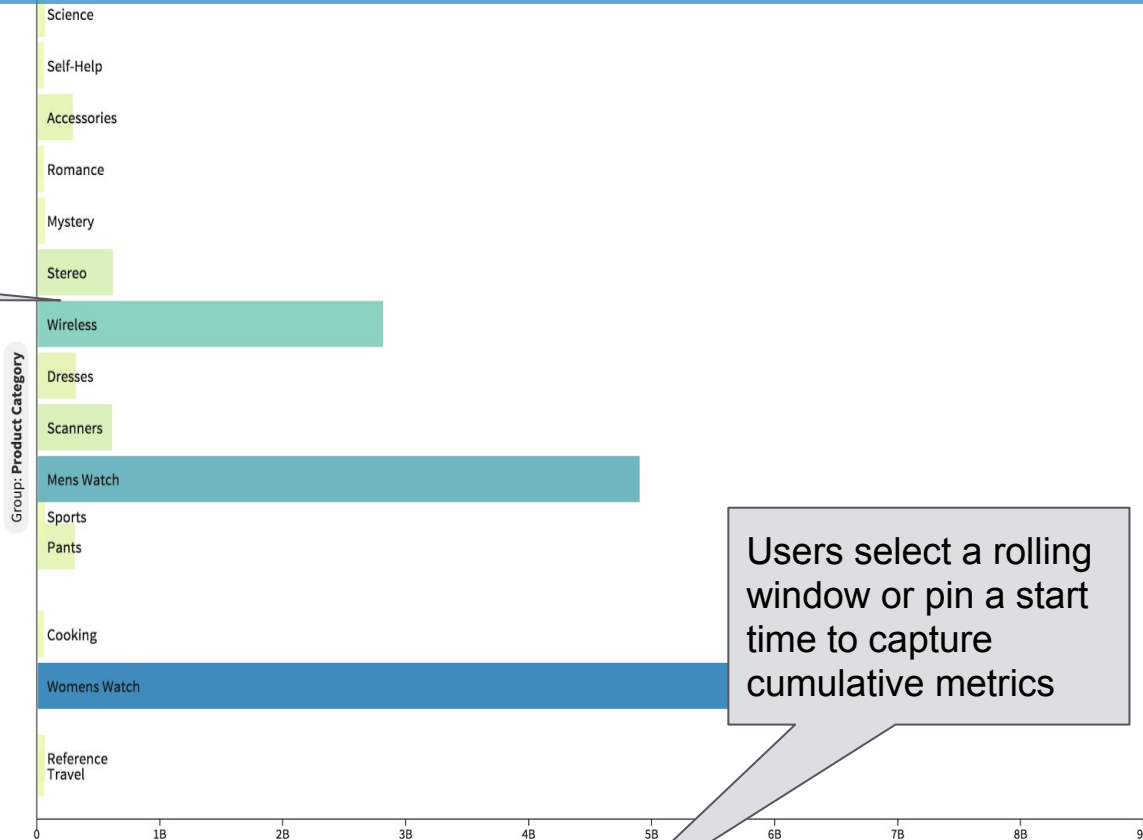
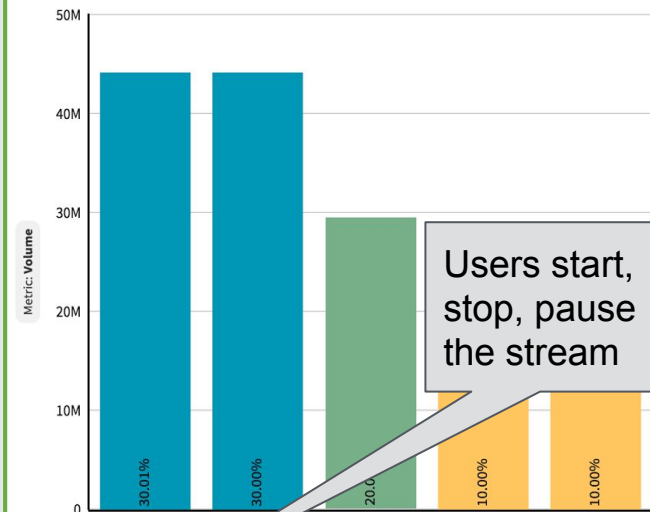
147,147,884  
Volume

Metric: Volume

Visualizations react to  
new data delivered

Users start,  
stop, pause  
the stream

Users select a rolling  
window or pin a start  
time to capture  
cumulative metrics



# Typical Streaming Architectures

Event

Kafka, JMS,  
RabbitMQ,  
etc...

Spark  
Streaming,  
Flink, etc..

Now What?

HDFS (what  
about query)

Cassandra (no  
aggregation)

Lambda (let's  
take a look at  
that for a sec)

# Lambda

- Stream data to hdfs
- Keep some in avro
- Do your own compactions to parquet / orc
- Expose via impala, sparksql, or other

**OR**

- Impala avro partition (speed)
- With history in parquet
- Union compatible schema
- Project as single table via view
- Works ok... still doing a lot of manual data management

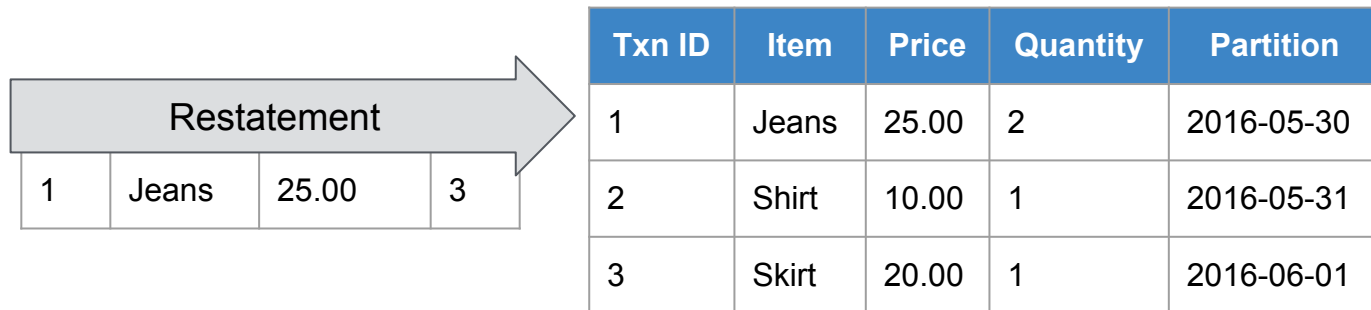
Oh.. and what happens to noncommutative operations like  
Distinct Count?

# Restatements ... yeah we went there

Restatement			
1	Jeans	25.00	3

Txn ID	Item	Price	Quantity	Partition
1	Jeans	25.00	2	2016-05-30
2	Shirt	10.00	1	2016-05-31
3	Skirt	20.00	1	2016-06-01

# Restatements ... how you do it



## General Algorithm

- Figure out which partition(s) are affected
- Recompute affected partition(s) with restated data
- Drop/replace existing partition(s) with new data

# Enter Kudu


## What is Kudu?

- **Kudu is an open source storage engine for structured data which supports low-latency random access together with efficient analytical access patterns.** (Source: <http://getkudu.io/kudu.pdf>)

## Why do you care?

- It makes management of streaming data for ad-hoc analysis **MUCH** easier
- Bridges the mental gap from random access to append only

## Why does Zoomdata care?





# Impala + Kudu: Performance

**Nearly the same performance as Parquet for many similar workloads**

**Simplified data management model**

**Can handle a new class of streaming use cases and workloads**

# Impala + Kudu: Performance

**Nearly the same performance as Parquet for many similar workloads**

**Simplified data management model**

**Can handle a new class of streaming use cases and workloads**

Great... let's just use Kudu from now on:

- We can ingest data with great write throughput
- Support analytic queries
- Support random access writes

What's not to love?

Ship It!



# There's a catch...

... it's your data model

Good news! If you have figured this out with HDFS and Parquet, you're not too far off.

Things to consider:

- **Access pattern and partition scheme (similar to partitioning data parquet)**
  - Has a big role to play in parallelism of your queries
- **Cardinality of your attributes**
  - Affects what type of column encoding you decide to use
- **Key structure**
  - You get only one, use it wisely

More on this can be found at : [http://getkudu.io/docs/schema\\_design.html](http://getkudu.io/docs/schema_design.html)

# Let's put it all together

**I have a fruit stand**

**I sell my fruits via phone order to remote buyers**

**My transactions look something like:**


**Orders(orderId,orderTS,fruit,price,customerID,  
customerPhone,customerAddress)**



# Impala DDL for Kudu

Key

```
CREATE EXTERNAL TABLE `strata_fruits_expanded` (  
  `_ts` BIGINT,  
  `_id` STRING,
```



```
TBLPROPERTIES(  
  'storage_handler' =  
  'com.cloudera.kudu.hive.KuduStorageHandler',  
  'kudu.table_name' = 'strata_fruits_expanded',  
  'kudu.master_addresses' = '10.xxx.xxx.xxx:7051',  
  'kudu.key_columns' = '_ts, _id'  
);
```

Key

# Impala DDL for Kudu



attributes

attributes

```
CREATE EXTERNAL TABLE `strata_fruits_expanded` (  
  `_ts` BIGINT,  
  `_id` STRING,  
  `fruit` STRING,  
  `country_code` STRING,  
  `country_area_code` STRING,  
  `phone_num` STRING,  
  `message_date` BIGINT,  
  `price` FLOAT,  
  `keyword` STRING  
)
```

**Low cardinality attributes -- things I want to group by -- are great candidates for dictionary encoding**

```
DISTRIBUTE BY HASH (_ts) INTO 60 BUCKETS  
TBLPROPERTIES(  
  'storage_handler' =  
  'com.cloudera.kudu.hive.KuduStorageHandler',  
  'kudu.table_name' = 'strata_fruits_expanded',  
  'kudu.master_addresses' = '10.xxx.xxx.xxx:7051',  
  'kudu.key_columns' = '_ts, _id'  
);
```

# Impala DDL for Kudu

```
CREATE EXTERNAL TABLE `strata_fruits_expanded` (  
  `_ts` BIGINT,  
  `_id` STRING,  
  `fruit` STRING,  
  `country_code` STRING,  
  `country_area_code` STRING,  
  `phone_num` STRING,  
  `message_date` BIGINT,  
  `price` FLOAT,  
  `keyword` STRING  
)
```

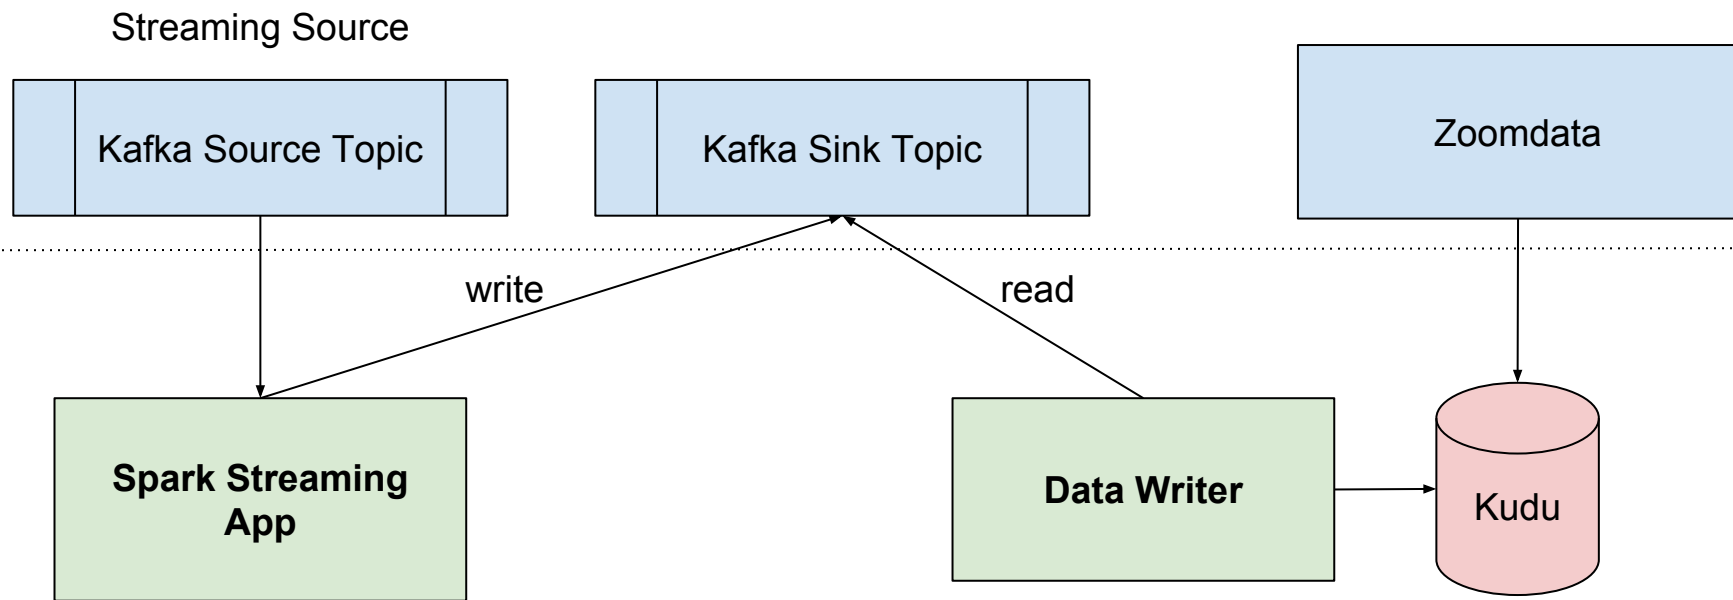
**How you distribute your data directly impacts your ability to process in parallel as well as any predicate push-down type of operations Kudu can perform**

Partition Scheme

```
DISTRIBUTE BY HASH (_ts) INTO 60 BUCKETS  
TBLPROPERTIES(  
  'storage_handler' =  
  'com.cloudera.kudu.hive.KuduStorageHandler',  
  'kudu.table_name' = 'strata_fruits_expanded',  
  'kudu.master_addresses' = '10.10.10.10:7000',  
  'kudu.key_columns' = '_ts, _id'  
);
```

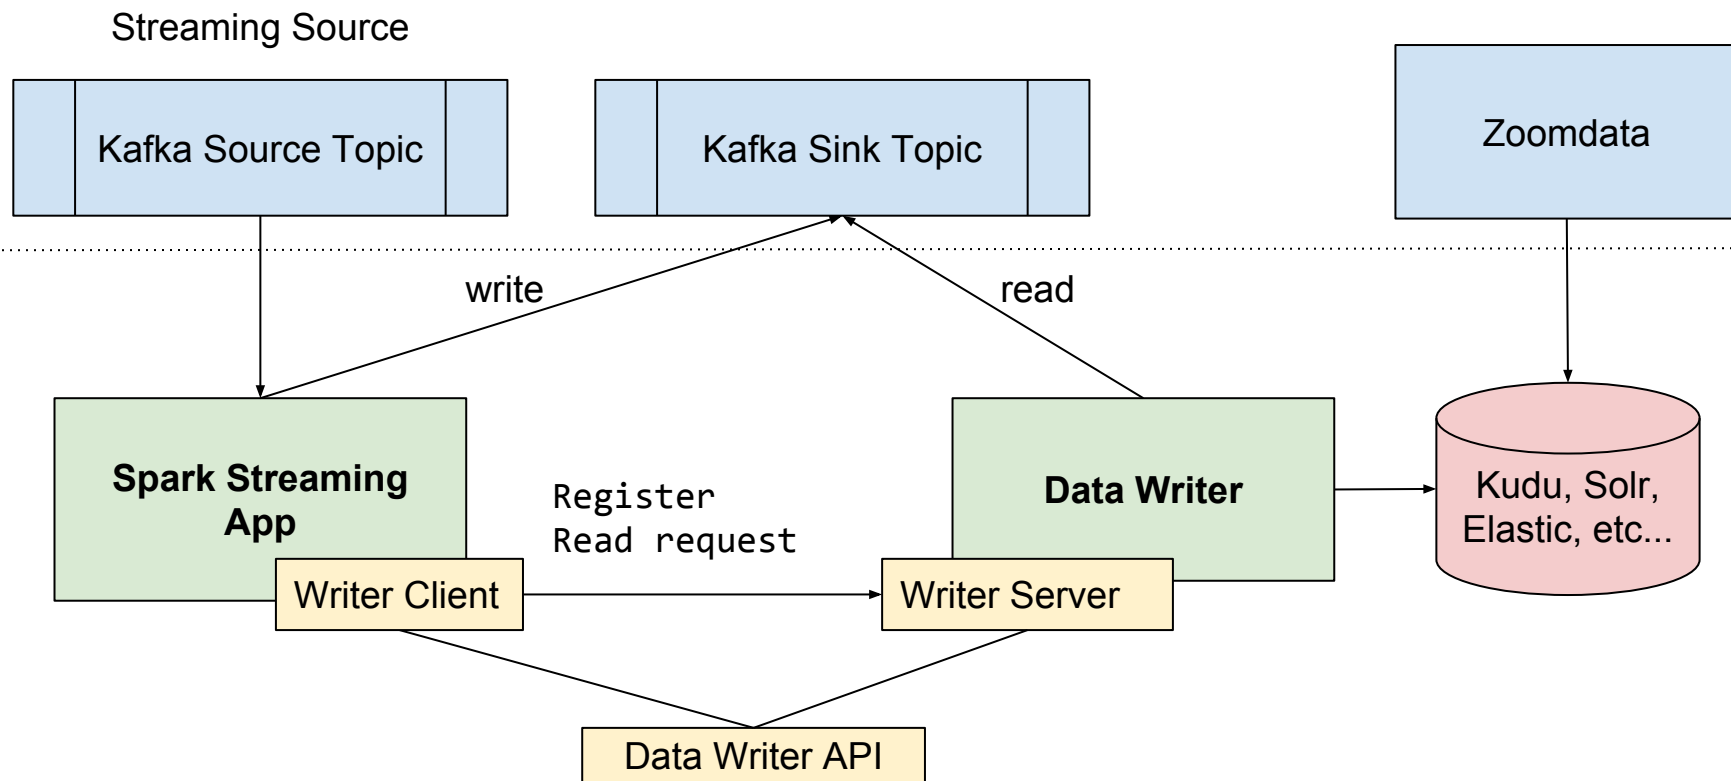
**For large tables, such as fact tables, aim for as many tablets as you have cores in the cluster -- but figure out what else you are running as well.**

# Let's see it in action....





# Let's see it in action... not actually that simple



# Special Thanks



**Anton Gorshkov:** For his original streaming with kafka fruit stand demo



**The Cloudera Kudu Team:** Specifically Todd Lipcon for all the insight into Kudu optimization



**Nexmo:** For use of their SaaS SMS service in this demo

# Thank You.

[www.zoomdata.com](http://www.zoomdata.com)

