



# SSR: Structured Streaming for R and Machine Learning

Felix Cheung

Principal Engineer & Spark Committer

Spark


+



Disclaimer:  
Apache Spark community contributions

# Agenda

- Structured Streaming
- ML Pipeline
- R - putting it all together
- Considerations



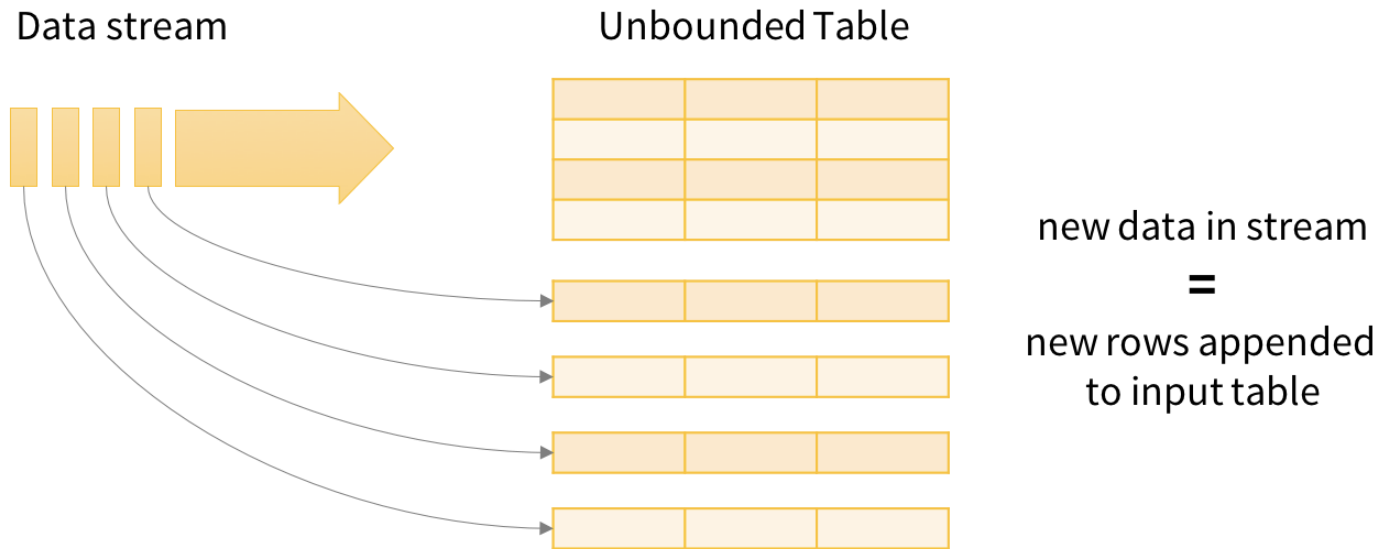
# Why Streaming?

- Faster insight at scale
- ETL
- Trends
- Latest data to static data
- Continuous Learning

# Spark Streaming

1. Receiver
2. Direct DStream
3. Structured Streaming

# Structured Streaming



Data stream as an unbounded Input Table



# Structured Streaming

- "Streaming Logical Plan"
  - Extending Dataset/DataFrame to include *incremental* execution of unbounded input
  - aka Rinse & Repeat



# Same

- Transformations:

map

filter

aggregate

window

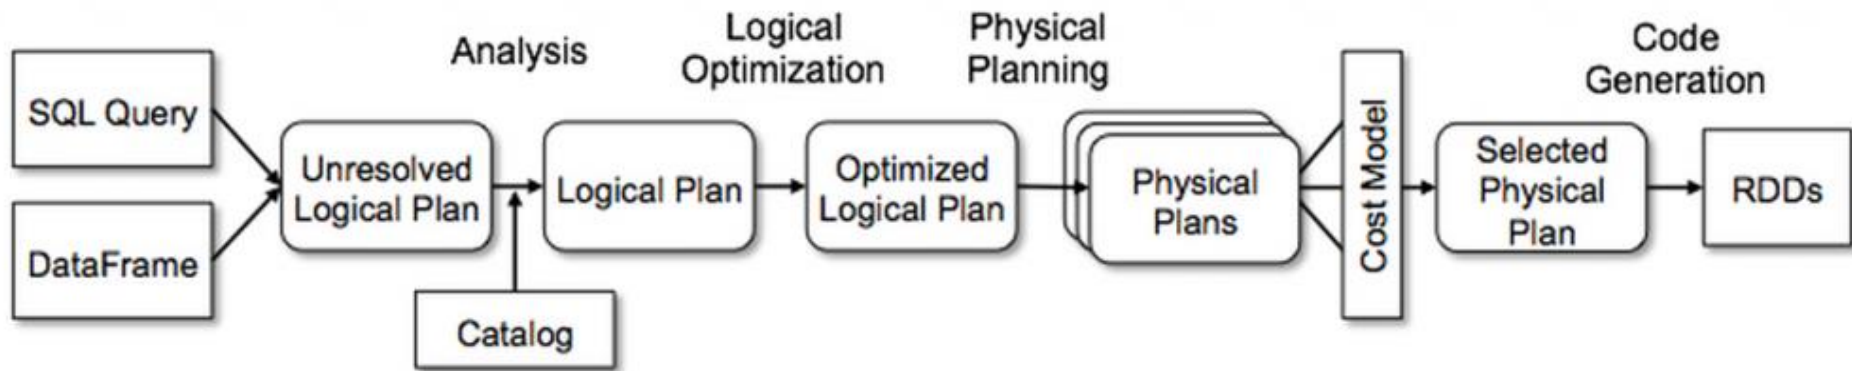
join\*

(\*some limitations)

# Better

- Trigger
- Consistency
- Fault Tolerance
- Event time – late data, watermark

# Execution Plan



# SS in a Circuit



DataFrame

DataFrame

Trigger

# Source

File

Kafka

Socket

MQTT

# Sink

File (new formats in 2.1+)

Console

Memory (aka Temp View)

Foreach

Kafka (new in 2.2)

# Output Mode


Append (default)

Complete

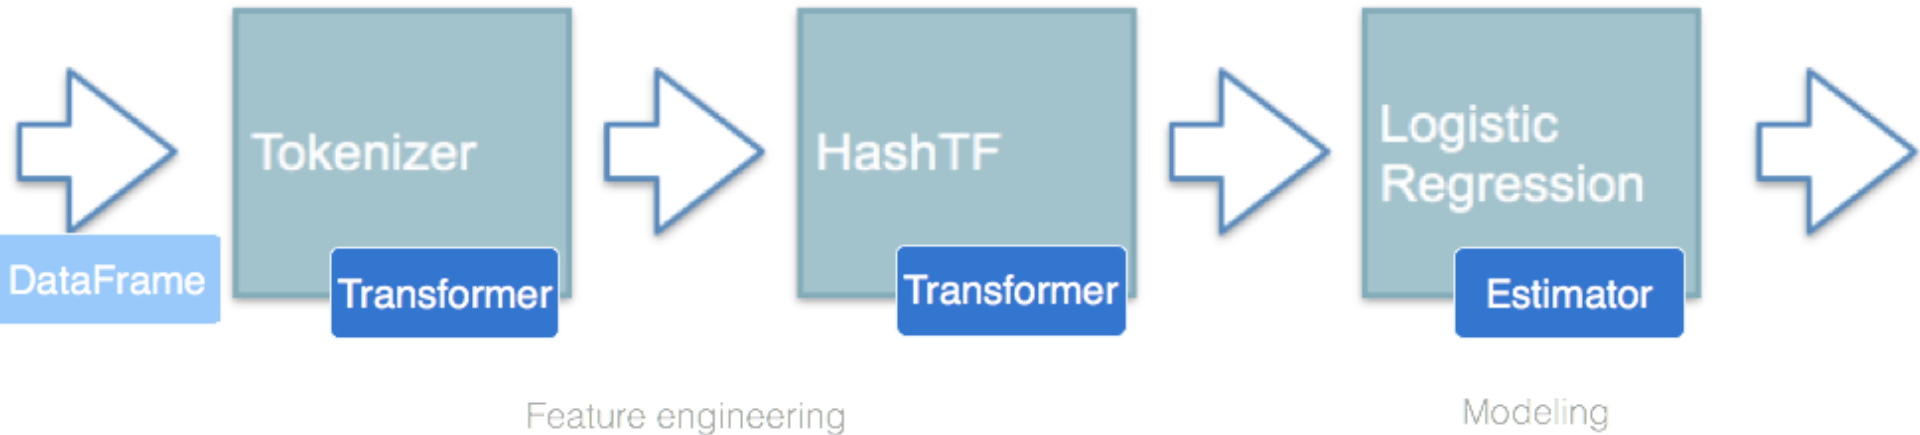
Update (new in 2.1.1)



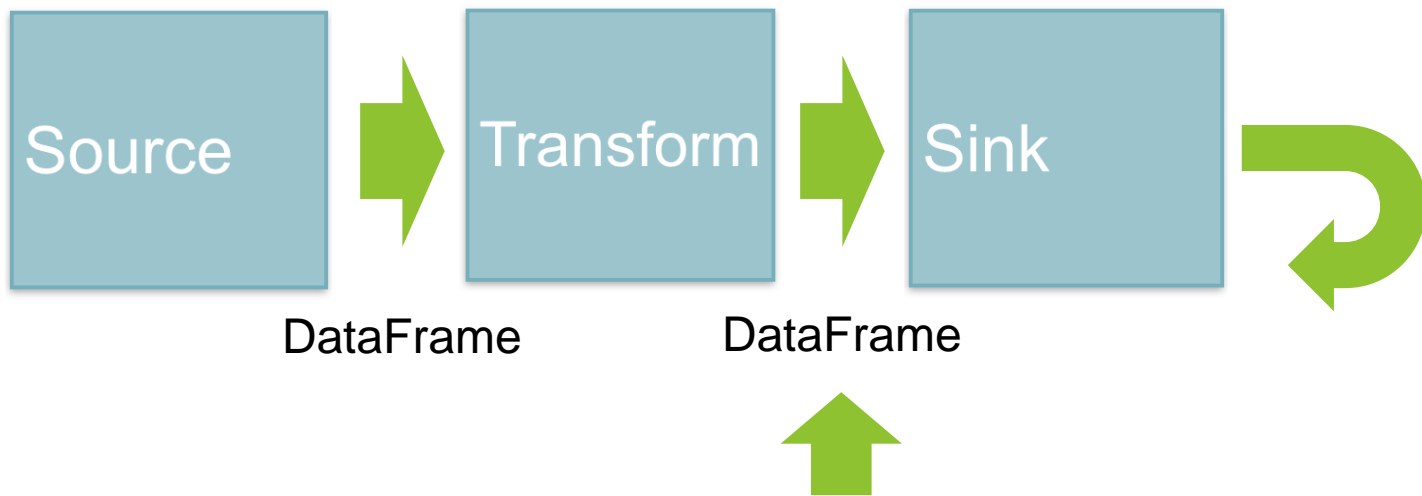
# Streaming & ML Don't Mix\*



# ML Pipeline Model



# Remember the SS Flow?



# ML Pipeline fit()

- Essentially an Action
- Results in a Model
- Sink start() also an Action
- Structured Streaming circuit must be completed with Sink start()

# R to the Rescue



# R

- Statistical computing and graphics
- 10.7k+ packages on [CRAN](https://cran.r-project.org/)

# Why Streaming in R

- Single integrated job for everything
  1. Ingest
  2. ETL
  3. Machine Learning
- Use your favorite packages - freedom to choose
- rkafka – last published 2015

*Spark*

+






# SparkR

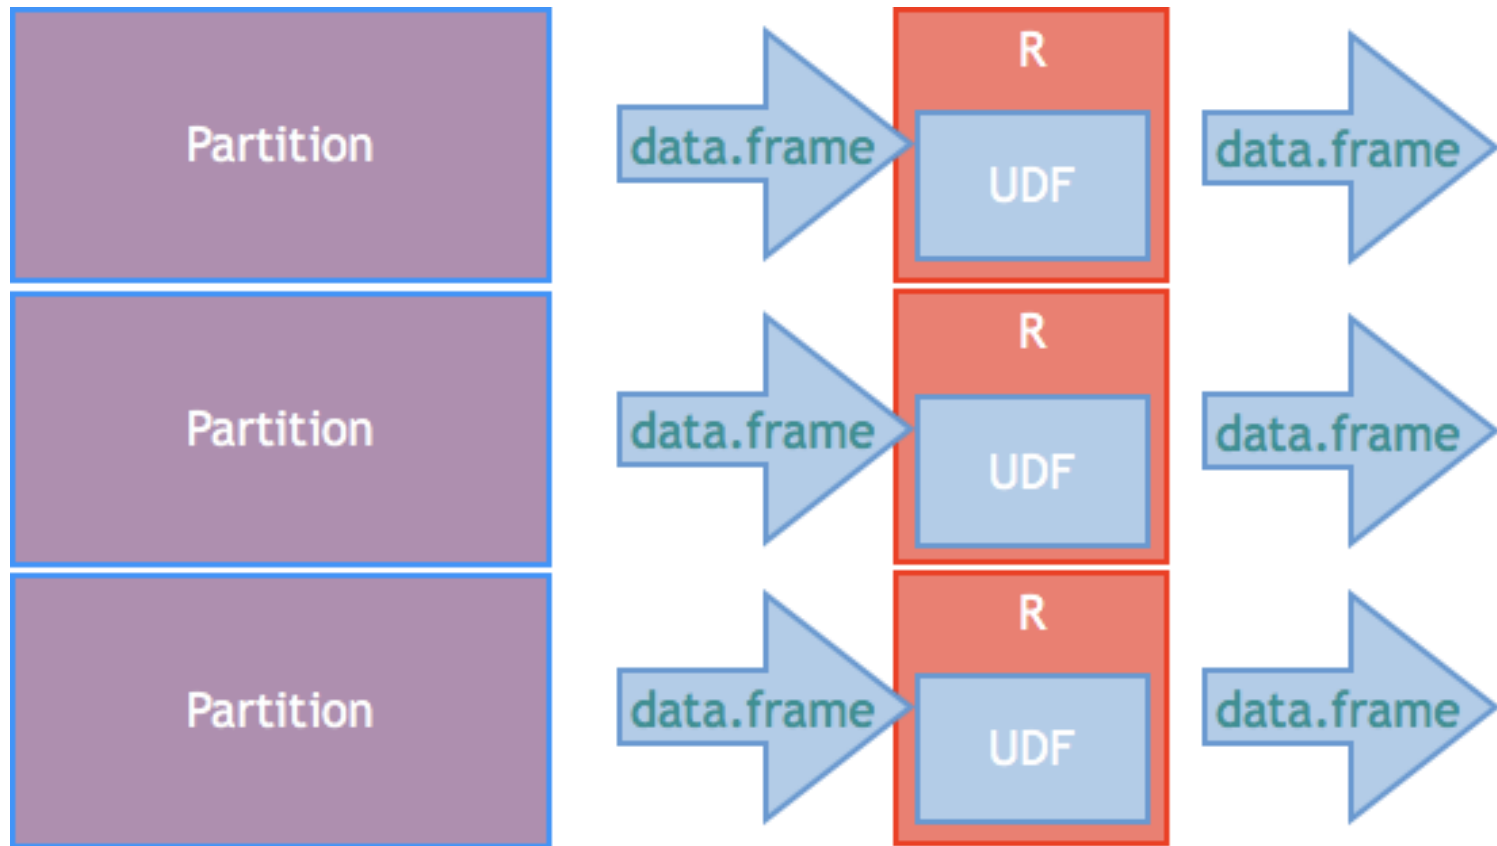
- DataFrame API like R data.frame, dplyr
  - Full Spark optimizations
- SQL, Session, Catalog
- “Spark Packages”
- ML
- R-native UDF
- SS

# Native R UDF

- User-Defined Functions - custom transformation
- Apply by Partition
- Apply by Group



# Parallel Processing By Partition



# SCALABLE DATA SCIENCE WITH SPARKR


Felix Cheung

Principal Engineer - *Spark @ Microsoft* & Apache Spark Committer



<https://spark-summit.org/east-2017/events/scalable-data-science-with-sparkr/>

# Native R UDF = DF Transform



# SS in R

1. DataStreamReader/Writer
2. StreamingQuery
3. Extending DataFrame (isStreaming)

# About Demo

- Create a job to discover trending news topics
  - Structured Streaming
  - Machine Learning with native R package in UDF

# Demo!

<https://goo.gl/0v6YxF>



# Demo

- SS – read text stream from Kafka
- R-UDF – a partition with lines of text
  - RTextTools – text vector into DTM – scrubbing
  - LDA
  - terms
- SQL – group by words, count
- SS – write to console

# Read DataFrame vs Stream

```
read.df(datapath, source = "parquet")
```

```
read.stream("kafka",  
            kafka.bootstrap.servers = servers,  
            subscribe = topic)
```

# Streaming WordCount in 1 line

```
library(magrittr)
kbsrvs <- "kafka-0.broker.kafka.svc.cluster.local:9092"
topic <- "test1"
```

```
read.stream("kafka", kafka.bootstrap.servers = kbsrvs, subscribe = topic) %>%
selectExpr("explode(split(value as string, ' ')) as word") %>%
group_by("word") %>%
count() %>%
write.stream("console", outputMode = "complete")
```

# Challenges



# Streaming and ML

- Streaming – small batch
- ML – *sometimes* large data to build model
  - => pre-trained model
  - => online machine learning
- Adopting to data schema, pattern changes
- Updating model (when?)

# Practical Implementation

- LSI – online training
- Online LDA
- kNN
- k-means with predict on new data

# SS Considerations

- Schema of DataFrame from Kafka:  
key (object), value (object), topic, partition,  
offset, timestamp, timestampType
- OutputMode requirements

# ML with R-UDF

- Native code UDF can break the job
  - eg. ML packages could be sensitive to empty row
  - more data checks In Real Life
- Debugging can be challenging – run separately first
- UDF must return that matches schema
- Model as state to distribute to each UDF instance



# Future – SSR

- Configurable trigger
- Watermark for late data



# Thank You.

<https://github.com/felixcheung>

linkedin: <http://linkd.in/1OeZDb7>

blog: <http://bit.ly/1E2z6OI>