



Improving Python and Spark Performance and Interoperability with Apache Arrow



Julien Le Dem
Principal Architect
Dremio

Li Jin
Software Engineer
Two Sigma Investments



About Us



Li Jin

@icexelloss

- Software Engineer at Two Sigma Investments
- Building a python-based analytics platform with PySpark
- Other open source projects:
 - Flint: A Time Series Library on Spark
 - Cook: A Fair Share Scheduler on Mesos



Julien Le Dem

@J_

- Architect at @DremioHQ
- Formerly Tech Lead at Twitter on Data Platforms
- Creator of Parquet
- Apache member
- Apache PMCs: Arrow, Kudu, Incubator, Pig, Parquet




Agenda

- Current state and limitations of PySpark UDFs
- Apache Arrow overview
- Improvements realized
- Future roadmap



Current state and limitations of PySpark UDFs



Why do we need User Defined Functions?

- Some computation is more easily expressed with Python than Spark built-in functions.
- Examples:
 - weighted mean
 - weighted correlation
 - exponential moving average



What is PySpark UDF

- PySpark UDF is a user defined function executed in **Python runtime**.
- Two types:
 - Row UDF:
 - `lambda x: x + 1`
 - `lambda date1, date2: (date1 - date2).years`
 - Group UDF (subject of this presentation):
 - `lambda values: np.mean(np.array(values))`

Row UDF

- Operates on a row by row basis
 - Similar to `map` operator

- Example ...

```
df.withColumn(  
    'v2',  
    udf(lambda x: x+1, DoubleType())(df.v1)  
)
```

- Performance:
 - **60x** slower than build-in functions for simple case

Group UDF

- UDF that operates on more than one row
 - Similar to ``groupBy`` followed by ``map`` operator
- Example:
 - Compute weighted mean by month



Group UDF

- Not supported out of box:
 - Need boiler plate code to pack/unpack multiple rows into a nested row
- Poor performance
 - Groups are materialized and then converted to Python data structures



Example: Data Normalization

$$(values - values.mean()) / values.std()$$



Example: Data Normalization

```
group_columns = ['year', 'month']
non_group_columns = [col for col in df.columns if col not in group_columns]
s = StructType([f for f in df.schema.fields if f.name in non_group_columns])
cols = list([F.col(name) for name in non_group_columns])

df_norm = df.withColumn('values', F.struct(*cols))
df_norm = (df_norm.groupBy('year', 'month')
           .agg(F.collect_list(df_norm.values).alias('values'))

s2 = StructType(s.fields + [StructField('v3', DoubleType())])
@udf(ArrayType(s2))
def normalize(values):
    v1 = pd.Series([r.v1 for r in values])
    v1_norm = (v1 - v1.mean()) / v1.std()
    return [values[i] + (float(v1_norm[i]),) for i in range(0, len(values))]

df_norm = (df_norm.withColumn('new_values', normalize(df_norm.values))
           .drop('values')
           .withColumn('new_values', F.explode(F.col('new_values'))))

for col in [f.name for f in s2.fields]:
    df_norm = df_norm.withColumn(col, F.col('new_values.{0}'.format(col)))

df_norm = df_norm.drop('new_values')
```



Example: Monthly Data Normalization

```
group_columns = ['year', 'month']
non_group_columns = [col for col in df.columns if col not in group_columns]
s = StructType([f for f in df.schema.fields if f.name in non_group_columns])
cols = list([F.col(name) for name in non_group_columns])

df_norm = df.withColumn('values', F.struct(*cols))
df_norm = (df_norm.groupBy('year', 'month')
           .agg(F.collect_list(df_norm.values).alias('values')))

s2 = StructType(s.fields + [StructField('v3', DoubleType())])
@udf(ArrayType(s2))
def normalize(values):
    v1 = pd.Series([r.v1 for r in values])
    v1_norm = (v1 - v1.mean()) / v1.std()
    return [values[i] + (float(v1_norm[i]),) for i in range(0, len(values))]

df_norm = (df_norm.withColumn('new values', normalize(df_norm.values))
           .drop('values')
           .withColumn('new_values', F.explode(F.col('new values'))))

for col in [f.name for f in s2.fields]:
    df_norm = df_norm.withColumn(col, F.col('new_values.{0}'.format(col)))

df_norm = df_norm.drop('new_values')
```

Useful bits

Example: Monthly Data Normalization

```
group_columns = ['year', 'month']
non_group_columns = [col for col in df.columns if col not in group_columns]
s = StructType([f for f in s.fields if f.name in non_group_columns])
cols = list([F.col(name) for name in non_group_columns])

df_norm = df.withColumn('values', F.struct(*cols))
df_norm = (df_norm.groupBy('year', 'month')
           .agg(F.collect_list(df_norm.values).alias('values'))

s2 = StructType(s.fields + [StructField('v3', DoubleType())])
@udf(ArrayType(s2))
def normalize(values):
    v1 = pd.Series([r.v1 for r in values])
    v1_norm = (v1 - v1.mean()) / v1.std()
    return [values[i] + (float(v1_norm[i]),) for i in range(0, len(values))]

df_norm = (df_norm.withColumn('new_values', normalize(df_norm.values))
           .drop('values')
           .withColumn('new_values', F.explode(F.col('new_values'))))

for col in [f.name for f in s2.fields]:
    df_norm = df_norm.withColumn(col, F.col('new_values.{0}'.format(col)))

df_norm = df_norm.drop('new_values')
```

Boilerplate

Boilerplate

Example: Monthly Data Normalization

- Poor performance - 16x slower than baseline

```
groupBy().agg(collect_list())
```




Problems

- Packing / unpacking nested rows
- Inefficient data movement (Serialization / Deserialization)
- Scalar computation model: object boxing and interpreter overhead



Apache
Arrow



Arrow: An open source standard

- Common need for in memory columnar
- Building on the success of Parquet.
- Top-level Apache project
- Standard from the start
 - Developers from 13+ major open source projects involved
- Benefits:
 - Share the effort
 - Create an ecosystem

Calcite
Cassandra
Deeplearning4j
Drill
Hadoop
HBase
Ibis
Impala
Kudu
Pandas
Parquet
Phoenix
Spark
Storm
R



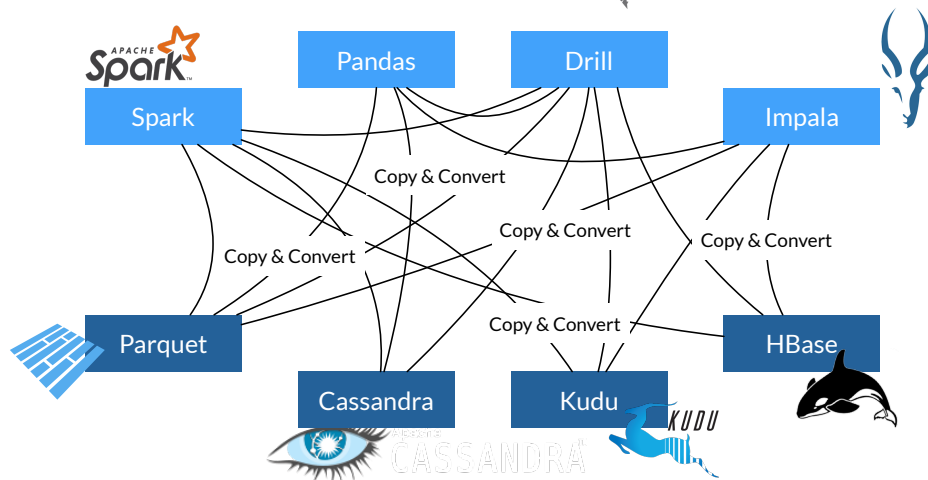
Arrow goals

- Well-documented and cross language compatible
- Designed to take advantage of modern CPU
- Embeddable
 - In execution engines, storage layers, etc.
- Interoperable



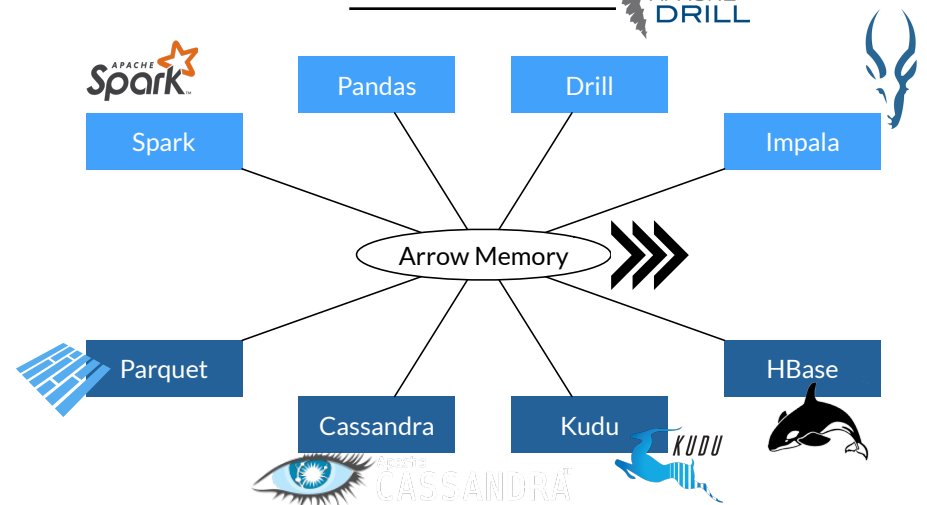
High Performance Sharing & Interchange

Before



- Each system has its own internal memory format
- 70-80% CPU wasted on serialization and deserialization
- Functionality duplication and unnecessary conversions

With Arrow



- All systems utilize the same memory format
- No overhead for cross-system communication
- Projects can share functionality (eg: Parquet-to-Arrow reader)



Columnar data

```

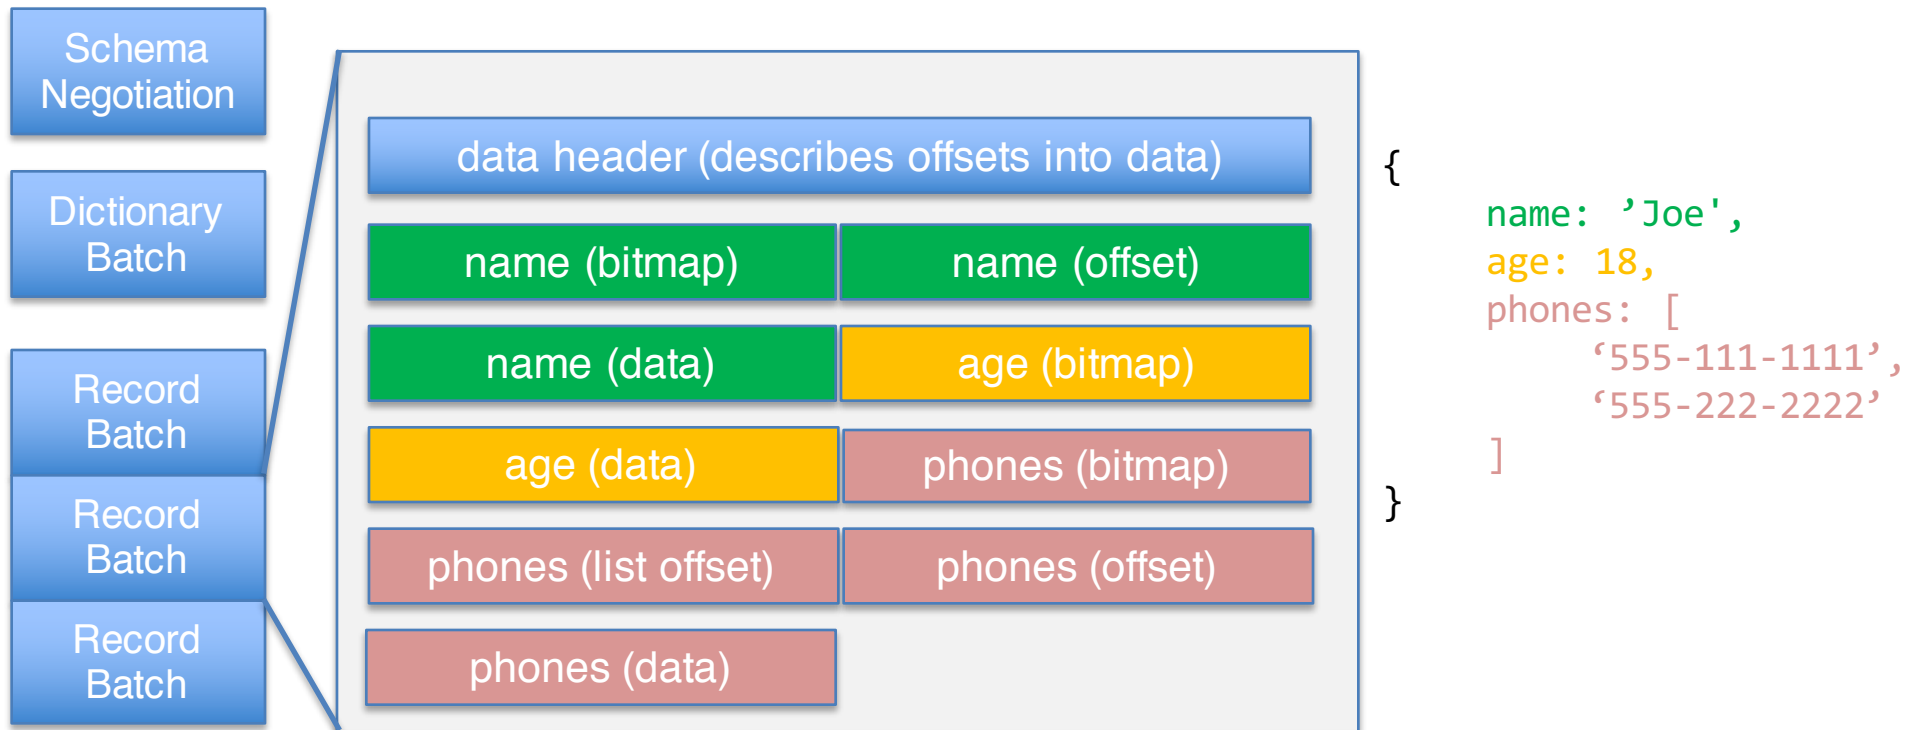
persons = [{
  name: 'Joe',
  age: 18,
  phones: [
    '555-111-1111',
    '555-222-2222'
  ]
}, {
  name: 'Jack',
  age: 37,
  phones: [ '555-333-3333' ]
}]

```

Name		Age	Phones		
Offset	Values	Values	List Offset	Str. Offset	Values
0	J	18	0	0	5
3	o	37	2	12	5
7	e		3	24	-
	J		4	36	1
	a				1
	c				-
	k				1
					1
					1
					5
					5
					5
					-
					2
					2
					2
					-
					...



Record Batch Construction



Each box (vector) is contiguous memory
The entire record batch is contiguous on wire

In memory columnar format for speed

- Maximize CPU throughput
 - Pipelining
 - SIMD
 - cache locality
- Scatter/gather I/O




Results


- PySpark Integration:
53x speedup (IBM spark work on SPARK-13534)
<http://s.apache.org/arrowresult1>
- Streaming Arrow Performance
7.75GB/s data movement
<http://s.apache.org/arrowresult2>
- Arrow Parquet C++ Integration
4GB/s reads
<http://s.apache.org/arrowresult3>
- Pandas Integration
9.71GB/s
<http://s.apache.org/arrowresult4>




Arrow Releases



Improvements to PySpark with Arrow



How PySpark UDF works



Current Issues with UDF

- Serialize / Deserialize in Python
- Scalar computation model (Python for loop)



Profile lambda x: x+1

Actual Runtime is 2s without profiling.
8 Mb/s

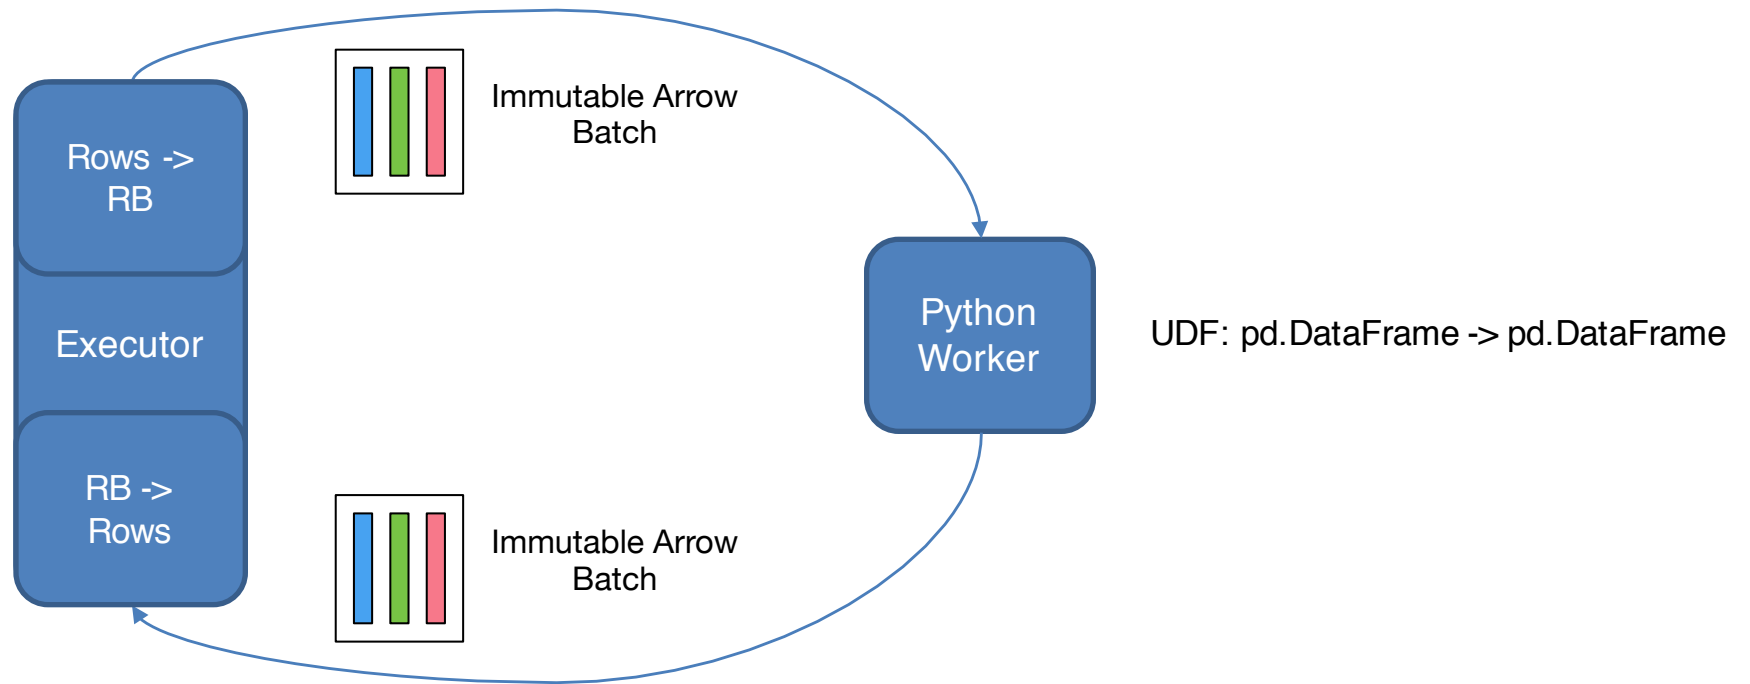
8787091 function calls in 4.084 seconds

Ordered by: internal time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
20973	1.296	0.000	3.820	0.000	serializers.py:223(_batched)
2097152	0.800	0.000	2.004	0.000	worker.py:107(<lambda>)
2097152	0.761	0.000	1.204	0.000	worker.py:72(<lambda>)
2097152	0.443	0.000	0.443	0.000	<ipython-input-2-853f857cd265>:14(<lambda>)
2097152	0.214	0.000	0.214	0.000	{method 'append' of 'list' objects}
20972	0.153	0.000	0.153	0.000	{built-in method _pickle.loads}
20972	0.086	0.000	0.086	0.000	{built-in method _pickle.dumps}
20972	0.045	0.000	0.045	0.000	{method 'write' of '_io.BufferedWriter' objects}
41944	0.044	0.000	0.287	0.000	serializers.py:161(_read_with_length)
41945	0.039	0.000	0.039	0.000	{method 'read' of '_io.BufferedReader' objects}
1	0.034	0.034	4.084	4.084	serializers.py:137(dump_stream)
20973	0.021	0.000	0.039	0.000	serializers.py:598(read_int)
20972	0.020	0.000	0.042	0.000	serializers.py:605(write_int)
20973	0.020	0.000	0.306	0.000	serializers.py:141(load_stream)
20972	0.019	0.000	0.172	0.000	serializers.py:474(loads)
20972	0.017	0.000	0.103	0.000	serializers.py:470(dumps)
62916	0.011	0.000	0.011	0.000	{built-in method builtins.len}
20972	0.009	0.000	0.009	0.000	{built-in method _struct.pack}
20973	0.008	0.000	0.008	0.000	{built-in method _struct.unpack}
1	0.000	0.000	0.000	0.000	serializers.py:246(load_stream)
1	0.000	0.000	4.084	4.084	serializers.py:243(dump_stream)
1	0.000	0.000	4.084	4.084	worker.py:217(process)
1	0.000	0.000	0.000	0.000	serializers.py:249(_load_stream_without_unbatching)
1	0.000	0.000	0.000	0.000	worker.py:121(<lambda>)
1	0.000	0.000	0.000	0.000	{built-in method builtins.hasattr}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}
1	0.000	0.000	0.000	0.000	{built-in method from_iterable}

91.8%

Vectorize Row UDF



Why pandas.DataFrame

- Fast, feature-rich, widely used by Python users
- Already exists in PySpark (toPandas)
- Compatible with popular Python libraries:
 - NumPy, StatsModels, SciPy, scikit-learn...
- Zero copy to/from Arrow



Scalar vs Vectorized UDF

Actual Runtime is **2s** without profiling

```
8787091 function calls in 4.084 seconds
Ordered by: internal time
ncalls  tottime  percall  cumtime  percall filename:lineno(function)
20973   1.296    0.000    3.820    0.000 serializers.py:223(<batched>)
2097152  0.800    0.000    2.004    0.000 worker.py:107(<lambda>)
2097152  0.761    0.000    1.204    0.000 worker.py:72(<lambda>)
2097152  0.443    0.000    0.443    0.000 <ipython-input-2-853f857cd265>:14(<lambda>)
2097152  0.214    0.000    0.214    0.000 {method 'append' of 'list' objects}
20972    0.153    0.000    0.153    0.000 {built-in method _pickle.loads}
20972    0.086    0.000    0.086    0.000 {built-in method _pickle.dumps}
20972    0.046    0.000    0.230    0.000 serializers.py:148(_write_with_length)
41944    0.045    0.000    0.045    0.000 {method 'write' of '_io.BufferedWriter' objects}
20973    0.044    0.000    0.287    0.000 serializers.py:161(_read_with_length)
41945    0.039    0.000    0.039    0.000 {method 'read' of '_io.BufferedReader' objects}
```

20x Speed Up

```
1245 function calls (1226 primitive calls) in 0.092 seconds
Ordered by: internal time
ncalls  tottime  percall  cumtime  percall filename:lineno(function)
3       0.013    0.004    0.013    0.004 {method 'read' of '_io.BufferedReader' objects}
2       0.012    0.006    0.012    0.006 {method 'write' of '_io.BufferedWriter' objects}
1       0.012    0.012    0.012    0.012 {built-in method _operator.add}
2       0.011    0.006    0.011    0.006 {method 'copy' of 'numpy.ndarray' objects}
1       0.011    0.011    0.012    0.012 {method 'to_pandas' of 'pyarrow._table.RecordBatch' objects}
1       0.009    0.009    0.009    0.009 {built-in method from_pandas}
1       0.006    0.006    0.006    0.006 {method 'get_result' of 'pyarrow._io.InMemoryOutputStream' objects}
1       0.006    0.006    0.006    0.006 {method 'to_pybytes' of 'pyarrow._io.Buffer' objects}
1       0.005    0.005    0.005    0.005 {method 'write_batch' of 'pyarrow._io._StreamWriter' objects}
1       0.003    0.003    0.003    0.003 internals.py:329(set)
```

Scalar vs Vectorized UDF

```
8787091 function calls in 4.084 seconds
Ordered by: internal time
```

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
20973	1.296	0.000	3.820	0.000	serializers.py:223(<batched>)
2097152	0.800	0.000	2.004	0.000	worker.py:107(<lambda>)
2097152	0.761	0.000	1.204	0.000	worker.py:72(<lambda>)
2097152	0.443	0.000	0.443	0.000	<ipython-input-2-853f857cd265>:14(<lambda>)
2097152	0.214	0.000	0.214	0.000	{method 'append' of 'list' objects}
20972	0.153	0.000	0.153	0.000	{built-in method '_pickle.loads'}
20972	0.086	0.000	0.086	0.000	{built-in method '_pickle.dumps'}
20972	0.046	0.000	0.230	0.000	serializers.py:148(write with length)
41944	0.045	0.000	0.045	0.000	{method 'write' of '_io.BufferedWriter' objects}
20973	0.044	0.000	0.287	0.000	serializers.py:161(read with length)
41945	0.039	0.000	0.039	0.000	{method 'read' of '_io.BufferedReader' objects}

Overhead
Removed

```
1245 function calls (1226 primitive calls) in 0.092 seconds
Ordered by: internal time
```

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
3	0.013	0.004	0.013	0.004	{method 'read' of '_io.BufferedReader' objects}
2	0.012	0.006	0.012	0.006	{method 'write' of '_io.BufferedWriter' objects}
1	0.012	0.012	0.012	0.012	{built-in method '_operator.add'}
2	0.011	0.006	0.011	0.006	{method 'copy' of 'numpy.ndarray' objects}
1	0.011	0.011	0.012	0.012	{method 'to_pandas' of 'pyarrow._table.RecordBatch' objects}
1	0.009	0.009	0.009	0.009	{built-in method 'from_pandas'}
1	0.006	0.006	0.006	0.006	{method 'get_result' of 'pyarrow._io.InMemoryOutputStream' objects}
1	0.006	0.006	0.006	0.006	{method 'to_pybytes' of 'pyarrow._io.Buffer' objects}
1	0.005	0.005	0.005	0.005	{method 'write_batch' of 'pyarrow._io._StreamWriter' objects}
1	0.003	0.003	0.003	0.003	internals.py:329(set)

Scalar vs Vectorized UDF

8787091 function calls in 4.084 seconds

Ordered by: internal time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
20973	1.296	0.000	3.820	0.000	serializers.py:223(_batched)
2097152	0.800	0.000	2.004	0.000	worker.py:107(<lambda>)
2097152	0.761	0.000	1.204	0.000	worker.py:72(<lambda>)
2097152	0.443	0.000	0.443	0.000	<ipython-input-2-853f857cd265>:14(<lambda>)
2097152	0.214	0.000	0.214	0.000	{method 'append' of 'list' objects}
20972	0.153	0.000	0.153	0.000	{built-in method _pickle.loads}
20972	0.086	0.000	0.086	0.000	{built-in method _pickle.dumps}
20972	0.046	0.000	0.230	0.000	serializers.py:148(write_with_length)
41944	0.045	0.000	0.045	0.000	{method 'write' of '_io.BufferedWriter' objects}
20973	0.044	0.000	0.287	0.000	serializers.py:161(_read_with_length)
41945	0.039	0.000	0.039	0.000	{method 'read' of '_io.BufferedReader' objects}

Less System Call
Faster I/O


1245 function calls (1226 primitive calls) in 0.092 seconds

Ordered by: internal time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
3	0.013	0.004	0.013	0.004	{method 'read' of '_io.BufferedReader' objects}
2	0.012	0.006	0.012	0.006	{method 'write' of '_io.BufferedWriter' objects}
1	0.012	0.012	0.012	0.012	{built-in method _operator.add}
2	0.011	0.006	0.011	0.006	{method 'copy' of 'numpy.ndarray' objects}
1	0.011	0.011	0.012	0.012	{method 'to_pandas' of 'pyarrow._table.RecordBatch' objects}
1	0.009	0.009	0.009	0.009	{built-in method from_pandas}
1	0.006	0.006	0.006	0.006	{method 'get_result' of 'pyarrow._io.InMemoryOutputStream' objects}
1	0.006	0.006	0.006	0.006	{method 'to_pybytes' of 'pyarrow._io.Buffer' objects}
1	0.005	0.005	0.005	0.005	{method 'write_batch' of 'pyarrow._io._StreamWriter' objects}
1	0.003	0.003	0.003	0.003	internals.py:329(set)



Scalar vs Vectorized UDF



Support Group UDF

- Split-apply-combine:
 - Break a problem into smaller pieces
 - Operate on each piece independently
 - Put all pieces back together
- Common pattern supported in SQL, Spark, Pandas, R ...



Split-Apply-Combine (Current)

- Split: groupBy, window, ...
- Apply: mean, stddev, collect_list, rank ...
- Combine: Inherently done by Spark



Split-Apply-Combine (with Group UDF)

- Split: groupBy, window, ...
- Apply: UDF
- Combine: Inherently done by Spark




Introduce `groupBy().apply()`

- UDF: `pd.DataFrame -> pd.DataFrame`
 - Treat each group as a pandas DataFrame
 - Apply UDF on each group
 - Assemble as PySpark DataFrame



Introduce `groupBy().apply()`



Previous Example: Data Normalization

$$(values - values.mean()) / values.std()$$



Previous Example: Data Normalization

Current:

```
group_columns = ['year', 'month']
non_group_columns = [col for col in df.columns if col not in group_columns]
s = StructType([f for f in df.schema.fields if f.name in non_group_columns])
cols = list([F.col(name) for name in non_group_columns])

df_norm = df.withColumn('values', F.struct(*cols))
df_norm = (df_norm.groupBy('year', 'month')
            .agg(F.collect_list(df_norm.values).alias('values'))

s2 = StructType(s.fields + [StructField('v3', DoubleType())])
@udf(ArrayType(s2))
def normalize(values):
    v1 = pd.Series([r.v1 for r in values])
    v1_norm = (v1 - v1.mean()) / v1.std()
    return [values[i] + (float(v1_norm[i]),) for i in range(0, len(values))]

df_norm = (df_norm.withColumn('new_values', normalize(df_norm.values))
            .drop('values')
            .withColumn('new_values', F.explode(F.col('new_values'))))

for col in [f.name for f in s2.fields]:
    df_norm = df_norm.withColumn(col, F.col('new_values.{0}'.format(col)))

df_norm = df_norm.drop('new_values')
```

Group UDF:

```
schema = StructType(df.schema.fields + [StructField('v3', DoubleType())])

def normalize(df):
    v1 = df.v1
    df['v3'] = (v1 - v1.mean()) / v1.std()
    return df

df_norm = (df.groupby('year', 'month')
            .apply(F.UserDefinedFunction(normalize, schema)))
```

5x Speed Up




Limitations

- Requires Spark Row <-> Arrow RecordBatch conversion
 - Incompatible memory layout (row vs column)
- (groupBy) No local aggregation
 - Difficult due to how PySpark works. See <https://issues.apache.org/jira/browse/SPARK-10915>



Future Roadmap



What's Next (Arrow)

- Arrow RPC/REST
- Arrow IPC
- Apache {Spark, Drill, Kudu} to Arrow Integration
 - Faster UDFs, Storage interfaces



What's Next (PySpark UDF)

- Continue working on SPARK-20396
- Support Pandas UDF with more PySpark functions:
 - `groupBy().agg()`
 - `window`



What's Next (PySpark UDF)

```
import numpy as np
@pandas_udf(Scalar, DoubleType())
def weighted_mean_udf(v1, w):
    return np.average(v1, weights=w)

df.groupBy('id').agg(weighted_mean_udf(df.v1, df.w).as('v1_wm'))
```

```
w = Window.partitionBy('id')

@pandas_udf(Series, DoubleType())
def rank_udf(v):
    return v.rank(pct=True)

df.withColumn('rank', rank_udf(df.v).over(w))
```

Get Involved

- Watch SPARK-20396
- Join the Arrow community
 - dev@arrow.apache.org
 - Slack:
 - <https://apachearrowslackin.herokuapp.com/>
 - http://arrow.apache.org
 - Follow @ApacheArrow



Thank you

- Bryan Cutler (IBM), Wes McKinney (Two Sigma Investments) for helping build this feature
- Apache Arrow community
- Spark Summit organizers
- Two Sigma and Dremio for supporting this work

