

Improving Spark's Reliability with DataSourceV2

Ryan Blue
Spark Summit 2019

NETFLIX

Data at Netflix



Cloud-native data warehouse

- YARN compute clusters are expendable
- Expendable clusters require architectural changes
 - **GENIE** is a job submission service that selects the cluster
 - **METACAT** is a cluster-independent metastore
 - **S3** is the source of truth for data

S3 is eventually consistent

- File list calls may be inaccurate
- Hive tables rely on accurate listing for correctness
- S3 queries may be incorrect, sometimes

S3 is eventually consistent

- File list calls may be inaccurate
- Hive tables rely on accurate listing for correctness
- S3 queries may be incorrect, **sometimes**

**At Netflix's scale,
sometimes is every day.**

A reliable S3 warehouse (in 2016)

- Requires consistent listing – **S3MPER**
- Requires in-place writes – **BATCH PATTERN**
- Requires atomic metastore changes – **METACAT**

Changes needed in Spark

- Integrate S3 batch pattern committers
- Spark versions
 - **1.6** – Hive path only
 - **2.0** – DataSource path for reads, not writes
 - **2.1+** – Use DataSource path for reads and writes

Problems and Roadblocks



DataFrameWriter

- Behavior is not defined
- What do `save` and `saveAsTable` do differently?
 - Create different logical plans . . .
that are converted to other logical plans
- When you use “overwrite” mode, what happens?
 - Depends on the data source

SaveMode

- Delegates behavior to the source when tables don't exist
- Overwrite might mean:
 - Replace table – data and metadata (Some code paths)
 - Replace all table data (Some code paths)
 - Replace static partitions (DataSource tables)
 - Replace dynamic partitions (Hive tables, SPARK-20236)

Validation

- What is “correct” for CTAS/overwrite when the table exists?
- `PreprocessTableCreation` vs `PreprocessTableInsertion`
 - Depends on the `DataFrameWriter` call
- Spark automatically inserts unsafe casts (e.g. string to int)
- Path tables have no schema validation on write


“[These] **should do the same thing,**
but as we've already published
these 2 interfaces and the
implementations **may have different**
logic, we have to keep these 2
different commands.”

“[These] **should do the same thing**, but as we’ve already published these 2 interfaces and the implementations **may have different logic**, we have to keep these 2 **different commands.**”



Commands

- RunnableCommand wraps a logical in a pseudo-physical plan
- Commands created inside run made it worse

 2.1.1-nflx-180

JobsStagesStorageEnvironmentExecutorsSQL

Details for Query 4

Submitted Time: 2019/04/22 22:40:53
Duration: 0.7 s

CreateDataSourceTableAsSelectCommand

▼ Details

```
== Parsed Logical Plan ==
'CreateTableAsSelect `rblue`.`test_ctas`, parquet, false
+- 'Project [*]
   +- 'UnresolvedRelation `rblue`.`test_table`

== Analyzed Logical Plan ==
CreateTable CatalogTable(Table: `rblue`.`test_ctas`), ErrorIfExists
+- Project [id#8L, data#9]
   +- SubqueryAlias test_table
      +- Relation[id#8L,data#9] parquet

== Optimized Logical Plan ==
CreateTable CatalogTable(Table: `rblue`.`test_ctas`), ErrorIfExists
+- Project [id#8L, data#9]
   +- SubqueryAlias test_table
      +- Relation[id#8L,data#9] parquet

== Physical Plan ==
CreateDataSourceTableAsSelectCommand CatalogTable(Table: `rblue`.`test_ctas`), ErrorIfExists
+- Project [id#8L, data#9]
   +- SubqueryAlias test_table
      +- Relation[id#8L,data#9] parquet
```

Community Roadblocks

- Substantial behavior changes for 2.0
 - Committed with no time to review
 - ... to the 2.0 release branch
- Behavior not up for discussion
- Parts of PRs merged without attribution

Iceberg and DataSourceV2



A reliable S3 warehouse (in 2019)

- **Iceberg:** tables without unpleasant surprises
- Fix tables, not the file system
- While fixing reliability and scale, fix usability:
 - Reliable schema evolution
 - Automatic partitioning
 - Configure tables, not jobs

Last year

- Need a way to plug in Iceberg cleanly
- Maintaining a separate write path takes time
- Spark's write path had solidified
- DataSourceV2 was proposed . . .

Why DataSourceV2?

- Isn't v2 just an update to the read/write API?
- Existing design problems also affect v2
 - No write validation – yet another logical plan
 - SaveMode passed to sources
- **Opportunity:** avoid needing v3 to fix behavior

What's different in DSv2

- Define a set of common logical plans
 - CTAS, RTAS, Append, OverwriteByExpression, etc.
 - Document user expectations and behavior
 - Implement consistent behavior in Spark for all v2 sources

- SPIP: Standardize SQL logical plans

<https://issues.apache.org/jira/browse/SPARK-23521>

Standard Logical Plans

- Specialize physical plans, not logical plans
 - No more `InsertIntoDataSourceTable` and `InsertIntoHiveTable`
 - No forgetting to apply rules to a new logical plan
- Apply validation rules universally
 - Same rules for Append and Overwrite
- Avoid using `RunnableCommand`

Consistent behavior

- Create, alter, and drop tables in Spark, not sources
 - CTAS when table exists: fail the query in Spark
 - Requires a catalog plugin API

- SPIP: Spark API for Table Metadata

<https://issues.apache.org/jira/browse/SPARK-27067>

Catalog API

- Multi-catalog support
 - Create tables in the source of truth
 - Avoiding this caused strange Spark behavior

- SPIP: Identifiers for multi-catalog support

<https://issues.apache.org/jira/browse/SPARK-27066>

Status

- **Goal:** working DSv2 in Spark 3.0
 - Independent of the v1 path
 - Default behavior to v1
- SPIPs have been adopted by community votes
- Append and overwrite plans are added and working
- Waiting on catalog API to add CTAS and DDL

Thank you!

Questions?

NETFLIX

Up next: **Migrating to Spark at Netflix**
At **11:50** today, in **Room 2006**